

# MongoDB – A document data store



# MongoDB in short

Feature	MongoDB
Model	Document-based
Query language	Supported, aggregation framework
Reference scenarios	Transactional (read/write intensive) & analytical (read intensive)
Partitioning (Sharding)	Hash-based & range-based (not all the collections are sharded)
Indexes	Primary, secondary, multiattr, fulltext
Replication	Master-slave, replica set
Consistency	Strong, eventual at replicas
Availability	Can be mediated with consistency, through r/w concerns
Fault tolerance	By replica set, system remains operational on failing nodes
Transactions	ACID transactions (multidocument since 4.0) (read concern)
CAP theorem	CP
Distributed by	MongoDB Inc.

# Mongo DB



- History:
  - Development started by 10gen in 2007
  - Open sourced in 2009
  - In 2013, 10gen became MongoDB Inc.
- Uses BSON format
  - Based on JSON –B stands for Binary
- Written in C++, C, Javascript
- Supports APIs (drivers) in many languages
  - JavaScript, Python, Ruby, Perl, Java, JavaScala, C#, C++, Haskell, Erlang, ...

# Who uses mongoDB?



Uber



Lyft



LaunchDarkly



Stack



Delivery Hero



ViaVarejo



Bepro  
Company



Accenture



CircleCI

# MongoDB

- MongoDB **Compass** (GUI to explore and manipulate documents)

- MongoDB **Atlas** (on cloud)



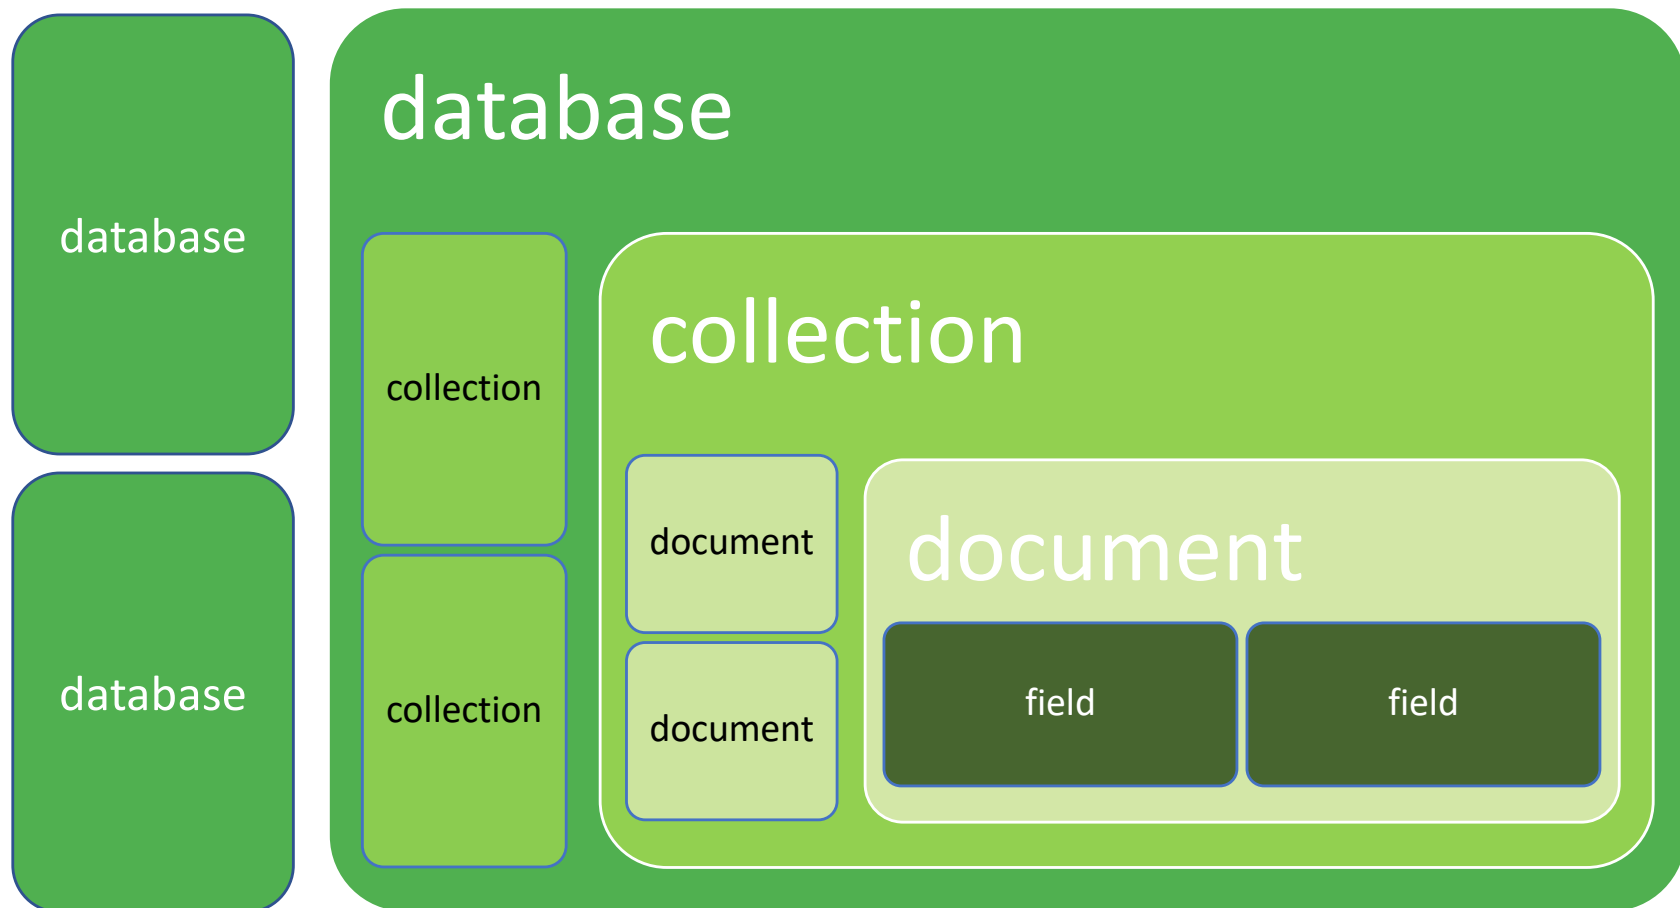
- Supported in Azure Cosmos DB



# Data Model and Interaction

# MongoDB: Basics

- Mongoddb instance



# MongoDB: Basics

Oracle	MongoDB
database instance	MongoDB instance
schema	database
table	collection
row	document
rowid	_id
join	embedded documents, \$lookup



# MongoDB: Basics

- Each MongoDB instance has multiple databases
  - Similar to a database schema in a RDBMS
- Each database can have multiple collections
  - Similar to a table in a RDBMS
- When we store a document, we have to choose which database and collection this document belongs

`db.collection.insertOne(document)`

- Document identifier `_id` will be created for each document, field name reserved by system

# Schema validation

A collection may be associated with a schema in JSON schema

```
db.createCollection("students", {
  validator: {
    $jsonSchema: {
      bsonType: "object",
      required: [ "name", "year", "major", "address" ],
      properties: {
        name: {
          bsonType: "string",
          description: "must be a string and is required"
        },
        year: {
          bsonType: "int",
          minimum: 2017,
          maximum: 3017,
          description: "must be an integer in [ 2017, 3017 ] and is required"
        },
        major: {
          enum: [ "Math", "English", "Computer Science", "History", null ],
          description: "can only be one of the enum values and is required"
        },
        gpa: {
          bsonType: [ "double" ],
```

# MongoDB query language

- `db.collection.find( <query>, <projection> )`
- Provides functionality similar to the SELECT command
  - `<query>` where condition
  - `<projection>` fields in result set

`<field>: <1 or true>`

Specifies the inclusion of a field.

---

`<field>: <0 or false>`

Specifies the exclusion of a field.

# MongoDB query language

```
// in videos
{"loc": 1234,
  "type": "dvd",
  "rentals": [{"rental": {"rentalDate": "15/10/2021",
                        "codCli": 375657}}],
  "title": "pulp fiction",
  "director": "quentin tarantino"}
```

```
db.videos.find()
```

# MongoDB query language

```
// in videos
{"loc": 1234,
  "type": "dvd",
  "rentals": [{"rental": {"rentalDate": "15/10/2021",
                        "codCli": 375657}}],
  "title": "pulp fiction",
  "director": "quentin tarantino"}
```

```
db.videos.find({"loc":1234},{title:1})
```

where  
(filtering -  
selection)

Select  
(projection)

# MongoDB query language

```
// in videos
{"loc": 1234,
  "type": "dvd",
  "rentals": [{"rental": {"rentalDate": "15/10/2021",
                        "codCli": 375657}}],
  "title": "pulp fiction",
  "director": "quentin tarantino"}
```

```
db.videos.find({}, {title:1})
```

Select  
(projection)

Projection only, `_id` is included in the result  
if not explicitly excluded (i.e., `{_id: 0}`)

# MongoDB query language

```
// in videos
{"loc": 1234,
  "type": "dvd",
  "rentals": [{"rental": {"rentalDate": "15/10/2021",
                        "codCli": 375657}}],
  "title": "pulp fiction",
  "director": "quentin tarantino"}
```

```
db.videos.find({"title": "pulp fiction"})
```

**where**  
(filtering -  
selection)

# MongoDB query language

```
// in videos
{"loc": 1234,
  "type": "dvd",
  "rentals": [{"rental": {"rentalDate": "15/10/2021",
                        "codCli": 375657}}],
  "title": "pulp fiction",
  "director": "quentin tarantino"}
```

```
db.videos.find({"rentals.rental.codCli": "375657"})
```



# MongoDB query language

```
// in videos
```

```
{ "loc": 1234,
```

```
  "type": "dvd",
```

```
  "rentals": [ { "rental": { "rentalDate": "15/10/2021",  
                           "codCli": 375657 } } ],
```

```
  "title": "pulp fiction",
```

```
  "director": "quentin tarantino" }
```

```
db.videos.find( { "type": { $in, [ "dvd", "vhs" ] } } )
```

# MongoDB query language

// in movies

```
{ "id": 12345678,
  "title": "pulp fiction",
  "director": "quentin tarantino",
  "year": 1994,
  "genre": ["drama", "crime"],
  "recommended_by": [],
  "contained_in": [ { "video": { "loc": 1234,
                                "type": "dvd" }
                    } ]
}
```

`db.movies.find({"year" : {$geq: 1990}})`

# MongoDB – Query Operators

Name	Description
\$eq	Matches value that are equal to a specified value
\$gt, \$gte	Matches values that are greater than (or equal to a specified value
\$lt, \$lte	Matches values less than or ( equal to ) a specified value
\$ne	Matches values that are not equal to a specified value
\$in	Matches any of the values specified in an array
\$nin	Matches none of the values specified in an array
\$or	Joins query clauses with a logical OR returns all
\$and	Join query clauses with a logical AND
\$not	Inverts the effect of a query expression
\$nor	Join query clauses with a logical NOR
\$exists	Matches documents that have a specified field

<https://docs.mongodb.org/manual/reference/operator/query/>

# MongoDB query language

```
// in videos
```

```
{ "loc": 1234,
```

```
  "type": "dvd",
```

```
  "rentals": [ { "rental": { "rentalDate": "15/10/2021",  
                           "codCli": 375657 } } ],
```

```
  "title": "pulp fiction",
```

```
  "director": "quentin tarantino" }
```

```
db.videos.find( { "title": "pulp fiction",  
                  "director" : "gabriele salvatores" } )
```

```
db.videos.find( { $or: [ { "title": "pulp fiction",  
                           "director" : "gabriele salvatores" } ] } )
```

# MongoDB query language

- `db.collection.find( <query>, <projection> )`
- Provides functionality similar to the SELECT command
  - `<query>` where condition
  - `<projection>` fields in result set

`<field>: <1 or true>`

Specifies the inclusion of a field.

---

`<field>: <0 or false>`

Specifies the exclusion of a field.

- Return a `cursor` to handle a result set
- Can modify the query to impose limits, skips, and sort orders
- Can specify to return the 'top' number of records from the result set

- `db.collection.findOne( <query>, <projection> )`

# Sort

- `db.movies.find({"year" : {$geq: 1990}}).sort ({"title": 1})`
- `db.movies.find({"director" : "quentin tarantino"}).sort ({"year": -1})`

# Count

- `count()` or `find().count()`
- Can have the same arguments as `find`
- `db.movies.count({"year" : {$geq: 1990}})`

# Lookup

```
{
  $lookup:
  {
    from: <collection to join>,
    localField: <field from the input documents>,
    foreignField: <field from the documents of the "from" collection>,
    as: <output array field>
  }
}
```

```
SELECT *, <output array field>
FROM collection
WHERE <output array field> IN (
  SELECT *
  FROM <collection to join>
  WHERE <foreignField> = <collection.localField>
);
```



# Lookup

```
db.orders.insert([
  { "_id" : 1, "item" : "almonds", "price" : 12, "quantity" : 2 },
  { "_id" : 2, "item" : "pecans", "price" : 20, "quantity" : 1 },
  { "_id" : 3 }
])
```

```
db.inventory.insert([
  { "_id" : 1, "sku" : "almonds", "description": "product 1", "instock" : 120 },
  { "_id" : 2, "sku" : "bread", "description": "product 2", "instock" : 80 },
  { "_id" : 3, "sku" : "cashews", "description": "product 3", "instock" : 60 },
  { "_id" : 4, "sku" : "pecans", "description": "product 4", "instock" : 70 },
  { "_id" : 5, "sku": null, "description": "Incomplete" },
  { "_id" : 6 }
])
```

```
db.orders.aggregate([
  {
    $lookup:
    {
      from: "inventory",
      localField: "item",
      foreignField: "sku",
      as: "inventory_docs"
    }
  }
])
```

# Lookup

```
db.orders.insert([
  { "_id" : 1, "item" : "almonds", "price" : 10, "quantity" : 2 },
  { "_id" : 2, "item" : "pecans", "price" : 20, "quantity" : 1 },
  { "_id" : 3 }
])
```

```
db.inventory.insert([
  { "_id" : 1, "sku" : "almonds", "description" : "product 1", "instock" : 120 },
  { "_id" : 2, "sku" : "bread", "description" : "product 2", "instock" : 50 },
  { "_id" : 3, "sku" : "cashews", "description" : "product 3", "instock" : 30 },
  { "_id" : 4, "sku" : "pecans", "description" : "product 4", "instock" : 70 },
  { "_id" : 5, "sku" : null, "description" : "product 5", "instock" : 40 },
  { "_id" : 6 }
])
```

```
db.orders.aggregate([
  {
    $lookup:
    {
      from: "inventory",
      localField: "item",
      foreignField: "sku",
      as: "inventory_docs"
    }
  }
])
```

```
{
  "_id" : 1,
  "item" : "almonds",
  "price" : 10,
  "quantity" : 2,
  "inventory_docs" : [
    { "_id" : 1, "sku" : "almonds", "description" : "product 1", "instock" : 120 }
  ]
}
{
  "_id" : 2,
  "item" : "pecans",
  "price" : 20,
  "quantity" : 1,
  "inventory_docs" : [
    { "_id" : 4, "sku" : "pecans", "description" : "product 4", "instock" : 70 }
  ]
}
{
  "_id" : 3,
  "inventory_docs" : [
    { "_id" : 5, "sku" : null, "description" : "product 5", "instock" : 40 },
    { "_id" : 6 }
  ]
}
```

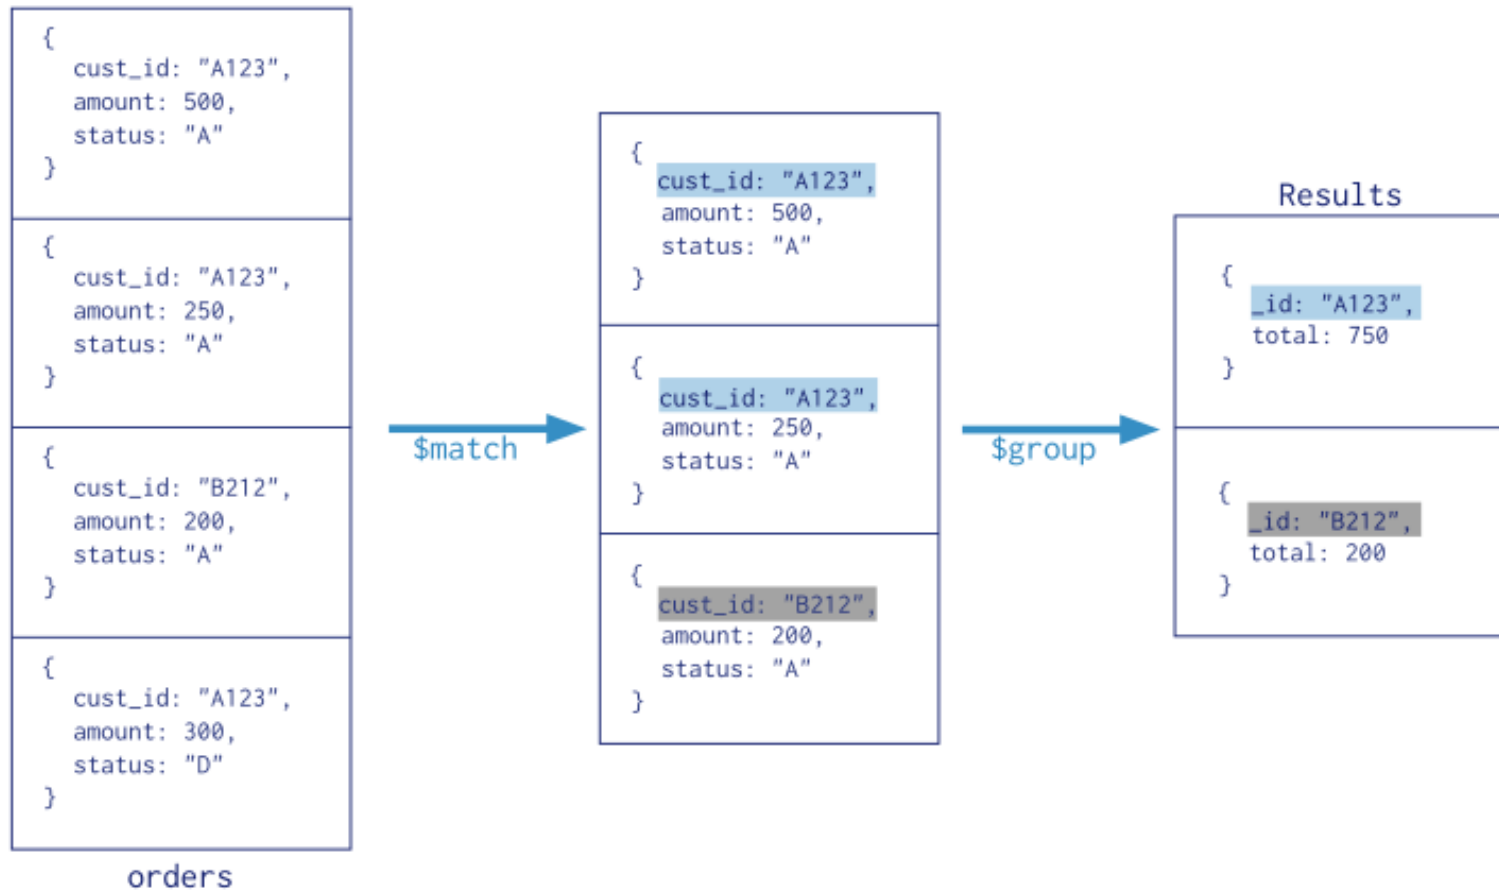
```
SELECT *, inventory_docs
FROM orders
WHERE inventory_docs IN (
  SELECT *
  FROM inventory
  WHERE sku = orders.item
);
```

# Aggregates

- Aggregation framework provides SQL-like aggregation functionality
- Documents from a collection pass through an aggregation pipeline, which transforms these objects as they pass through
- Expressions produce output documents based on calculations performed on input documents
- Example
- ```
db.orders.aggregate( { $group : { _id: type,
                                totalquantity: { $sum: quantity } } } )
```

# Aggregates

Collection  
↓  
db.orders.aggregate( [  
 \$match stage → { \$match: { status: "A" } },  
 \$group stage → { \$group: { \_id: "\$cust\_id", total: { \$sum: "\$amount" } } }  
] )



REMARK: partition\* is called  
shard\* in mongodb terminology

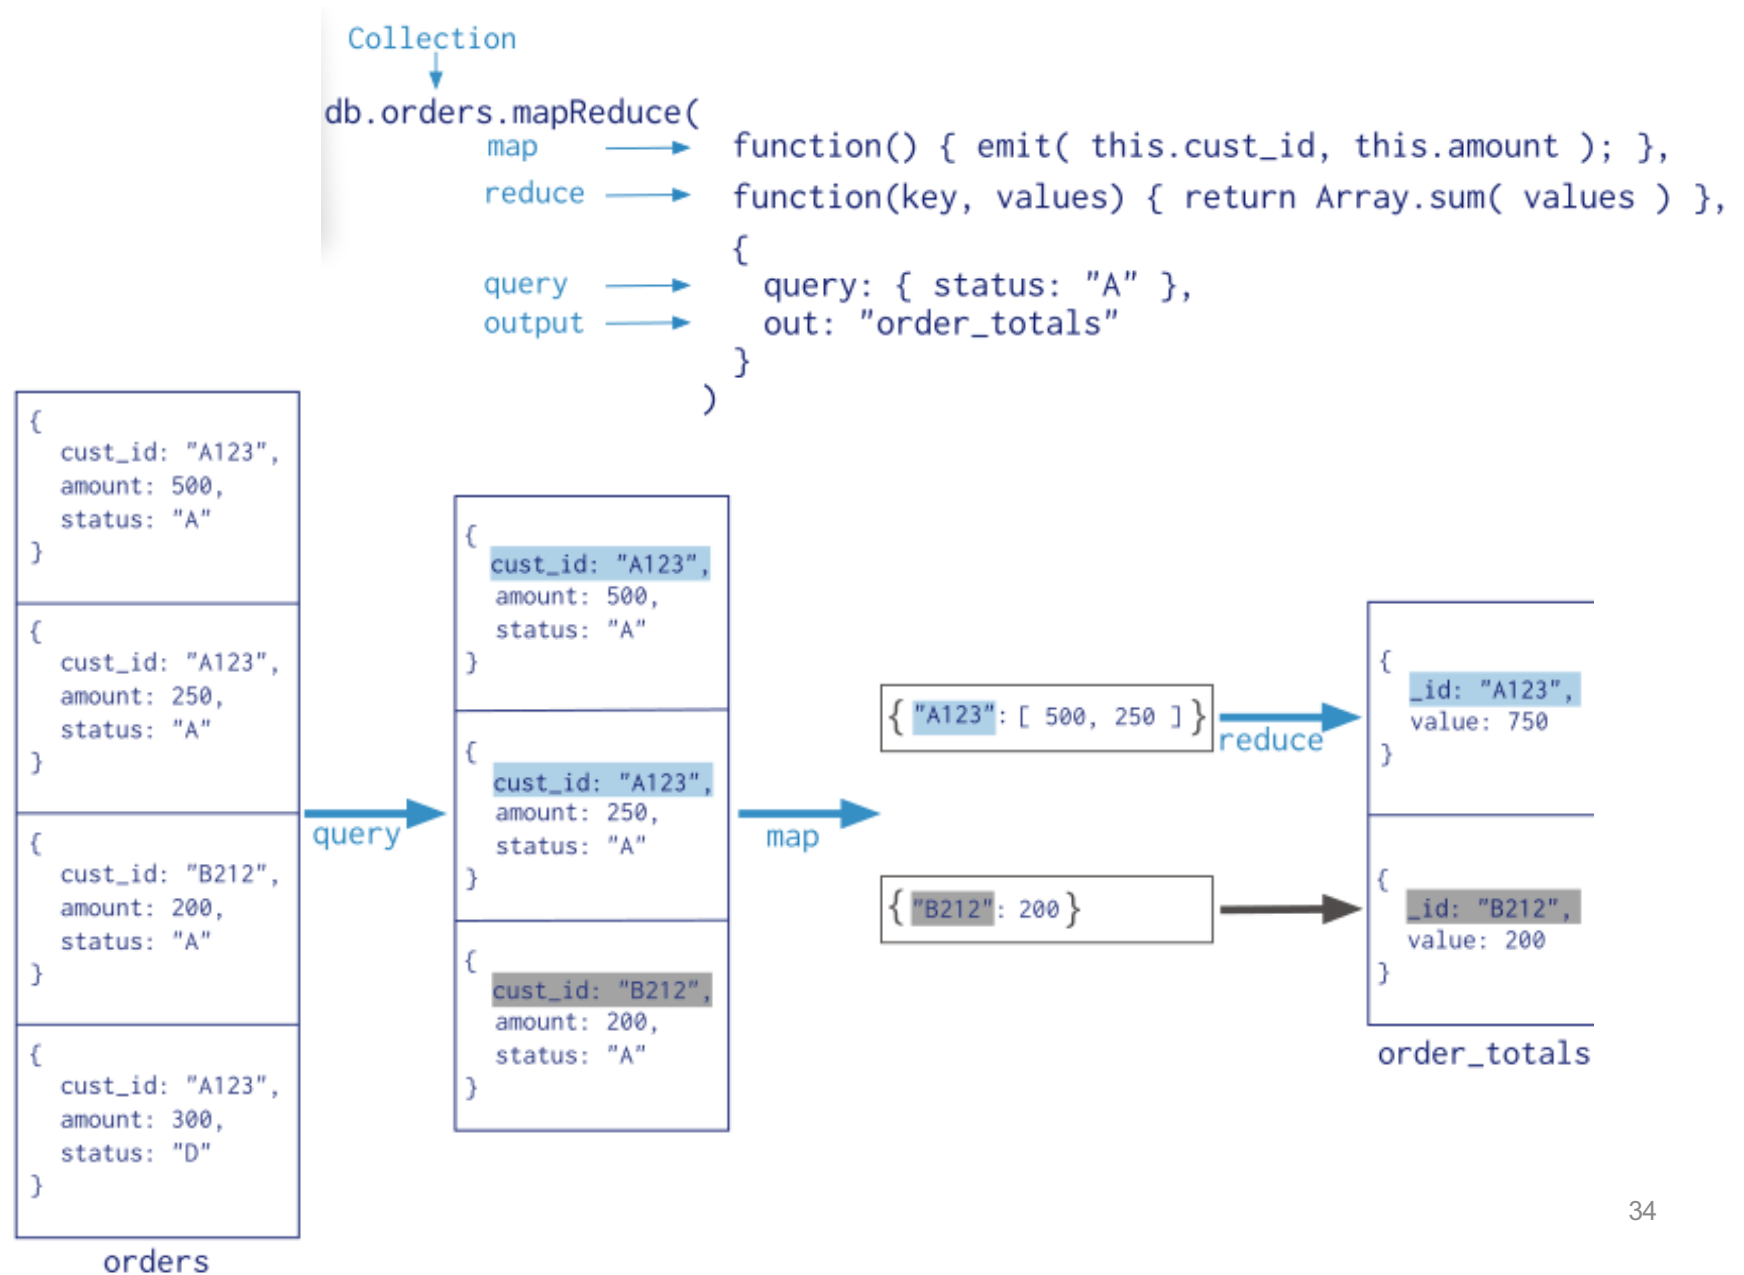
# Map Reduce

- Performs complex aggregator functions given a collection of keys, value pairs
- Must provide at least a map function, reduction function and a name of the result set

```
db.collection.mapReduce( <mapfunction>,  
                          <reducefunction>,  
                          { out: <collection>,  
                            query: <document>,  
                            sort: <document>,  
                            limit: <number>} )
```

- If the input is a partitioned collection, mongos will automatically dispatch the map-reduce job to each partition in parallel
- if the out field for mapReduce has the partitioning value, the output collection is partitioned using the `_id` field as the partition key

# Map Reduce



# CRUD Operations

All write operations in MongoDB are **atomic** on the level of a single **document**.

- Create

```
db.collection.insertOne( <document> )  
db.collection.insertMany( [<document>, <document>, ...] )
```

- Update

```
db.collection.update( <query>, <update>, <options> )  
... updateOne and updateMany
```

- Delete

```
db.collection.deleteMany( <query> )  
db.collection.deleteOne( <query> )
```

- Read

```
db.collection.find( <query>, <projection> )
```

# Indexing



# MongoDB - Indexes

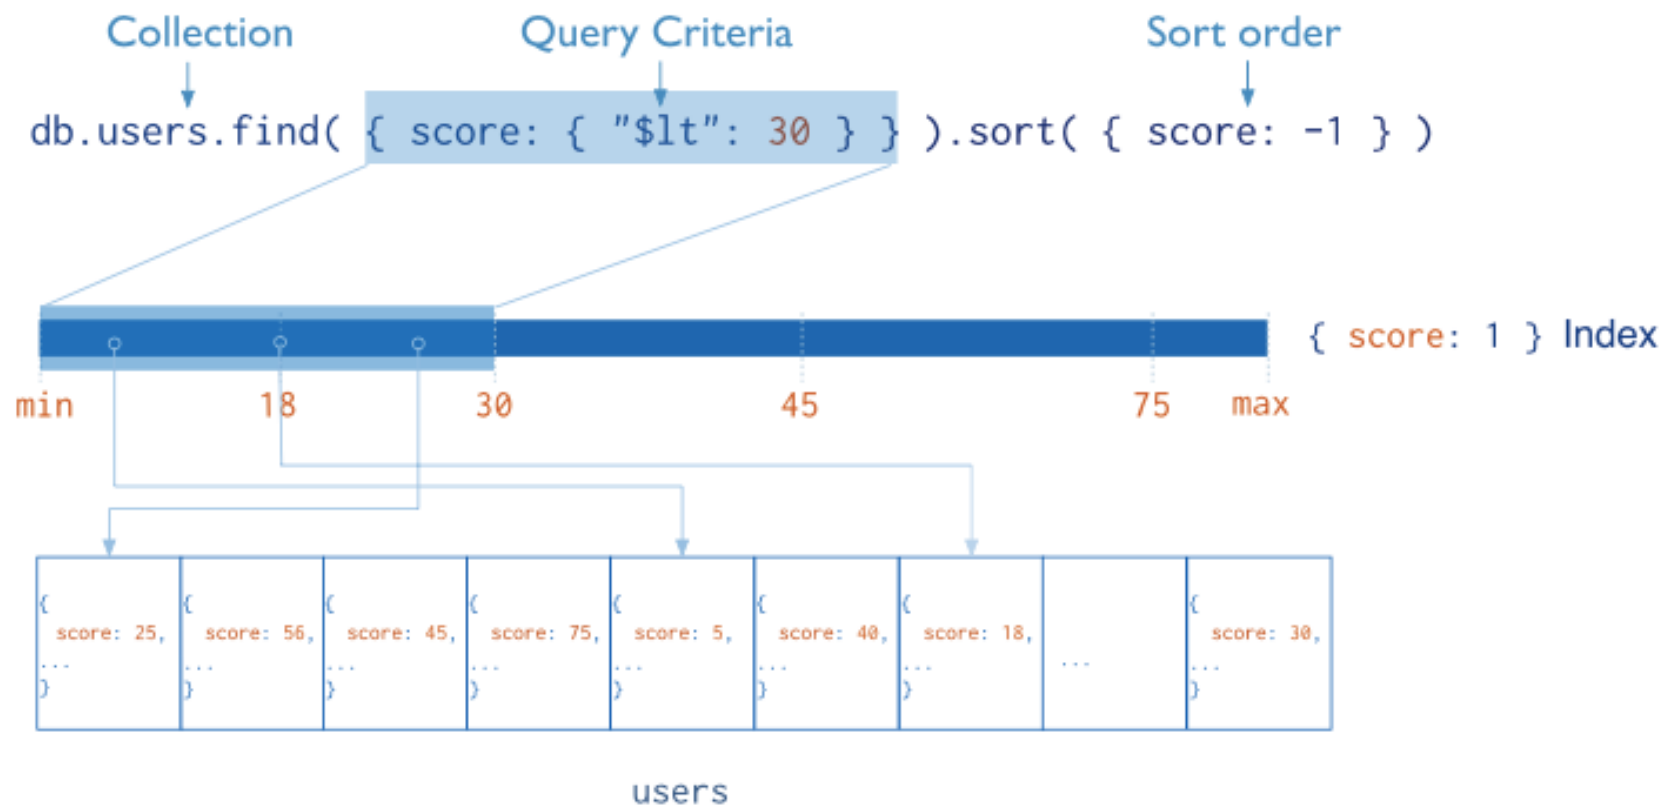
```
> db.clients.find(firstname: "alice")).explain()
{
  "cursor" : "BasicCursor",
  "nscanned" : 1000000,
  "nscannedObjects" : 1000000,
  "n" : 1,
  "millis" : 721,
  "nYields" : 0,
  "nChunkSkips" : 0,
  "isMultiKey" : false,
  "indexOnly" : false,
  "indexBounds" : {}
}
```

# MongoDB - Indexes

```
>db.users.ensureIndex({"firstname" : 1})
```

```
> db.clients.find(firstname: "alice").explain()
{
  "cursor" : "BtreeCursor username_1",
  "nscanned" : 1,
  "nscannedObjects" : 1,
  "n" : 1,
  "millis" : 3,
  "nYields" : 0,
  "nChunkSkips" : 0,
  "isMultiKey" : false,
  "indexOnly" : false,
  "indexBounds" : {"firstname" : [ ["alice", "alice" ] ] }
}
```

# Index functionalities



# Index functionalities

- An index is automatically created on the `_id` field (the primary key)
- Users can create other indexes
  - to improve query performance (filter conditions, sorting on the field) or
  - to enforce unique values for a particular field (`unique`)

```
db.users.ensureIndex({"username" : 1}, {"unique" : true})
```

# Index functionalities

- Supports single field index as well as compound index  
`db.users.ensureIndex({"age" : 1, "username" : 1})`
  - Like SQL: order of the fields in a compound index matters
  - If you index a field that holds an array value, MongoDB creates separate index entries for *every* element of the array
  - Also supports hash indexes
- 
- <https://docs.mongodb.com/manual/indexes/>

# Full-text Indexes and Search

- MongoDB has a special type of index for searching for text within documents
  - built-in support for multi-language stemming and stop words
- Heavyweight, be cautious
- `db.stores.createIndex( { name: "text", description: "text" } )`
- Text search
- `db.stores.find( { $text: { $search: "java coffee shop" } } )`
- It is possible to exclude words, and to sort on relevance ranking `textScore`

Architecture

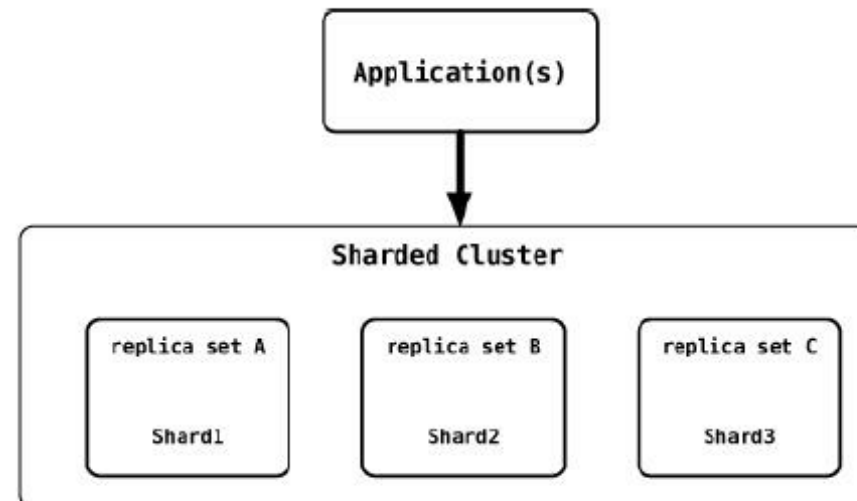
# MongoDB in short

| Feature                          | MongoDB                                                            |
|----------------------------------|--------------------------------------------------------------------|
| Model                            | Document-based                                                     |
| Query language                   | Supported, aggregation framework                                   |
| Reference scenarios              | Transactional (read/write intensive) & analytical (read intensive) |
| Partitioning ( <b>Sharding</b> ) | Hash-based & range-based (not all the collections are sharded)     |
| Indexes                          | Primary, secondary, multiattr, fulltext                            |
| Replication                      | <b>Master-slave</b> , replica set                                  |
| Consistency                      | Strong, eventual at replicas                                       |
| Availability                     | Can be mediated with consistency, through r/w concerns             |
| Fault tolerance                  | By replica set, system remains operational on failing nodes        |
| Transactions                     | ACID transactions (multidocument since 4.0)<br>(read concern)      |
| CAP theorem                      | CP                                                                 |
| Distributed by                   | MongoDB Inc.                                                       |



# Collection sharding/data partitioning

- By sharding/partitioning the data is split by certain field and moved to different nodes

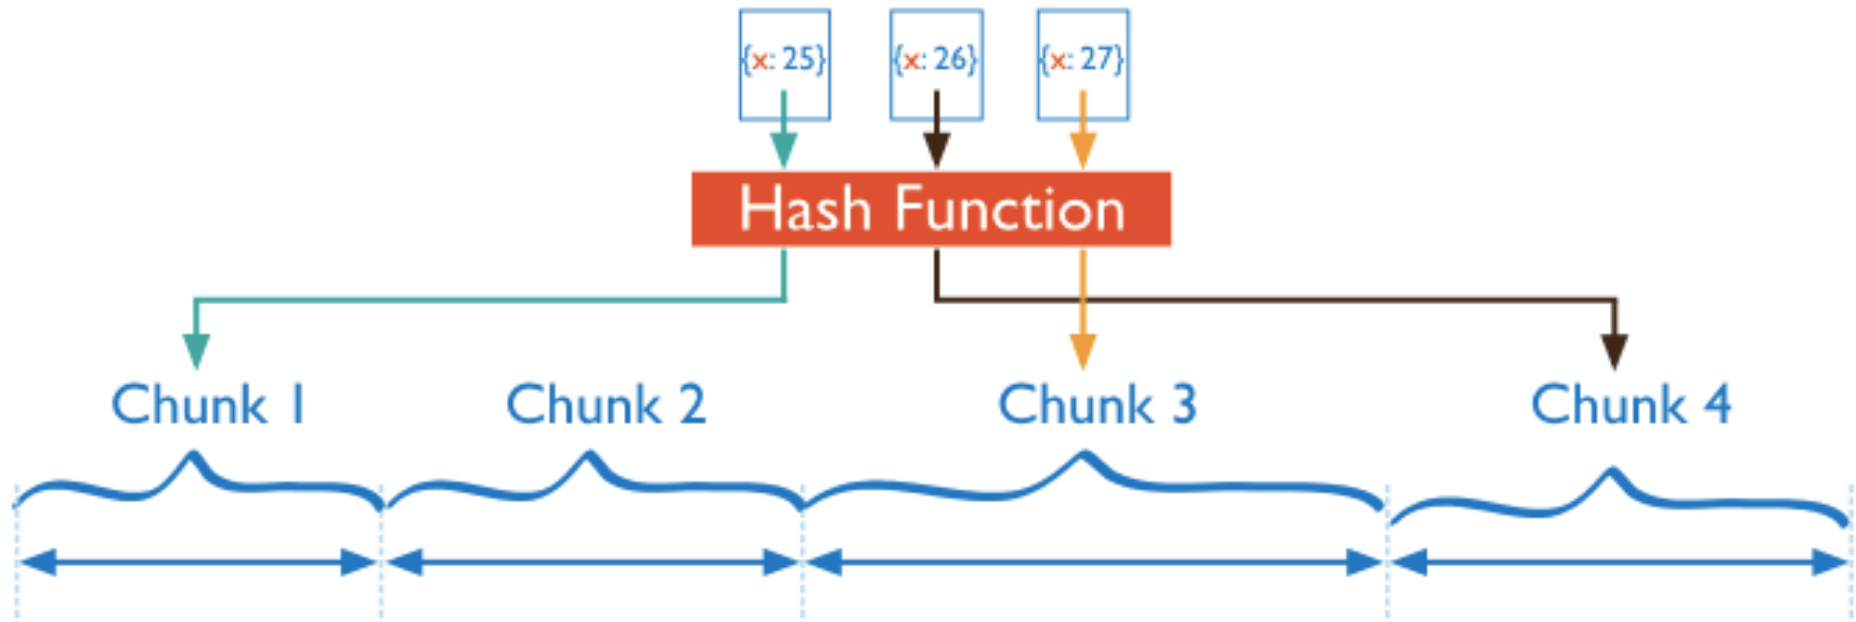


- Partitioning on the first name of the customer:

```
db.runCommand({shardcollection:"videorental.clients",  
               key:{firstname:1}})
```

# Partitioning / Sharding

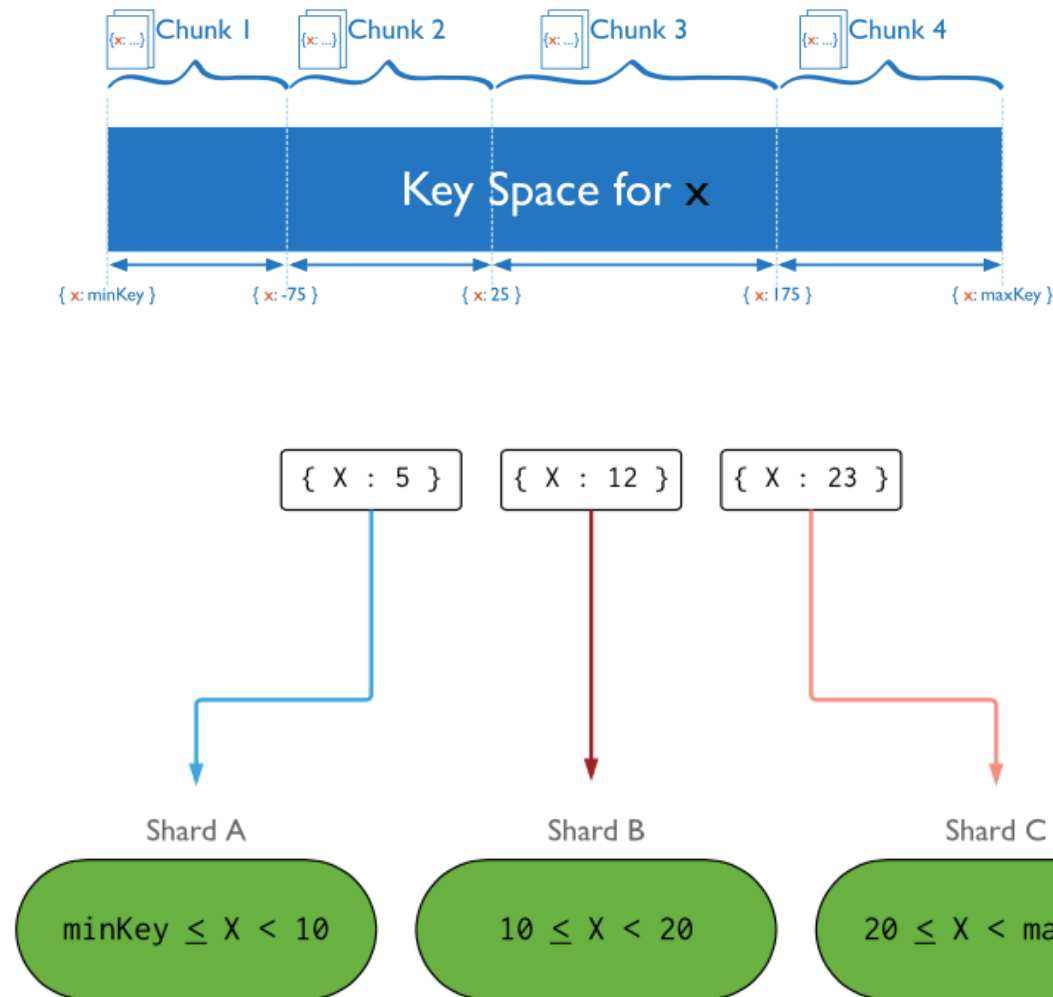
- Hash-based sharding



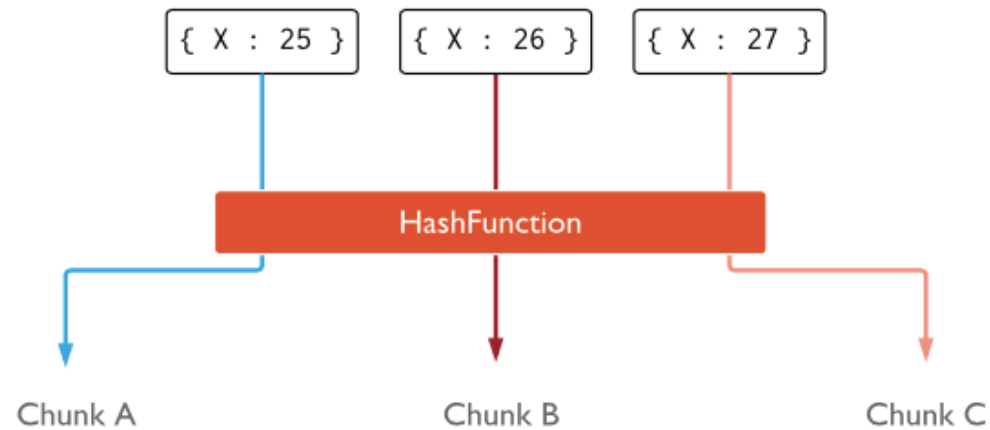
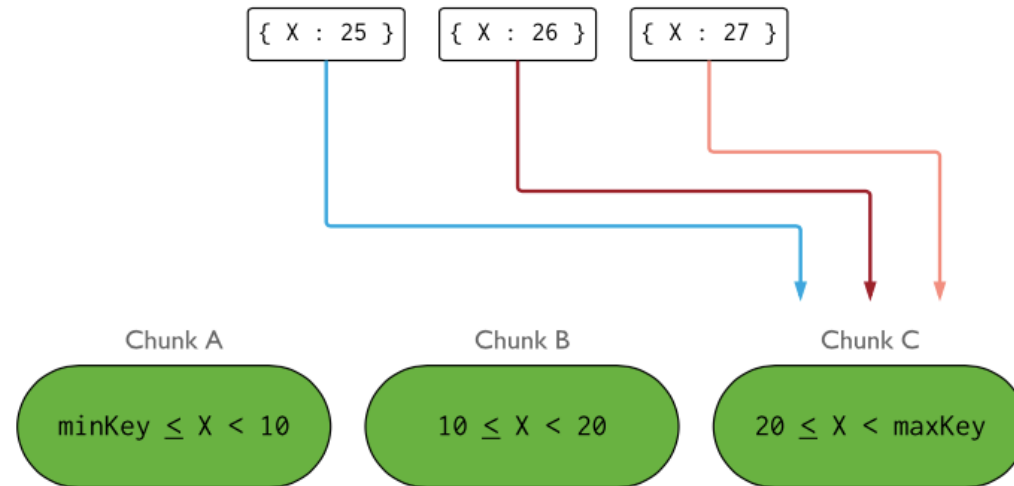
# Partitioning / Sharding

- Ranged sharding

Ranged sharding is most efficient  
when the shard key is  
Large Cardinality  
Low Frequency  
Non-Monotonically Changing




# Partitioning with monotonically increasing keys



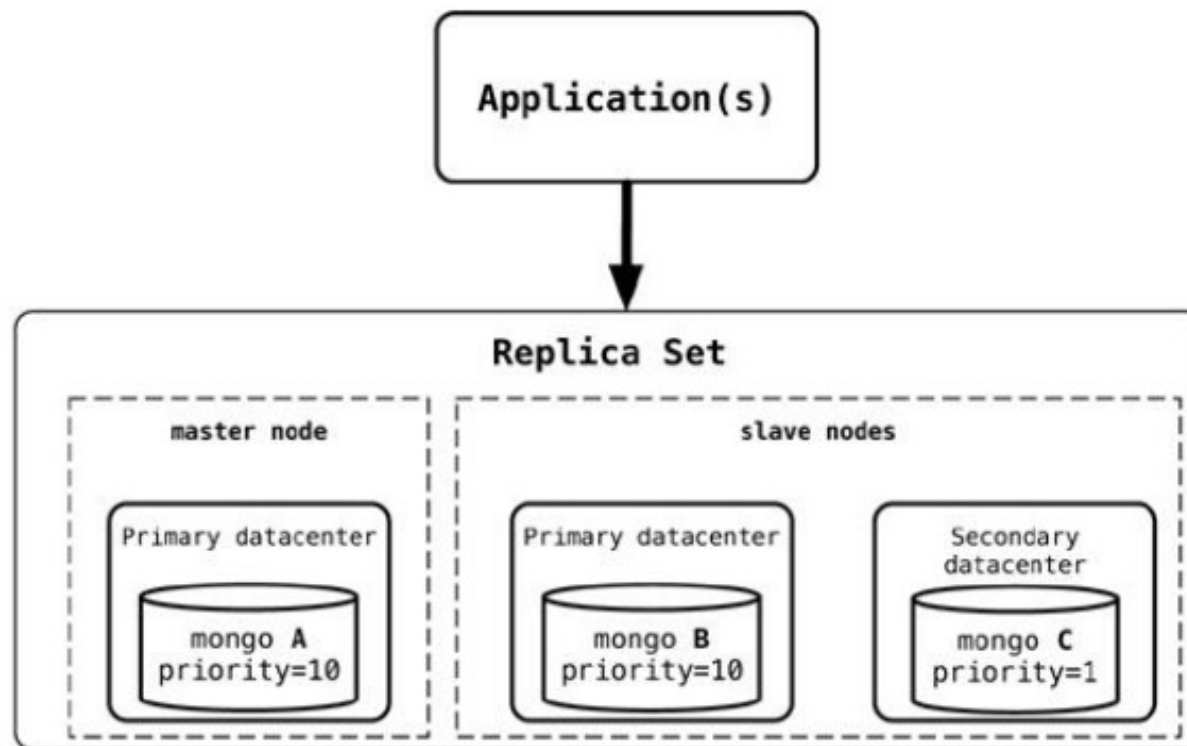
# CAP in MongoDB

- **Focus on Consistency and Partition tolerance**
- Consistency
  - all replicas contain the same version of the data
- Partition tolerance
  - multiple entry points
  - system remains operational on system split
- Availability
  - system remains operational on failing nodes
  - traded off with consistency



# Replication

- A MongoDB database makes use of **replica sets** for consistency and availability following a **master-slave** approach

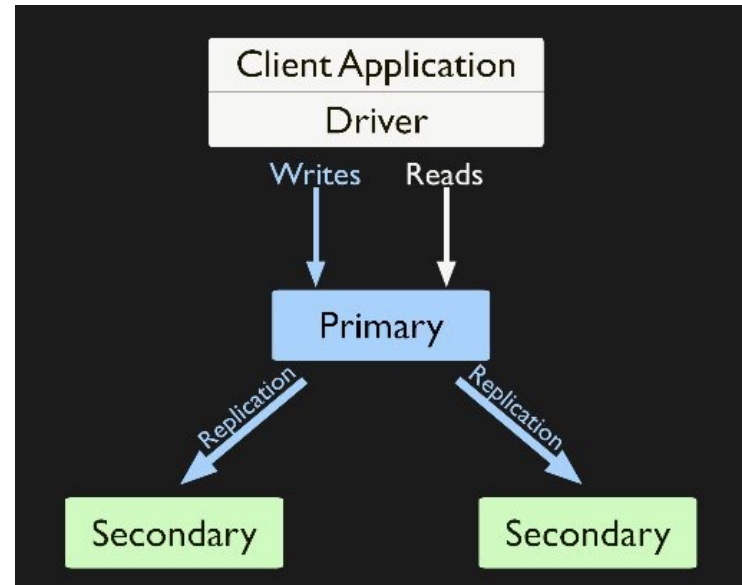


# Replica sets

- The replica-set nodes elect the master, or primary, among themselves
  - The one closer to the other servers or having more RAM
  - Users can affect this by assigning a priority to a node
- All requests go to the master node
- Data is replicated to the slave nodes and the clients can get to the data even when the primary node is down
- If the master node goes down, the remaining nodes in the replica set vote among themselves to elect a new master

# Replication & fault tolerance

- Primary accepts **all read and write** operations
- Secondaries **only accepts read operations** not write



- Secondaries replicate the primary's oplog and apply the operations to their data sets **asynchronously**
  - With a **replication lag**
- By having the secondaries' data sets reflect the primary's data set, the replica set can continue to function despite the failure of one or more members



# Consistency vs availability

- Through the effective use of write concerns and read concerns, the level of consistency and availability, can be adjusted such as
  - waiting for stronger consistency guarantees, or
  - loosening consistency requirements to provide higher availability

# Transactions

- All write operations in MongoDB are **atomic** on the level of a single **document**
- When a single write operation modifies multiple documents (e.g. [`db.collection.updateMany\(\)`](#)), the modification of each document is atomic, but the operation as a whole is not atomic.
- For situations that require atomicity of reads and writes to multiple documents (in a single or multiple collections), MongoDB supports multi-document transactions

# Transactions

- Atomic transactions are possible at the multi-document level since version 4.0 (2018)
- All transactions that contain read operations must use read preference **primary**
  - All operations in a given transaction must route to the same member
- Until a transaction commits, the data changes made in the transaction are not visible outside the transaction

Use cases

# Suitable use cases



- **Event Logging**

- Storing logs of events, acting as a central data store for event storage
- Events can be sharded by
  - the application that generated the event
  - the type of the event (e.g., order\_processed, customer\_logged)

- **Content Management Systems, Blogging Platforms**

- content management systems or applications for publishing websites, managing user comments, user registrations, profiles, web-facing documents.

- **Web Analytics or Real-Time Analytics**

- store data for real-time analytics; since parts of the document can be updated, it's very easy to store page views or unique visitors, and new metrics can be easily added without schema changes

- **E-Commerce Applications**

- E-commerce applications often need to have flexible schema for products and orders, as well as the ability to evolve their data structure without expensive database refactoring or data migration

# When not to use



- **Complex Transactions Spanning Different Documents**
  - Document data stores are not suited for atomic cross-document operations.
- **Queries against Varying Aggregate Structure**
  - In document databases data is saved as an aggregate in the form of application entities.
  - If the design of the aggregate is constantly changing, you need to save the aggregates at the lowest level of granularity.

# References & Credits

- References:
  - Kristina Chodorow, MongoDB – The definitive guide, 3rd Ed., O'Reilly, 2019
  - <https://docs.mongodb.com/>
- Credits:
  - Riccardo Torlone, Big Data, Università di Roma Tre
  - Kathleen Durant, CS 3200, Northeastern University