

From the aggregate-oriented  
logical schema to MongoDB  
logical/physical schema

# Input

- The aggregate-oriented logical schema in meta-notation
- The annotated ER diagram
- One NoSQL system  $S$  (for today, MongoDB)

# Output

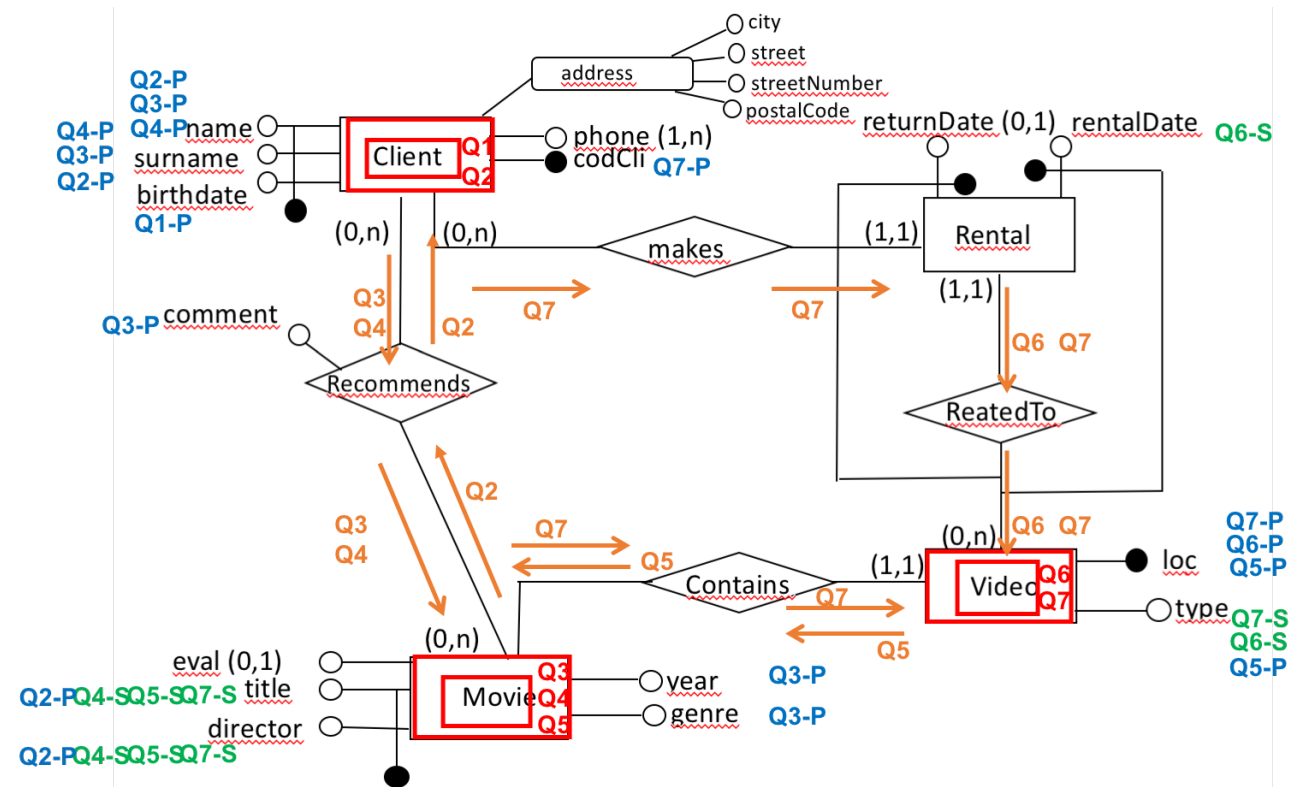
- A schema for  $S$  (in metanotation +  
+ partition keys + indexes)

**ASSUMPTION: all selection conditions are equality based**

# Issues

- From the aggregate-oriented logical schemas in meta-notation, the corresponding aggregation entity and the set of associated queries to a set of collections for S
- Each collection allows one subset of the queries to be executed
- From the set of selection attributes and the identifiers of the aggregation entity to the partition key and indexes

# Input



- client: {name, surname, birthdate, recommends: [{title, director}]}
- movie: {title, director, year, genre, recommended\_by: [{name, surname, comment]}, contained\_in: [ {loc, type } ]}
- video: {loc, type, rentals: [{rentalDate, codCli]}, title, director}

# Design in MongoDB

# Remarks

- The identifier in MongoDB always is the `_id` field (ObjectId type)
  - Automatically assigned
  - Monotonically increasing
  - Can be assigned at document insertion, but cannot be updated later
- The `_id` field is automatically indexed
- Partition/shard keys must be indexed

(in what follows, we use the term partition key for uniformity with aggregates, in MongoDB read shard key)

# Aggregation entity Client

client: {name, surname, birthdate, recommends: [{title, director}]}

Queries associated with Client: Q1, Q2

Selection attributes for Q1: { }

Selection attributes for Q2: { }

Projection attributes for Q1: {birthdate}

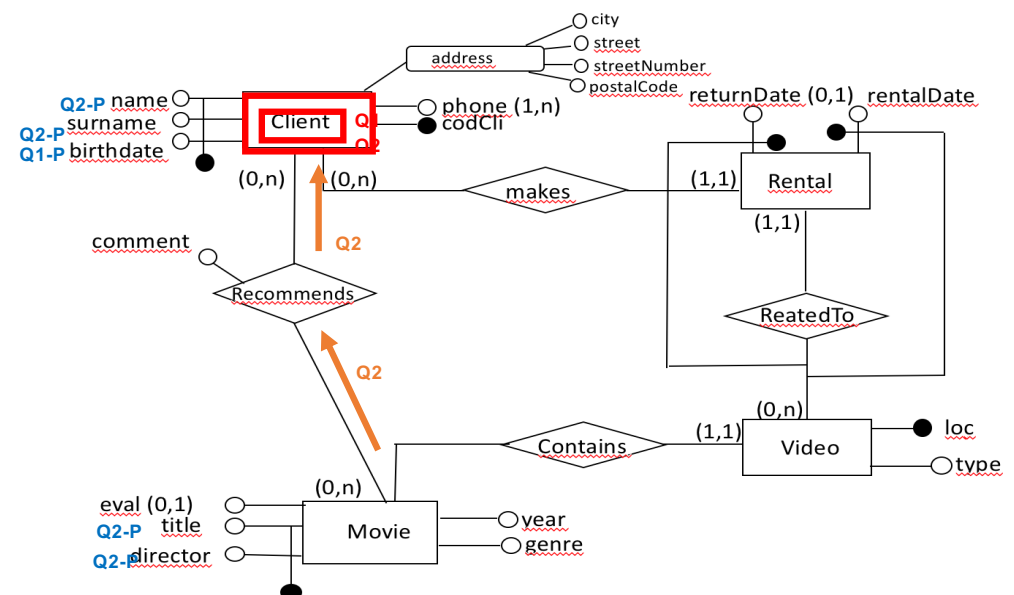
Projection attributes for Q2: {name, surname, title, director}

No selection attribute → no  
need for a specific partition key

\_id field automatically assigned  
could be partition key as well

Other partition keys: whatever  
you want (name)

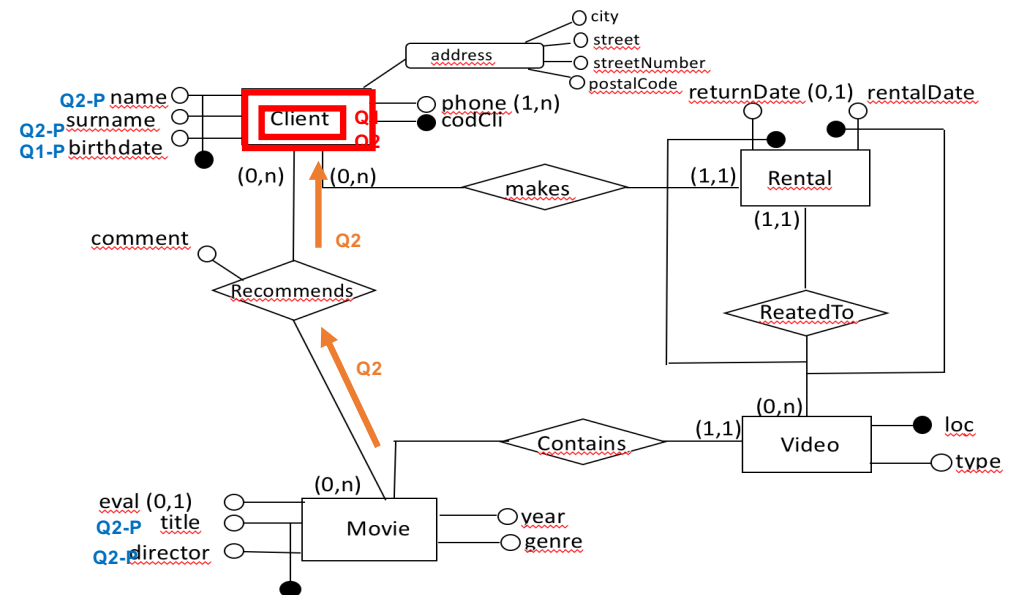
**Q1. Average age of clients**  
**Q2. Name and surname of clients**  
**and related recommended movies**



# Aggregation entity Client

client: { id, name, surname, birthdate, recommends: [{title, director}]}

**Q1.** Average age of clients  
**Q2.** Name and surname of clients and related recommended movies



```
db.clients.createIndex( {"name": 1, "surname":1, "birthdate":1}, {unique: true})
```

```
db.clients.createIndex( {"name": 1}, {unique: false})
```

```
db.adminCommand( { shardCollection: "db.clients", key: { name: 1 },  
                  field: "hashed" } )
```

(non unique index on "name" automatically created if clients is empty)



# Aggregation entity Movie

movie: {title, director, year, genre, recommended\_by: [{name, surname, comment}],  
contained\_in: [ {loc, type} ] }

Queries associated with Movie: Q3, Q4, Q5

Selection attributes for Q3: { }

Selection attributes for Q4: { title, director }

Selection attributes for Q5: { title, director }

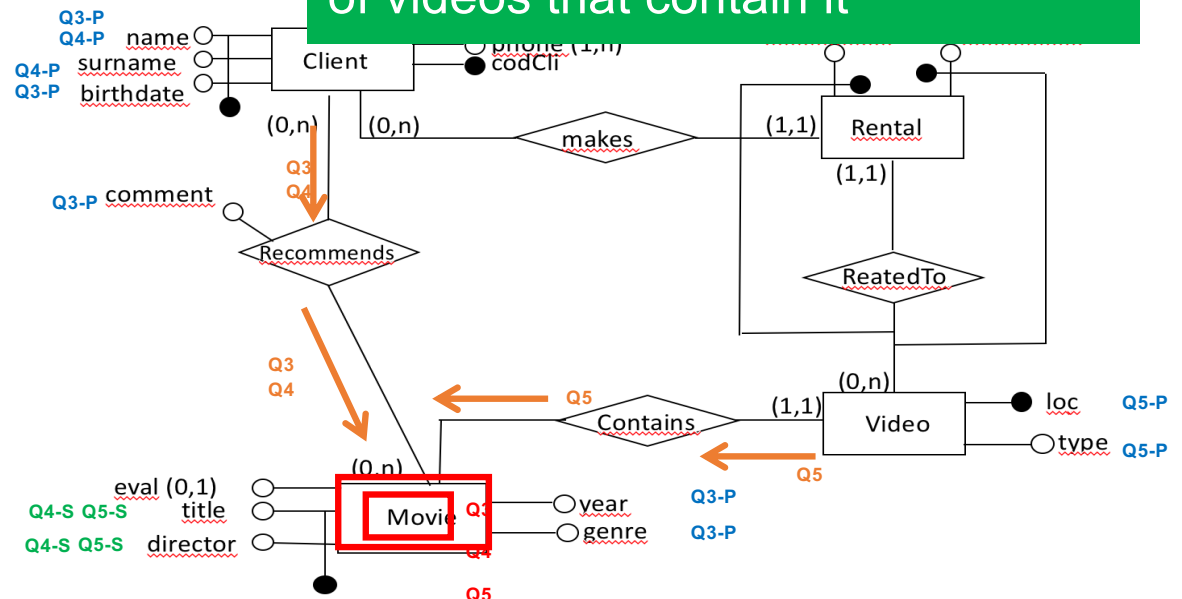
Partition key { title, director }  
(with unique index)  
hashed

\_id field automatically assigned  
could be partition key as well

**Q3.** Genre and year of the movies  
and their related  
recommendations, together with  
the name and the surname of the  
client who made them

**Q4.** Name and surname of clients  
who recommended the movie 'pulp  
fiction' by 'quentin tarantino'

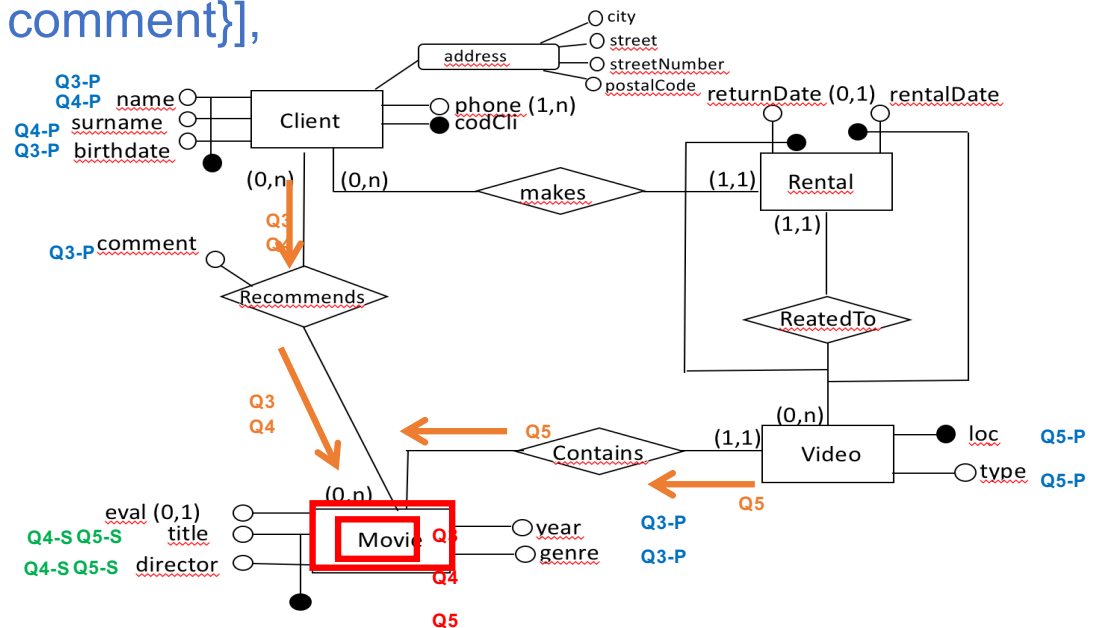
**Q5.** Given a movie, all information  
of videos that contain it



# Aggregation entity Movie

movie: { id, title, director, year, genre,  
recommended\_by: [{name, surname, comment}],  
contained\_in: [ {loc, type} ] }

**Q3.** Genre and year of the movies and their related recommendations, together with the name and the surname of the client who made them  
**Q4.** Name and surname of clients who recommended the movie 'pulp fiction' by 'quentin tarantino'  
**Q5.** Given a movie, all information of videos that contain it



```
db.movies.createIndex( {"title": 1, "director":1}, {unique: true})
```

```
db.adminCommand( { shardCollection: "db.movies",
                    key: {"title": 1, "director":1},
                    field: "hashed" } )
```

(non unique index on {title, director} automatically created if movies is empty)

# Aggregation entity Video

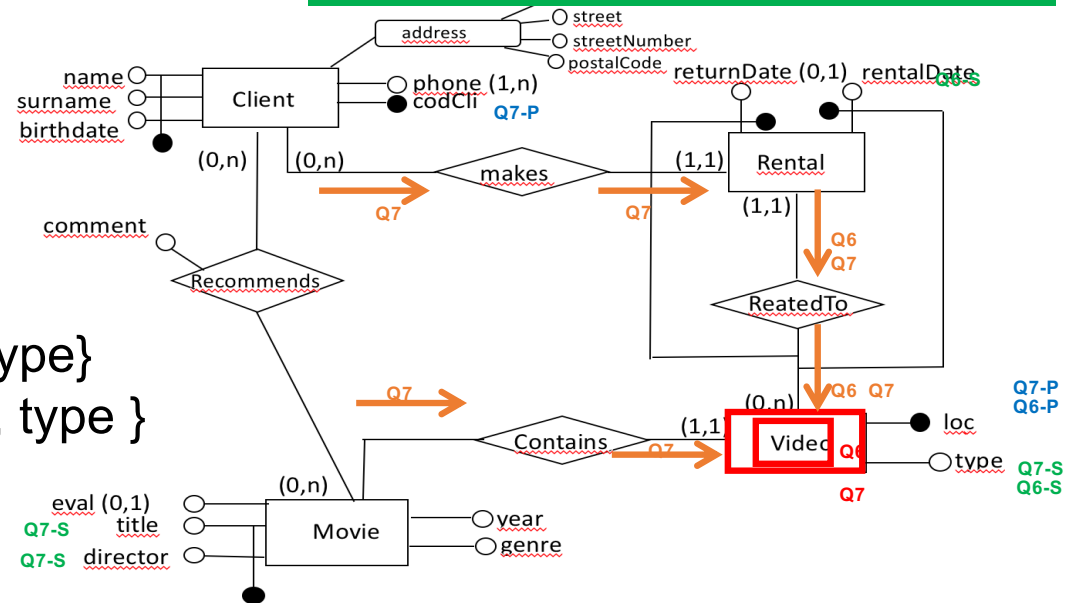
video: {loc, type, rentals: [{rentalDate, codCli}], title, director}

Q6. Videos of type 'DVD', rented from a certain date  
Q7. The videos of type 'VHS' containing the movie 'pulp fiction' by 'quentin tarantino' and the clients that rented them

Queries associated with Video: Q6, Q7

Selection attributes for Q6: { rentalDate, type }

Selection attributes for Q7: { title, director, type }



type appears in both → it becomes the partition key (an equality is always specified)

Partition key = { type }

A non-unique index has to be created on the partition key

# Aggregation entity Video

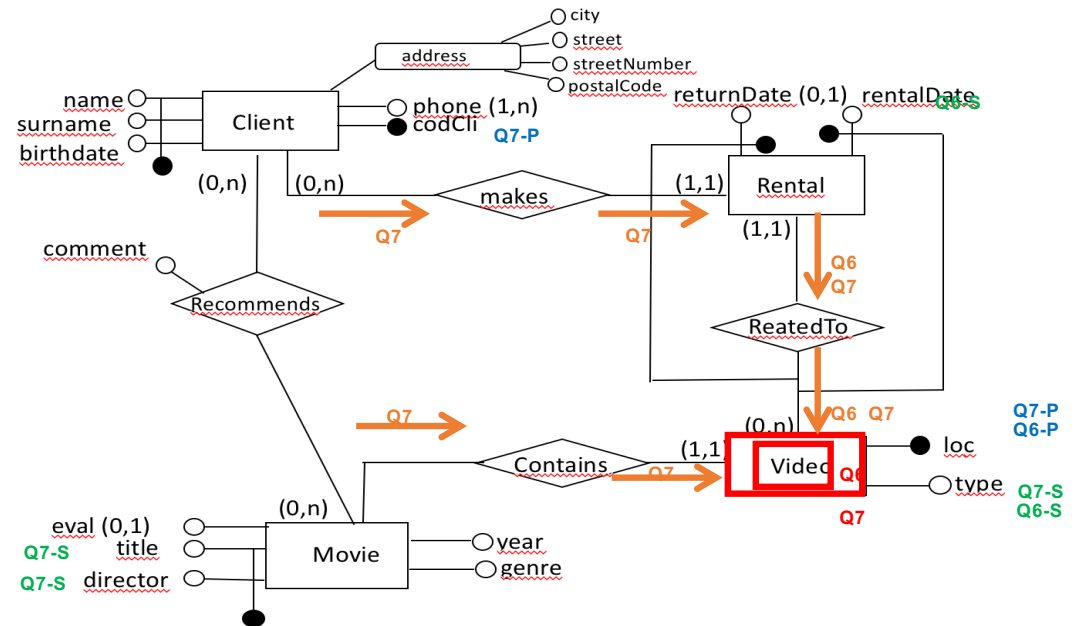
video: {\_id, loc, type, rentals: [{rentalDate, codCli}], title, director}

*What about the unique index on loc?*

MongoDB does not support unique indexes across shards, except when the unique index contains the full shard key as a prefix of the index.

If you don't have an index on the shard key, you should at least have a compound index that starts with the shard key.

Unique index on type, loc



# Index rules

- A unique index on `_id` is always created by the system
- All sharded collections must have an index that supports the shard key. The index can be:
  - an index on the shard key or
  - a compound index where the shard key is a prefix of the index
- If sharding is executed on an empty collection, a non-unique index is automatically created on the shard key
- For a ranged sharded collection, only the following indexes can be unique:
  - the index on the shard key
  - a compound index where the shard key is a prefix
  - the default `_id` index; however, the `_id` index only enforces the uniqueness constraint per shard if the `_id` field is not the shard key or the prefix of the shard key.

# Index rules

- Although you can have a unique compound index where the shard key is a prefix, if using unique parameter, the collection must have a unique index that is on the shard key.
- You cannot specify a unique constraint on a hashed index
- The decision on whether an index must be unique or non-unique depends on the identifiers detected during the aggregate modeling