

Column-family NoSQL data stores

Column-family data model at a glance

- Each data instance is represented in the form (key, value)
 - key is an identifier
 - value is the aggregate, corresponding to a set of pairs (column-key, column-value)
- Column-value can be either an atomic value or a complex one
- In this second case, the value structure is visible only at the application level, it is opaque (a blob) at the logical value
- Data instances can be grouped into logical collections

Logical level

client: { codCli,

name, surname, birthdate,

address: [city, street, street, postalCode], BLOB

movies: [{movie: {comment, director, ...}}, BLOB

}

KEY

VALUE

Column-family data model at a glance

- Each data instance is represented in the form (key, value)
 - key is an identifier
 - value is the aggregate, corresponding to a set of pairs (column-key, column-value)
- Column-value can be either an atomic value or a complex one
- In this second case, the value structure is visible only at the application level, it is opaque (a blob) at the logical value
- Data instances can be grouped into logical collections

client: { codCli,
name, surname, birthdate,
address: {city, street, streetNumber, postalCode},
movies: [{movie: {comment, title, director,...}}]
}

Application level

KEY

VALUE

Column-family data model at a glance

- Columns can be organised into **families**: sets of related data (columns) that are often accessed together

Logical level

Client

Family: info: {codCli,

name, surname, birthdate,

address: [city, street, streetNumber, postalCode],

}

Column family associated with client personal information

Client

Family: movies: {codCli,

movies: [[movie: {comment, title, director, ...}]]

}

Column family associated with client movie recommendations

Column-family data model at a glance

- Columns can be organised into **families**: sets of related data (columns) that are often accessed together

Application level

Client

Family: info: {codCli,

name, surname, birthdate,

address: {city, street, streetNumber, postalCode},

}

Column family associated with
client personal information

Client

Family: movies: {codCli,

movies: [{movie: {comment, title, director,...}}]

}

Column family associated with
client movie recommendations

Instance structure

- **Limited schema information**: column-family names + column-names inside column-families
- Nested values are not visible at the NoSQL system level but only at the application level

Collections

- Pairs can be grouped into logical **collections/namespaces**: sets of aggregates, i.e., sets of (key, value) pairs sharing the same column-family
- Each collection becomes a sort of table

Client

Family: info: {codCli,

name, surname, birthdate,

address: {city, street, si BLOB ber, postalCode},

}

Client

Family: info

name	surname	birthdate	address

Collections and schema information

- Schema information can be provided either when inserting the data instances (DML) or before (DDL)
- Inside the same collection, the chunk of data corresponding to a column family **might** correspond (depending on the specific system) to different sets of column-names for each row
- **Flexibility**
- Columns **must** however belong to the same column family

Client
Family: info: {codCli,

name, surname, birthdate,
address: {city, street, street, BLOB, number, postalCode}
}

client
Family: info

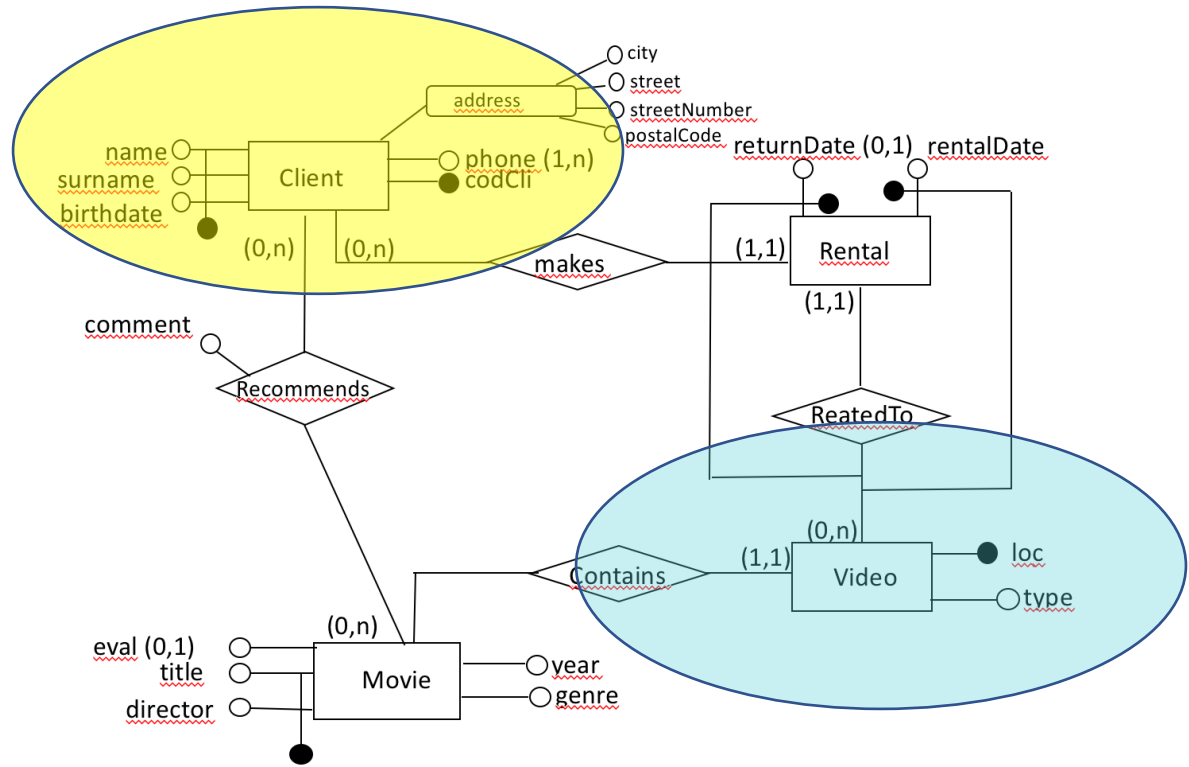
Customer Details				
Row 1	1	Name: "John"	Surname: 'Black'	
Row 2	2	Name: 'Sarah'	Birthdate: "15/10/2000"	
Row 3	3	Name: Philippe"		

Keys

- (key, value) pair
- At the logical level: identification + data retrieval
 - the key **needs** to assume **unique values** in the collection (i.e., it is an **identifier**)
- At the physical level, the key tells the system **how to partition the data** and where to store the data
 - **partition key**
- Usually, the partition key is a subset of or is equal to the primary key
- Aggregates can be directly retrieved **only by specifying values for attributes in the partition key**
 - primary key attributes can be used for retrieval together with the partition key, under some restrictions

Keys

- Usually, more than one attribute
- *Partition key subset of (or equal to) primary key (in bold in the examples)*
- The type and the content of the key is identified starting from the designed aggregates, taking into account the features of the system at hand



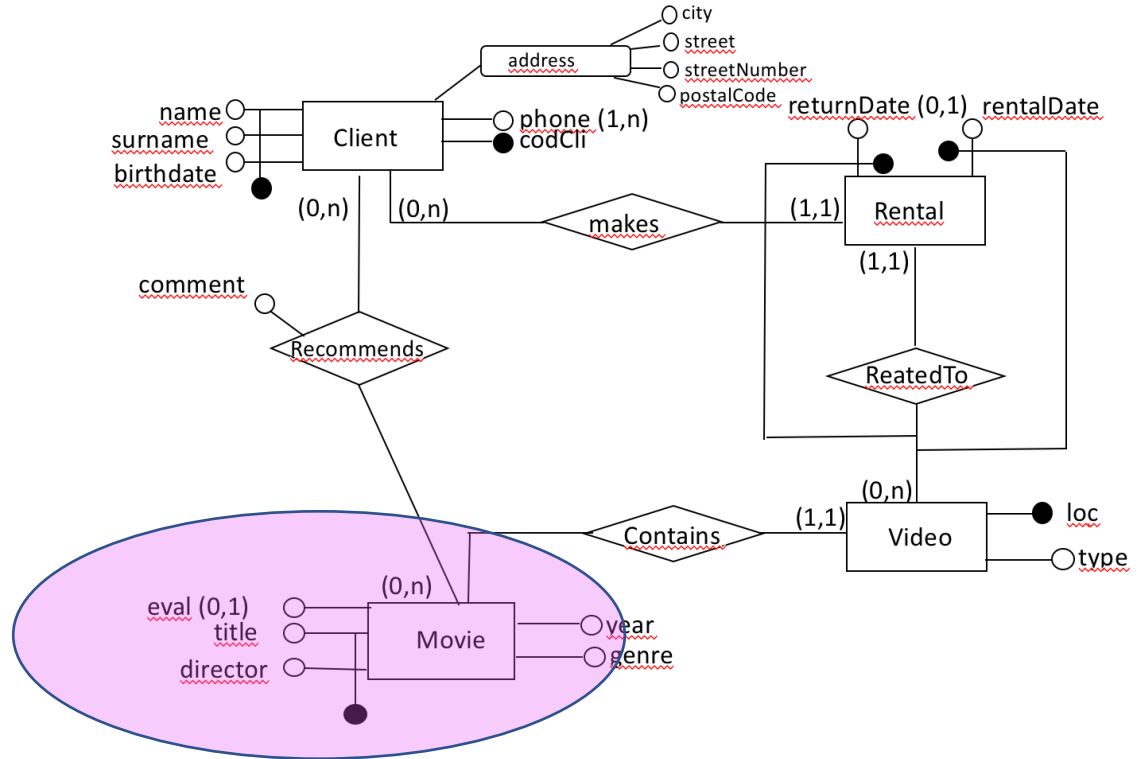
Client: {**codCli**, ...}

primary key is codCli
(partition) key is codeCli

Video: {**loc**, ...}

primary key is loc
(partition) key is loc

Keys



Movie : {**title**, **director**, ...}


(partition) key could be either
(title, director)

title

director

→ we select (title, director) (in bold)

Keys



Movie : {**title**, **director**, ...}

(partition) key could be either
(title, director)

title

director

→ we selected (title, director) (in bold)

Movie aggregate from our modeling example:

movie: { **title**, **director**,
year, genre, recommended_by: [{name, surname, comment}],
contained_in: [{video: {loc, type}}] }

Identifier vs partition/sharding key

- In the Videos collection, videos are partitioned by loc
- In the Clients collection, clients are partitioned by codCli
- In the Movies collection, movies are partitioned by (title, director)
- This impacts the way data are stored
- This impacts the way data can be retrieved

Video Rental Example

Logical level
(references
exists but they
are opaque in
the system)

```
client: { codCli,
```

```
  name, surname, birthdate,
```

```
  address: {city, street, streetNumber, postalCode},
```

```
  movies: [{movie: {comment, id, {title, director}, ...}}]
```

```
]
```

```
}
```

REFERENCE

REFERENCE

```
movie: { title, director,
```

```
  year, genre, recommended_by: {name, surname, comment},
```

```
  contained_in: [{video: {loc, time}}]
```

```
video: {loc,
```

```
  type, rentals: [{rental: {rentalDate, codCli}}], {title, director}}
```

Video Rental Example

Application
level
(references
can be used at
this level, by
accessing the
structure of
opaque
components)

```
client: { codCli,
```

```
    name, surname, birthdate,  
    address: {city, street, streetNumber, postalCode},  
    movies: [ {movie: {comment, id:{title, director},...}}
```

```
]
```

```
}
```

REFERENCE

REFERENCE

```
movie: { id: {title,director},
```

```
    year, genre, recommended_by: [ {name, surname, comment} ],  
    contained_in: [ {video: {loc, type}} ] }
```

```
video: {loc,
```

```
    type, rentals: [ {rental: {rentalDate, codCli}} ], title, director}}
```

Interaction

- Basic **lookup** based on the **key** and **column name** (of a given column family)
 - `void define(family)`
 - `void insert(key, family, columns)`
 - `columns get(key, family)`
 - `value get(key, family, column)`
 - `value get(key, family, column, value)`
 - `void put(key, family, column, value)`
- It is not possible to arbitrarily **join** data contained in different column-families (always possible **at the application level**)
- It is not possible to navigate the nested structure of the column-values
- In case collections are supported, collection name is also specified

Example 1 – Find the title of the movie in a certain video

Column family: All

video: {loc, type, rentals: [{rentalDate, codCli}], title, director}

```
{"loc": 1234,  
  "type": "dvd"  
  "rentals": [{"rental": {"rentalDate": "15/10/2021",  
                           "codCli": 375657}}],  
  "title": "pulp fiction",  
  "director": "quentin tarantino"}
```

value get(collection, key, family, column)

get(Videos, '1234', 'all', [title])

- loc is the partition key, we can retrieve data starting from it
- the title of the movie contained in the video can be directly retrieved from the data store

Example 2 – Find the videos containing a certain movie

Collection family: All

video: {loc, type, rentals: [{rentalDate, codCli}], title, director}

```
{"loc": 1234,
```

```
"type": "dvd",  
"rentals": [{"rental": {"rentalDate": "15/10/2021",  
                        "codCli": 375657}}],  
"title": "pulp fiction",  
"director": "quentin tarantino"}
```

value get(collection, key, column)

- Even if the title and director values are «visible» at the data store level, we have to retrieve data starting from the key
- We can only get all the videos, and
- At the application level, filter them by title and director

Example 3 – Find the videos rented by a certain client

Collection family: All

video: {loc, type, rentals: [{rentalDate, codCli}], title, director}

```
{"loc": 1234,
```

```
  "type": "dvd",  
  "rentals": [{"rental": {"rentalDate": "15/10/2021",  
                           "codCli": 375657}}],  
  "title": "pulp fiction",  
  "director": "quentin tarantino"}
```

```
value get(collection, key, family, column)
```

- Client data is «not visible» at the data store level
- At the data store level, we can only get all the videos, and
- At the application level, go inside video rentals and filter them by codCli

Example 2 – Find the videos containing a certain movie - with the movies collection

Collection family: All

movie: { **title**, **director**,

year, genre, recommended_by: [{name, surname, comment}],
contained_in: [{video: {loc, type}}] }

value get(collection, key, family, column)

- `get(Movies, 'pulp fiction:quentin tarantino', 'all', contained_in)`
- At the data store level, retrieve the information about the videos containing a certain movie (the whole blob)

REMARK: the relationship is part of the aggregate, a single data access to retrieve all the relevant information (contained in a single data node)

Example 4 - Relationships

Find the age(s) of customer(s) that rented a given video

```
{"loc": 1234,
```

```
"type": "dvd",
```

```
"rentals": [{"rental": {"rentalDate": "15/10/2021",  
                        "codCli": 375657}}],
```

```
"title": "pulp fiction",
```

```
"director": "quentin tarantino"}
```

```
{  
  "codcli": 375657,  
  "name": "John",  
  "surname": "Black",  
  "birthdate": "15/10/2000",  
  "address": {"city": "Genoa", "street": "Via XX Settembre",  
              "streetNumber": 15, "postalCode": 16100},  
  "movies": [{"movie": {"comment": "very nice", "title": "pulp fiction",  
                        "director": "quentin tarantino"}},  
              {"movie": {"comment": "very nice", "title": "pulp fiction",  
                        "director": "quentin tarantino"}}]  
}
```

`get(Videos, 1234, 'all')`

- At the data store level, find Video 1234 in collection Videos and retrieve all the information
- at the *application level*, we go inside the aggregate value and we discover that it was rented by Client 375657
- at the *application level*, we can execute `get(Clients, 375657, 'all', birthdate)` to retrieve information about a/the customer birthdate that rented video 1234 (thus navigating the customer reference stored inside the video), and then the age is computed

Example 4 - Relationships

```
{ "loc": 1234,  
  "type": "dvd",  
  "rentals": [ { "rental": { "rentalDate": "15/10",  
                           "codCli": 375657,  
                           "rentalCode": "R1" } } ],  
  "title": "pulp fiction",  
  "director": "quentin tarantino" }
```

Find the age(s) of client(s) that rented a given video

- REMARKS:
1. Who knows that inside **rentals** there is the code of the client that rented the video?
And where to look for the client information and how to match?
- The application, the data store is completely unaware of that!**
- 2. Since the relationship is not part of the aggregate, navigating it requires a sort of join, no system support for it
 - 3. Since the relationship is not part of the aggregate, navigating it (at application level) requires two distinct data accesses, at two possibly distinct data nodes
- bin can execute `get(Clients, 375657, 'all', ...)` to retrieve information about a/the customer birthdate that rented video 1234 (thus navigating the customer reference stored inside the video), and then the age is computed

Advanced interaction

- Some column-family systems support SQL-like languages
- Joins are never supported
- *Only queries for which the set of nodes containing relevant data can be determined by the system through the partition key in advance are executed*
 - A value for the partition key **has always to be** specified
- The execution of other queries (but not joins) can however be forced or admitted through the creation of indexes

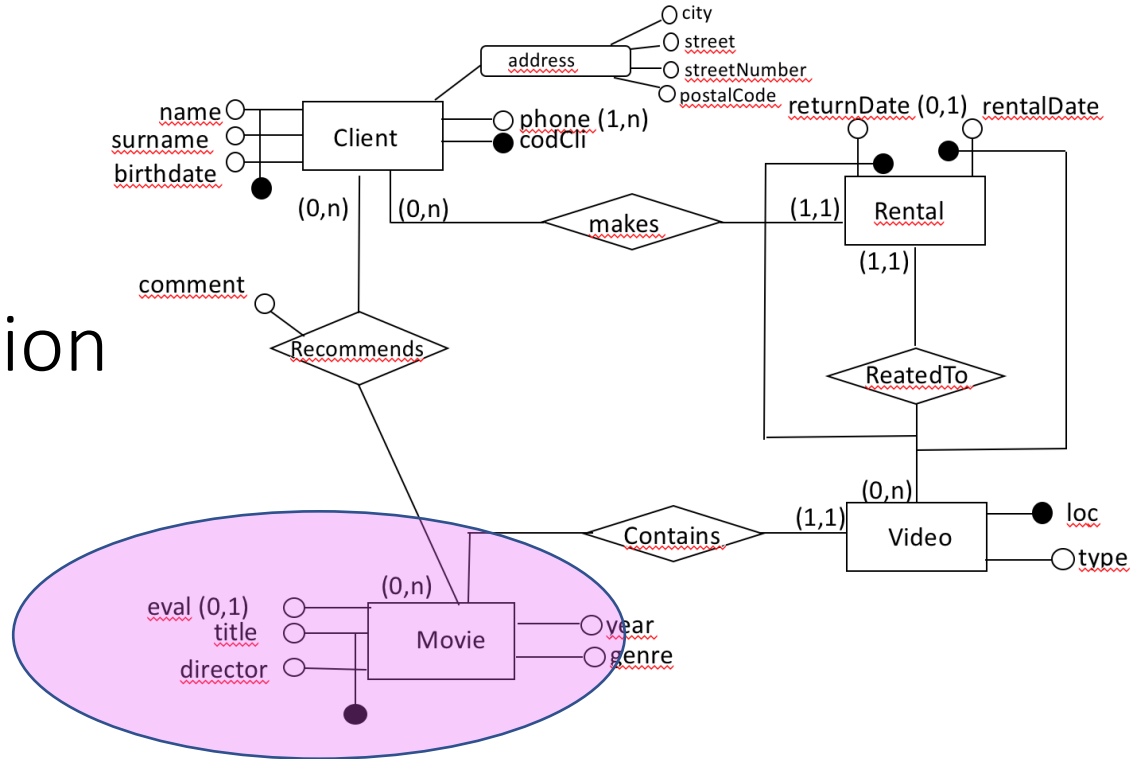
Keys physical representation

Movie : { title, director ... }

primary key is (title, director)
(partition) key is title

movie: { title,

director, year, genre, recommended_by: [{name, surname, comment}],
contained_in: [{video: {loc, type}}] }




Keys physical representation

movie: { **title**,

director, year, genre, recommended_by: [{name, surname, comment}],
contained_in: [{video: {loc, type}}] }

<u>title</u>	<u>director</u>	year	genre	Recommended_by	Contained_in
A	B	1975	comedy
A	C	2010	comedy



<u>title</u>	B:year	B:genre	...	C:year	C:genre	C:...
A	1975	comedy		2010	comedy	

Example – Find the videos containing a certain movie - using the movies collection

movie: { title,

director, year, genre, recommended_by: [{name, surname, comment}],
contained_in: [{video: {loc, type}}] }

(partition) key is title
primary key is (title, director)

```
SELECT contained_in  
FROM Movies  
WHERE title = 'pulp fiction' AND director = 'quentin tarantino'
```

- An equality condition with respect to the partition key has been specified, the query is admitted
- Single operation at the system level retrieves information about the videos containing a certain movie

Example – Find the videos containing a movie of a certain director - using the movies collection

movie: { title,

director, year, genre, recommended_by: [{name, surname, comment}],
contained_in: [{video: {loc, type}}] }

(partition) key is title
primary key is (title, director)

```
SELECT contained_in  
FROM Movies  
WHERE director = 'quentin tarantino'
```

- This query is not admitted because no partition key value (title) is specified

Popular column-family data stores



HYPERTABLE

