



A column-family data store

# Cassandra

- Used by more than 1500 companies
- Developed at Facebook
  - Initial release: 2008 (stable release: 2013)
  - now, Apache Software License
- Written in: Java
- Operating System: cross-platform
- Java, Scala, Ruby, C#, Python, Perl, PHP, C++  
Cassandra clients

# Who uses Cassandra?



- Some of the largest production deployments:
  - **Apple**: over 75,000 nodes storing over 10 PB of data
  - **Netflix**: 500 nodes, 420 TB, over 1 trillion requests per day
  - **eBay**: over 100 nodes, 250 TB

# Who *else* uses Cassandra?



# Data model and interaction

# Cassandra data model


- The **original** Cassandra **data model** is similar to that discussed for generic column-family data stores, with a basic interaction protocol - Thrift
- More recently, a **new** high level data **model (with primary keys)** and related (**SQL-like**) language has been proposed
- We will discuss only the high level data model and language since the interaction protocol for the original data model is now deprecated

# Data model and interaction

- Concepts closer to a traditional relational database with tables, columns and rows
- Cassandra Query Language (CQL) – similar to SQL
  - Still the same data underneath, just abstracted away
  - reintroduction of schema so that you don't have to read code to understand the data model
  - CQL creates a common language so that details of your data can be easily communicated (but it is not a standard)
  - best-practice Cassandra interface that hides the details of the original data model
- CQL is one of the strongest reasons why Cassandra has become so widely used today

# Data model


- Keyspace – Database in an RDBMS





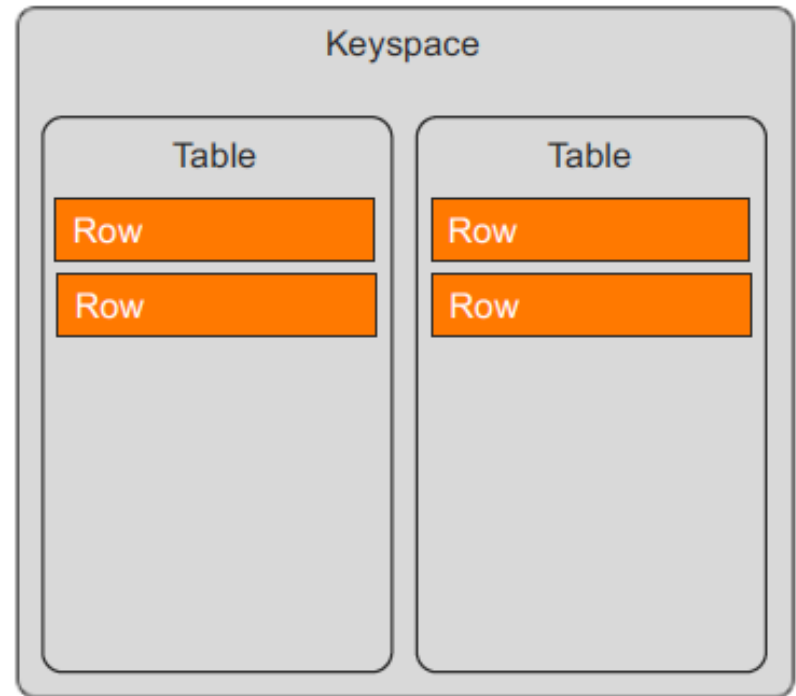
# Data model

- Keyspace – Database in an RDBMS
- Table – Table in an RDBMS



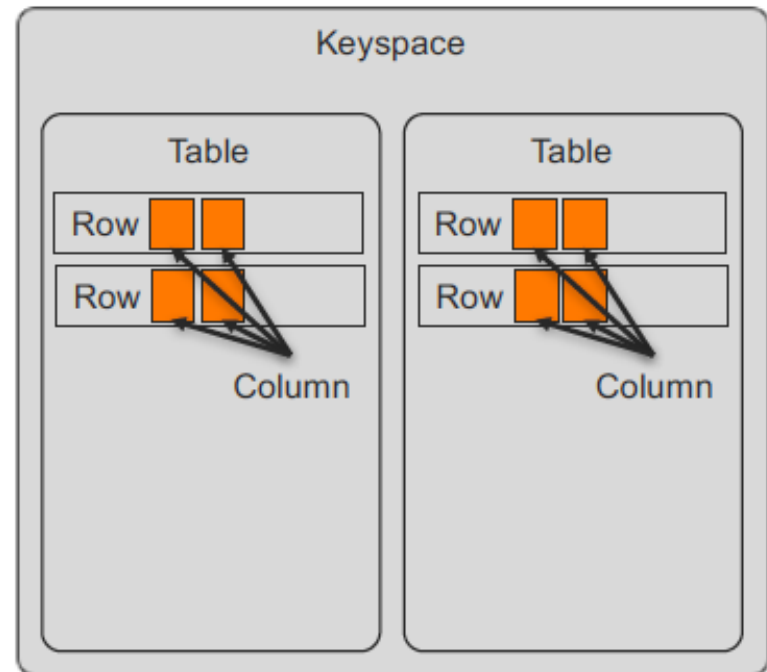
# Data model

- Keyspace – Database in an RDBMS
- Table – Table in an RDBMS
- Row – Row in an RDBMS



# Data model

- Keyspace – Database in an RDBMS
- Table – Table in an RDBMS
- Row – Row in an RDBMS
- Column – Column in an RDBMS



# CQL at the glance

```
CREATE TABLE clients (  
    codCli int PRIMARY KEY,  
    name text,  
    surname text,  
    birthdate date );
```

It looks like SQL but the original Cassandra model is still there, even if hidden by the CQL data model

```
INSERT INTO clients (codCli, name, surname, birthdate)  
VALUES (3, 'john', 'black', '15/10/2000');
```

```
SELECT * FROM users;
```

```
    codCli | name | surname | birthdate  
-----+-----+-----+-----  
        3 | john | black   | 15/10/2000
```

# CQL atomic data types

- Strings
  - Ascii, varchar, text
- Numbers
  - Bigint, decimal, double, float, int,...
- Time
  - Date, time, timestamp
- ...

# CQL – CREATE TABLE and partition key

- Similar to SQL CREATE TABLE

Partition key **codCli**  
Primary key **codCli**

```
CREATE TABLE clients (  
    codCli int PRIMARY KEY,  
    name text,  
    surname text,  
    birthdate date );
```

- By default, the **partition key** is the first column of a composite key (thus, partition key is not unique but do not forget the match with the old Cassandra model)

```
CREATE TABLE movies (  
    title text,  
    director text,  
    year int,  
    genre text,  
    PRIMARY KEY (title, director));
```

Partition key **title**  
Primary key (title, director)

Partition key (title, director)  
Primary key (title, director)

```
CREATE TABLE movies (  
    title text,  
    director text,  
    year int,  
    genre text,  
    PRIMARY KEY ((title, director)));
```

# CQL – CREATE TABLE and partition key

Partition key =

first attribute of the primary key

or

a prefix of the primary key, specified inside ( )

# CQL – CREATE TABLE

- No foreign keys

```
CREATE TABLE movie (  
    title text,  
    director text,  
    Year int,  
    genre text,  
    PRIMARY KEY ((title, director));
```

```
CREATE TABLE video (  
    loc int PRIMARY KEY,  
    type text,  
    title text,  
    director text);
```

There is no way to specify the foreign key

Side effect:

Table video can refer movies not included in table movie



# CQL - INSERT

- Syntax similar to SQL

```
INSERT INTO t2 (a, b, c, d) VALUES (0,0,0,0);
```

- **No primary key violation**: if you insert a row with the same primary key of another row in the table, the first will be simply updated
  - INSERT becomes an UPDATE!

# CQL – CREATE TABLE and partition key

- All the attributes that follow the partition key attributes in a primary key definition are called **clustering columns**
- In a node, all rows in a collection with the same partition key value are **locally sorted with respect to clustering columns** (default: ascending order)
- Consequences on how to query data

# Example

```
CREATE TABLE T
(
    A int,
    B int,
    C int,
    D int,
    PRIMARY KEY (A, B, C)
);
```

```
INSERT INTO T (A, B, C, D) VALUES (0,0,0,0);
INSERT INTO T (A, B, C, D) VALUES (0,1,2,2);
INSERT INTO T (A, B, C, D) VALUES (0,0,1,1);
INSERT INTO T (A, B, C, D) VALUES (0,1,3,3);
INSERT INTO T (A, B, C, D) VALUES (1,1,4,4);
```

Two partitions  
(physical level)

P1

A	B	C	D
1	1	4	4

P2

A	B	C	D
0	0	0	0
0	0	1	1
0	1	2	2
0	1	3	3

# Example

Two partitions  
(physical level)

	A	B	C	D
P1	(1, 1, 4, 4)			
	A	B	C	D
P2	(0, 0, 0, 0)			
	(0, 0, 1, 1)			
	(0, 1, 2, 2)			
	(0, 1, 3, 3)			

If the partitions are physically stored on two nodes N1 and N2...

*Node N1*

**A 1:4:D**  
1 4

*Node N2*

**A 0:0:D 0:1:D 1:2:D. 1:3:D**  
0 0 1 2 3

If the partitions are physically stored on a single node N...

*Node N*

**A 1:4:D**  
1 4  
**A 0:0:D 0:1:D 1:2:D. 1:3:D**  
0 0 1 2 3

# CQL - UPDATE

```
CREATE TABLE T
(
    a int,
    b int,
    c int,
    d int,
    PRIMARY KEY (a, b, c)
);
```

```
INSERT INTO T (a, b, c, d) VALUES (0,0,0,0);
```

- UPDATE is similar to INSERT!

```
UPDATE T SET a = 0 AND b = 0 AND c = 0 AND d = 1;
```

- is equivalent to

```
INSERT INTO T (a, b, c, d) VALUES (0,0,0,1);
```

- It can contain a WHERE clause following the same rules of the WHERE clause in SELECT (see later)

# CQL - DELETE

- Marks data for removal from a table
- Data is removed later through a process called **compaction**
- Unlike SQL it MUST contain a WHERE clause, following the same rules than the WHERE clause in SELECT (see later)

```
DELETE FROM T  
WHERE a = 0 AND b = 0 AND c = 0;
```

# CQL - SELECT

- Similar to SELECT statement in SQL

- Retrieve all the rows from a table

```
SELECT * FROM T;
```

- Project all the rows from a table

```
SELECT a FROM T;
```

- Select some rows from a table, with many limitations

```
SELECT * FROM T WHERE a = 0 AND b = 0 AND c = 1;
```

# CQL – SELECT, restrictions

- Many restrictions on what you can do in the WHERE clause
  - no join
  - restrictions on selection conditions
- Some queries cannot be executed
- Solutions
  - Query-based design
  - Indexes
  - Forcing the execution of non-admitted queries



# CQL – SELECT, restrictions: no join

- Queries must refer a single table
- Joins are not supported
  - As we know, joining (portions of) tables stored in different nodes leads to a high data communication
  - Impact on performance
- Two options to overcome the problem
  1. Go back to the logical design and restructure your data (query-based design)
  2. Execute joins by writing specific applications, accessing data stored in Cassandra
    - Store data in Cassandra and use MapReduce frameworks or Spark

## CQL – SELECT, restrictions: no join

- If you design your aggregates using the methodology, all workload queries can be executed over your data without the need of join execution
- However, join execution might be needed for executing queries that do not belong to the workload

# Example

```
CREATE TABLE movie (  
    title text,  
    director text,  
    year int,  
    genre text,  
    PRIMARY KEY ((title, director));
```

```
CREATE TABLE video (  
    loc int PRIMARY KEY,  
    type text,  
    title text,  
    director text);
```

```
SELECT year  
FROM movie M, video V  
WHERE M.title = V.title AND  
M.director = V.director AND loc  
= 1234;
```

Non admitted

# Example – approach 1

```
CREATE TABLE movie (  
  title text,  
  director text,  
  year int,  
  genre text,  
  PRIMARY KEY ((title, director));
```

```
CREATE TABLE video (  
  loc int PRIMARY KEY,  
  type text,  
  title text,  
  director text);
```

```
CREATE TABLE movie_video (  
  title text,  
  director text,  
  year int,  
  genre text,  
  loc int PRIMARY KEY,  
  type text);
```

```
SELECT year  
FROM movie_video  
WHERE loc = 1234;
```

Admitted

Query-based design: go back to the design

# Example – approach 2

- Write a program that do the following:

1. executes

```
SELECT title, director
FROM videos
WHERE loc = 1234;
```

2. Let `t` represent the returned tuple

3. Execute

```
SELECT year
FROM movies
WHERE title = t.title AND director = t.director;
```

# CQL – SELECT, restrictions: selection conditions

- The WHERE clause **must** contain an equality-based selection condition on each attribute of the partition key
  - The system is able to identify the nodes contributing to the result
  - IN clause is also allowed (a set of nodes is identified)
- The WHERE clause **can** contain selection conditions on clustering columns, following the ordering provided in the primary key, with some additional restrictions
  - On the identified nodes, rows to be returned must be sequentially stored

# Example

```
SELECT*  
FROM movies  
WHERE title = 'pulp fiction'
```

Admitted

```
CREATE TABLE movies (  
  title text,  
  director text,  
  year int,  
  genre text,  
  PRIMARY KEY (title, director);
```

```
SELECT*  
FROM movies  
WHERE title = 'pulp fiction' AND  
      director= 'quentin tarantino'
```

Admitted

Non admitted

```
SELECT*  
FROM movies  
WHERE director= 'quentin tarantino'
```

```
SELECT*  
FROM movies  
WHERE title = 'pulp fiction' AND  
      year = 2000
```

Non admitted

# CQL – SELECT, restrictions: selection conditions

- Data to be retrieved from a partition has to be sequentially stored
- Clustering columns support =, IN, >, >=, <=, <, CONTAINS operators
  - CONTAINS only for collection types, see later
- The WHERE clause can contain conditions over any prefix of the clustering column list, as defined in the primary key
- If more than one clustering column has been defined, range restrictions are allowed only on the last clustering column being restricted in the WHERE clause



```
CREATE TABLE movies (  
  title text,  
  director text,  
  year int,  
  genre text,  
  PRIMARY KEY (title, director, year);
```

# Example

```
SELECT*  
FROM movies  
WHERE title = 'pulp fiction' AND  
       director= 'quentin tarantino' AND  
       year > 2000
```

Admitted

Non admitted

```
SELECT*  
FROM movies  
WHERE title = 'pulp fiction' AND  
       year > 2000
```

```
CREATE TABLE movies (  
    title text,  
    director text,  
    year int,  
    genre text,  
    PRIMARY KEY (title, year, director);
```

# Example

```
SELECT*  
FROM movies  
WHERE title = 'pulp fiction' AND  
       director= 'quentin tarantino' AND  
       year > 2000
```

Non admitted

Admitted

```
SELECT*  
FROM movies  
WHERE title = 'pulp fiction' AND  
       year > 2000
```

# Indexes

- Exceptions to the previous rules are possible if indexes are created
- An index can be created on any column
  - B+-Tree indexes
  - Local at any node
- WHERE conditions on any attributes upon which an index has been created are admitted

```
CREATE TABLE movies (  
    title text,  
    director text,  
    year int,  
    genre text,  
    PRIMARY KEY (    title,  
                    director,  
                    year);
```

```
CREATE INDEX ON movies(year);
```

```
SELECT*  
FROM movies  
WHERE  year > 2000;
```

Admitted

# Enforcing query execution

- Executing a non admitted query, you will get the following message

*Bad Request: Cannot execute this query as it might involve data filtering and thus may have unpredictable performance. If you want to execute this query despite the performance unpredictability, use ALLOW FILTERING.*

- The system tells you that the execution of your request **might be** inefficient
- You can force the query execution by using clause **ALLOW FILTERING**  
enforcing a full scan of data (quite inefficient)

# Example

```
CREATE TABLE movies (  
    title text,  
    director text,  
    year int,  
    genre text,  
    PRIMARY KEY ( title,  
                  director,  
                  year);
```

```
CREATE INDEX ON movies(year);
```

```
SELECT*  
FROM movies  
WHERE  year > 2000;
```

Admitted

```
SELECT*  
FROM movies  
WHERE  year > 2000  
ALLOW FILTERING ;
```

Admitted

# Other clauses

- Aggregate functions: COUNT, SUM, MAX, MIN,...
- GROUP BY
  - Over primary key columns in defined order as arguments
- ORDER BY
  - Depending on the clustering order

# CQL – Collection types

- Extension of basic features of a column-family data mode with nested atomic values: **sets, lists, maps**
- Sets hold list of unique elements **that are visible to the system**
- Lists hold ordered, possibly repeating elements
- **Maps hold a list of key-value pairs**
- Collection types cannot be nested
- Collection field cannot be used in primary keys

```
CREATE TABLE mytable(  
    row text,  
    Y text,  
    myset set<text>,  
    mylist list<int>,  
    mymap map<text, text>,  
    PRIMARY KEY (row)  
);
```

# CQL – Collection types - INSERT

- Inserting

```
INSERT INTO mytable (row, myset)  
VALUES (123, { 'apple', 'banana'});
```

```
INSERT INTO mytable (row, mylist)  
VALUES (123, ['apple', 'banana', 'apple']);
```

```
INSERT INTO mytable (row, mymap)  
VALUES (123, {1:'apple', 2:'banana'});
```

```
CREATE TABLE mytable(  
    row text,  
    Y text,  
    myset set<text>,  
    mylist list<text>,  
    mymap map<int, text>,  
    PRIMARY KEY (row)  
);
```



# CQL – Collection types - UPDATE

- Updating

```
UPDATE mytable SET myset = myset + {'apple', 'banana'};
```

```
WHERE row = 123;
```

```
UPDATE mytable SET myset = myset - {'apple'};
```

```
WHERE row = 123;
```

```
UPDATE mytable SET mylist = mylist + ['apple', 'banana']
```

```
WHERE row = 123;
```

```
UPDATE mytable SET mylist = ['banana'] + mylist
```

```
WHERE row = 123;
```

```
UPDATE mytable SET mymap['fruit'] = 'apple'
```

```
WHERE row = 123
```

```
UPDATE mytable SET mymap = mymap + { 'fruit': 'apple'}
```

```
WHERE row = 123
```

```
CREATE TABLE mytable(  
    row text,  
    Y text,  
    myset set<text>,  
    mylist list<text>,  
    mymap map<text, text>,  
    PRIMARY KEY (row)
```

# Example – collection types

```
CREATE TABLE clients (  
  codCli int PRIMARY KEY,  
  name text,  
  surname text,  
  birthdate date,  
  emails set<text>,  
  phones map<text, text>,  
  hobbies list<text>  
);
```

```
CREATE INDEX ON clients (name);  
CREATE INDEX ON clients (emails);  
CREATE INDEX ON clients (hobbies);  
CREATE INDEX ON clients (keys(phones));  
CREATE INDEX ON clients (values(phones));  
CREATE INDEX ON clients (entries(phones));
```

```
SELECT * FROM clients  
WHERE name = 'Benjamin';
```

```
SELECT * FROM clients  
WHERE emails CONTAINS  
      'Benjamin@oops.com';
```

```
SELECT * FROM clients  
WHERE hobbies CONTAINS tennis';
```

```
SELECT * FROM clients  
WHERE phones CONTAINS KEY 'office';
```

```
SELECT * FROM clients  
WHERE phones CONTAINS '0108567586';
```

```
SELECT * FROM clients  
WHERE phones['office'] = '0108567586';42
```

# CQL - User Defined Types

- User defined types allow a compact usage of complex data
- Nesting
- An alternative approach for dealing with associations and joins

# Example

```
client: {  
    codCli,  
  
    name, surname, birthdate,  
    address: {city, street, streetNumber, postalCode},  
    movies: [ {comment, title, director,...}]  
}
```

```
CREATE TYPE address_t(  
    city text,  
    street text,  
    streetNumber int,  
    postalCode int);
```

```
CREATE TYPE movie_t(  
    comment text,  
    title text,  
    director int,  
    ...);
```

```
CREATE TABLE clients (  
    codCli int PRIMARY KEY,  
    name text,  
    surname text,  
    birthdate date,  
    address frozen<address_t>,  
    movies set<frozen<movie_t>>  
);
```

Values for `address_t` treated as a blob (no way to select or update components, you should do that at the application level)

Values for `movie_t` treated as a blob (no way to select components, you should do that at the application level)

# Frozen types

- Frozen <> makes a component **opaque** for the system
- Frozen<> can be used with both user defined types and collection types, **it is mandatory in most situations**
- With frozen <>, we can overcome most limitations related to collection types and UDTs
- Collection types and UDTs nesting **is allowed by using frozen types**
- Frozen types can be used inside either frozen or non frozen collection types
- Collection fields with frozen types can be used as primary key

# Architecture

# Cassandra in short

Feature	In Cassandra
Model	Column-family
Architecture	P2P
Query language	Supported (CQL)
Reference scenarios	Transactional (read/write intensive)
Partitioning	Consistent hashing
Indexes	Local secondary index
Replication	Leader-less
Consistency	Eventual consistency
Availability	High, tunable
Fault tolerance	High (no master, P2P ring)
Transactions	No ACID transactions
CAP theorem	AP
Distributed by	Apache third-parties

# CAP theorem

## Visual Guide to NoSQL Systems

**A**  
Availability:  
Each client can  
always read  
and write.

**Data Models**

Relational (comparison)  
Key-Value  
Column-Oriented/Tabular  
Document-Oriented

**CA**

RDBMSs  
(MySQL,  
Postgres,  
etc)

Aster Data  
Greenplum  
Vertica

**AP**

Dynamo  
Voldemort  
Tokyo Cabinet  
KAI

Cassandra  
SimpleDB  
CouchDB  
Riak

**Pick Two**

**C**

**Consistency:**  
All clients always  
have the same view  
of the data.

**CP**

BigTable  
Hypertable  
Hbase

MongoDB  
Terrastore  
Scalaris

Berkeley DB  
MemcacheDB  
Redis

**P**

**Partition Tolerance:**  
The system works  
well despite physical  
network partitions.

Cassandra

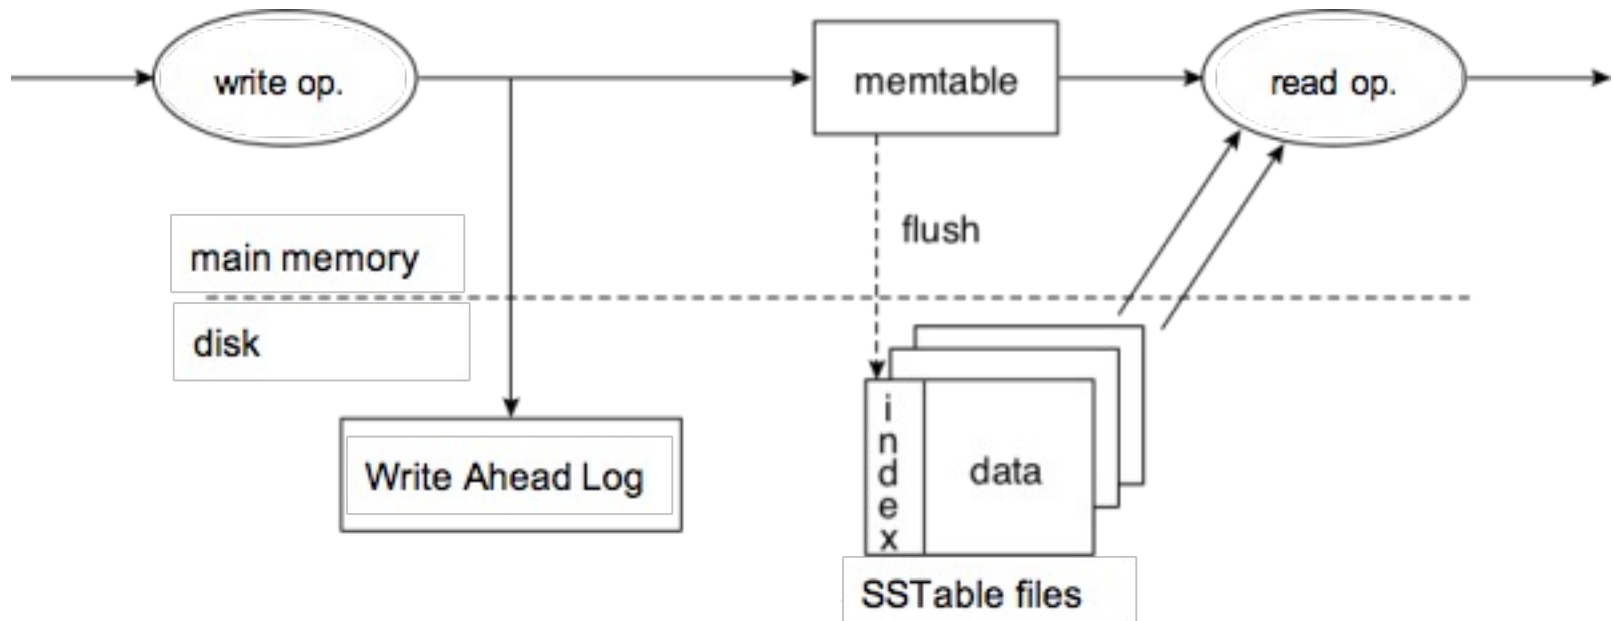
AP: Highly  
available system



# Local Persistence

- Organization of local data store at nodes relies on **local file system**
- Write operations happens in three steps
  - Write to **commit log** in local disk of the node
  - Update **in-memory data structure** (write-back cache of data rows that can be looked up by key)
  - When the in-memory structure is full, it is written (flushed) on disk as an **SSTable**
- Read operation
  - Looks up in-memory data structures first before looking up files on disk

# Local persistence



Persistency + **durability**  
but also high **throughput** of **write** operations

# Configurable read/write consistency

- A **consistency level** can be specified for a given query session or for each read and write query
- It allows to set parameters **w** and **r** in the formula  $w + r > n$  ( $n$  number of replicas)
- Common consistency levels

**ALL** – All nodes in the cluster must confirm (even non-local nodes)

**QUORUM** – A quorum of nodes (half the replication factor plus one) in the cluster must confirm

**ONE / TWO / THREE** – One, two or three nodes in the cluster must confirm

WRITE		
Level	Description	Usage
ALL	A write must be written to the <a href="#">commit log and memtable</a> on all replica nodes in the cluster for that partition.	Provides the highest consistency and the lowest availability of any other level.
QUORUM	A write must be written to the <a href="#">commit log and memtable</a> on a quorum of replica nodes across <i>all</i> datacenters.	Use in single or multiple datacenter clusters to maintain strong consistency across the cluster. Use if you can tolerate some level of failure.
ONE	A write must be written to the <a href="#">commit log and memtable</a> of at least one replica node.	Satisfies the needs of most users because consistency requirements are not stringent.
TWO	A write must be written to the <a href="#">commit log and memtable</a> of at least two replica nodes.	Similar to ONE.
THREE	A write must be written to the <a href="#">commit log and memtable</a> of at least three replica nodes.	Similar to TWO.

READ		
Level	Description	Usage
ALL	Returns the record after all replicas have responded. The read operation will fail if a replica does not respond.	Provides the highest consistency of all levels and the lowest availability of all levels.
QUORUM	Returns the record after a quorum of replicas from all <a href="#">datacenters</a> has responded.	Used in single or multiple datacenter clusters to maintain strong consistency across the cluster. Ensures strong consistency if you can tolerate some level of failure.
ONE	Returns a response from the closest replica, as determined by the <a href="#">snitch</a> . By default, a <a href="#">read repair</a> runs in the background to make the other replicas consistent.	Provides the highest availability of all the levels if you can tolerate a comparatively high probability of stale data being read. The replicas contacted for reads may not always have the most recent write.
TWO	Returns the most recent data from two of the closest replicas.	Similar to ONE.
THREE	Returns the most recent data from three of the closest replicas.	Similar to TWO.

# Transactions

- No ACID transactions with rollback or locking mechanisms
- Cassandra offers Atomic, Isolated, and Durable (AID) transactions with eventual/tunable consistency (no classical consistency, no integrity constraints)

Column-family stores use cases

# Suitable use cases



- *Applications with a high number of writes*
- Event Logging
  - Great choice to store event information
- Content Management Systems, Blogging Platforms
  - You can store blog entries with tags, categories, links, and trackbacks in different columns
  - Comments can be either stored in the same row or moved to a different keyspace
  - Similarly, blog users and the actual blogs can be put into different column families
- Counters
  - Often, in Web applications you need to count and categorize visitors of a page to calculate analytics
  - Counter can be maintained in Cassandra

# When not to use



- Systems that require ACID transactions for writes and reads
- Prototypes
  - during the early stages query patterns may change and this requires to change the column family design
  - in Cassandra, the cost may be higher for query change as compared to schema change



# Cassandra, MapReduce, Hadoop

- Integration with Hadoop MapReduce (MapReduce I/O from/to Cassandra)
- Integration with Pig (Pig I/O from/to Cassandra)