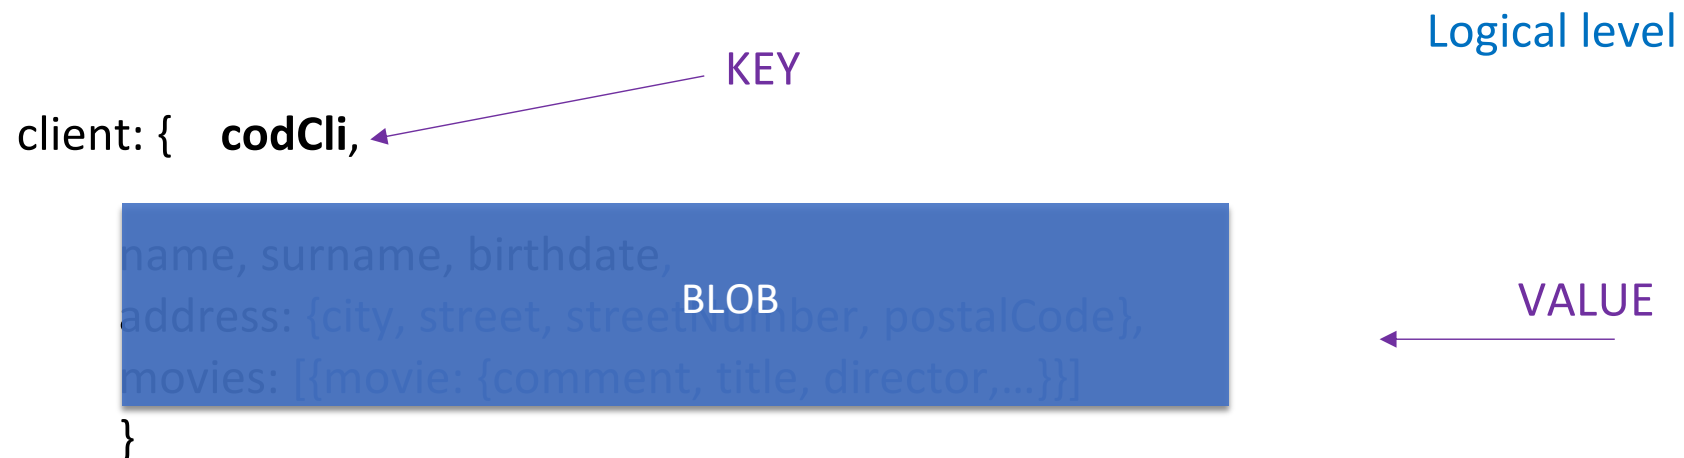


Key-value NoSQL data stores

Key-value data model at a glance

- Each data instance is represented in the form (key, value)
 - key is an identifier
 - value is the aggregate
- The aggregate is totally opaque at the logical level
 - just a big blob of mostly meaningless bit, that the system just stores
 - values do not have a known structure
 - very high flexibility: arbitrary aggregate content
- The aggregate is visible at the application level
- Data instances can be grouped into logical collections



Key-value data model at a glance

- Each data instance is represented in the form (key, value)
 - key is an identifier
 - value is the aggregate
- The aggregate is totally opaque at the logical level
 - just a big blob of mostly meaningless bit, that the system just stores
 - values do not have a known structure
 - very high flexibility: arbitrary aggregate content
- The aggregate is visible at the application level
- Data instances can be grouped into logical collections

Logical level

client: { **codCli**, ← KEY

name, surname, birthdate,
address: {city, street, streetNumber, postalCode},
movies: [{movie: {comment, title, director,...}}]
}

← VALUE

Instance structure

- No schema information
- No nested values are visible at the NoSQL system level but only at the application level

Collections

- Pairs can be grouped into logical **collections/namespaces**: sets of aggregates, i.e., sets of (key, value) pairs

| | |
|-----|-------|
| key | value |
| key | value |
| key | value |
| key | value |

In RDBMS world: A table with two columns:

ID column (primary key)

DATA column storing the value
(unstructured BLOB)

A

| | |
|---|----|
| 1 | V1 |
| 2 | V2 |
| 3 | V3 |
| 4 | V4 |

```
"type": "dvd",  
"rentals": [{"rental": {"rentalDate": "15/10/2021",  
                        "codCli": 375657}}],  
"title": "pulp fiction",  
"director": "quentin tarantino"
```

...

```
"name": "John",  
"surname": "Black",  
"birthdate": "15/10/2000",  
"address": {"city": "Genoa", "street": "Via XX Settembre",  
            "streetNumber": 15, "postalCode": 16100},  
"movies": [{"movie": {"comment": "very nice", "title": "pulp fiction",  
                      "director": "quentin tarantino"}},  
            {"movie": {"comment": "very nice", "title": "pulp fiction",  
                      "director": "quentin tarantino"}}]
```

(see buckets in Riak)

Collections and schema information

- The only «schema» information the system is aware of is the collection
- There is no structure specification nor schema constraint in the data model
- But at application-level, we can take advantage of some regularities in data, if any

Collections

- Pairs can be grouped into logical **collections/namespaces**: sets of aggregates, i.e., sets of (key, value) pairs

| | |
|-----|-------|
| key | value |
| key | value |
| key | value |
| key | value |

Videos

| | |
|------|-------|
| 1234 | value |
| 5678 | value |
| 9999 | value |
| 0000 | value |

Clients

| | |
|--------|-------|
| 375657 | value |
| 375658 | value |
| 375659 | value |
| 375660 | value |

(see buckets in Riak)

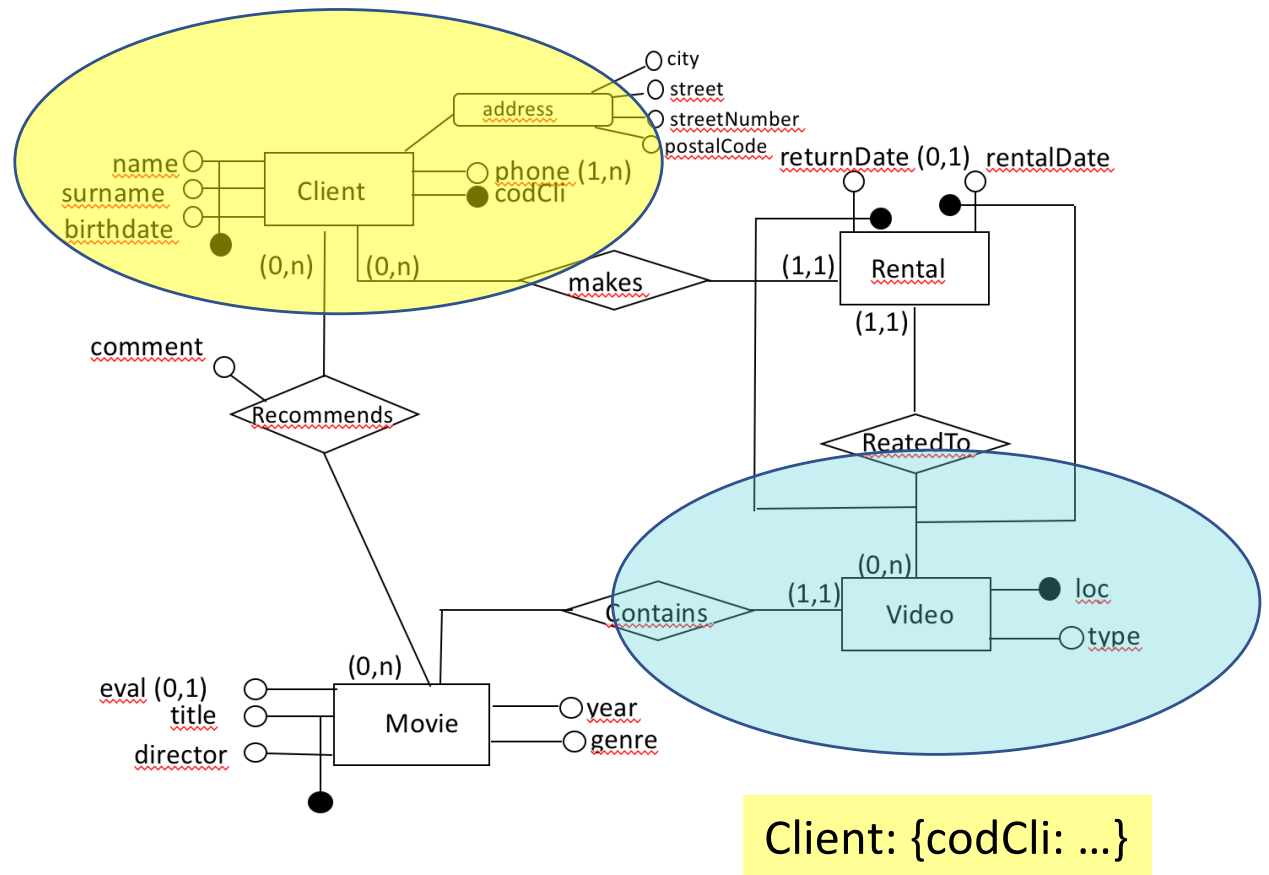
Keys

- (key, value) pair
- At the logical level: identification + data retrieval
 - the key **needs** to assume **unique values** in the collection (i.e., it is an **identifier**)
- At the physical level, the key tells the system **how to partition the data** and where to store the data
 - **partition key**
- The partition key and the identifier/primary key coincide
- Aggregates can be directly retrieved **only by specifying values for attributes in the partition key**
- *A simple **hash table** (map), primarily used when all accesses to the data are via the **key***

Keys

- How to **design** the key?
 - **Provided** by the user (natural unique key): userID, e-mail,...
 - **Generated** by some algorithm
 - **Derived** from time-stamps (or other data)
- Use cases with **natural keys**:
 - user profiles (user ID), ...
 - shopping cart data (**user ID**)
 - web session data (with the **session ID** as the key)
- **Expiration** of keys
 - After a certain **time interval**
 - e.g. for caches, session/shopping cart objects,...

Keys



Videos

| | |
|------|-------|
| 1234 | value |
| 5678 | value |
| 9999 | value |
| 0000 | value |

Key is loc

Clients

| | |
|--------|-------|
| 375657 | value |
| 375658 | value |
| 375659 | value |
| 375660 | value |

Key is codeCli

Keys

There is no single unique attribute, different options are possible

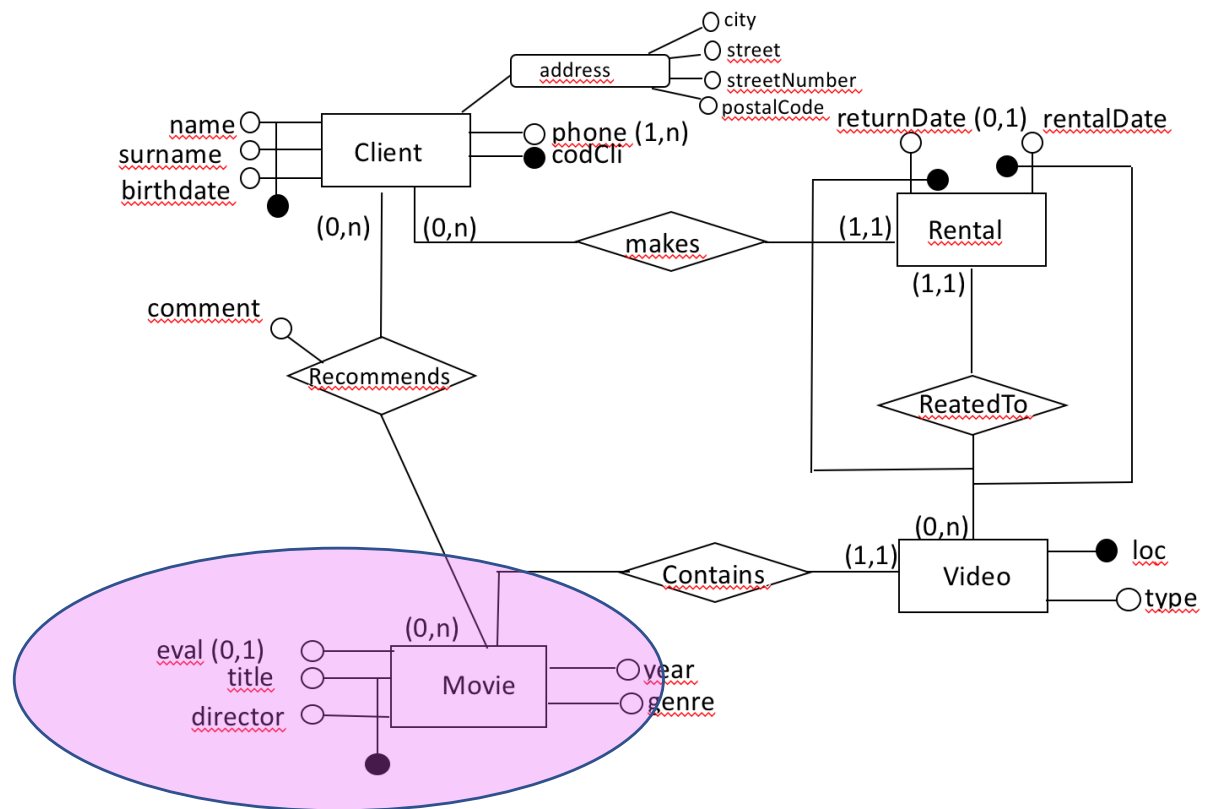
1. Id
2. (title and director are included in the value)

Movie : {id: ...}

Movies

| | |
|---|-------|
| 1 | value |
| 2 | value |
| 3 | value |
| 4 | value |

Key is id



Keys

There is no single unique attribute, different options are possible

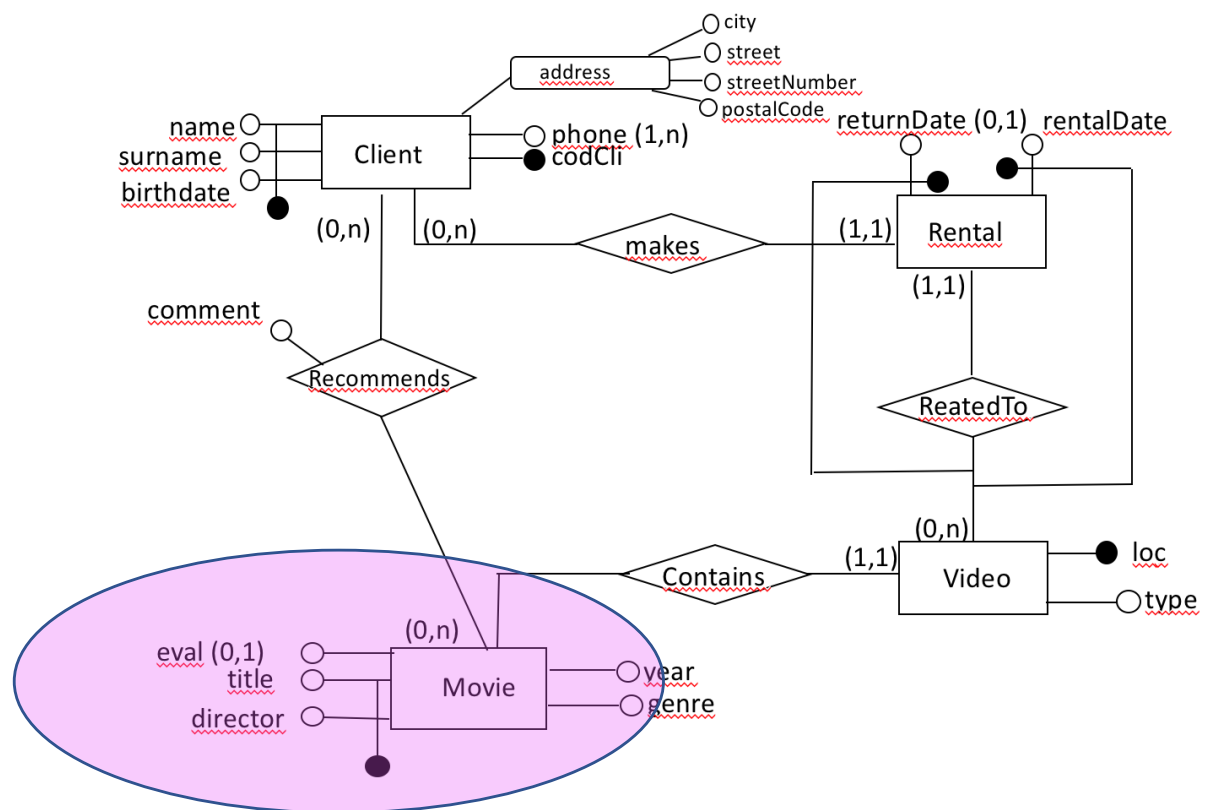
2. title:director

Movie : {'title:director': ...}

Movies

| | |
|--------------------------------|-------|
| pulp fiction:quentin tarantino | value |
| kill bill:quentin tarantino | value |
| le iene:quentin tarantino | value |
| dumbo : gabriele salvatores | value |

Key is title:director

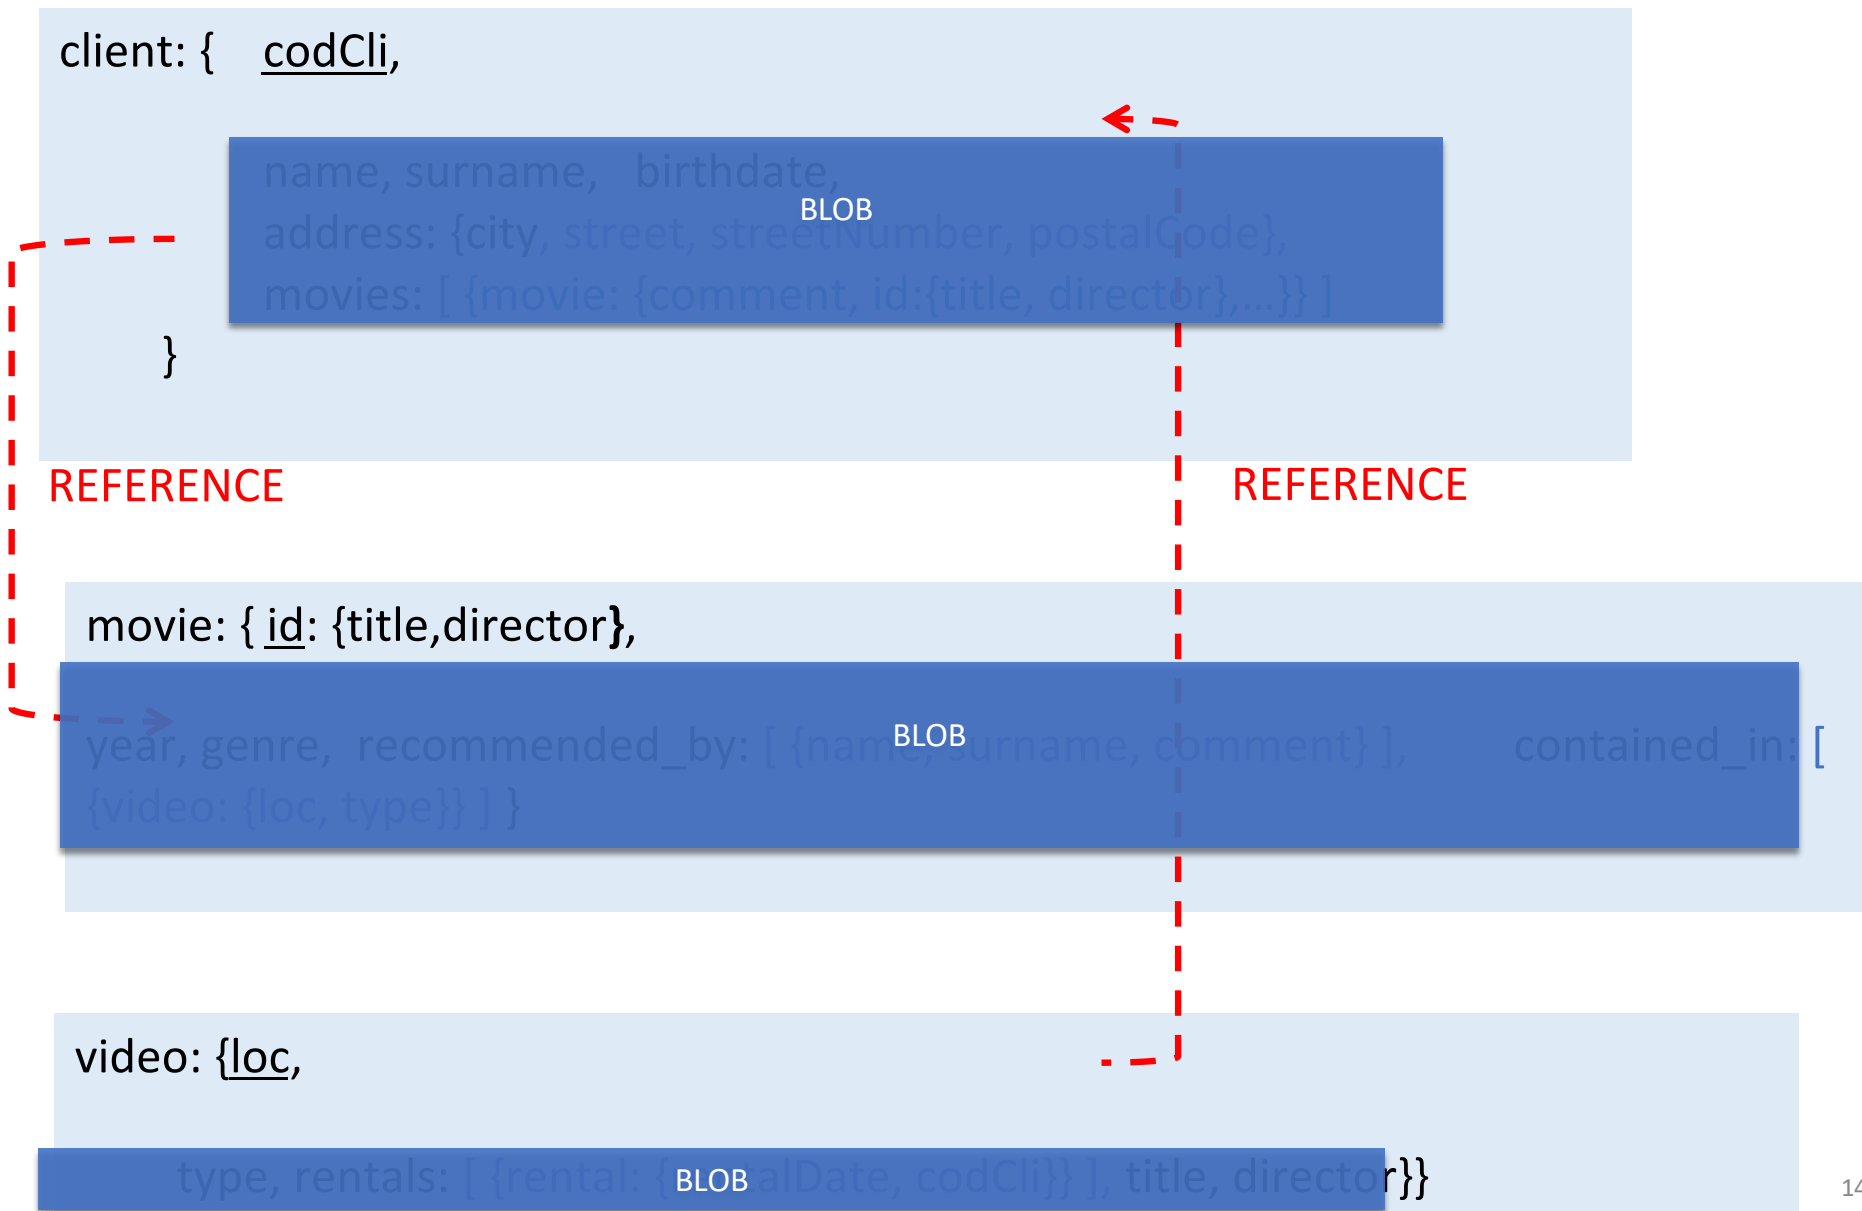


Identifier vs partition/sharding key

- In the Videos collection, videos are partitioned by loc
- In the Clients collection, clients are partitioned by codCli
- In the Movies collection, clients are partitioned by
 - Id
 - Title:director
- This impacts the way data are stored
- This impacts the way data can be retrieved

Video Rental Example

Logical level



Video Rental Example

Application level

```
client: { codCli,  
  name, surname, birthdate,  
  address: {city, street, streetNumber, postalCode},  
  movies: [ {movie: {comment, id:{title, director},...}} ]  
}
```

REFERENCE

REFERENCE

```
movie: { id: {title,director},  
  year, genre, recommended_by: [ {name, surname, comment} ],  
  contained_in: [ {video: {loc, type}} ] }
```

```
video: {loc,  
  type, rentals: [ {rental: {rentalDate, codCli}} ], title, director}}
```

Interaction

- Basic **operations**:

- **Put** a value for a key
(if the key already exists
the corresponding value
is overwritten)

```
put(key, value)
```

- **Get** the value for the key

```
value := get(key)
```

- **Delete** a key (and the
corresponding value)

```
delete(key)
```


Interaction

- **Lookup** based on the **key**
- The application can read the entire aggregate by using the key
- Queries with respect to specific aggregate fields are not supported (in general): we need to read the whole aggregate and check **query conditions at the application level**
- **Associations navigated by sequence of lookups**
- In case collections are supported, key values are preceded by the collection name in each operation

Example 1 – Find the title of the movie in a video

```
{"loc": 1234,
```

```
  "type": "dvd",  
  "rentals": [{"rental": {"rentalDate": "15/10/2021",  
                        "codCli": 375657}}],  
  "title": "pulp fiction",  
  "director": "quentin tarantino"}
```

`get(Videos, 1234)`

- at the data store level, find Video 1234 in namespace Videos, no detailed information about its content
- at the *application level*, we go inside the aggregate value and we discover that it contains «Pulp Fiction»

Example 2 – Find the videos containing a certain movie

```
{"loc": 1234,
```

```
  "type": "dvd",  
  "rentals": [{"rental": {"rentalDate": "15/10/2021",  
                           "codCli": 375657}}],  
  "title": "pulp fiction",  
  "director": "quentin tarantino"}
```

`get(Videos,???)`

- We cannot filter on other attributes than the key
- At the data store level, can only get all the videos, and
- At the application level, filter them by titles and director

Example 2 – Find the videos containing a certain movie - we should use the movies collection instead

`get(Movies,pulp fiction:quentin tarantino)`

- At the data store level, get the value
- At application level, find *contained_in* and get the location

Movies collection

Key: title:director,

Value contains {year, genre,

recommended_by: [{name, surname, comment}],

contained_in: [{video: {loc, type}}]}

REMARK: the relationship is part of the aggregate, a single data access to retrieve all the relevant information (=a single data node)

Example 2 – Find the videos containing a certain movie - we should use the movies collection instead, what if we chose the other key?

Movies collection

Key:MovieID,

Value contains {title, director, year, genre,
recommended_by: [{name, surname, comment}],
contained_in: [{video: {loc, type}}]}

get(Movies,???)

- We cannot filter on other attributes than the key
- At the data store level, can only get all the movies, and
- At the application level, filter them by titles and director

REMARK: still likely more efficient than starting from Videos, since Movies are less than videos and all the videos containing a certain movie are stored inside the same aggregate

Example 3 - Relationships

Find the age(s) of customer(s) that rented a given video

```
{"loc": 1234,
```

```
  "type": "dvd",  
  "rentals": [{"rental": {"rentalDate": "15/10/2021",  
                        "codCli": 375657}}],  
  "title": "pulp fiction",  
  "director": "quentin tarantino"}
```

```
{  
  "codcli": 375657,  
  "name": "John",  
  "surname": "Black",  
  "birthdate": "15/10/2000",  
  "address": {"city": "Genoa", "street": "Via XX Settembre",  
             "streetNumber": 15, "postalCode": 16100},  
  "movies": [{"movie": {"comment": "very nice", "title": "pulp fiction",  
                      "director": "quentin tarantino"}},  
             {"movie": {"comment": "very nice", "title": "pulp fiction",  
                      "director": "quentin tarantino"}}]  
}
```

get(Videos, 1234)

- At the data store level, find Video 1234 in namespace Videos, no detailed information the clients that rented it
- at the *application level*, we go inside the aggregate value and we discover that it was rented by Client 375657
- at the *application level*, we can execute get(Clients, 375657) to retrieve information about a/the customer that rented video 1234 (thus navigating the customer reference stored inside the video)

Example 3 - Relationships

```
{"loc": 1234,
```

Find the age(s) of customer(s) that rented a given video?

```
"type": "dvd",  
"rentals": [{"rental": {"rentalDate": "15/10/2020",  
                        "codCli": 375657,  
                        "title": "pulp fiction",  
                        "director": "quentin tarantino"}}
```

REMARKS:

1. Who knows that inside BLOB1 there is the code of the client that rented the video? And where to look for the client information and how to match?
- The application, the data store is completely unaware of that!**
- 2. Since the relationship is not part of the aggregate, navigating it requires a sort of join, no system support for it
 - 3. Since the relationship is not part of the aggregate, navigating it (at application level) requires two distinct data accesses, at two possibly distinct data nodes
- ... Videos, no detailed
- ... aggregate value and we discover
- ... 375657
- ... we can execute get(Clients, 375657) to retrieve
- ... at a/the customer that rented video 1234 (thus
- ... the customer reference stored inside the video)

Advanced interaction

-
- Some systems support additional functionality
 - Use of indexes: The data must be indexed first
 - Some kind of additional index (e.g. full text) can be used
 - Example: Riak search

Key-value stores use cases


K-V Stores: Suitable Use Cases

- Storing **Web Session** Information
 - Every web session is assigned a **unique session_id** value
 - Everything about the session can be **stored by** a **single** PUT request or retrieved using a **single** GET
 - **Fast**, everything is stored in a **single object**
- **User Profiles**, Preferences
 - Every user has a unique **user_id/user_name** + preferences (language, time zone, design, access rights, ...)
 - As in the previous case: Fast, single object, single GET/PUT
- **Shopping Cart** Data
 - Similar to the previous cases

K-V Stores: When Not to Use

- **Relationships** among Data
 - Relationships between **different sets** of data
 - **Some** key-value stores (Riak) provide **link-walking** features
- Multi-operation **Transactions**
 - **Saving multiple** keys
 - Failure to save any of them → revert or roll back the rest of the operations
- Query by Data
 - **Search the keys** based on something found in the value part
 - **Additional indexes** needed (some stores provide them)
- Operations by **Key Sets**
 - Operations are limited to one key at a time
 - **No way** to operate upon **multiple keys** at the same time

Popular key-value data stores



Ranked list: <http://db-engines.com/en/ranking/key-value+store>

Key-value stores are quite diverse

- **Dozens** of key-value stores – are all of them the same?
- **Embedded** local storages
 - LevelDB
 - **Local storage** for many systems, Log-structured Merge Tree
- **Distributed** key-value Stores
 - Riak, Infinispan
- Memory **caches**
 - Redis, Memcached



CAP theorem: key-value stores

