

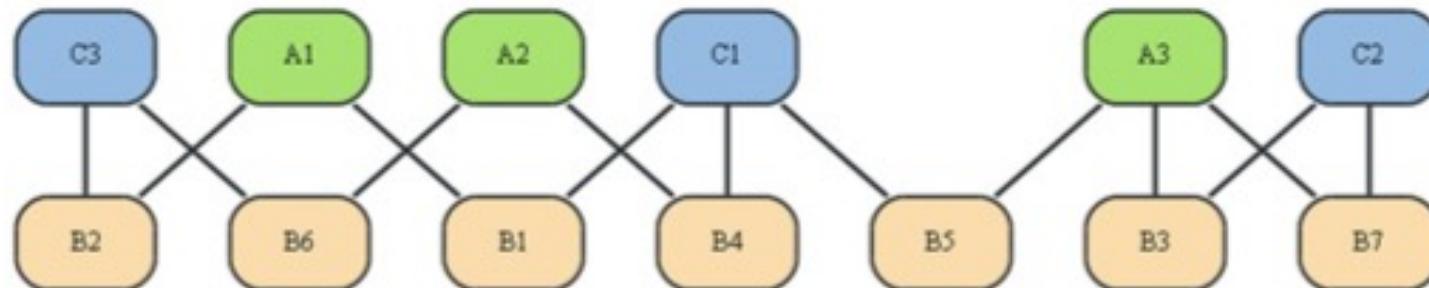
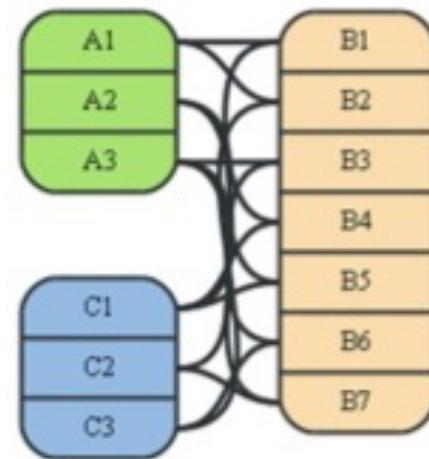
Graph Databases - neo4j

Outline

- Short intro to graph databases
- Neo4j
 - Introduction
 - Data model & interaction
 - Architecture
- Use cases

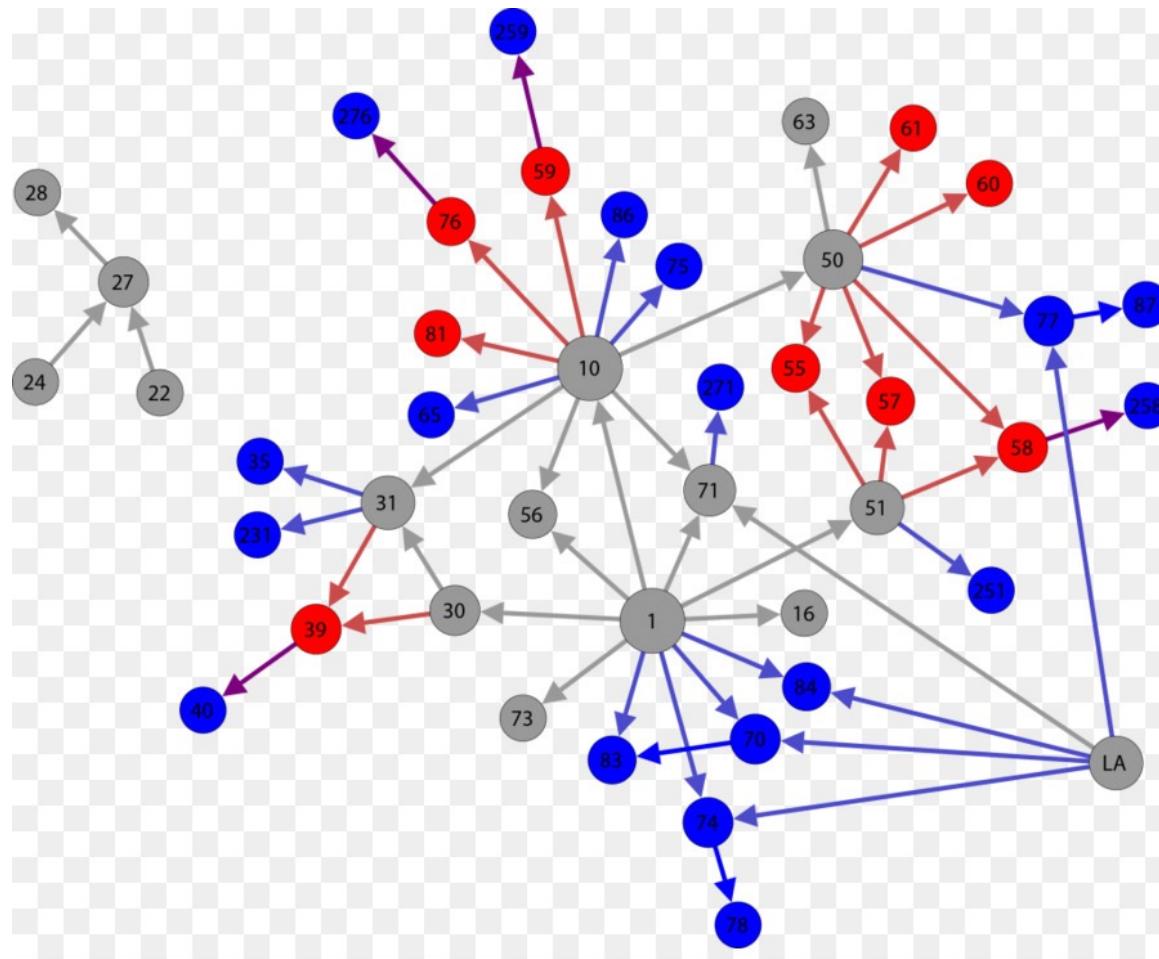
Graph databases

- Graph databases are motivated by the issues of relational databases in coping with complex relationships
- Complex relationships require complex join



Graph databases

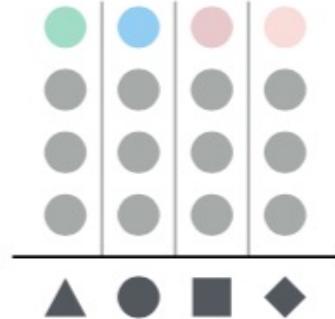
- Capture data consisting of complex relationships
 - Focus on data inter-connection and topology



Why Graphs? (recall)

- Well-known generic data structure
- One way to address the impedance mismatch
(objects in your Java don't match the structure of a DB)
- Maths and algorithms are well understood
- «blackboard friendly»
- Data naturally modelled as graphs

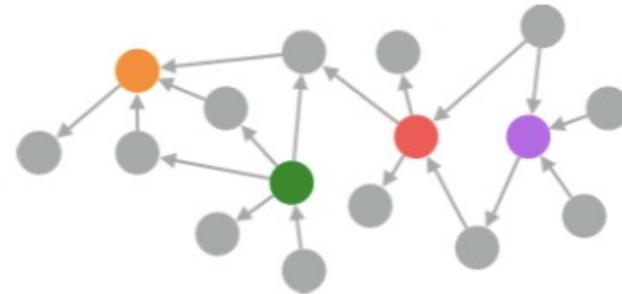
Which is Good at What?



Relational
Database



Graph
Database



Good for:

- Well-understood data structures that don't change too frequently
- Known problems involving discrete parts of the data, or minimal connectivity

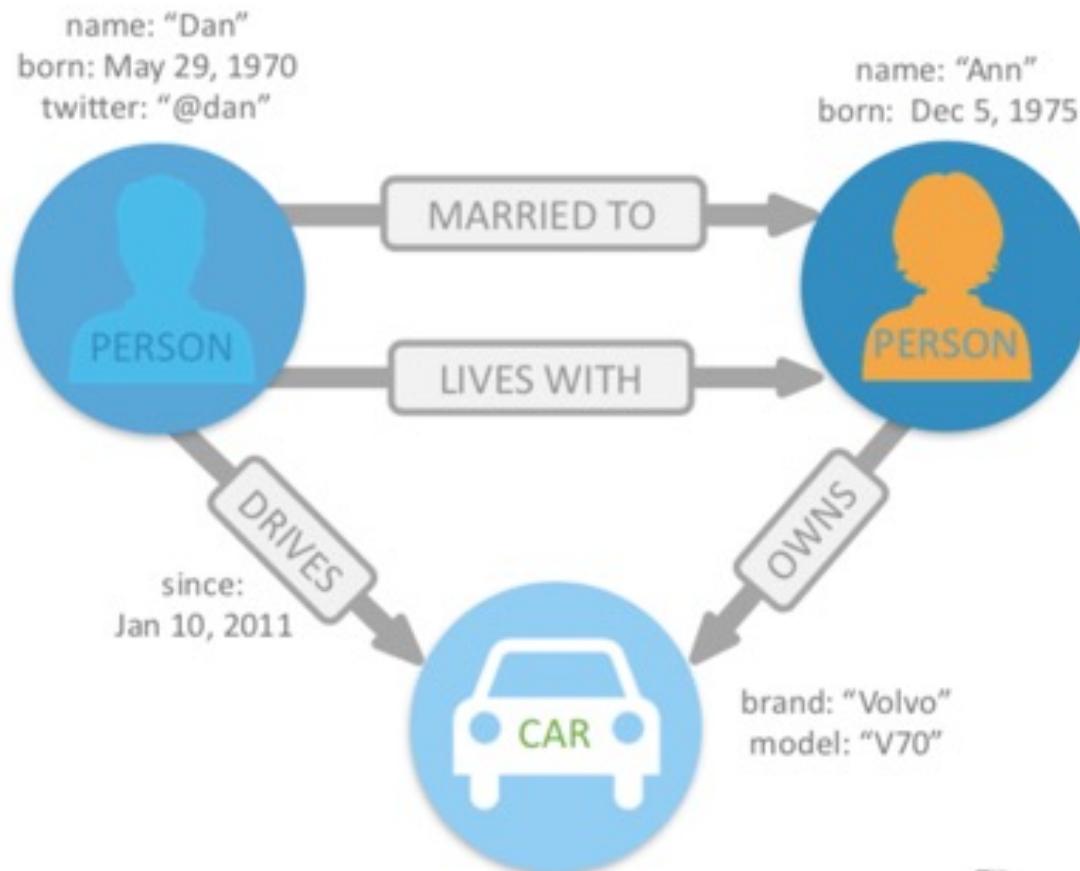
Good for:

- Dynamic systems: where the data topology is difficult to predict
- Dynamic requirements: that evolve with the business
- Problems where the relationships in data contribute meaning & value

Graph Model

- Complex, densely-connected domains
 - Lots of join tables? Relationships
 - Lots of sparse tables? Semi-structure
- Messy data
 - Ad hoc exceptions
- Relationships as **first-class elements**
 - Semantic clarity: named, directed
 - Not simply constraints
 - Can have properties

Graph Data Model (property graph)



NoSQL databases

graph vs aggregate-oriented databases

Graph databases

- graph-based query languages
- **partitioning is difficult:**
more likely to run on a single server
- transactions maintain **consistency** over multiple nodes and edges

Aggregate-oriented

- simple query languages
- distributed across clusters
- no ACID guarantees

Popular graph dbs/systems



Property graph model: Neo4j, Titan, InfiniteGraph, ...

Triple store model (RDF): AllegroGraph, Virtuoso, ...

Neo4j - Introduction

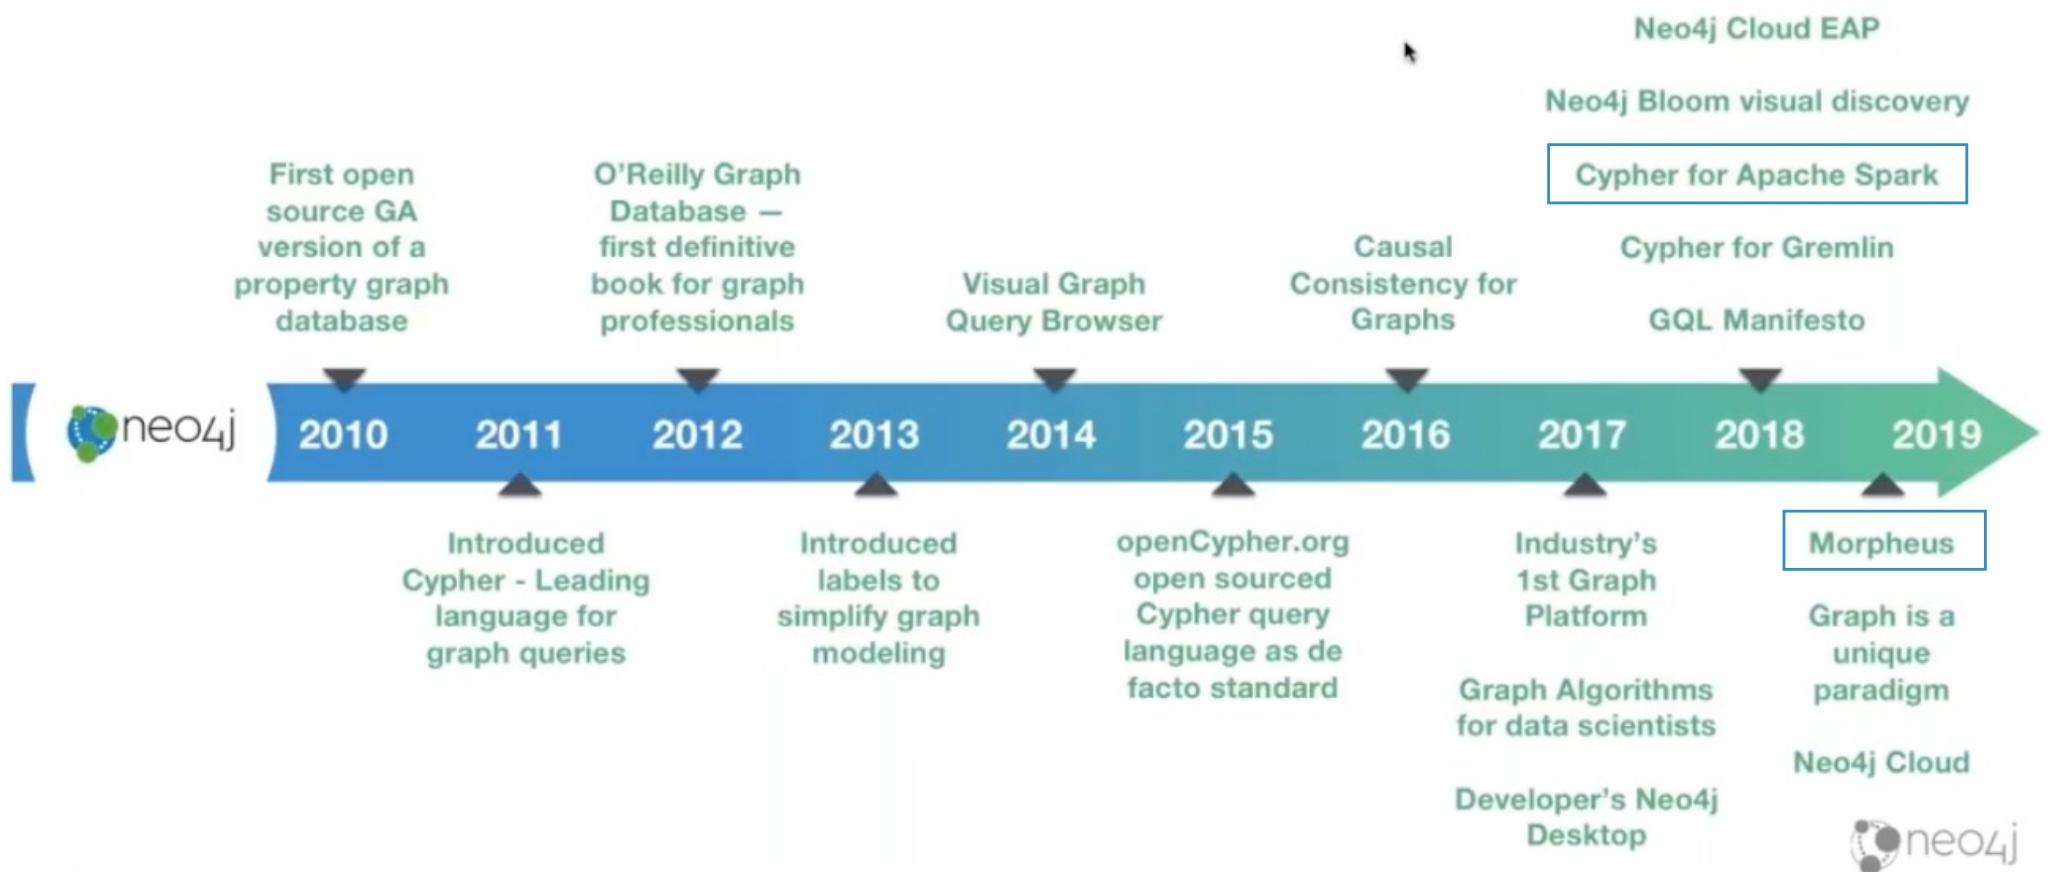


neo4j in short

Feature	neo4j
Model	Graph-based
Query language	Supported, Cypher
Reference scenarios	transactional (read intensive) & analytical
Partitioning	Difficult (application-level)
Indexes	on properties (simple and composite), full-text
Replication	Master-slave, single-leader
Consistency	Strong
Availability	Load balancing among read replicas
Fault tolerance	By re-electing a master in case it goes down
Transactions	ACID (maintain consistency over multiple nodes and edges)
CAP theorem	CA
Distributed by	Neo4j Inc.

neo4j

- First released in 2010
- Developed in Java
- Community and Enterprise versions, dual-licensed: GPL v3 and a commercial license
 - the Community Edition is free but is limited to running on one node only due to the lack of clustering
- Query language: Cypher
- Java traversal API



Who uses neo4j?



v o l v o



<https://neo4j.com/customers>

Data model and interaction

Property Graph, Traversals, and Cypher Query Language

Property Graph Model

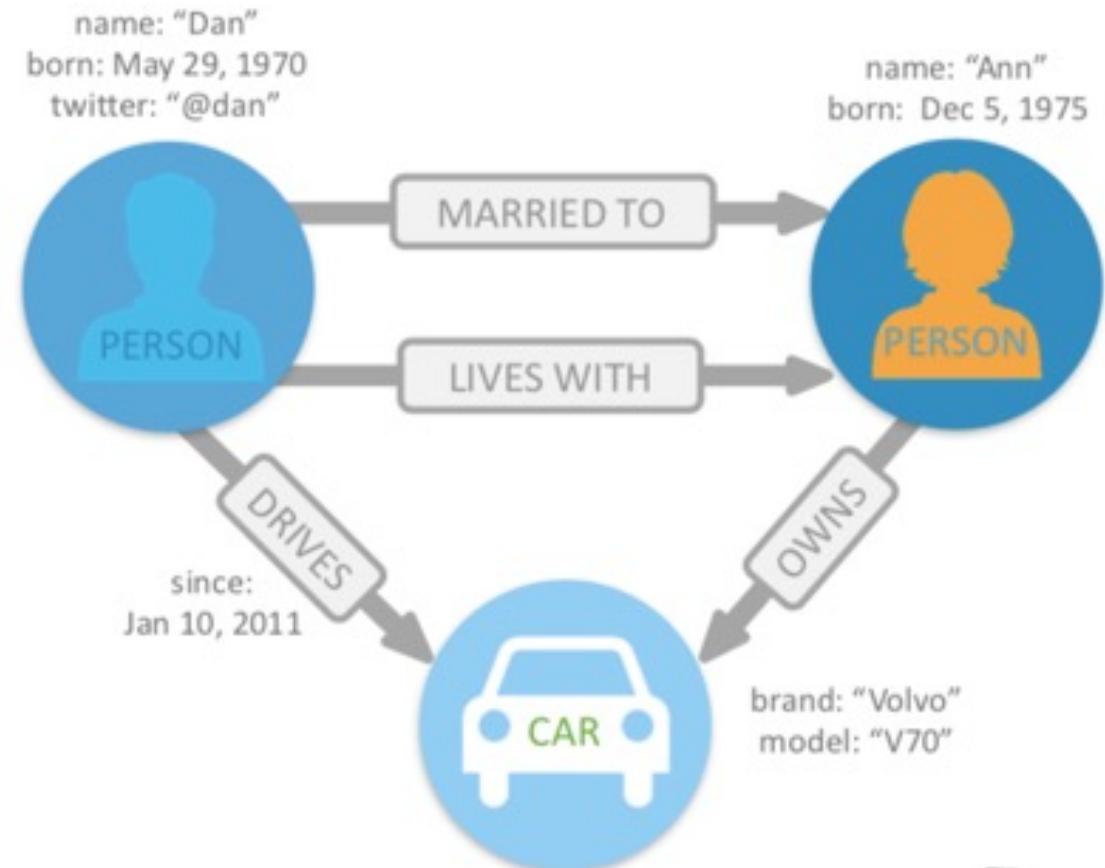
A property graph $G = (V, E)$ is a **directed multigraph** where:

- every node $v \in V$ and every edge $e \in E$ can be associated with a set of pairs $\langle \text{key}, \text{value} \rangle$, called **properties**
- value can be of primitive type or an array of primitive type
- every node $v \in V$ can be tagged with one or more **label(s)**
- every edge $e \in E$ is associated with a **type**

Property Graph Model

Edges have incoming/outgoing edges but traversal is equally efficient in both directions

Schema-less model



Graph vs Relational Data Model

- The **relational model** is designed for a single type of **relationship**
 "
- **Adding** another relationship usually means a lot of **schema changes**
- In RDBMS **we model** the graph beforehand based on the **traversal** we want
 - If the traversal changes, the data will have to change

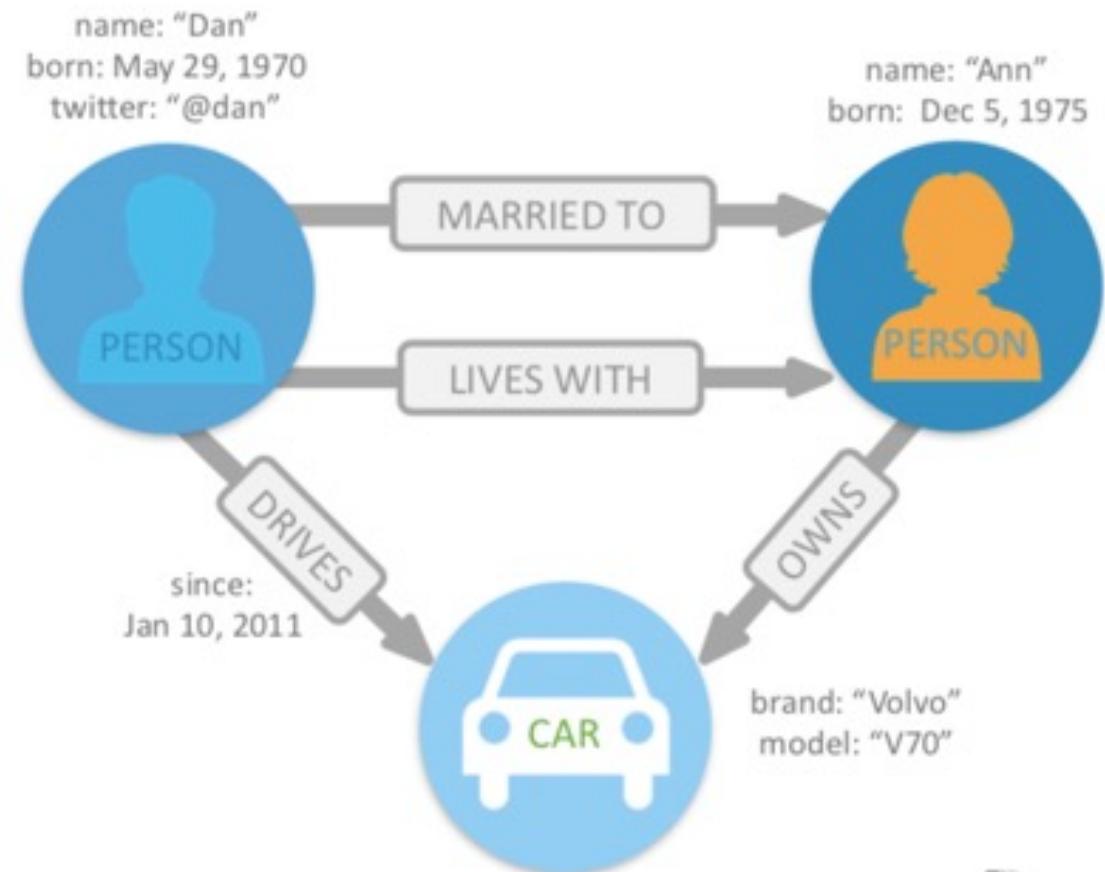
Schemaless Model: Flexibility

- Flexible property graph schema
- Neither node labels nor relationship types prescribe properties
 - freely add new edges
 - freely add properties to nodes and edges
- Schema can adapt over time
 - easy to add new nodes and relationships when the business needs change
 - changing existing nodes and their relationships is costly! (similar to data migration)

Queries as Traversals

Queries = Traversals

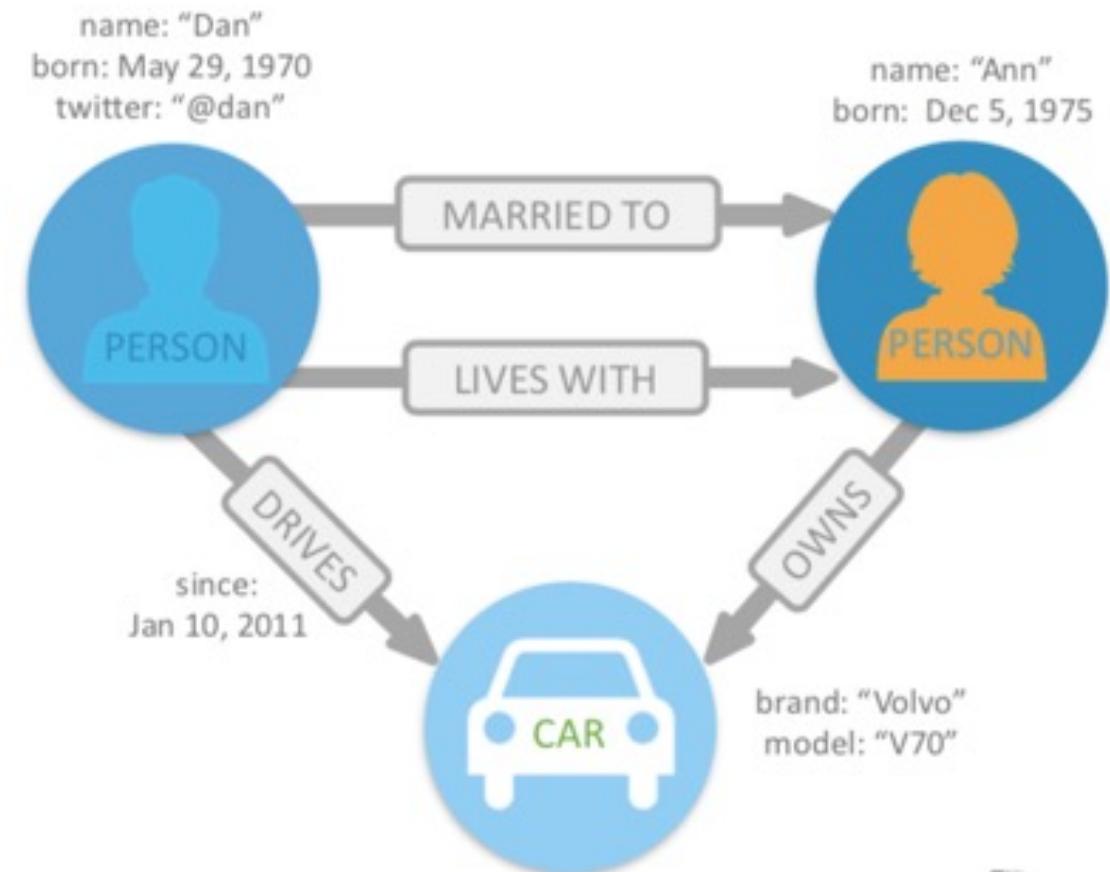
get the spouse of the owner of a Volvo



Queries = Traversals

get the spouse of the owner of a Volvo

get the owner of a car a person drives

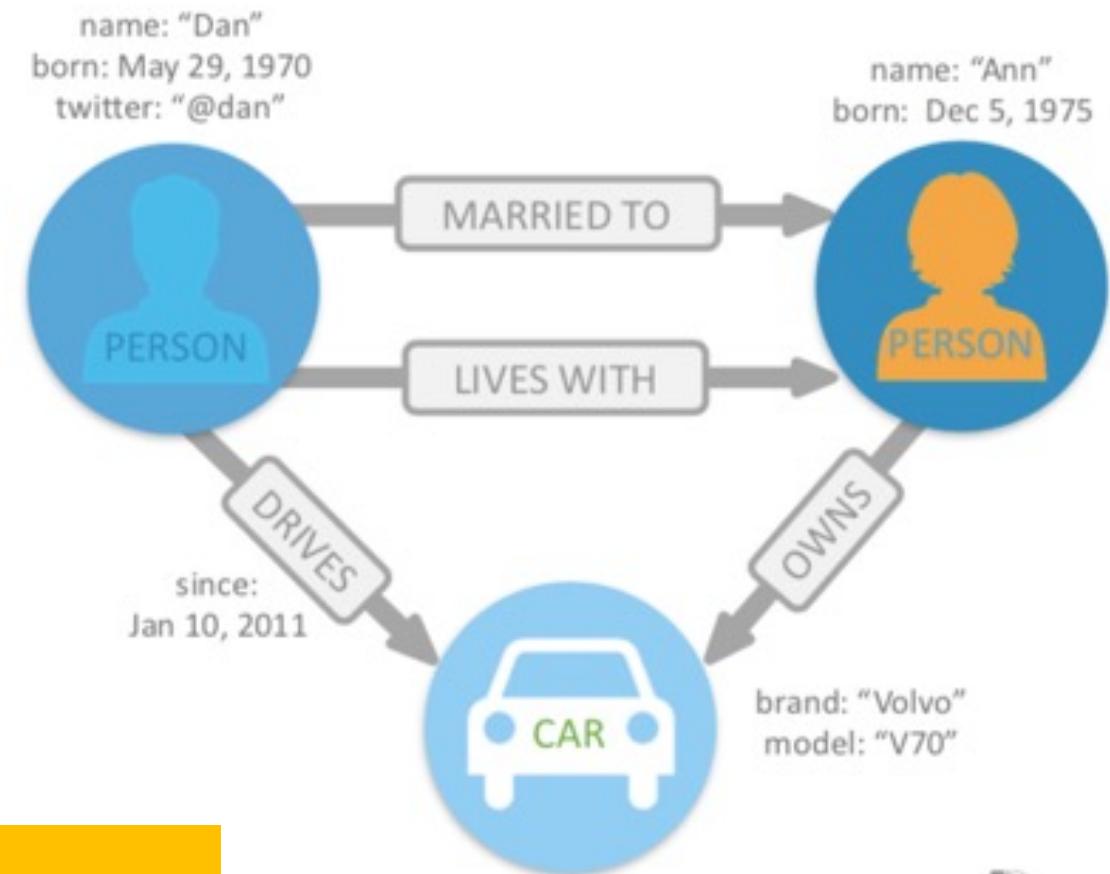


Queries = Traversals

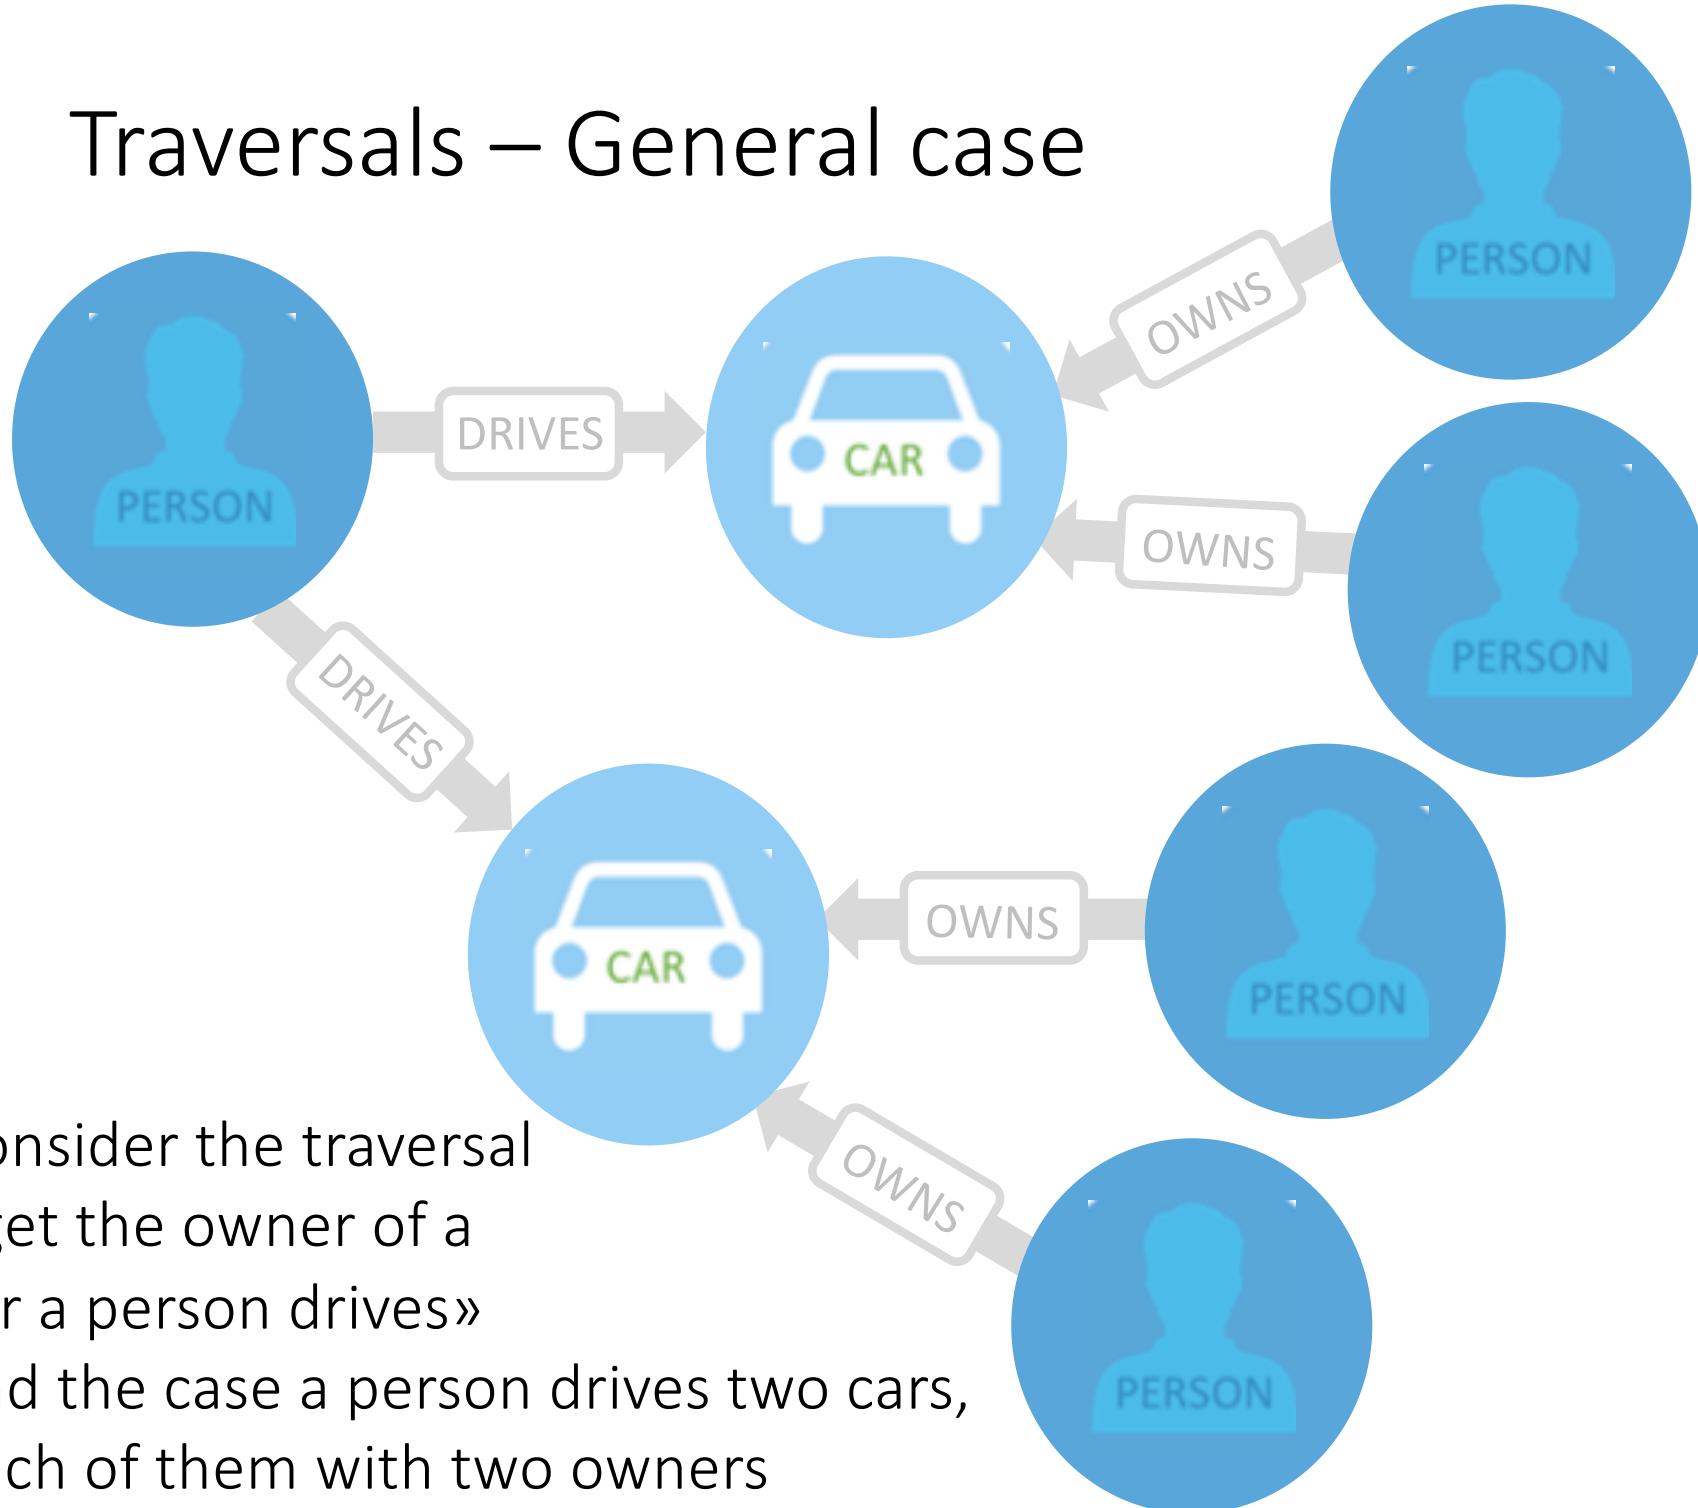
get the spouse of the owner of a Volvo

get the owner of a car a person drives

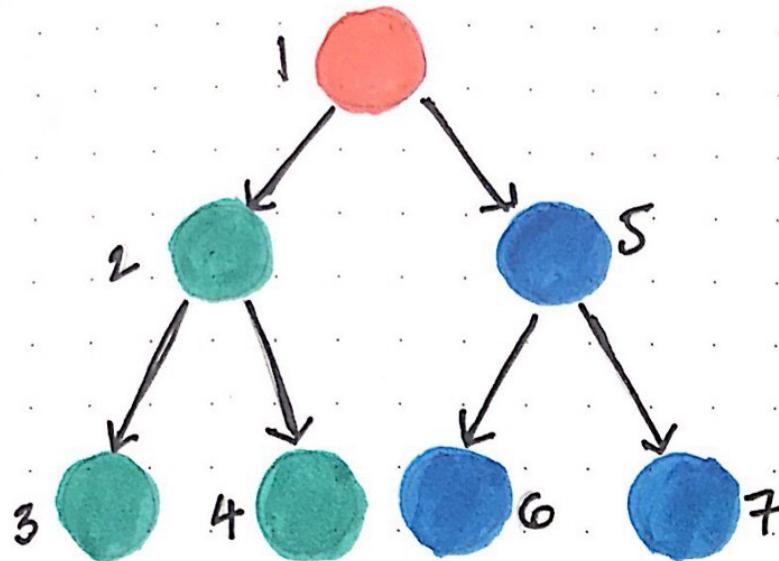
get the spouse of the owner of a car a person drives



Traversals – General case

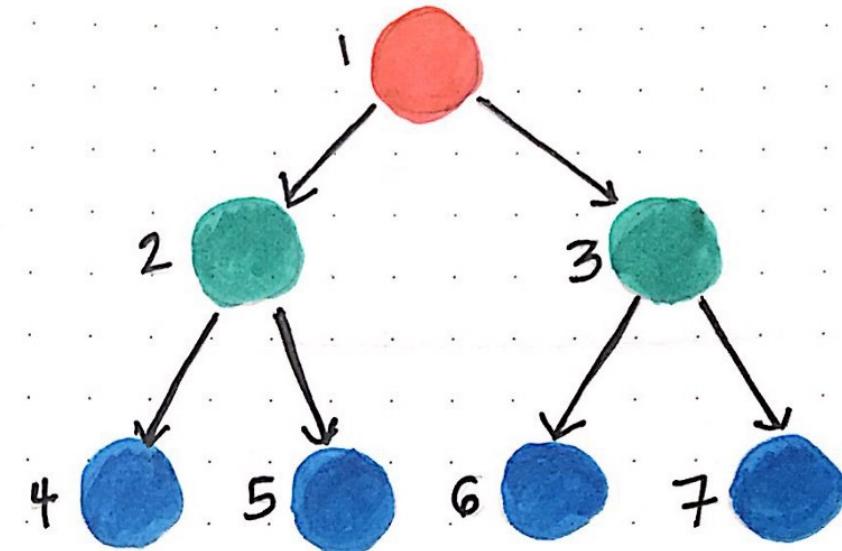


Graph Traversals



Depth-first search

- Traverse through left subtree(s) first, then traverse through the right subtree(s).



Breadth-first search

- Traverse through one level of children nodes, then traverse through the level of grandchildren nodes (and so on...).

Constant Time Traversal

- Constant time traversals in big graphs for both depth and breadth due to efficient representation of nodes and relationships
- Enables scale-up to billions of nodes on moderate hardware

Graph traversal

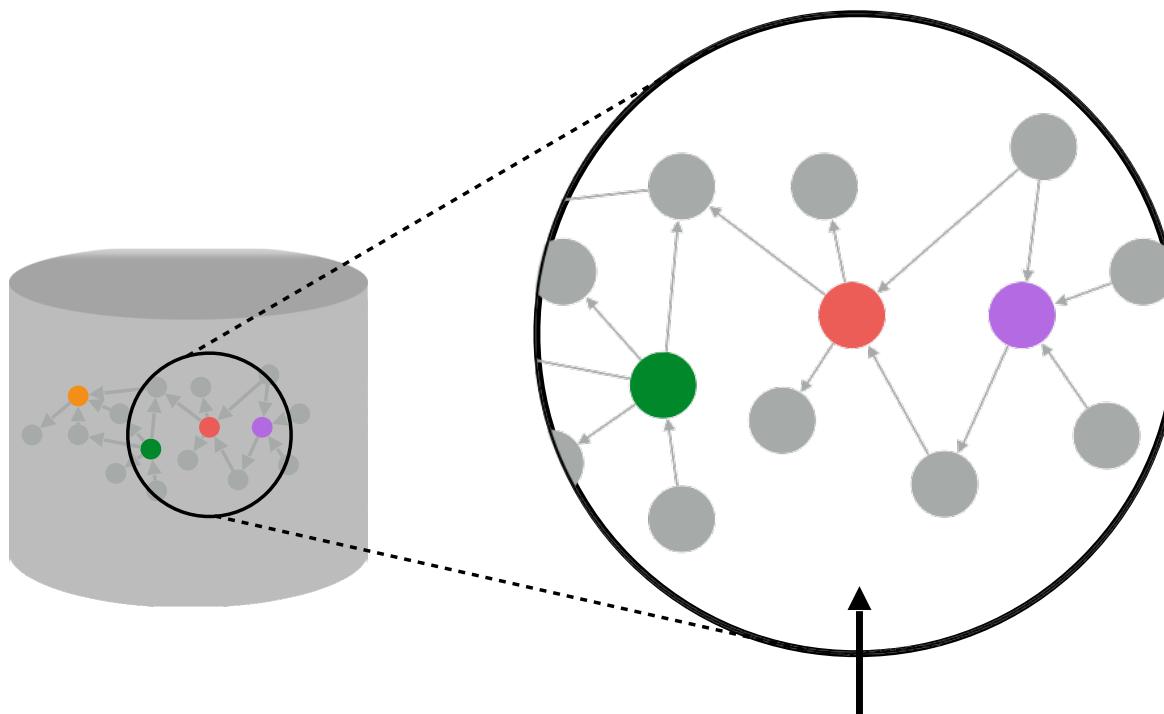
- Traversing relationships is very fast
- The relationship between nodes is not calculated at query time but it is actually persisted
- Traversing persisted relationships is faster than calculating them for every query

Index-free adjacency

We say that a (graph) database g satisfies the index-free adjacency if the existence of an edge between two nodes n_1 and n_2 in g can be tested on those nodes and does not require to access an external, global, index.

The cost of a basic traversal is independent of the size of the database

Index-free adjacency



At Write Time:
Data is *connected*
as it is stored

At Read Time:
Lightning-fast retrieval of data and relationships
via pointer chasing

Fast Graph Traversal

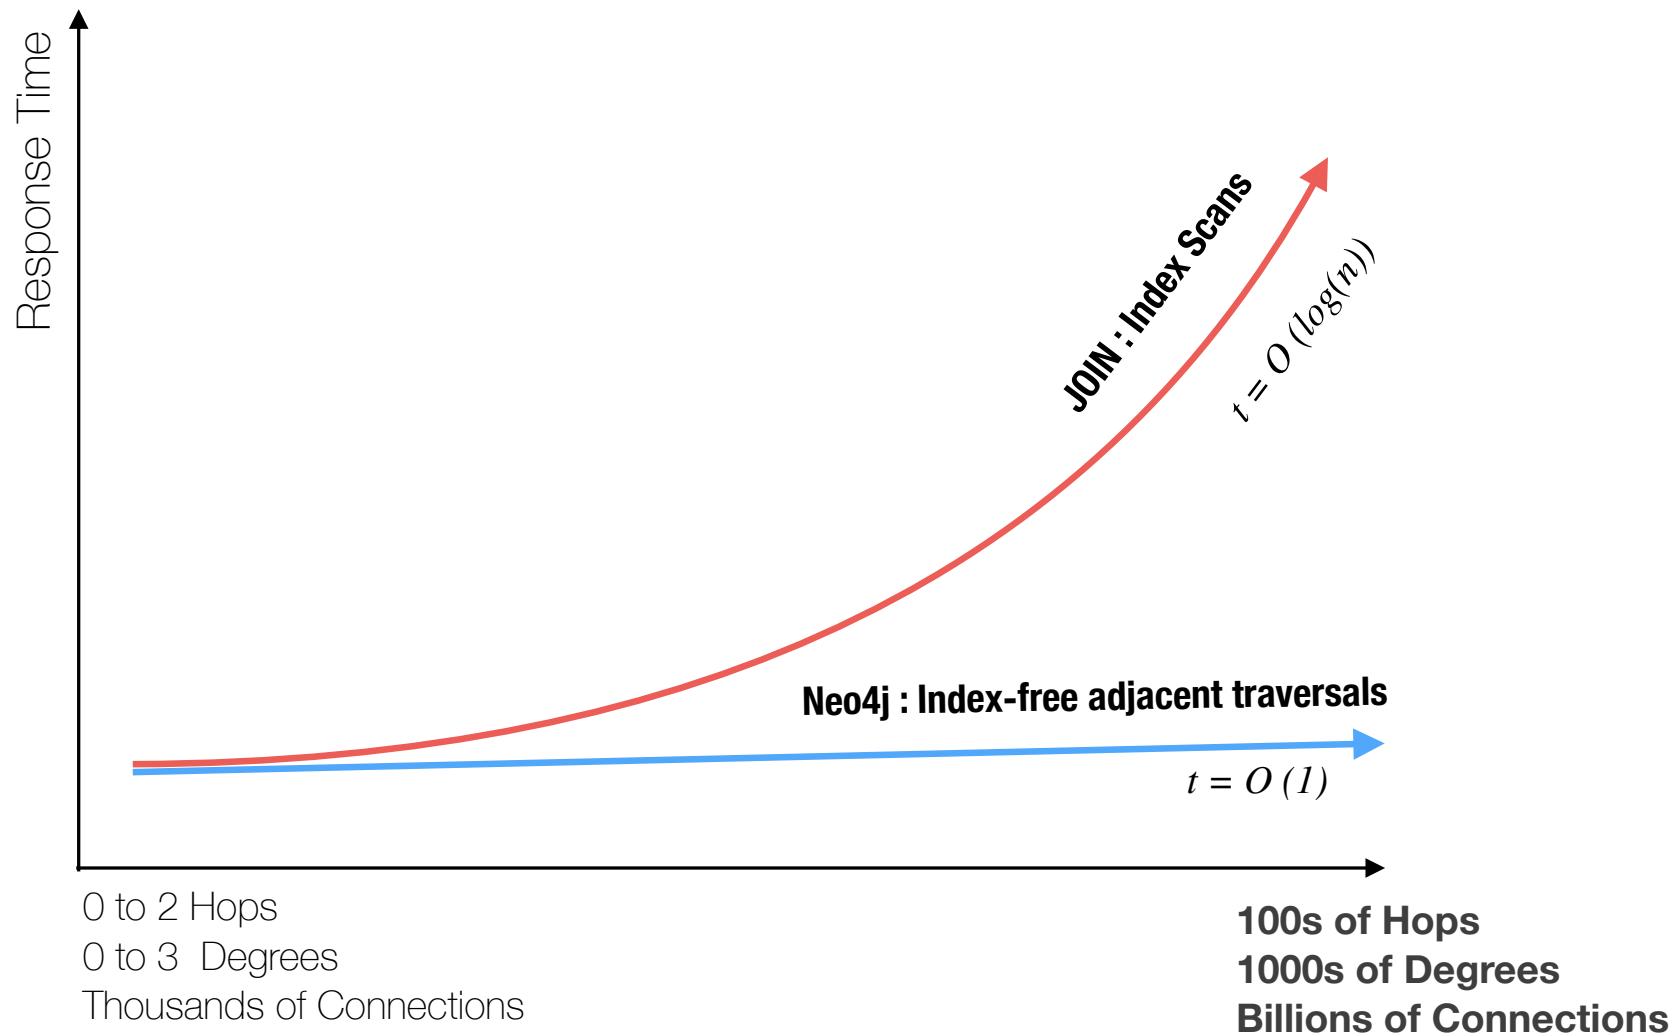
How Fast is Fast?

- Sample Social Graph with roughly 1,000 persons
- On average each person has 50 friends
- `pathExists(a,b)` limited to depth 4
- Caches warmed up to eliminate disk I/O

DATABASE	# OF PERSONS	QUERY TIME
MySQL	1,000	2,000 ms
Neo4j	1,000	2 ms
Neo4j	10,000,000	2 ms

Index-Free Traversal

Connected-Data Query Performance



Graph vs relational data model

- Relational databases
 - implement relationships using foreign keys
 - joins require to navigate around and can get quite expensive
- Graph databases
 - make traversal along the relationships very cheap
 - performance is better for highly connected data
 - shift most of the work from query time to insert time
 - good when querying performance is more important than insert speed

Cypher Query Language

Declarative

- Most of the time, Neo4j knows better than you

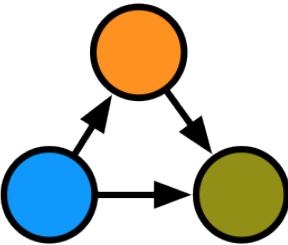
Imperative

- follow relationship
breadth-first vs depth-first
- explicit algorithm

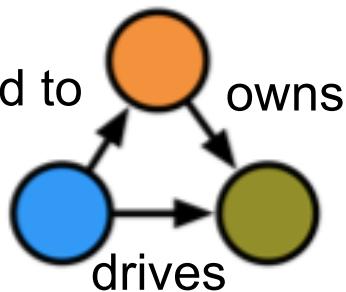
Declarative

- specify starting point
specify desired outcome
- algorithm adaptable based
on query

Pattern Matching



Pattern

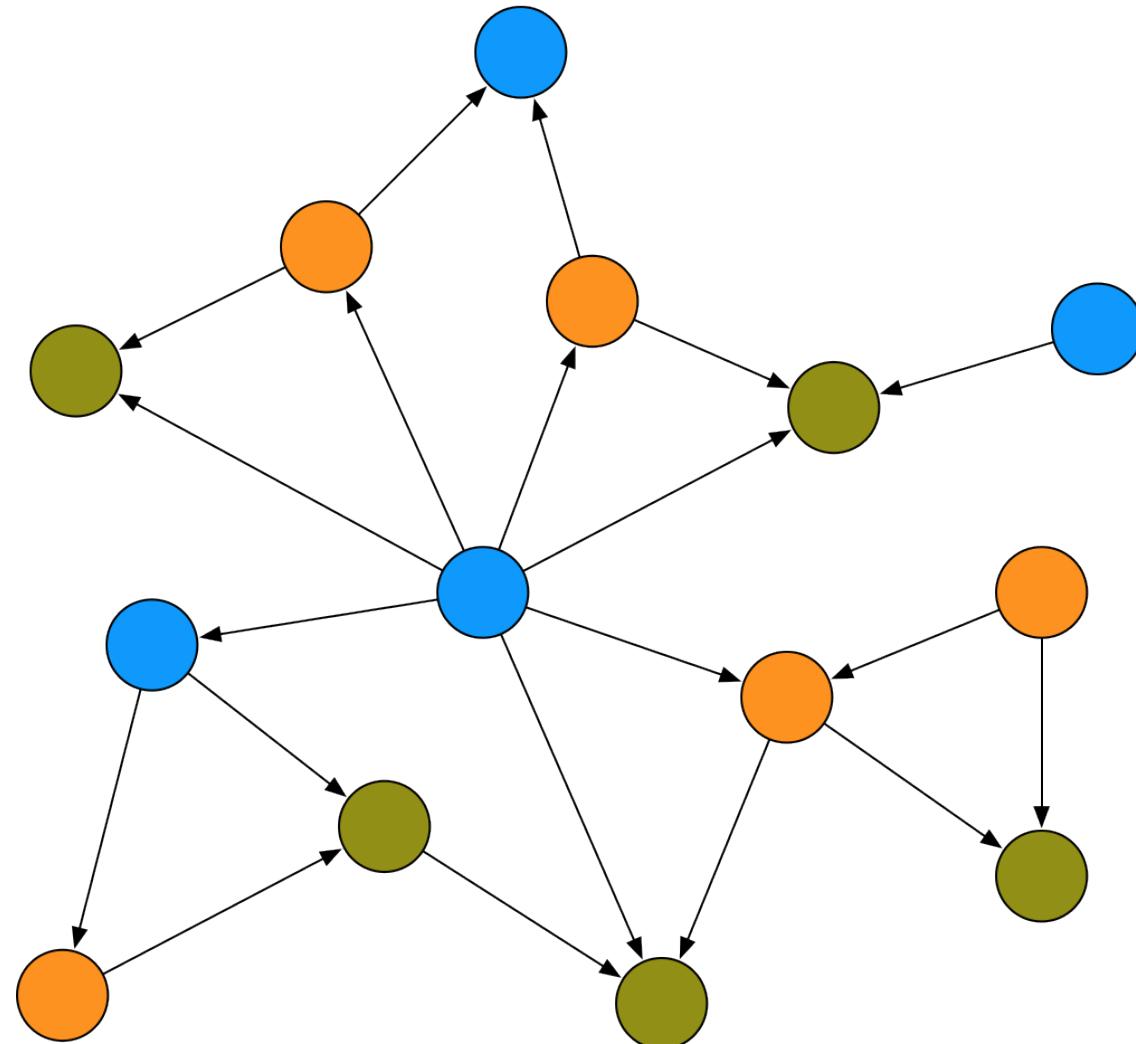


drives

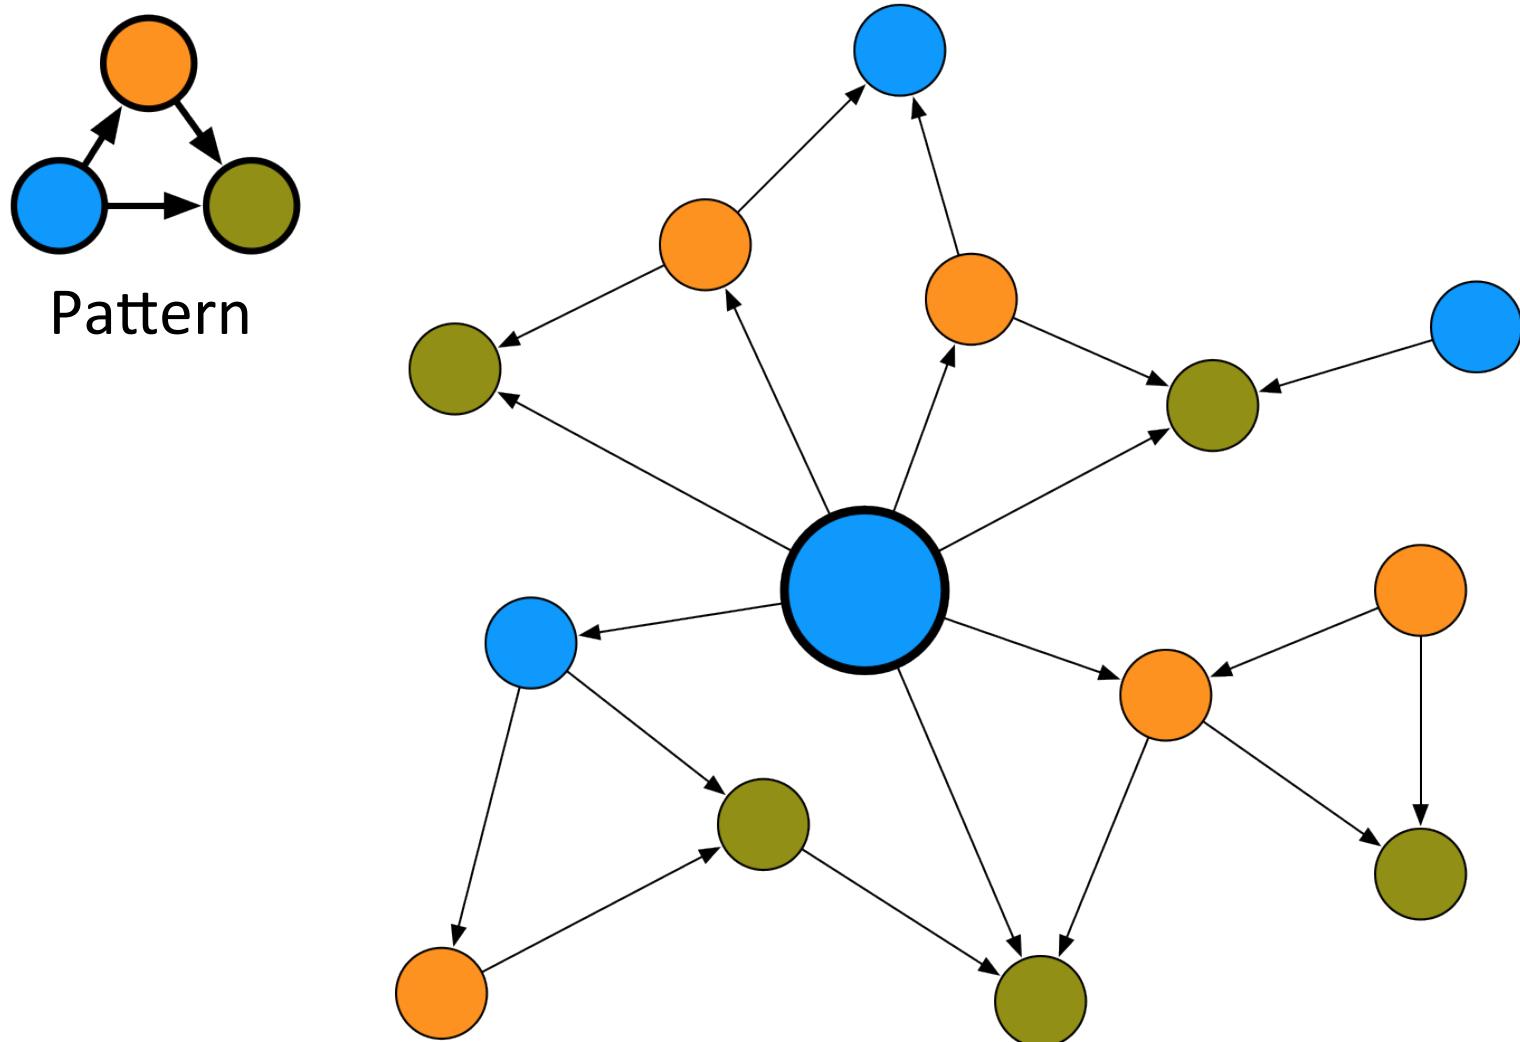
owns

married to

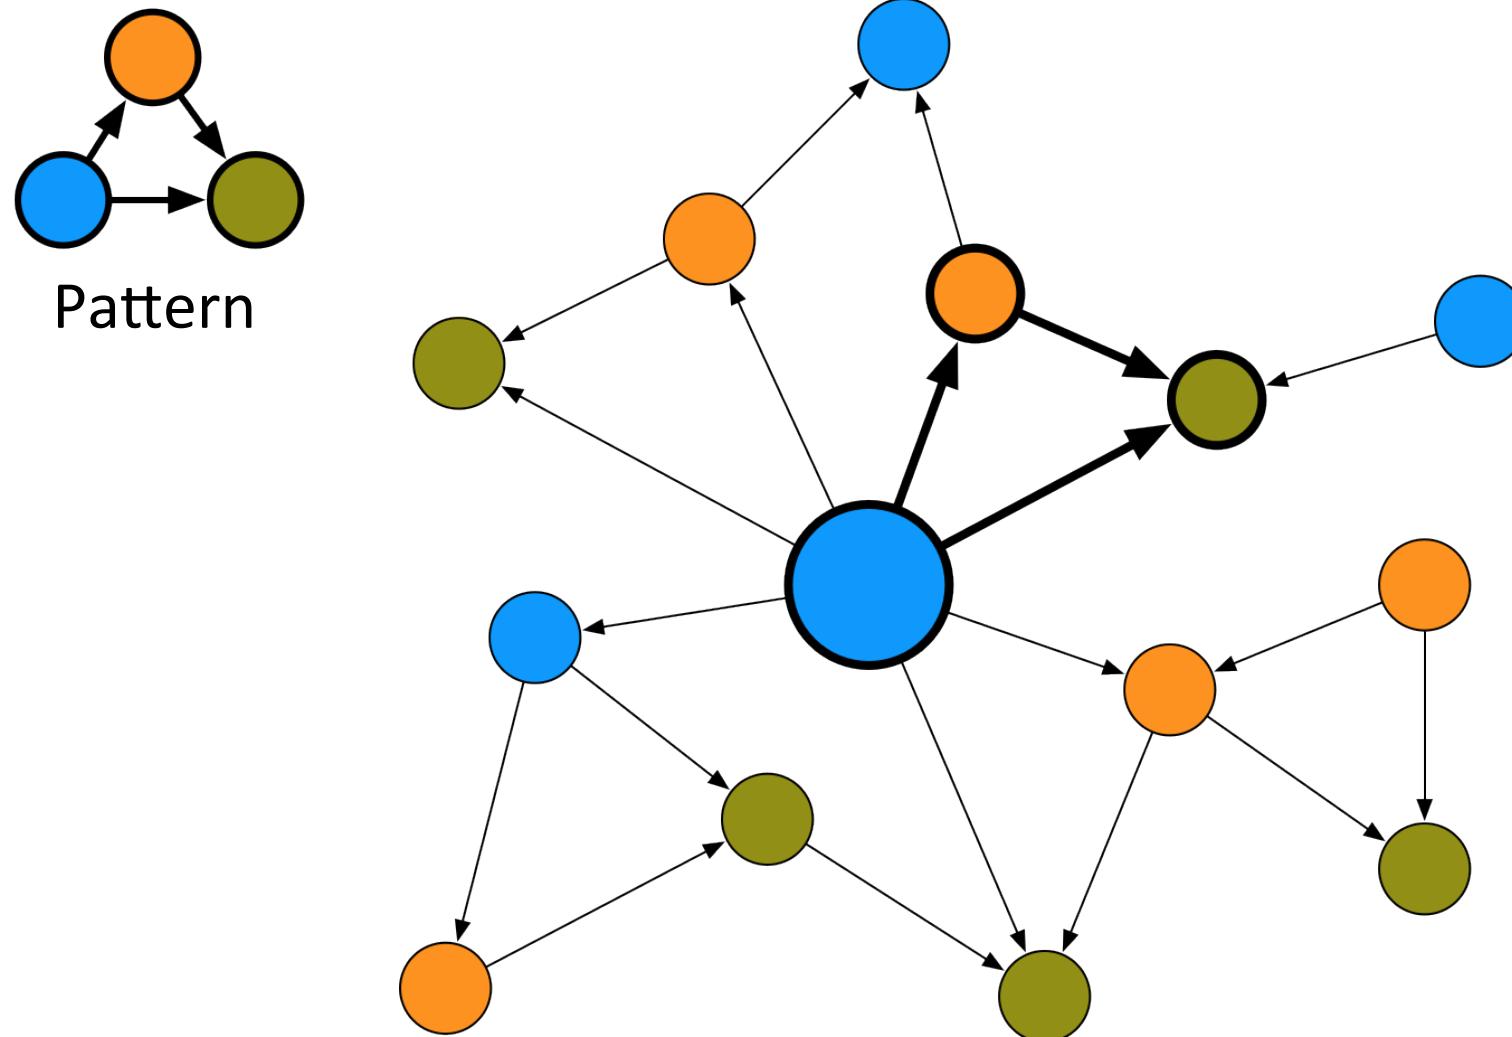
e.g.,



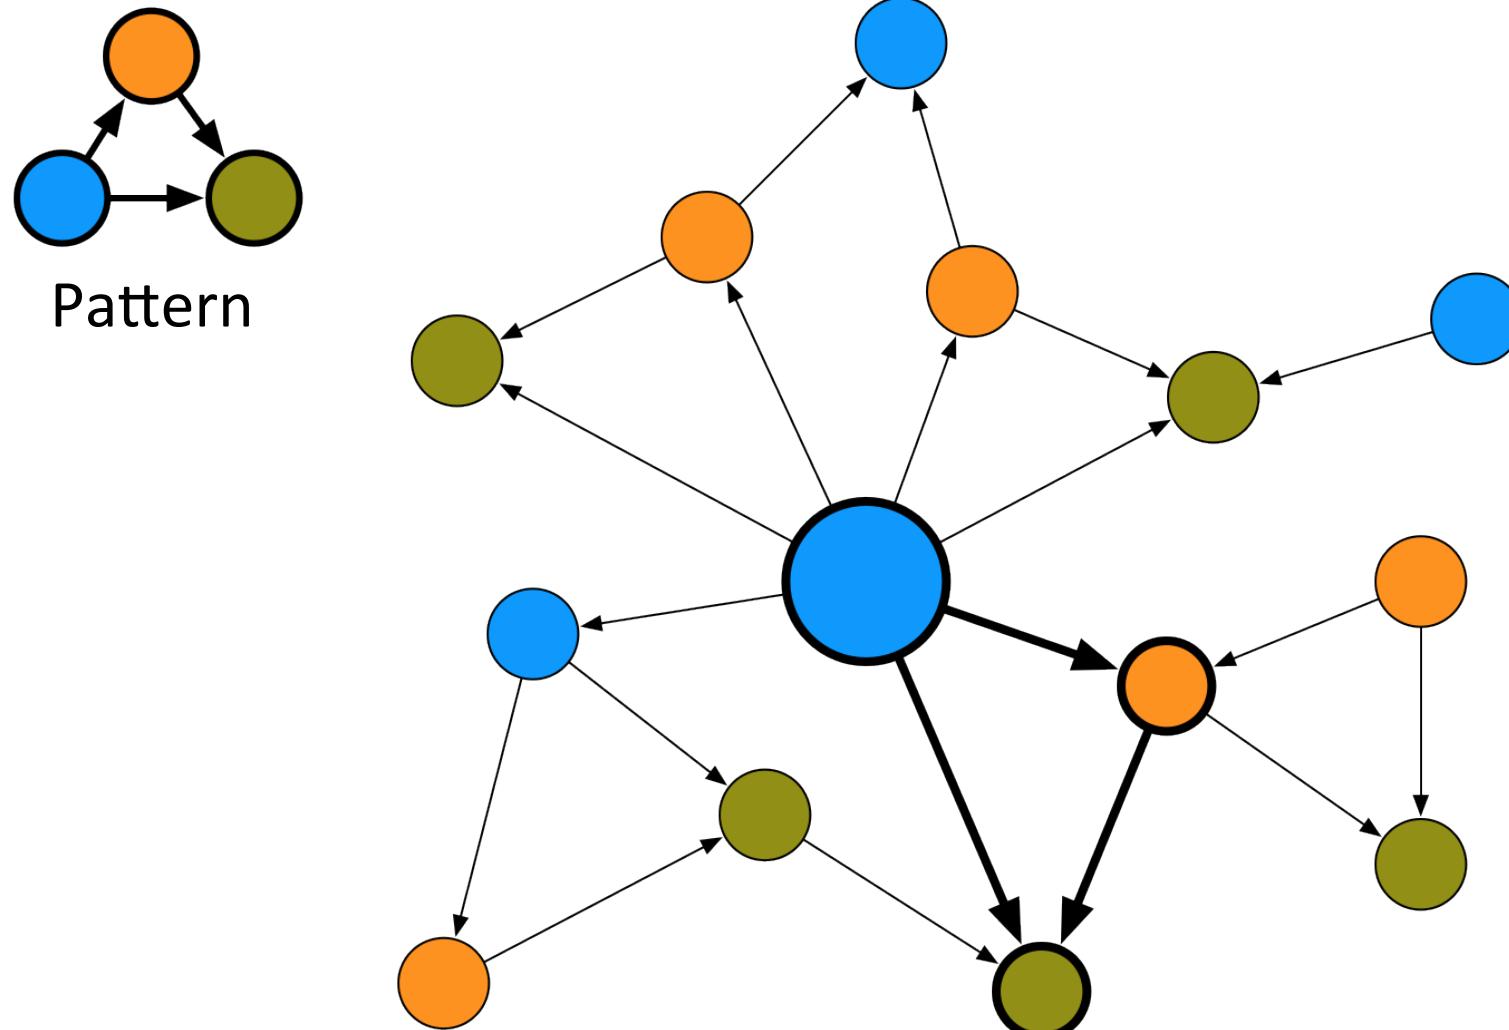
Pattern Matching – Start Node



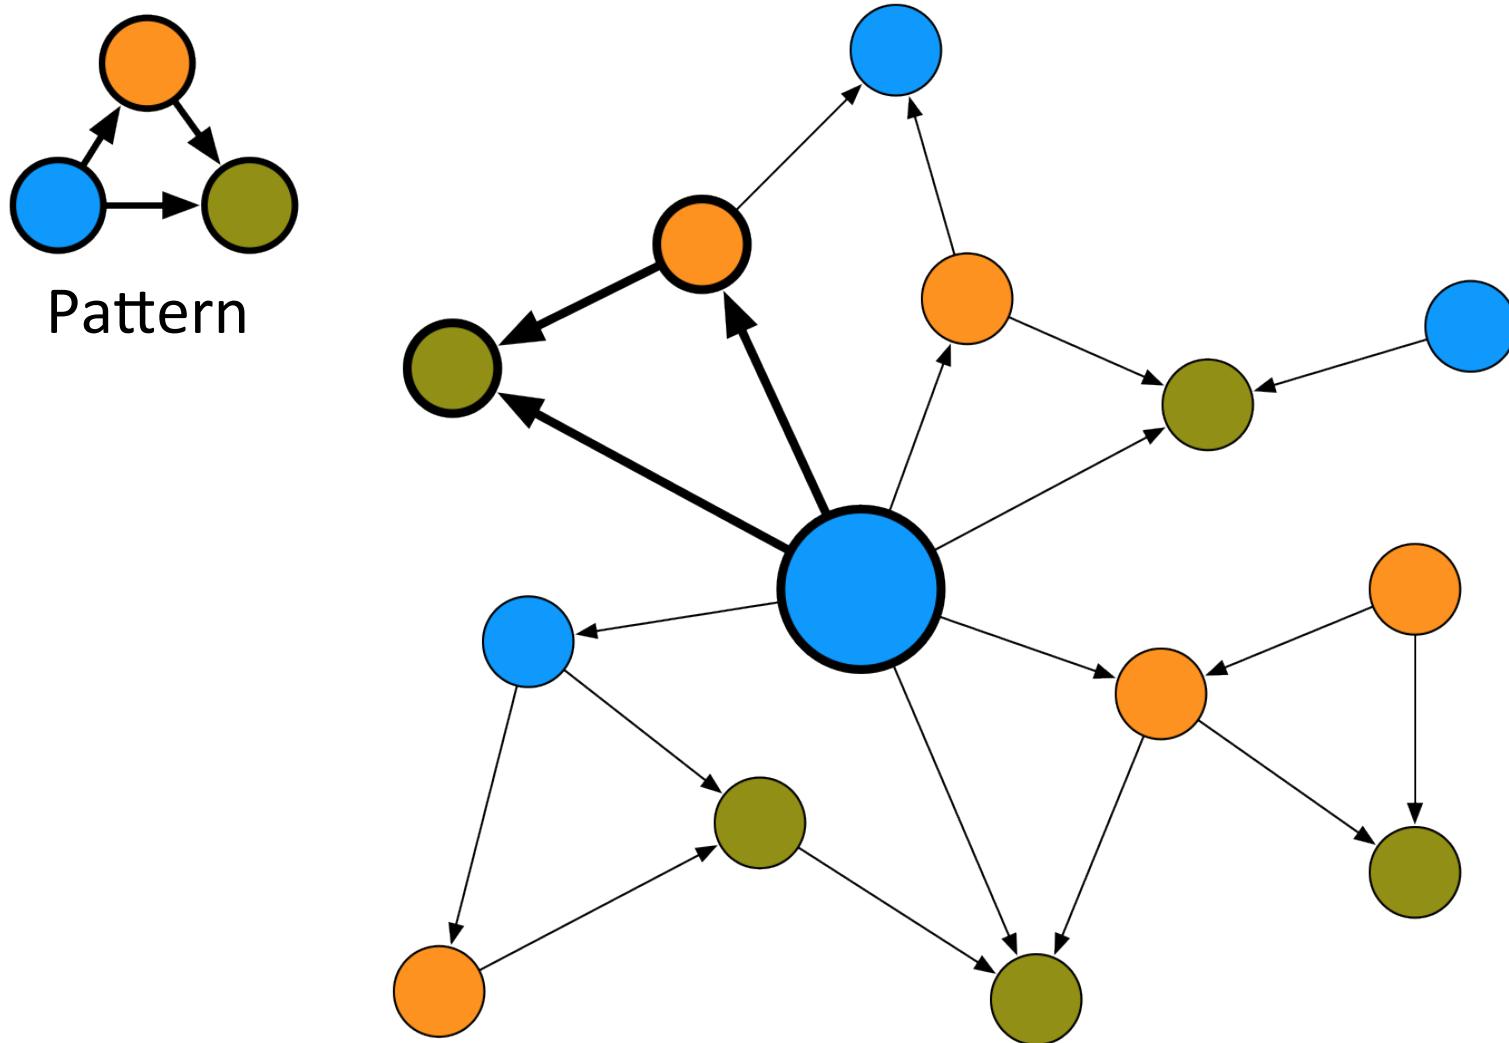
Pattern Matching - Match



Pattern Matching - Match



Pattern Matching - Match



ASCII Art Patterns

Pattern: Persons who live in San Francisco

● Person

● Location



ASCII Art Patterns

Pattern: Persons who live in San Francisco

● Person

● Location



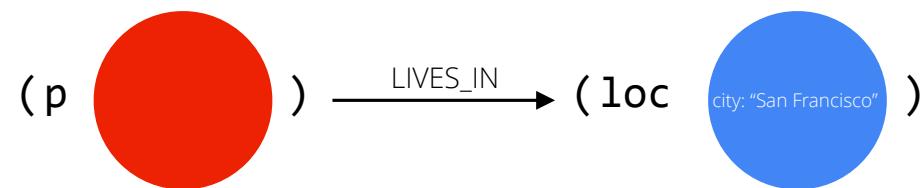
node ()

ASCII Art Patterns

Pattern: Persons who live in San Francisco

● Person

● Location



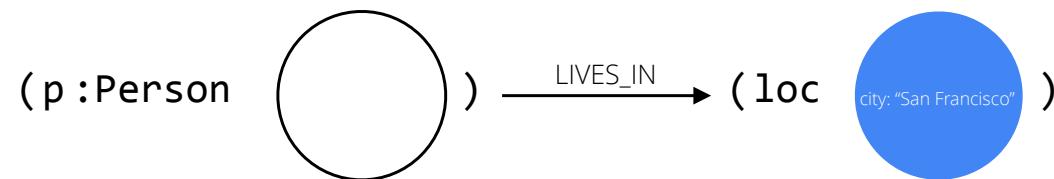
node variable (p)

ASCII Art Patterns

Pattern: Persons who live in San Francisco

● Person

● Location



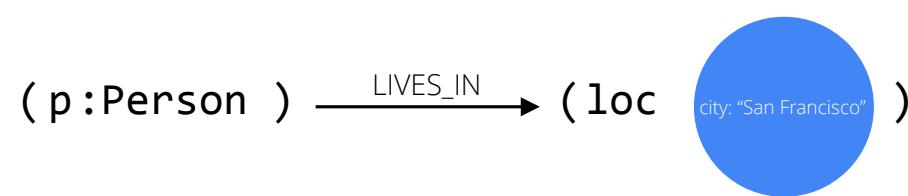
node label (:Person)

ASCII Art Patterns

Pattern: Persons who live in San Francisco

● Person

● Location



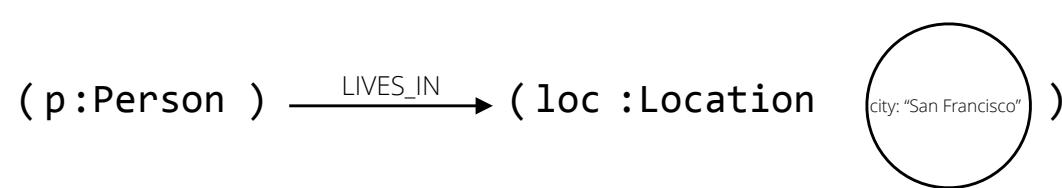
node label (:Person)

ASCII Art Patterns

Pattern: Persons who live in San Francisco

● Person

● Location



node label

(:Location)

ASCII Art Patterns

Pattern: Persons who live in San Francisco

- Person
- Location

(p :Person) —————→ (loc :Location {city: “San Francisco”})

property (city: ' 'San Francisco' ')

ASCII Art Patterns

Pattern: Persons who live in San Francisco

- Person
- Location

(p:Person) $\overset{\text{LIVES_IN}}{\dashrightarrow}$ (loc :Location {city: "San Francisco"})

edge \dashrightarrow

ASCII Art Patterns

Pattern: Persons who live in San Francisco

● Person

● Location

```
( p:Person ) - [:LIVES_IN] -> ( loc :Location {city: "San Francisco"} )
```

edge type - [:LIVES_IN] ->

ASCII Art Patterns

Pattern: Persons who live in San Francisco

● Person

● Location

```
MATCH (p:Person) -[:LIVES_IN] -> (loc:Location {city: "San Francisco"})  
RETURN p
```

```
MATCH (node1:Label1 {property:value})-->(node2:Label2)  
RETURN node1.property
```

```
MATCH (n1:Label1)-[rel:TYPE {property:value}]->(n2:Label2)  
RETURN rel.property, type(rel)
```

Structural Patterns

“ASCII art”

variable
:node label
property values {k:v}

() - - ()

Conditions on the relationship may be specified inside []

() <- - ()
() - -> ()

variable
:edge type
property values

Structural Patterns - Nodes

- ()
- (matrix)
- (:Movie)
- (matrix:Movie)
- (matrix {title: "The Matrix"})
- (:Movie {title: "The Matrix", released: 1999})
- (matrix:Movie {released: 1999})
- ...

Structural Patterns - Relationships

More “ASCII art” (combines with <, >, :)

variable length path

() – [*] – ()

constrained length path

() – [*min..max] – ()

Structural Patterns - Paths

- $(a)-[*]->(b)$ traverse any depth
- $(a)-[*\ depth]->(b)$ exactly \textit{depth} steps long
- $(a)-[*1..4]->(b)$ from one to four levels deep
- $(a)-[:KNOWS*3]->(b)$ relationships of type KNOWS at 3 levels distance
- $(a)-[:KNOWS*..5]->(b)$ relationships of type KNOWS at most 5 levels distance
- $(a)-[:KNOWS|:LIKES*2..]->(b)$ relationships of type KNOWS or LIKES from 2 levels distance

Structural Patterns

- friend-of-a-friend

(user)-[:KNOWS]-(friend)-[:KNOWS]-(foaf)

- shortest path:

path = shortestPath((user)-[:KNOWS*..5]-(other))

- collaborative filtering

(user)-[:PURCHASED]->(product)<-[:PURCHASED](otherUser)

NOTE we need to impose that otherUser <> user !

- tree navigation

(root)<-[:PARENT*]-(leaf:Category)-[:ITEM]->(data:Product)

Structural Patterns

- collaborative filtering

(user)-[:PURCHASED]->(product)<- [:PURCHASED] (otherUser)

otheruser <> user

:Product

:Customer

What if we are looking for a product bought by (at least) 3 users?

We may have multiple paths in the same MATCH, comma separated, sharing variables

Structural Patterns

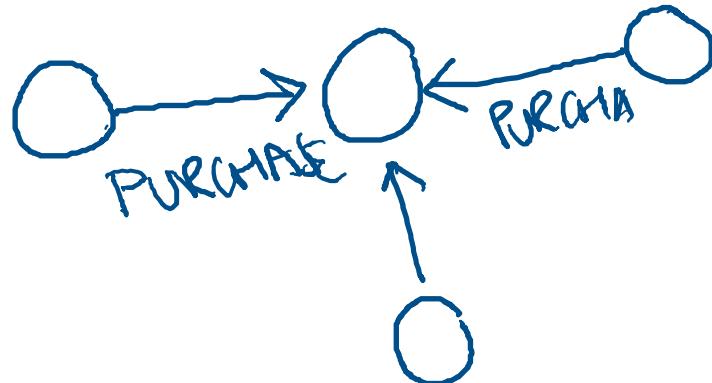
- collaborative filtering

(user)-[:PURCHASED]->(product)<- [:PURCHASED] (otherUser)
otheruser <> user

:Product

:Customer

What if we are looking for a product bought by (at least) 3 users?



We may have multiple paths in the same MATCH, comma separated, sharing variables

```
MATCH (user)-[ :PURCHASED ]->(product)<-[ :PURCHASED ](otherUser),  
(product)<- [ :PURCHASED ]-(thirdUser)
```

otheruser <> user, thirduser <> user, thirduser <> otheruser

Structural Patterns

- collaborative filtering

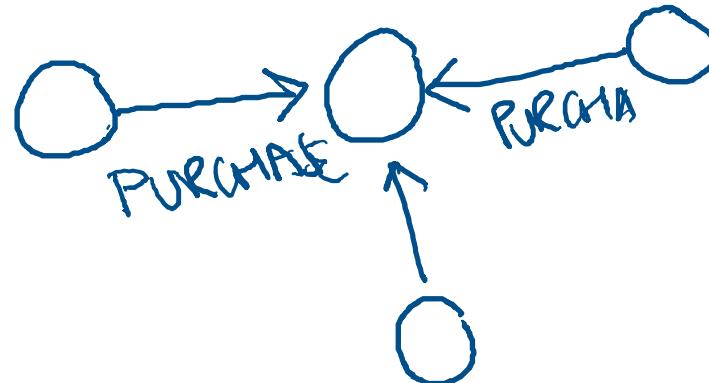
(user)-[:PURCHASED]->(product)<- [:PURCHASED] (otherUser)

otheruser <> user

:Product

:Customer

What if we are looking for a product bought by (at least) 3 users?



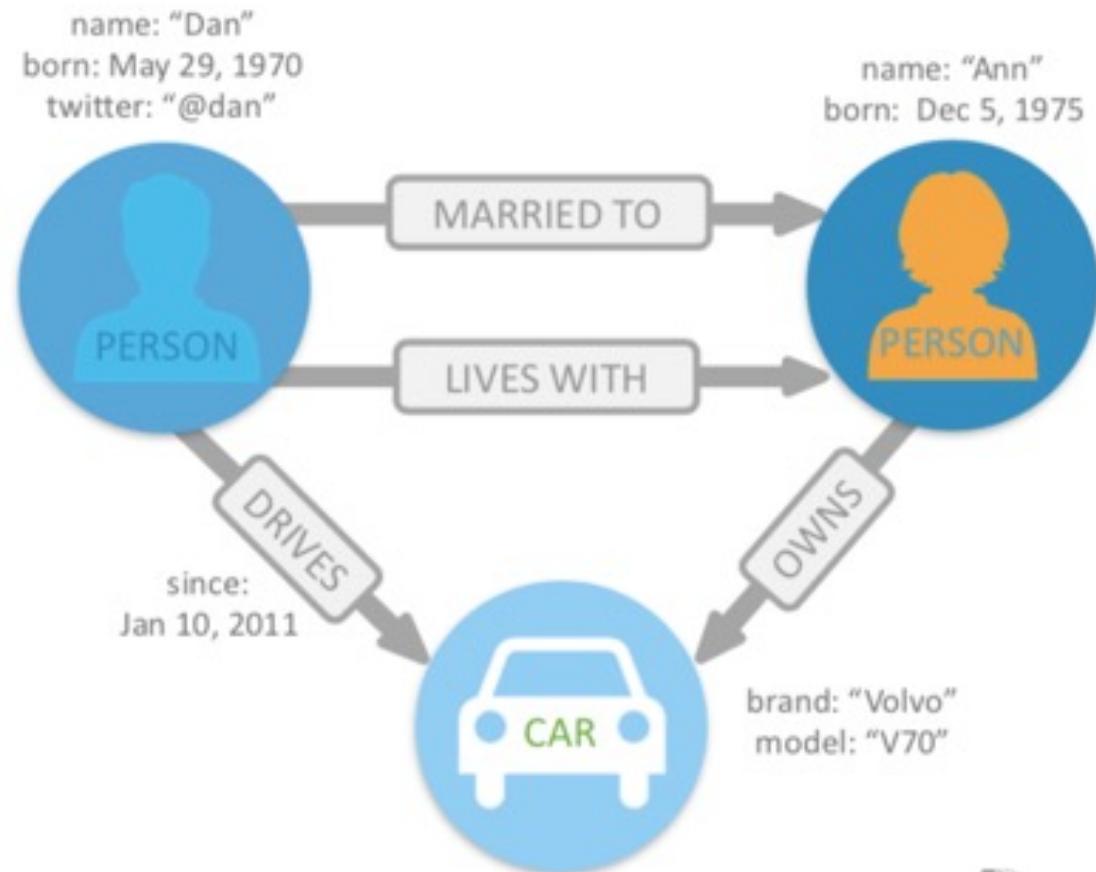
We may have multiple paths in the same MATCH, comma separated, sharing variables

```
MATCH (user)-[ :PURCHASED ]->(product),  
(product) <- [ :PURCHASED ]-(otherUser),  
(product)<- [ :PURCHASED ]-(thirdUser)
```

otheruser <> user, thirduser <> user, thirduser <> otheruser

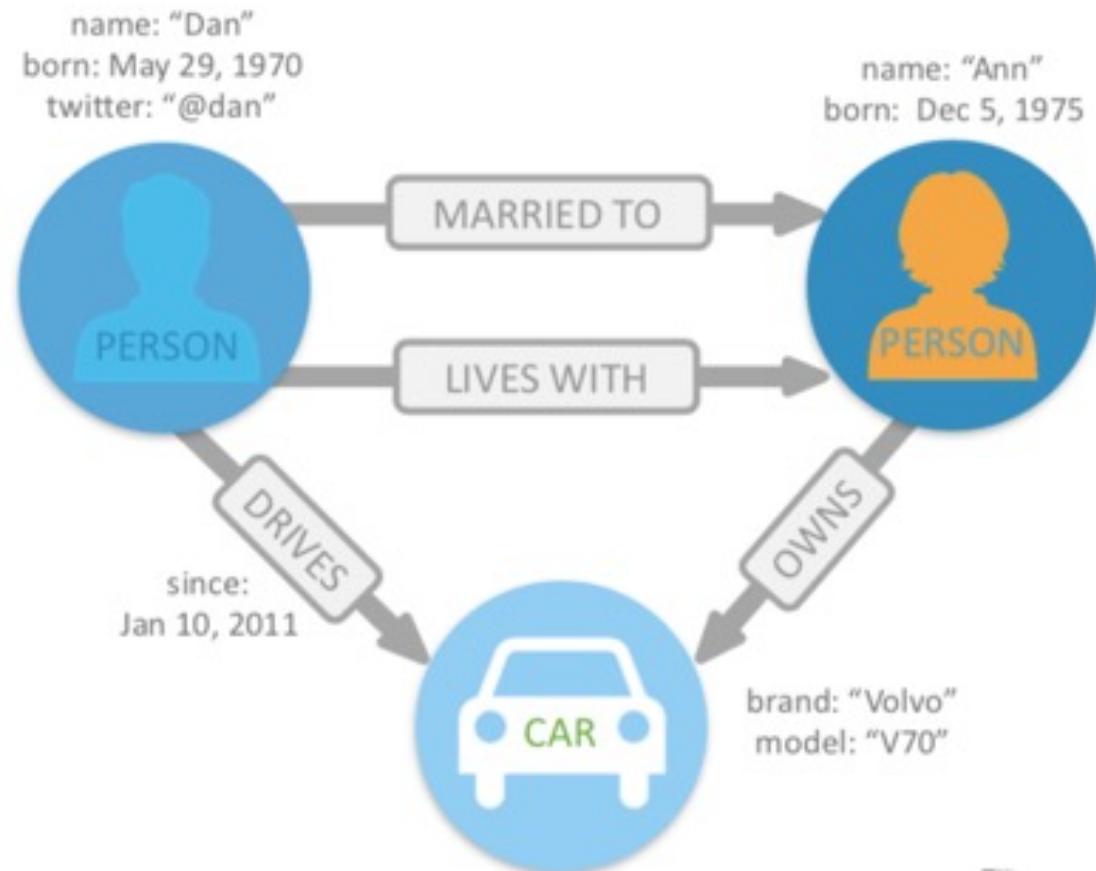
Queries = Traversals

get the spouse of the owner of a Volvo



Queries = Traversals

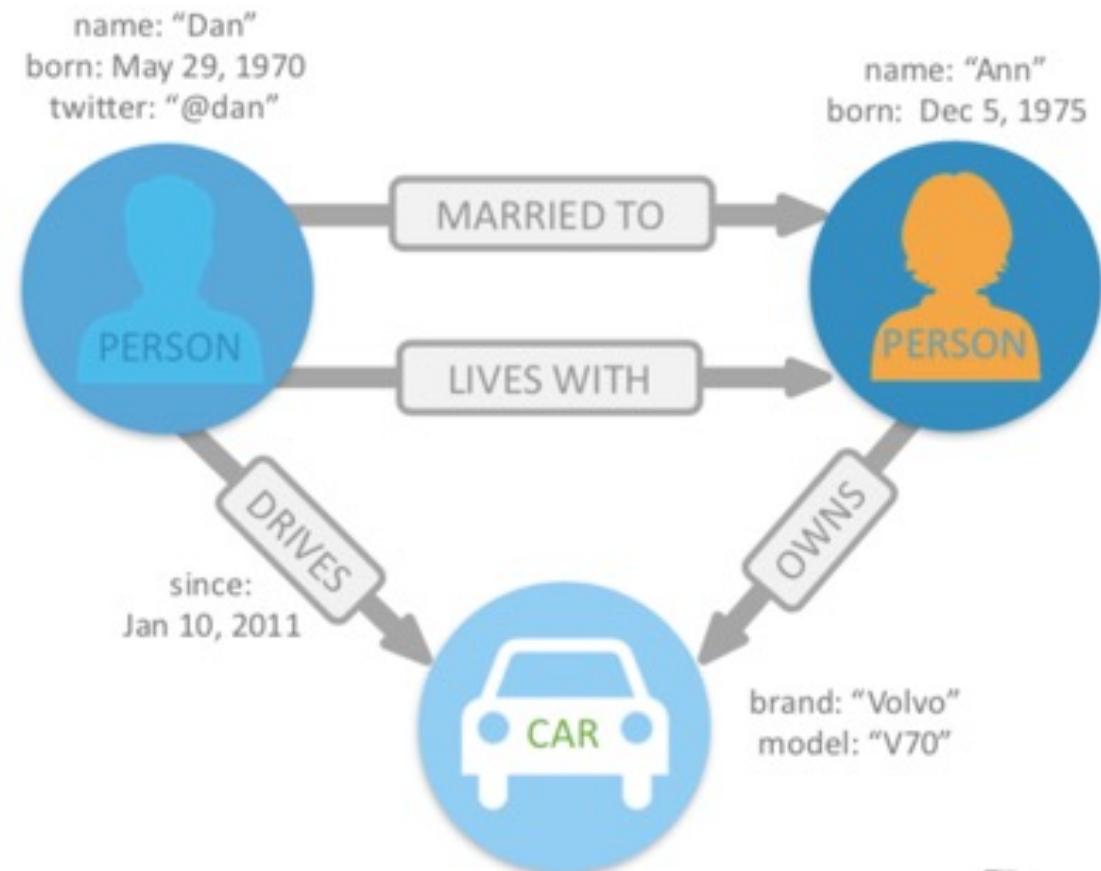
get the spouse of the owner of a Volvo



```
MATCH (:Car {brand:'Volvo'})<-[:OWNS]-(:Person)-[:MARRIED_TO]-(p:Person)  
RETURN p
```

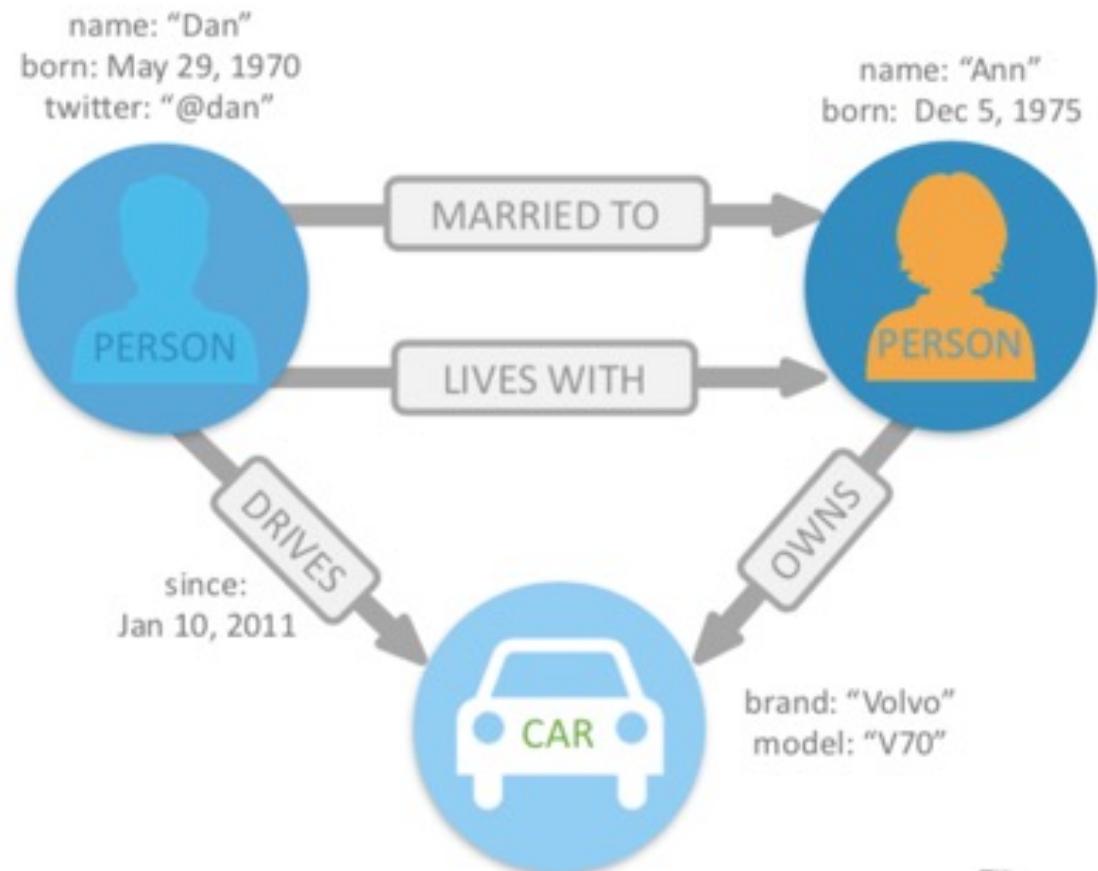
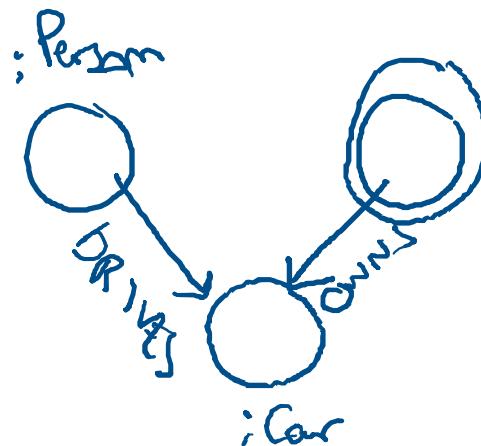
Queries = Traversals

get the owner of a car a person drives



Queries = Traversals

get the owner of a car a person drives

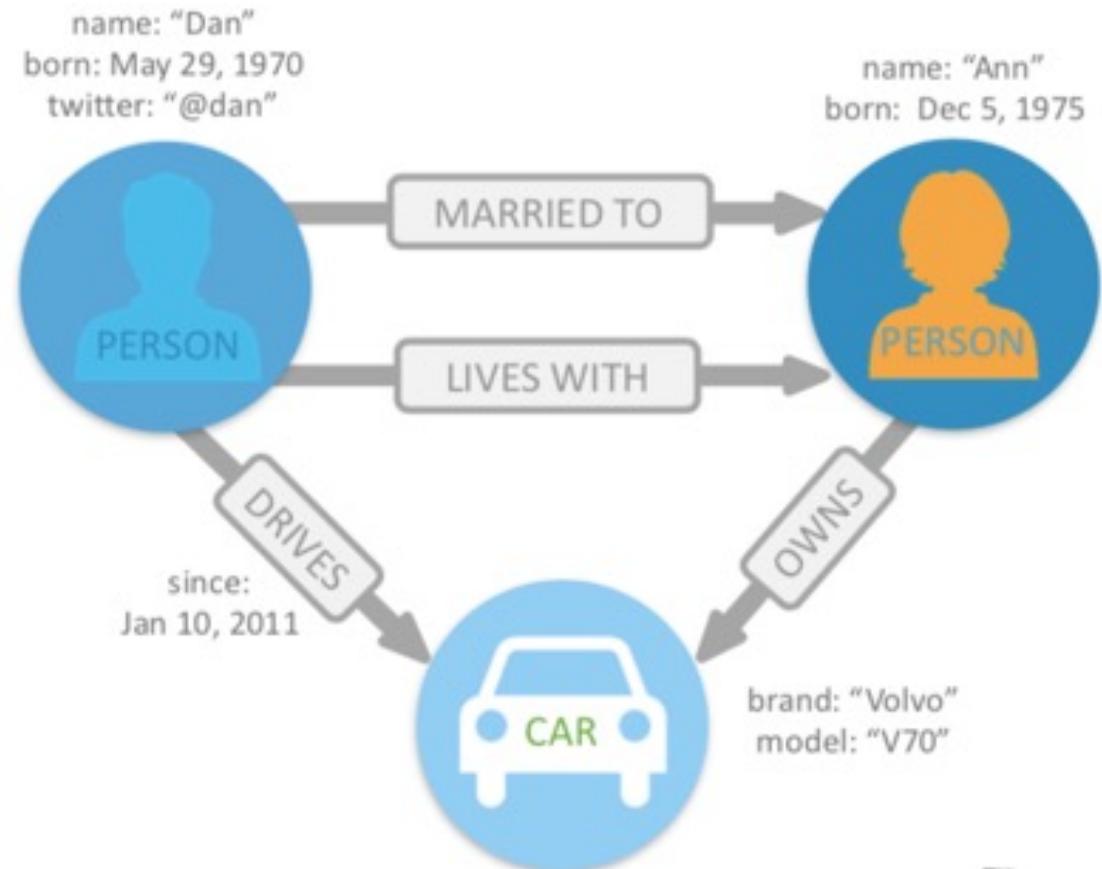


```
MATCH (:Person)-[:DRIVES]->(:Car)<-[:OWNS]-(p:Person)  
RETURN p
```

```
MATCH (c:Car)<-[:DRIVES]-(:Person), (p:Person)-[:OWNS]->(c)  
MATCH (p:Person)-[:OWNS]->(:Car)<-[:DRIVES]-(:Person)
```

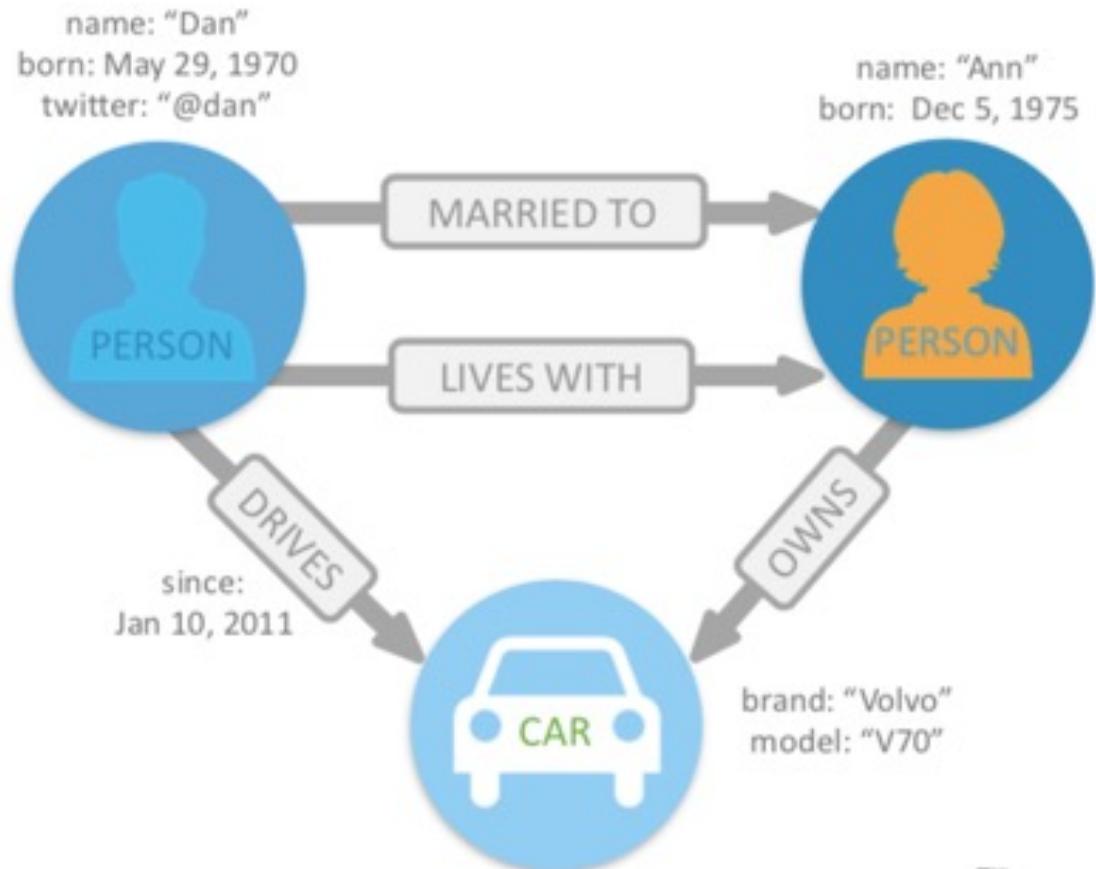
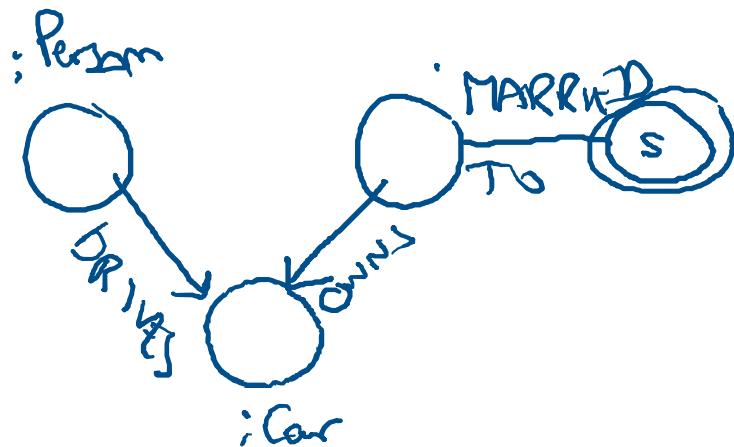
Queries = Traversals

get the spouse of the owner of a car a person drives



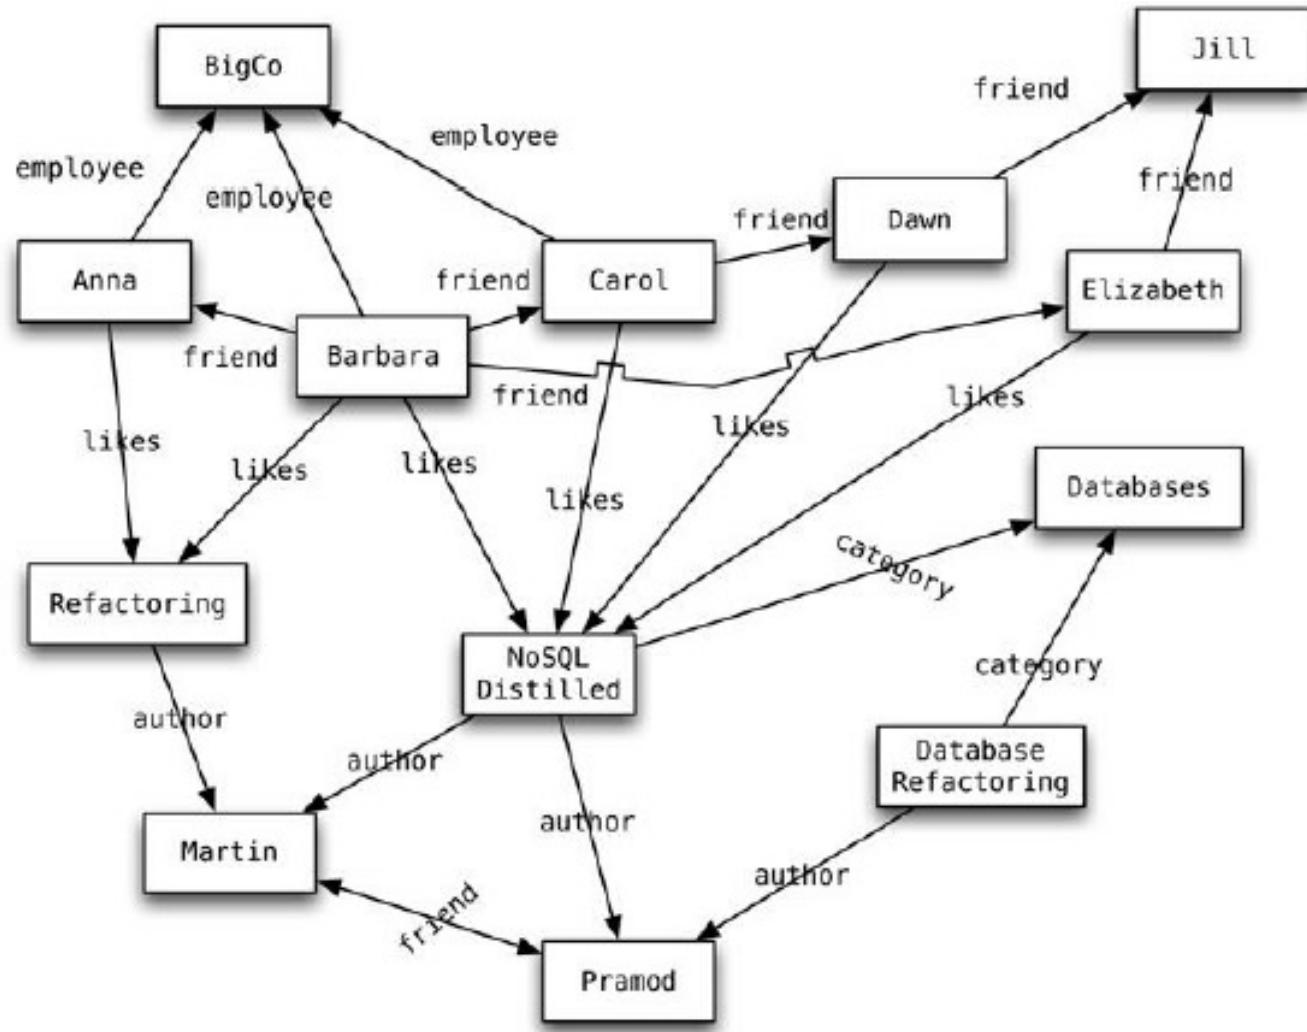
Queries = Traversals

get the spouse of the owner of a car a person drives

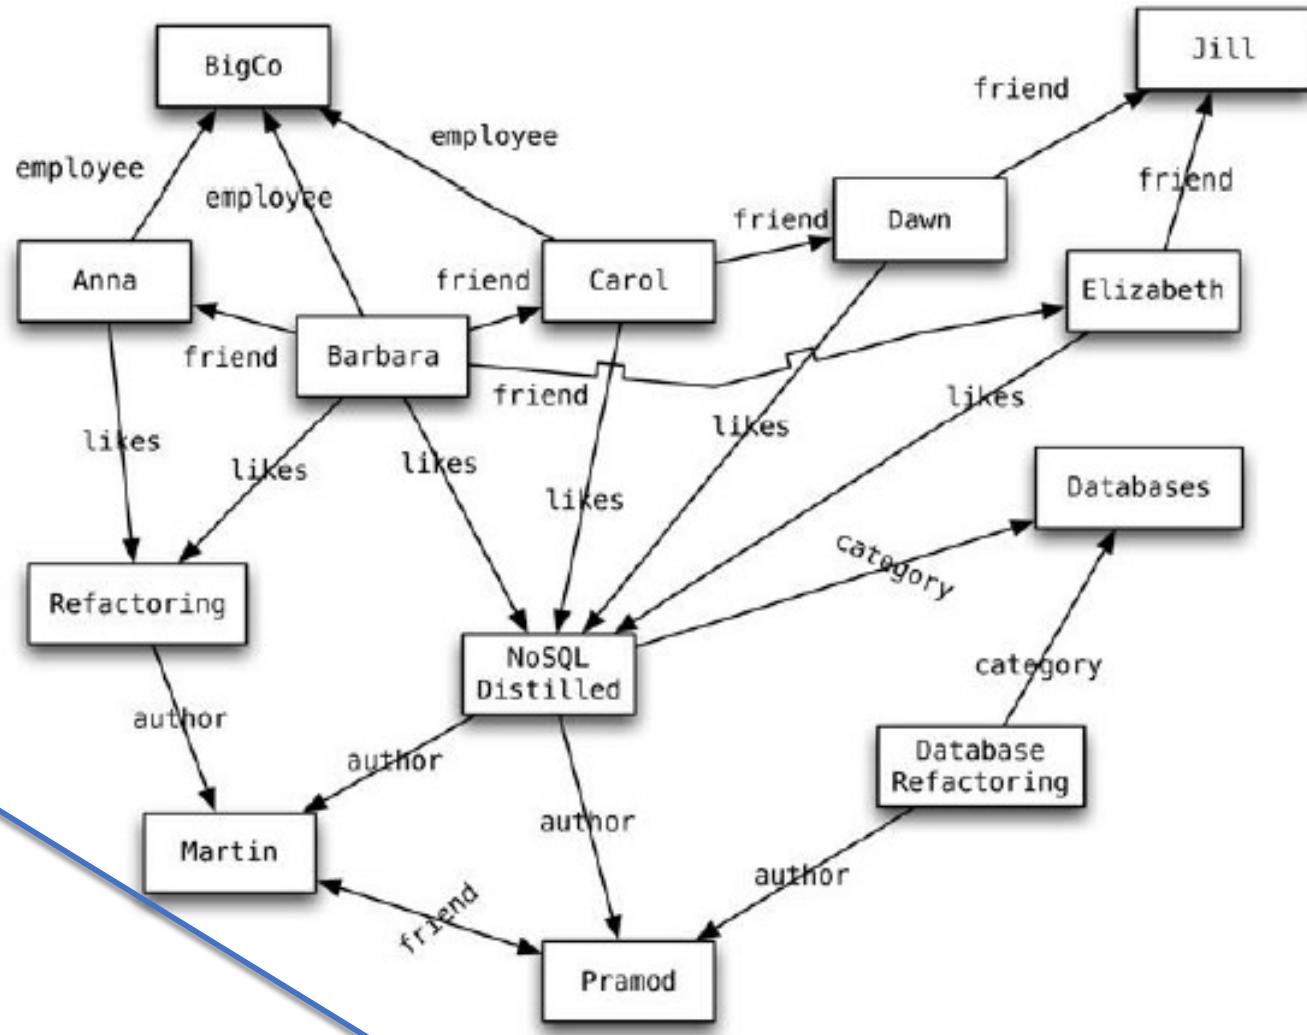


```
MATCH (:Person)-[:DRIVES]->(:Car)<-[:OWNS]-  
(:Person)-[:MARRIED_TO]-(s:Person)  
RETURN s
```

“find the friends of Barbara that like a node of database category.”

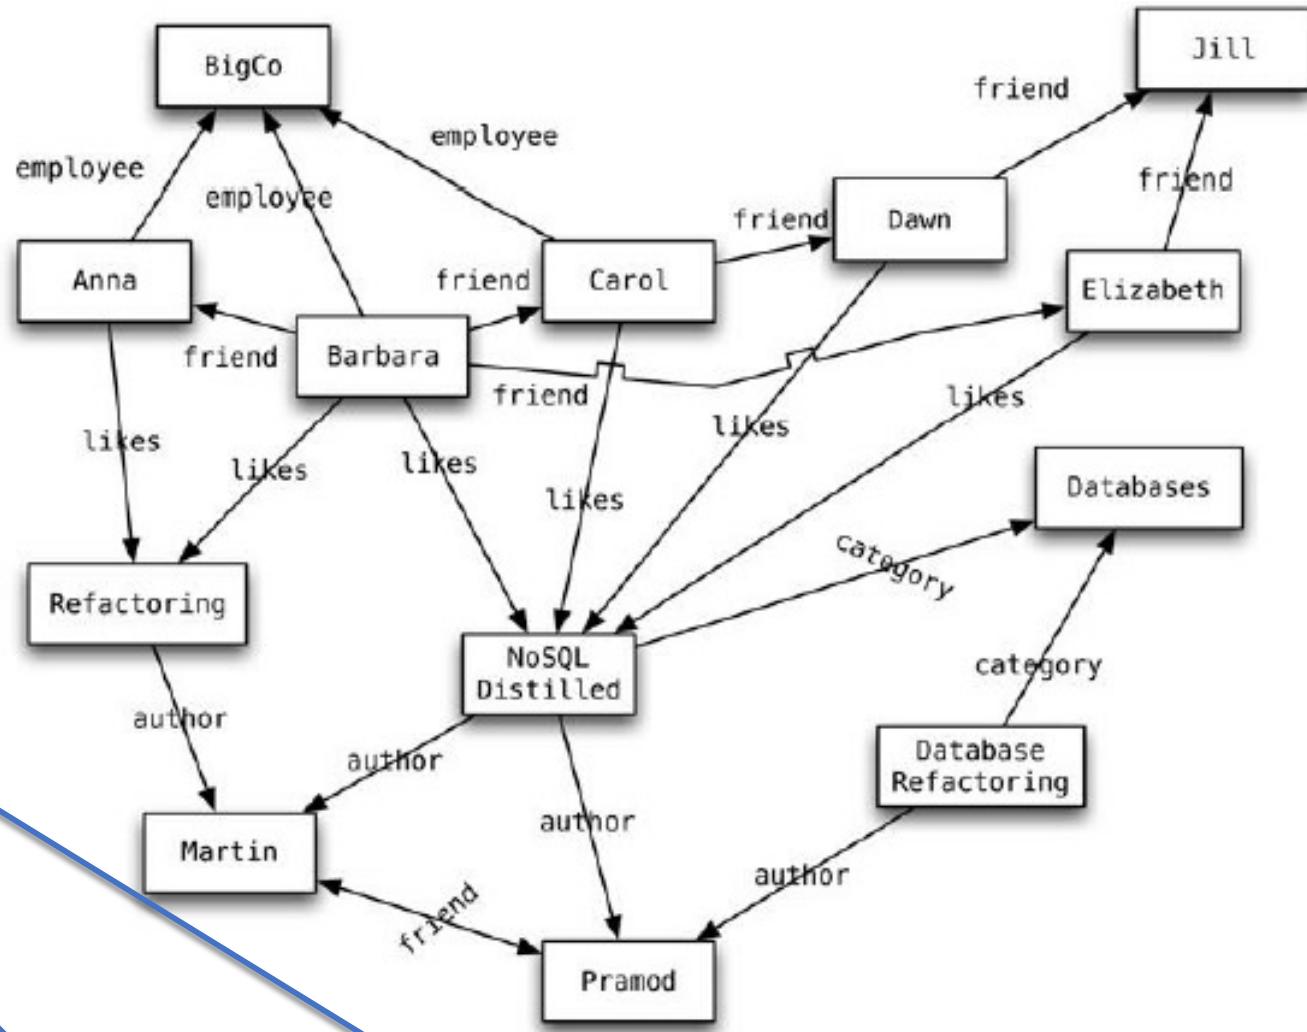


“find the friends of Barbara that like a node of database category.”



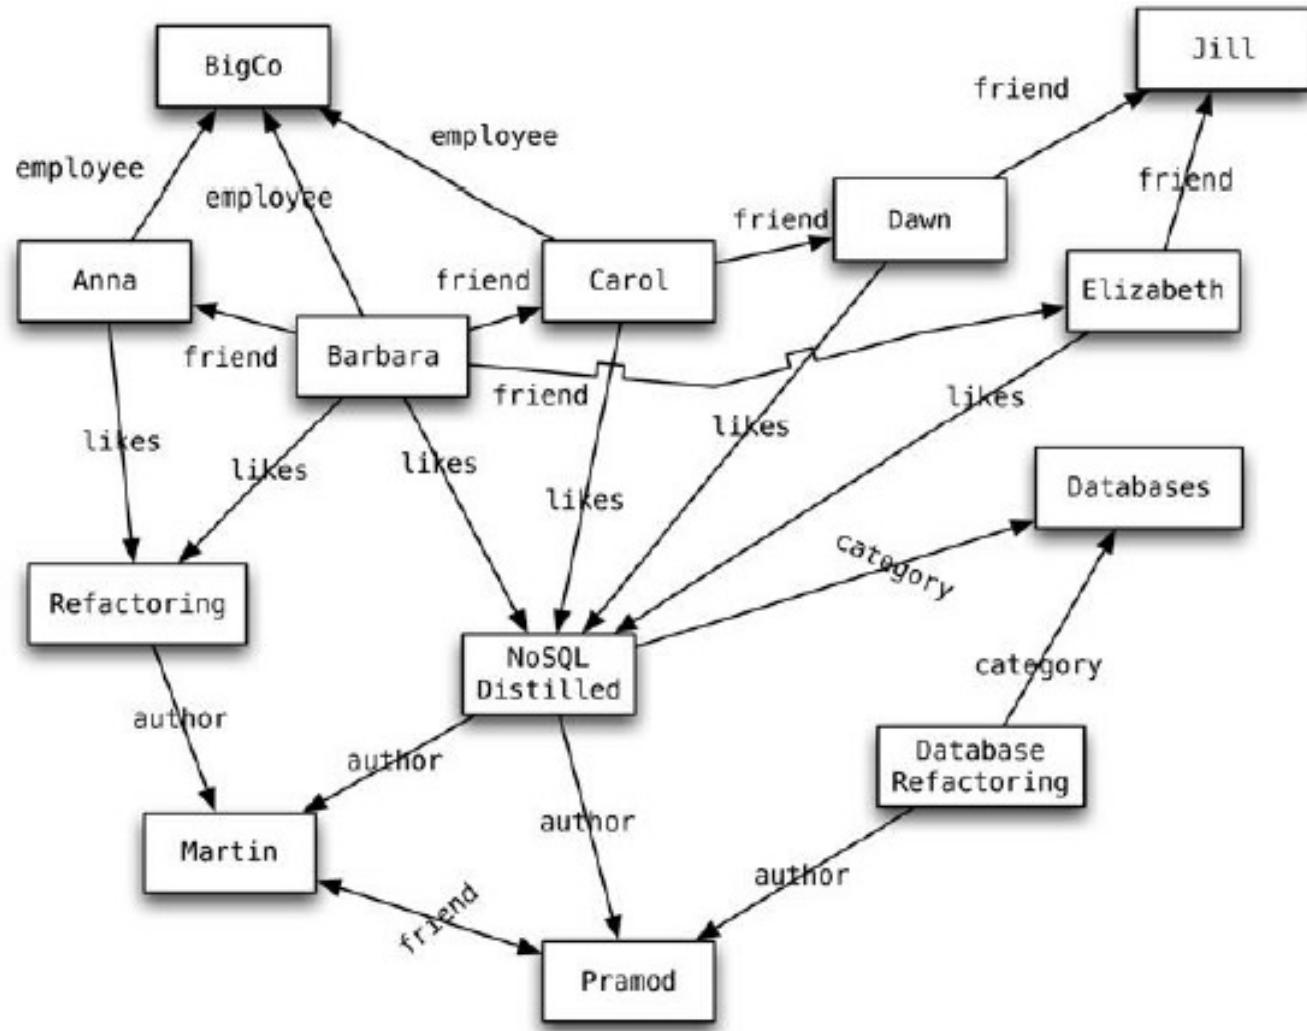
```
MATCH (:Barbara)-[:friend]-(node)-[:likes]->()-[:category]->(:Databases)  
RETURN node ;
```

“find the friends of Barbara that like a node of database category.”



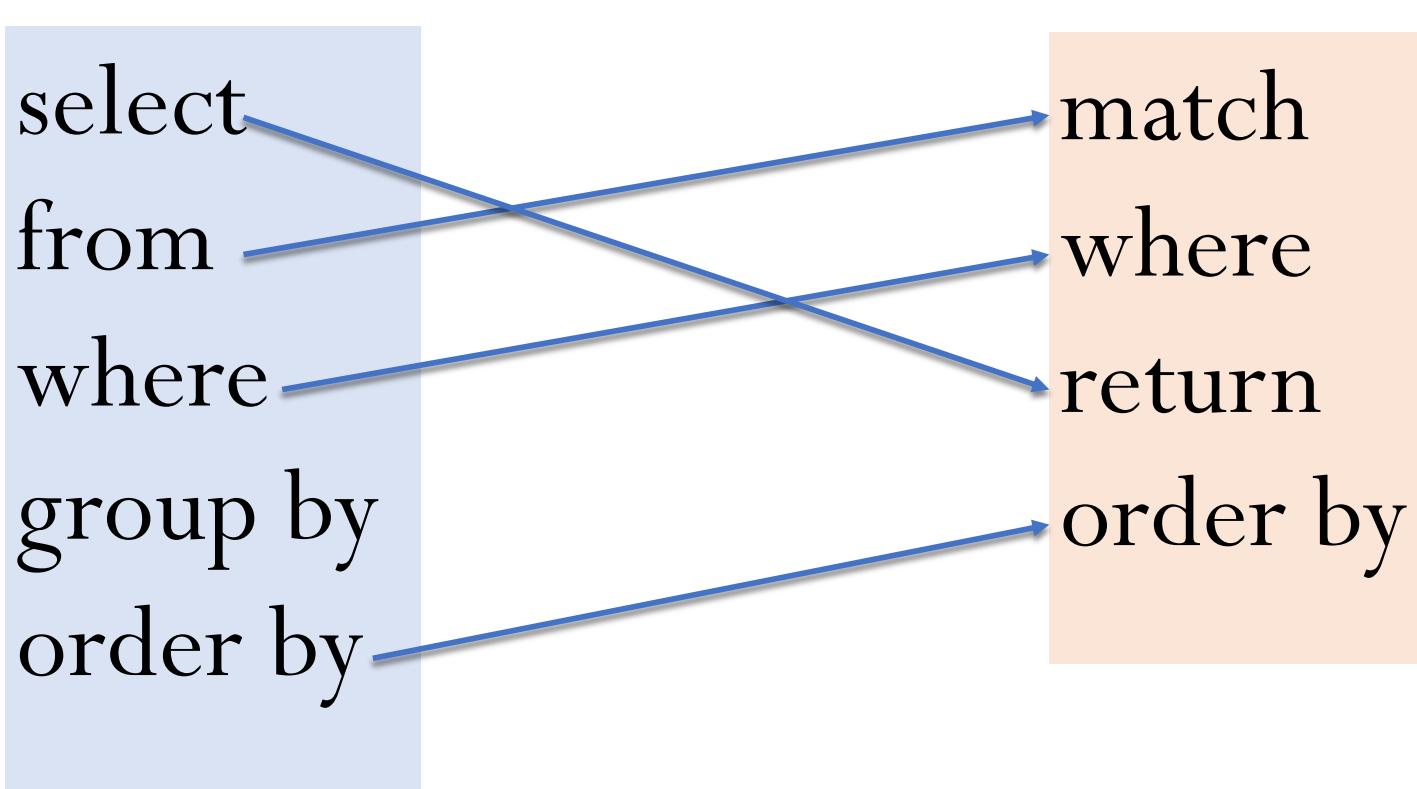
MATCH (:Barbara)-[:friend]-(node)-[:likes]->()-[:category]->(:Databases)
RETURN node ;

MATCH (:Databases)<-[:category]-()<-[:likes]-(node)-[:friend]-(:Barbara)
RETURN node ;



Cypher queries – Structure

Familiar for SQL Users



Where

```
MATCH (node:Label {property:value})  
MATCH (node:Label)  
WHERE property=value
```

- usual big 6 (=, <, >, <=,>=, <>)
- Usual Boolean (AND, OR, NOT)
- Property existence checking: exists
 - MATCH (n)
 - WHERE exists(n.marriedTo)
 - RETURN n.name, n.marriedTo
- Collection membership: IN
- ...

Returning nodes and values

```
MATCH (:Barbara)-[:friend]-(node)-  
[:likes]->()- [:category]-> (:Databases)  
RETURN node
```

```
MATCH (n)  
WHERE exists(n.marriedTo)  
RETURN n.name, n.marriedTo
```

Returning Paths

- `MATCH p=(:Barbara)-[:friend]-(node)-[:likes]->()-[:category]->(:Databases)`
`RETURN p`
all of the nodes and relationships for each path,
including all of their properties
- Just nodes in the path:
`RETURN nodes(p)`
- Just relationships in the path:
`RETURN rels(p)`

Cypher: Outputs (from yesterday)

A node

```
MATCH (p:Person {name:"Tom"}) RETURN p
```

A value

```
MATCH (p:Person {name:"Tom"}) RETURN p.age
```

A list of values

```
MATCH (p:Person) RETURN p.name LIMIT 5
```

An array

```
MATCH p=shortestPath((a)-[*]->(b)) WHERE a.name="Axel" AND b.name="Tom" RETURN p
```

A list of arrays

```
MATCH p=((a)-[*]->(b)) WHERE a.name="Axel" AND b.name="Frank" RETURN p
```

Return – aggregate functions

- MATCH (n:Person)

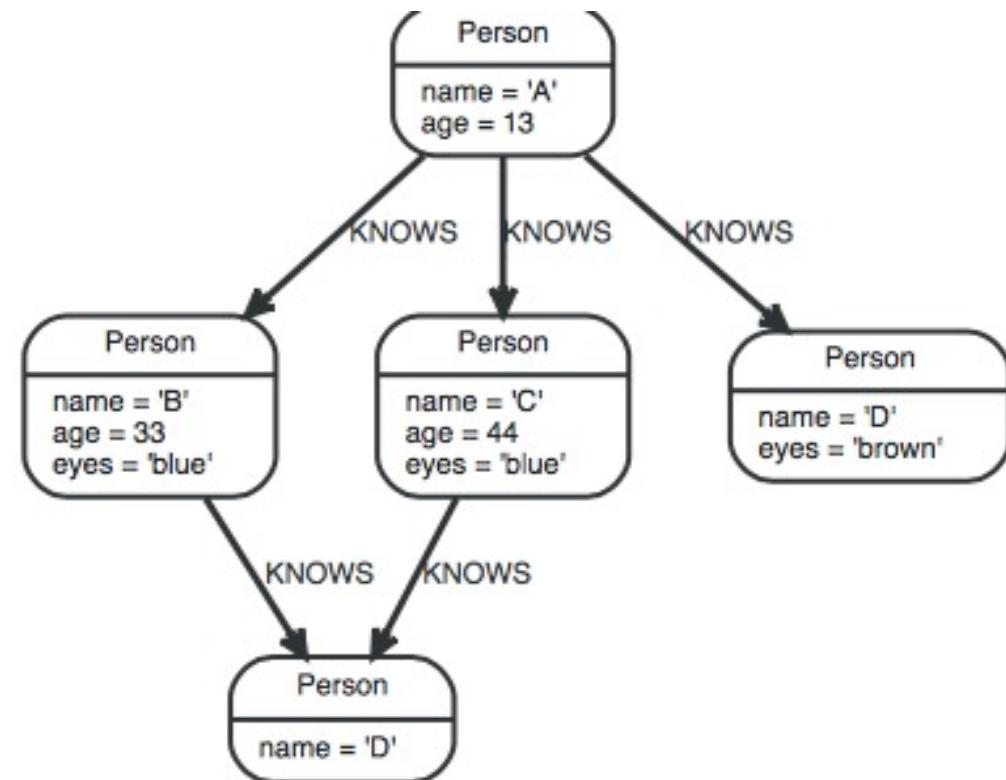
RETURN avg(n.age)

30

- MATCH (n:Person)

RETURN collect(n.age)

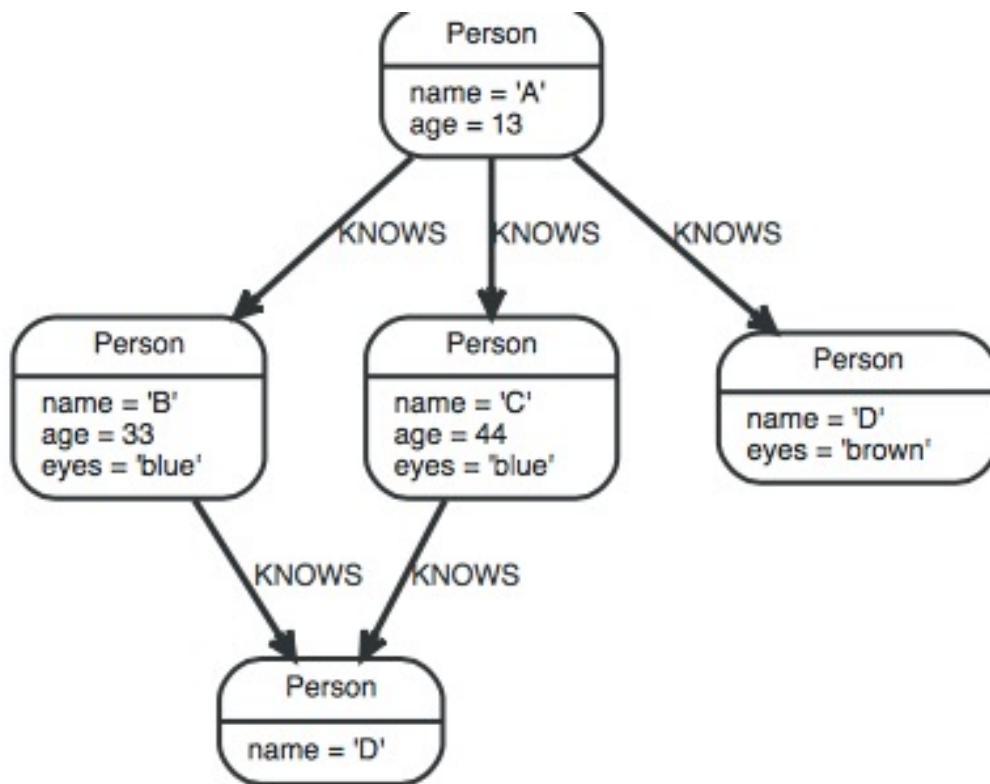
(13, 33, 44)



avg(), collect(), count(), max(), min(), sum(), ...

GROUP BY

- In Cypher, GROUP BY is done implicitly by all of the aggregate functions
- In a RETURN clause (as in a WITH, see next slides), any columns not part of an aggregate will be the GROUP BY key



- MATCH (n:Person)
RETURN COUNT(n), n.eyes
2,'blue' 1,'brown'
- MATCH (n:Person)
RETURN COUNT(n), n.eyes,
AVG(age)
2,'blue',38.5 1,'brown'

GROUP BY – how many friends?

```
MATCH (p:Person)-[:Knows]->(friend)  
WHERE p.age = 20  
RETURN p.name, count(friend)
```

What if we only want people with at least 10 friends?
(a sort of SQL HAVING)
We need a WITH clause ...

WITH – to divide a query into multiple parts

```
MATCH (p:Person)-[:Knows]->(friend)
```

```
WHERE p.age = 20
```

```
WITH p, count(friend) as friends
```

```
WHERE friends > 10
```

```
RETURN p.name, friends
```

Ex: Reformulate the 3 users queries

Cumulative

```
reduce(accumulator = initial,  
variable IN list | expression)
```

1. iterate through each element e in the given list
2. run the expression on e and
3. store the new partial result in the accumulator

```
MATCH p =(a:{name:'Alice'})-[ :KNOWS*..3 ]->  
      (c: {name:'Bob'})  
RETURN reduce(totalAge = 0, n IN nodes(p) |  
totalAge + n.age) AS reduction
```

A few more Cypher queries

Return the pairs of nodes that are related by two different relationships

```
MATCH(n1)-[r1]-(n2)-[r2]-(n1) WHERE type(r1)<>type(r2)  
RETURN n1, n2
```

Return the three nodes with the highest number of relationships.

```
MATCH (a)-[r]->()  
RETURN a, COUNT(r) ORDER BY COUNT(r) DESC LIMIT 3
```

Return the length of the shortest path between two nodes of your choice.

```
MATCH p = shortestPath((p1:Person{name:'Alice'})-[*]-  
                      (p2:Person{name:'Bob'}))  
RETURN length(p)
```

Cypher: Clauses

- **MATCH**: The graph **pattern** to match
- **WHERE**: **Filtering** criteria
- **RETURN**: What to return
- **CREATE**: Creates nodes and relationships
- **DELETE**: Remove nodes, relationships, properties
- **SET**: Set values to **properties**
- **WITH**: Divides a query into multiple parts

Cypher vs SQL

```
(SELECT T.directReportees AS directReportees, sum(T.count) AS count
FROM (
    SELECT manager.pid AS directReportees, 0 AS count
        FROM person_reportee manager
       WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
UNION
    SELECT manager.pid AS directReportees, count(manager.directly_manages) AS count
        FROM person_reportee manager
       WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
      GROUP BY directReportees
UNION
    SELECT manager.pid AS directReportees, count(reportee.directly_manages) AS count
        FROM person_reportee manager
       JOIN person_reportee reportee
          ON manager.directly_manages = reportee.pid
       WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
      GROUP BY directReportees
UNION
    SELECT manager.pid AS directReportees, count(L2Reportees.directly_manages) AS count
        FROM person_reportee manager
       JOIN person_reportee L1Reportees
          ON manager.directly_manages = L1Reportees.pid
       JOIN person_reportee L2Reportees
          ON L1Reportees.directly_manages = L2Reportees.pid
       WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
      GROUP BY directReportees
) AS T
     GROUP BY directReportees)
UNION
(SELECT T.directReportees AS directReportees, sum(T.count) AS count
FROM (
    SELECT manager.directly_manages AS directReportees, 0 AS count
        FROM person_reportee manager
       WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
UNION
    SELECT reportee.pid AS directReportees, count(reportee.directly_manages) AS count
        FROM person_reportee manager
       JOIN person_reportee reportee
          ON manager.directly_manages = reportee.pid
       WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
      GROUP BY directReportees
UNION
```

```
SELECT depth1Reportees.pid AS directReportees,
       count(depth2Reportees.directly_manages) AS count
  FROM person_reportee_manager
    JOIN person_reportee L1Reportees
      ON manager.directly_manages = L1Reportees.pid
    JOIN person_reportee L2Reportees
      ON L1Reportees.directly_manages = L2Reportees.pid
     WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
   GROUP BY directReportees
 ) AS T
 GROUP BY directReportees)
UNION
 (SELECT T.directReportees AS directReportees, sum(T.count) AS count
  FROM(
    SELECT reportee.directly_manages AS directReportees, 0 AS count
  FROM person_reportee_manager
    JOIN person_reportee reportee
      ON manager.directly_manages = reportee.pid
     WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
   GROUP BY directReportees
UNION
    SELECT L2Reportees.pid AS directReportees, count(L2Reportees.directly_manages) AS count
  FROM person_reportee_manager
    JOIN person_reportee L1Reportees
      ON manager.directly_manages = L1Reportees.pid
    JOIN person_reportee L2Reportees
      ON L1Reportees.directly_manages = L2Reportees.pid
     WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
   GROUP BY directReportees
 ) AS T
 GROUP BY directReportees)
UNION
 (SELECT L2Reportees.directly_manages AS directReportees, 0 AS count
  FROM person_reportee_manager
    JOIN person_reportee L1Reportees
      ON manager.directly_manages = L1Reportees.pid
    JOIN person_reportee L2Reportees
      ON L1Reportees.directly_manages = L2Reportees.pid
     WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
  )
```

Cypher vs SQL

```
MATCH (boss) - [:MANAGES*0..3] -> (sub),  
      (sub) - [:MANAGES*1..3] -> (report)  
WHERE boss.name = "John Doe"  
RETURN sub.name AS Subordinate,  
count(report) AS Total;
```

Cypher vs SQL

```
(SELECT T.directReportees AS directReportees, sum(T.count) AS count
FROM (
  SELECT manager.pid AS directReportees, 0 AS count
  FROM person_reportee manager
  WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
UNION
  SELECT manager.pid AS directReportees, count(manager.directly_manages) AS count
  FROM person_reportee manager
  WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
GROUP BY directReportees
UNION
  SELECT manager.pid AS directReportees, count(reportee.directly_manages) AS count
  FROM person_reportee manager
  JOIN person_reportee reportee
  ON manager.directly_manages = reportee.pid
  WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
GROUP BY directReportees
UNION
  SELECT manager.pid AS directReportees, count(L2Reportees.directly_manages) AS count
  FROM person_reportee manager
  JOIN person_reportee L1Reportees
  ON manager.directly_manages = L1Reportees.pid
  WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
  GROUP BY directReportees
) AS T
GROUP BY directReportees)
UNION
(SELECT depth1Reportees.pid AS directReportees,
count(depth2Reportees.directly_manages) AS count
FROM person_reportee manager
JOIN person_reportee L1Reportees
ON manager.directly_manages = L1Reportees.pid
JOIN person_reportee L2Reportees
ON L1Reportees.directly_manages = L2Reportees.pid
WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
GROUP BY directReportees
) AS T
GROUP BY directReportees)
UNION
(SELECT T.directReportees AS directReportees, sum(T.count) AS count
FROM(
  SELECT reportee.directly_manages AS directReportees, 0 AS count
  FROM person_reportee manager
  JOIN person_reportee reportee
  ON manager.directly_manages = reportee.pid
  WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
  GROUP BY directReportees
  ) AS T
GROUP BY directReportees)
UNION
(SELECT L2Reportees.pid AS directReportees, count(L2Reportees.directly_manages) AS count
FROM person_reportee manager
JOIN person_reportee L1Reportees
ON manager.directly_manages = L1Reportees.pid
JOIN person_reportee L2Reportees
ON L1Reportees.directly_manages = L2Reportees.pid
WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
GROUP BY directReportees
) AS T
GROUP BY directReportees)
UNION
(SELECT L2Reportees.directly_manages AS directReportees, 0 AS count
FROM person_reportee manager
JOIN person_reportee L1Reportees
ON manager.directly_manages = L1Reportees.pid
JOIN person_reportee L2Reportees
ON L1Reportees.directly_manages = L2Reportees.pid
WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
) AS T
GROUP BY directReportees)
UNION
SELECT reportee.pid AS directReportees, count(reportee.directly_manages) AS count
FROM person_reportee manager
JOIN person_reportee reportee
ON manager.directly_manages = reportee.pid
WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
GROUP BY directReportees
UNION
```

```
MATCH (boss)-[:MANAGES*0..3]->(sub),
      (sub)-[:MANAGES*1..3]->(report)
WHERE boss.name = "John Doe"
RETURN sub.name AS Subordinate,
       count(report) AS Total
```

Less time writing queries

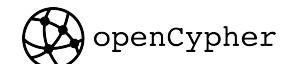
- More time understanding the answers
- Leaving time to ask the next question

Less time debugging queries:

- More time writing the next piece of code
- Improved quality of overall code base

Code that's easier to read:

- Faster ramp-up for new project members
- Improved maintainability & troubleshooting



Indexes and Query Tuning

- Indexes can be added to improve search performance
`CREATE INDEX [index_name]
FOR (n:LabelName)
ON (n.propertyName)`
 - Different indexes, recommended BTREE
 - Composite (multiproperty) indexes can be defined as well
 - Full-text indexes can be defined as well (Apache Lucene)
- The use of indexes by the query optimizer can be checked by looking at the query plan (`EXPLAIN`) or suggested by `USING`

Other querying mechanisms

- Java API
- Graph algorithm library

