

# Decentralized Systems

**IPFS & ENS**

# **InterPlanetary File System**

(thanks to Marina Ribaudó)

# A Centralized Internet


- We do **not own our data**
- Storing all information in central servers can cause problems, like having **single points of failure** and **censorship**
- Nowadays, the Internet is **way less centralized** than one may imagine, and that even causes catastrophic failures...

# A Centralized Internet

- The current way in which information is provided to users is **location-based**
  - Users refer to resources by URL
  - <http://gallery.com/meme1.png>



# InterPlanetary File System



Upgrading the web





# A Centralized Internet

- The current way in which information is provided to users is **location-based**
  - Users refer to resources by URL
  - `http://gallery.com/meme1.png`
- Problem: images can change behind the same URL



# IPFS

- The InterPlanetary File System (**IPFS**) is a protocol and P2P network for **storing** and **sharing data** in a distributed file system
  - **Ethereum** does **computation**
  - **IPFS** does **storage**
- IPFS allows users to not only receive but host content, in a similar manner to BitTorrent

# IPFS

- IPFS lets users **back up files and websites** by hosting them across **numerous nodes**
- Content is **resistant to censorship** and **centralized points of failure**, such as server issues or **coordinated attacks**
- Rather than using file locations, **IPFS points towards contents**, which could be stored on any number of computers around the world



# Content-Based Addressing

- Addresses content by **what** it is, rather than **where** it is stored
  - Computer → file://path-to-file/index.html
  - Internet → https://domain.com/path-to-file/index.html
  - IPFS → ipfs://[CID]/path-to-file/index.html
- Example CID:  
QmcniBv7UQ4gGPQQW2BwbD4ZZHzN3o3tPuNLZCbBchd1zh
  - Long binary string encoded to be **printable**

# Content-Based Addressing

- For each file, and for each folder, a Content Identifier, or **CID**, is computed using a hash function
- This hash is used to store and retrieve all files in IPFS



Example of a CIDv1:

bafybeigdyrzt5sfp7udm7hu76uh7y26nf3efuylqabf3oclgty55fbzdi



<multibase>(cid-version || multicodec || multihash)

v1	dag-pb	sha2-256	32 bytes	SHA256 hash
00000001	01110000	00010010	00100000	110010010...
CID-Version	Multicodec	Multicodec	Length	Actual Hash

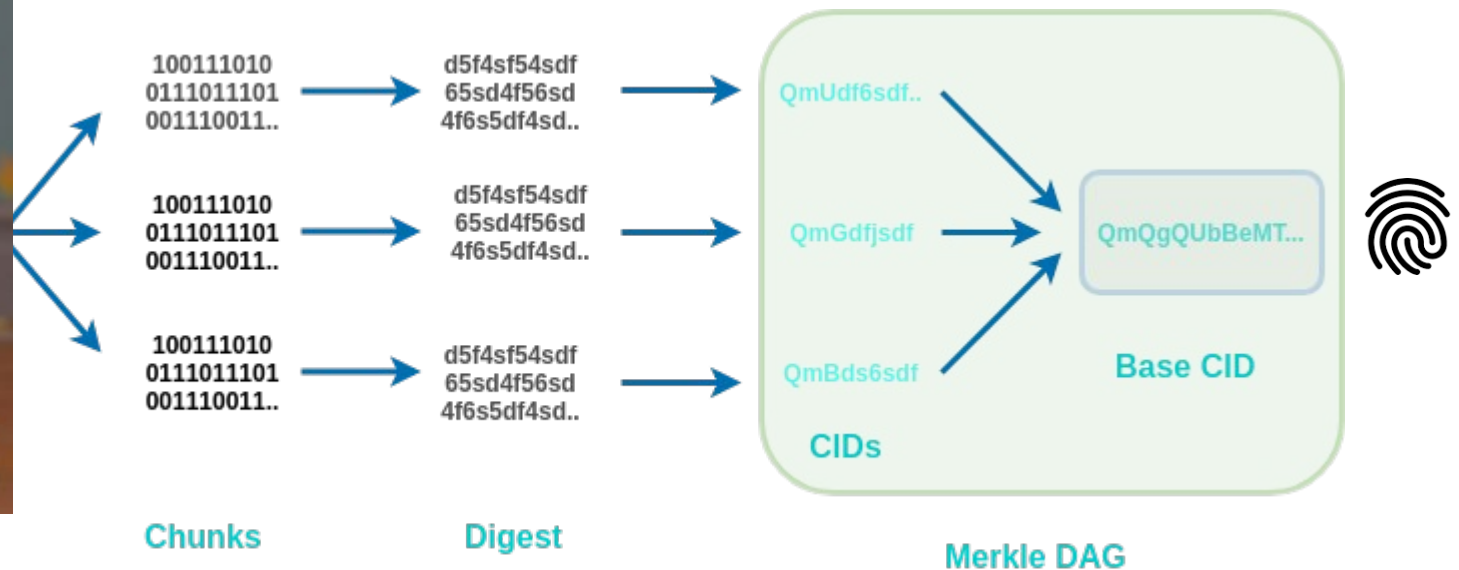
Multihash

# Content-Based Addressing

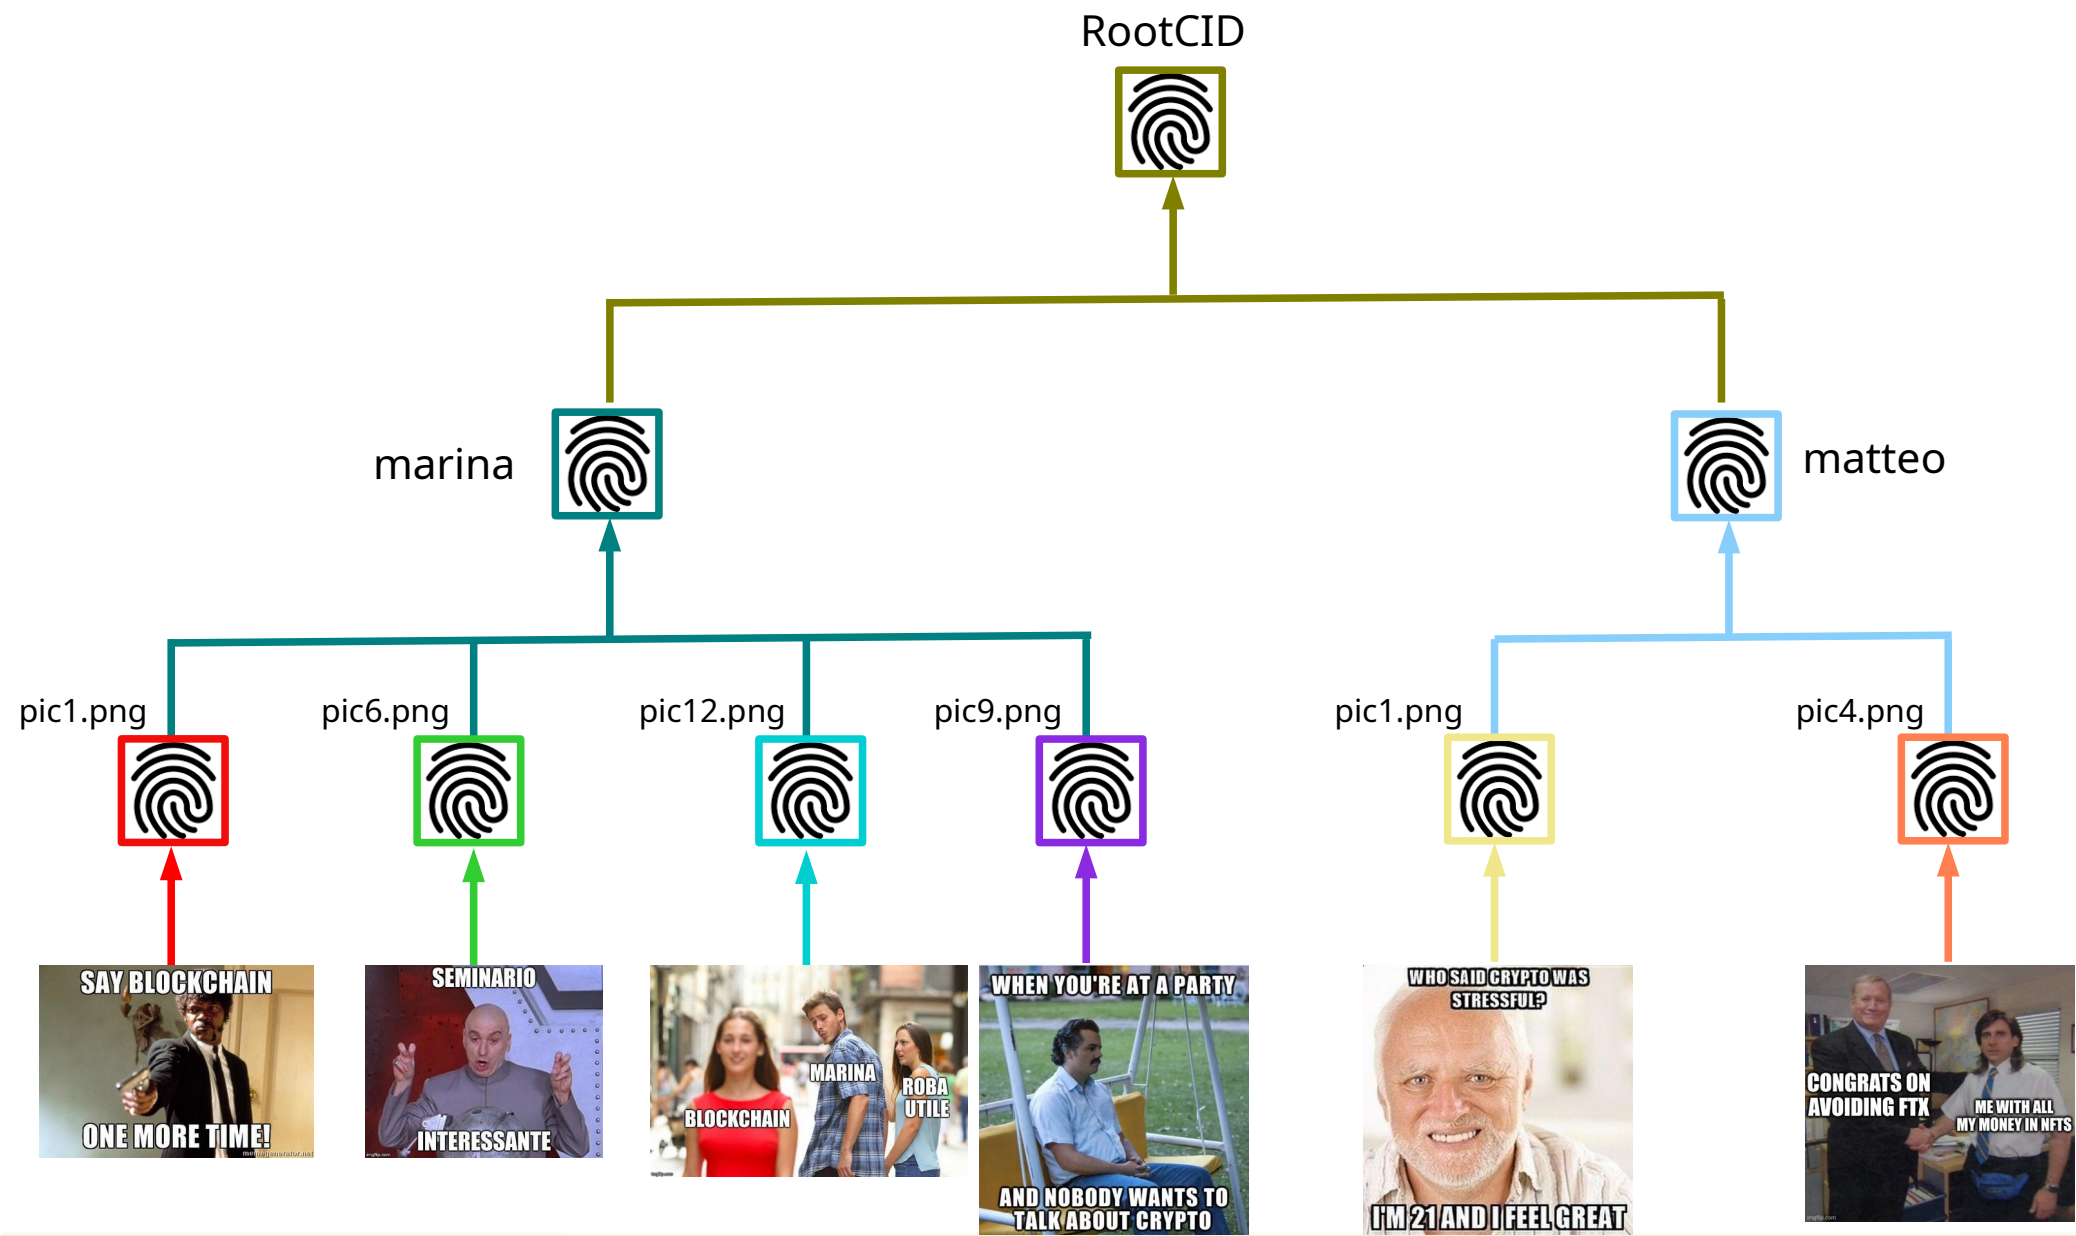
- Large files are **chunked**, **hashed**, and **organized** into **Merkle DAG** objects <https://docs.ipfs.tech/concepts/merkle-dag/>
- Blocks equal or smaller than 256 kB



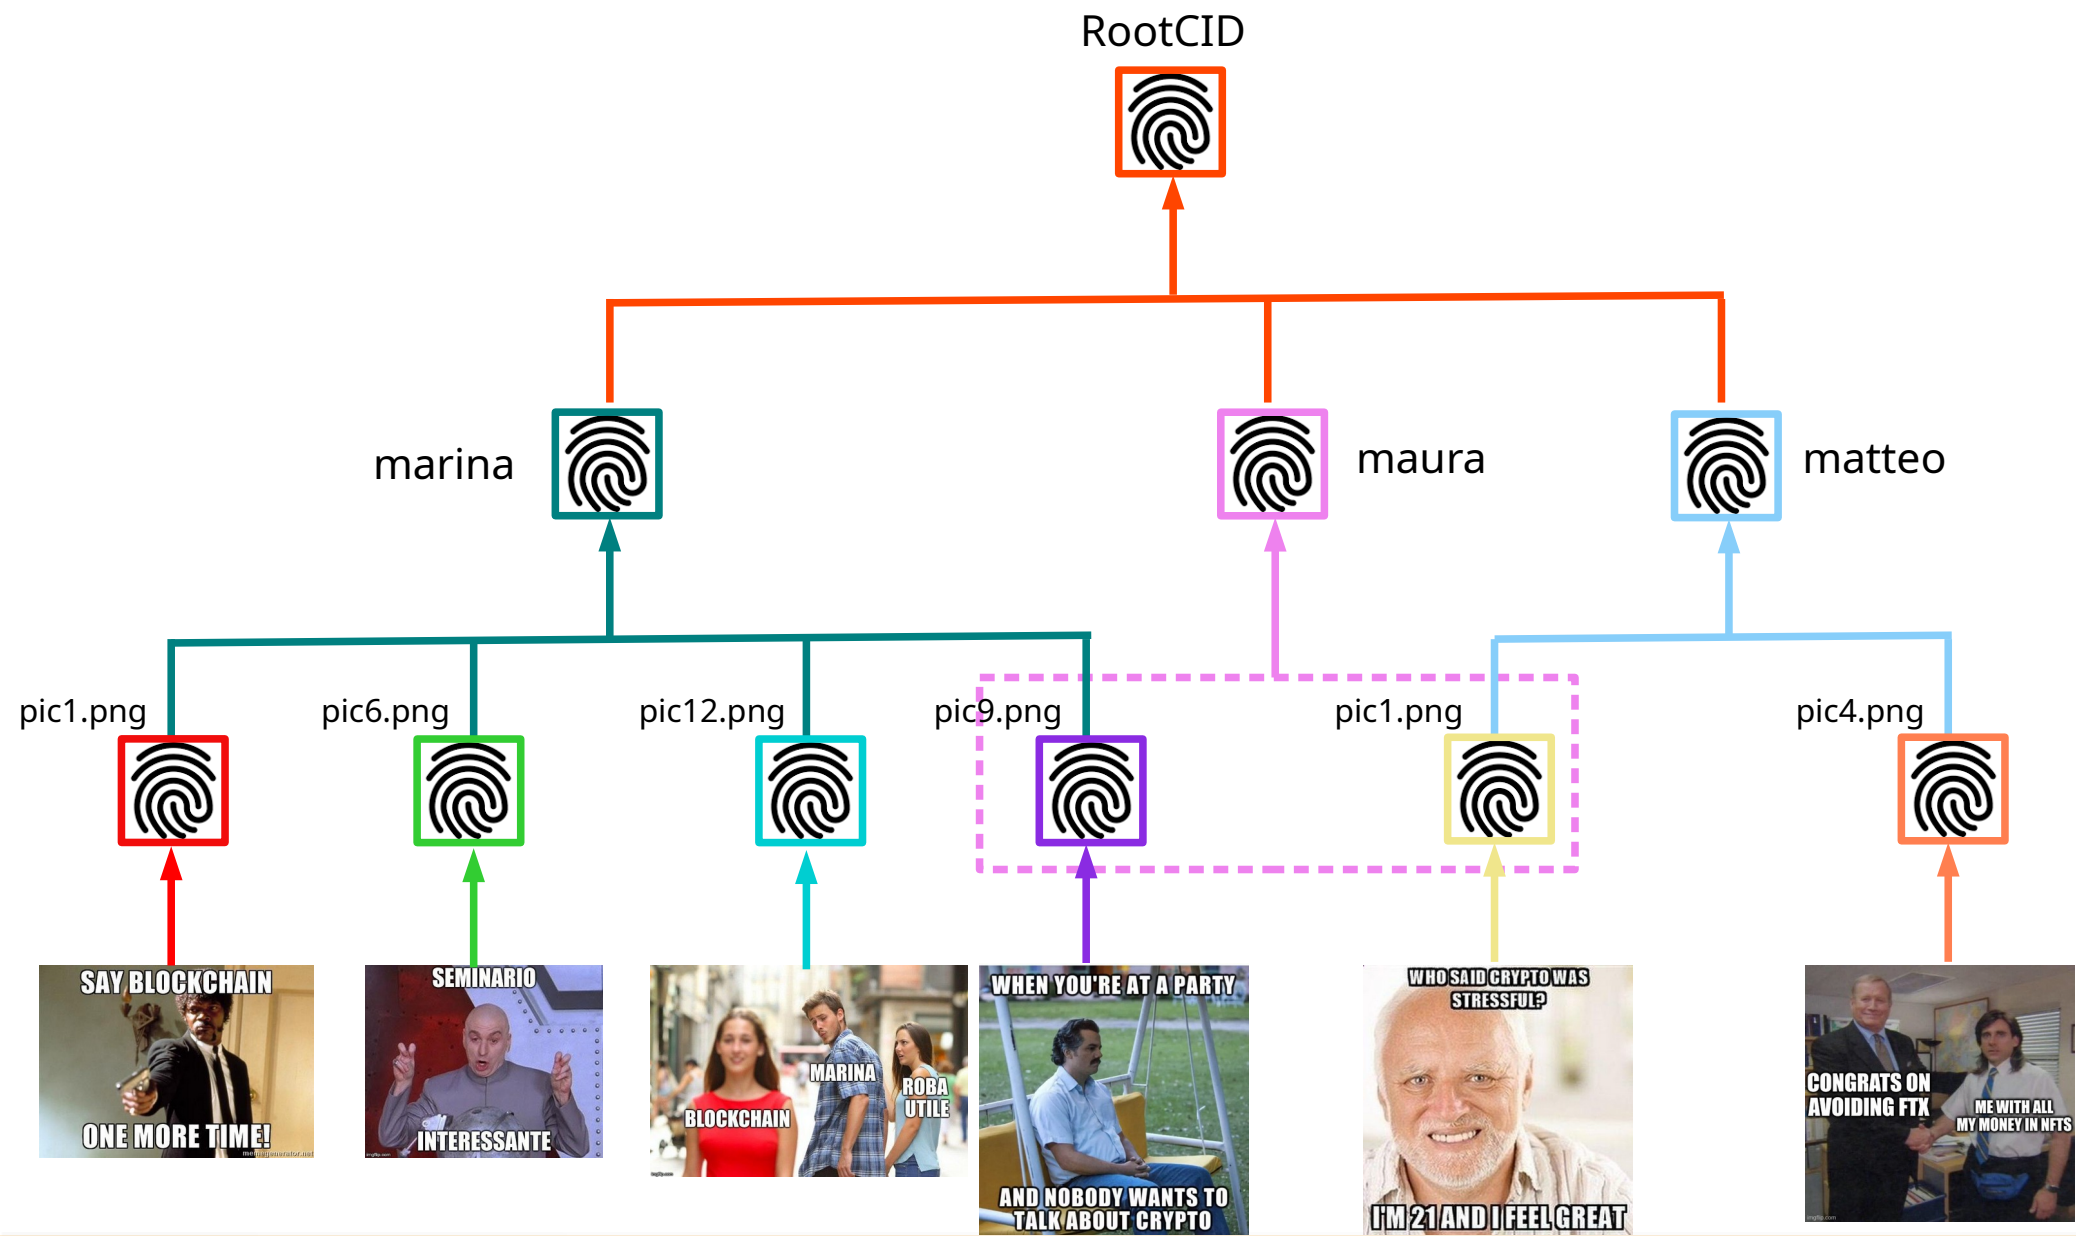
Photo



# IPFS: files and folders



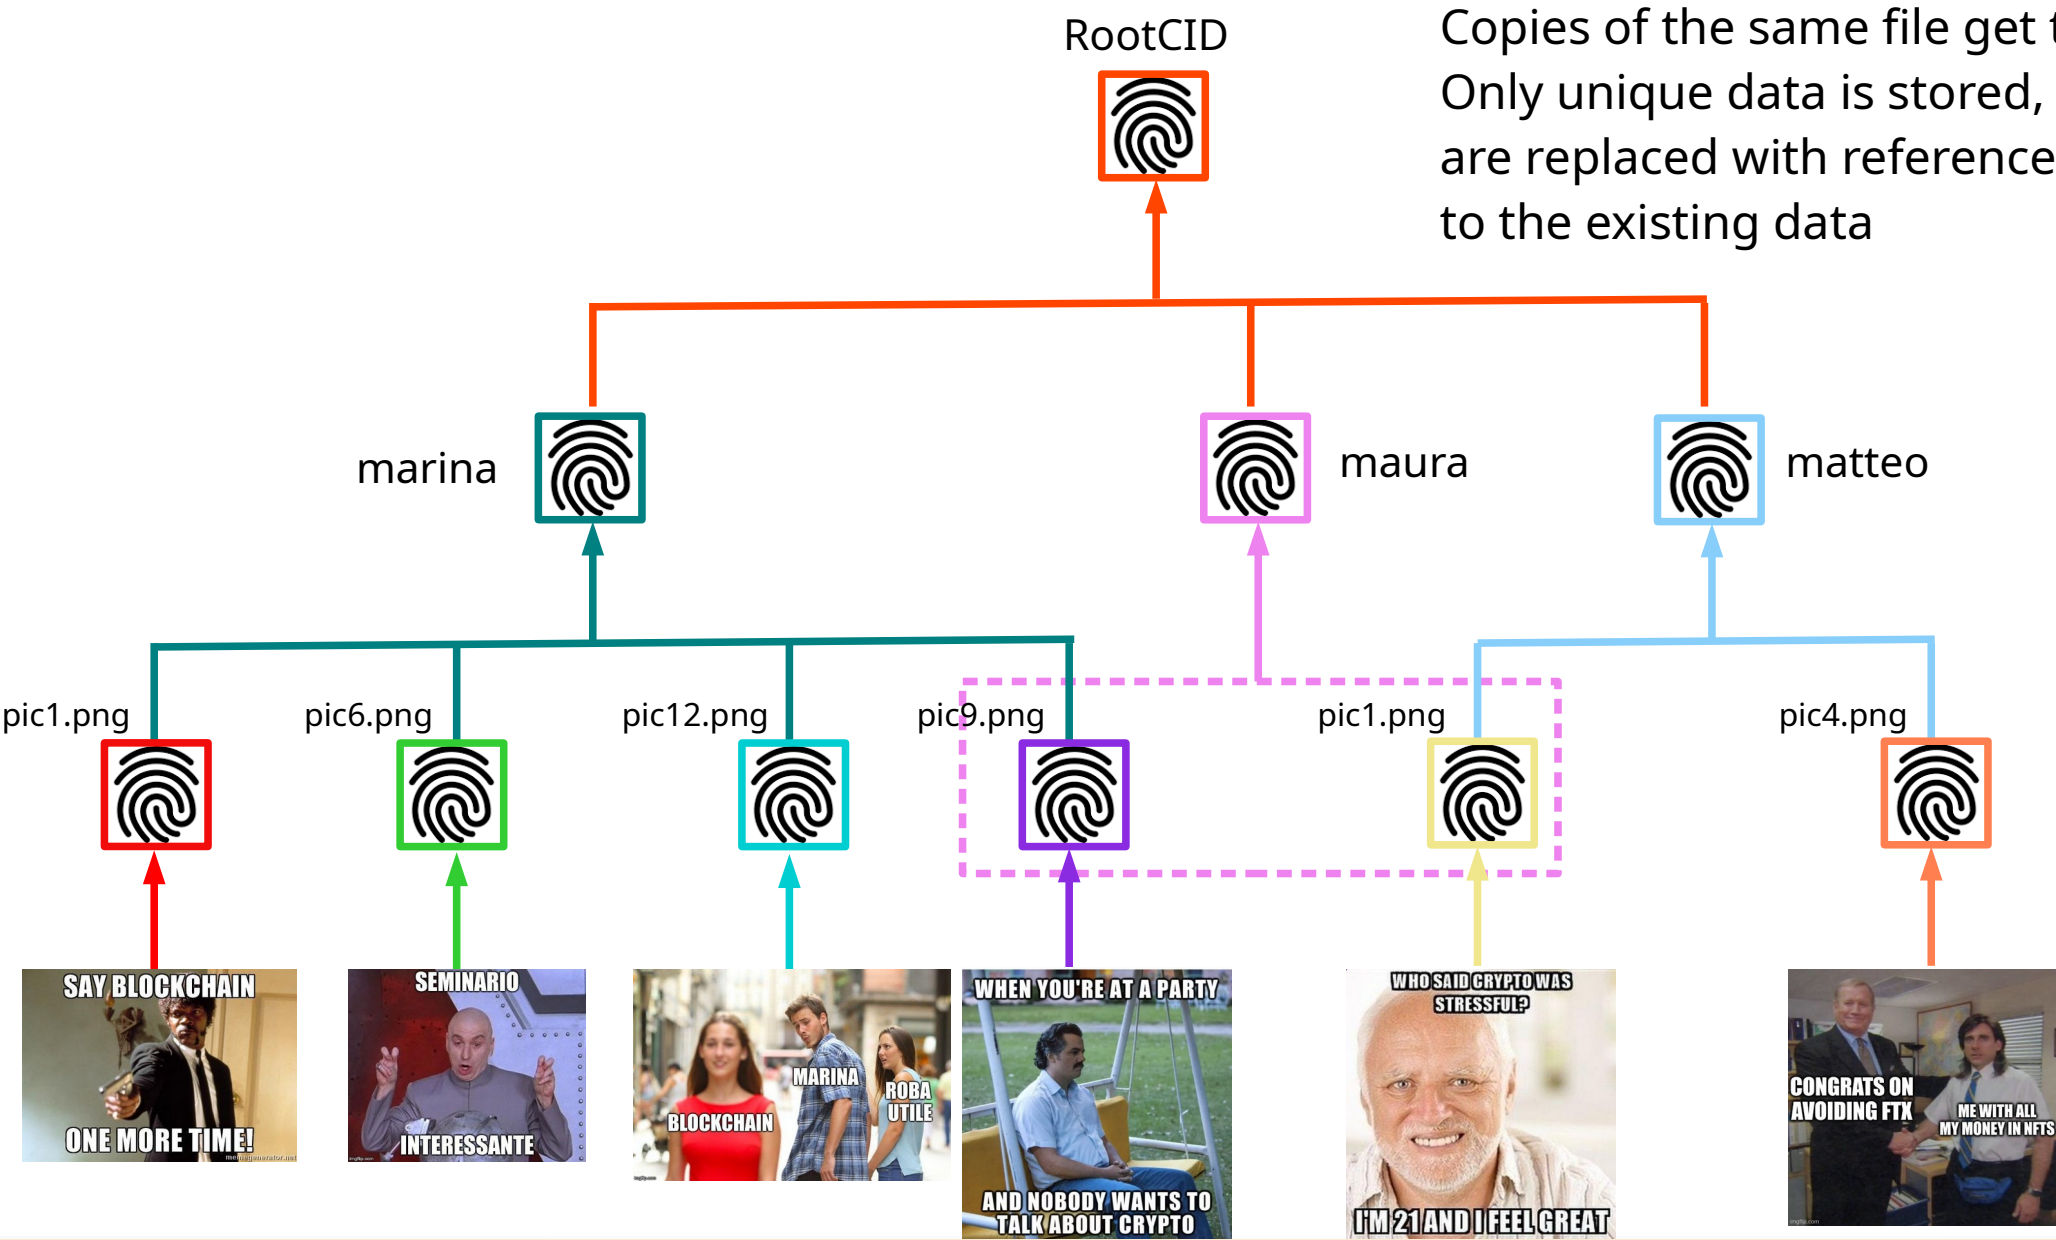
# IPFS: deduplication





# IPFS: deduplication

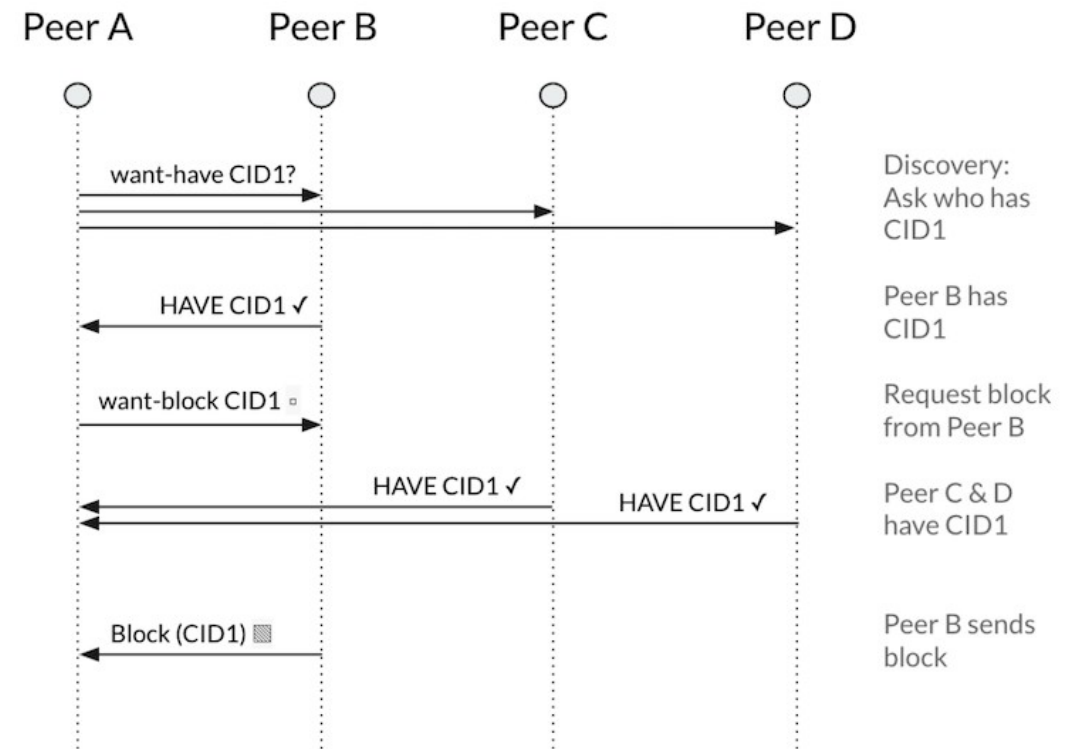
Copies of the same file get the same CID  
Only unique data is stored, and duplicates are replaced with references or pointers to the existing data





# Upload/Download

- Uses the BitSwap protocol, inspired by BitTorrent
- Remember? A **swarm** of peers, exchanging a file made of **small blocks**
  - Tit-for-tat incentives
- Here, a single swarm for **all of IPFS**



<https://docs.ipfs.tech/concepts/bitswap/>

# Content Storage

- After you download a file/block, you **cache** it and **serve it to others**
  - Similarly to **seeding** in BitTorrent
- When disk space is over, you will erase **old pieces of data** you didn't access anymore
- You can decide you'll keep storing a piece of data by **pinning** it
  - It won't be evicted from the cache
- A file remains available as long as **somebody is still keeping it**

# Versioning: IPNS

- IPFS supports **versioning** through IPNS (InterPlanetary Naming System)
  - Instead of a CID, you use the **hash of a public key**
  - As content, you have a file with
    - The full public key
    - The **version number**
    - The CID of the **current version**
    - Signature
- Storage nodes will store the **last correctly signed version**

# DNSLink

- You can have a **human-readable** IPFS address by putting **an IPFS link in your DNS record**
- It's going to be automatically substituted to your DNS name


```
$ dig +noall +answer TXT _dnslink.docs.ipfs.tech
_dnslink.docs.ipfs.tech. 34 IN TXT "dnslink=/ipfs/QmVMxjouRQCA2Qy
```

<https://docs.ipfs.tech/concepts/dnslink/>

- Here, `/ipns/docs.ipfs.tech/introduction/` will be converted to `/ipns/QmVMxjouR.../introduction/`

# Lookup: DHT

- Current connections may be missing some CIDs
- IPFS uses the **Kademlia DHT** (remember?) to support routing and discovery of content and peers on the network
- Peer Ids are directly mapped to CIDs



# Content Lookup

- IPFS implements content lookup through a distributed hash table (**DHT**) and **CIDs**
  - Each **node** in the IPFS network is assigned a **unique identifier**, typically a **256-bit value**
  - When a node wants to find a particular piece of content, it queries the DHT with the **content's CID (256-bit value)**
- Lookup is performed on Kademlia; at last one will get the **addresses of nodes having that CID** (if any)
  - They're added to the list of local neighbors, and download proceeds from there



# Notable Properties



**Nodes can cache content** they have recently accessed, making it available to other nodes in the network, this encourages content replication across multiple nodes, enhancing fault tolerance and availability



IPFS uses **deduplication** to help **reducing redundancy at the storage level** by ensuring that identical content is only stored once



IPFS uses **caching** to help **improving access times** by keeping frequently accessed content closer to the requester. The decision to use caching and the specific caching policies can vary among different IPFS implementations

# IPFS Nodes (1)

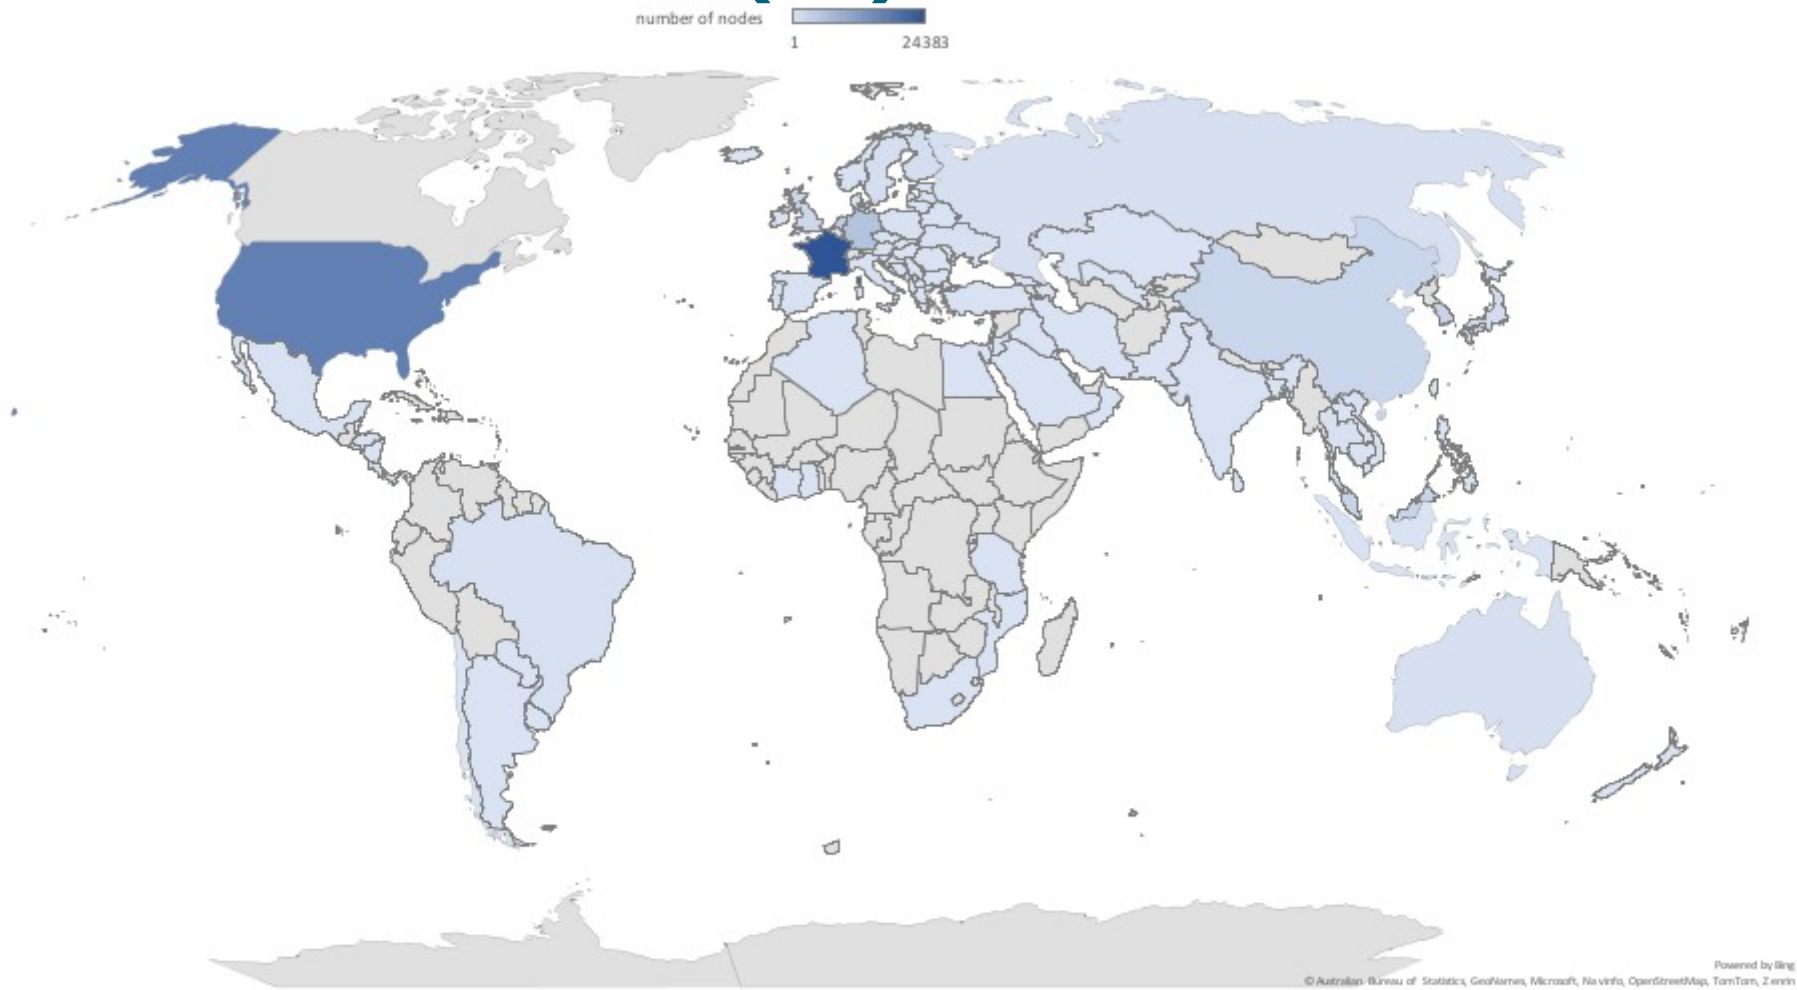




Fig. 5: Number of IPFS nodes per country.

Source: [Confais et al., 2023](#)

# IPFS Nodes (2)



- Source: [Trautwein et al.](#), ACM SIGCOMM '22
- Something I don't understand in the plots, if anybody wants to download the data and try to reproduce them as their seminar...

# Filecoin




- Ad-hoc blockchain complementing IPFS by providing **incentives to store data**
- To store data persistently, you can pay some filecoin to ensure “miners” store your data for some time
- Based on crypto-magic (ZK-SNARKs, seminar opportunity for adventurous students):
  - Proof of space-time
  - Proof of replication

# **Ethereum Name Service**

**(thanks to Federico Fontana)**

# Zooko's Triangle

- A 2001 [conjecture](#) about the impossibility of having names that are at once **meaningful, decentralized** and **secure**
- Q: What about CIDs?  
Ethereum/Bitcoin addresses?  
DNSLink?
- It has been **proven false**
  - ENS is one of the solutions





# Ethereum Name Service



- A way to associate **human-readable names** to **Ethereum addresses**
- Solves Zooko's triangle: decentralized, secure, meaningful
- Based on Ethereum smart contracts
- Contracts are updateable thanks to the proxy pattern (you'll see)

# As a User: Registration Fee

fedfontana.eth

**Before we start**

Registering your name takes three steps

- 1 Complete a transaction to begin the timer
- 2 Wait 60 seconds for the timer to complete
- 3 Complete a second transaction to secure your name

2.88 Gwei

ETH USD

10 year registration	0.0319 ETH
Est. network fee	0.0014 ETH
Estimated total	0.0333 ETH

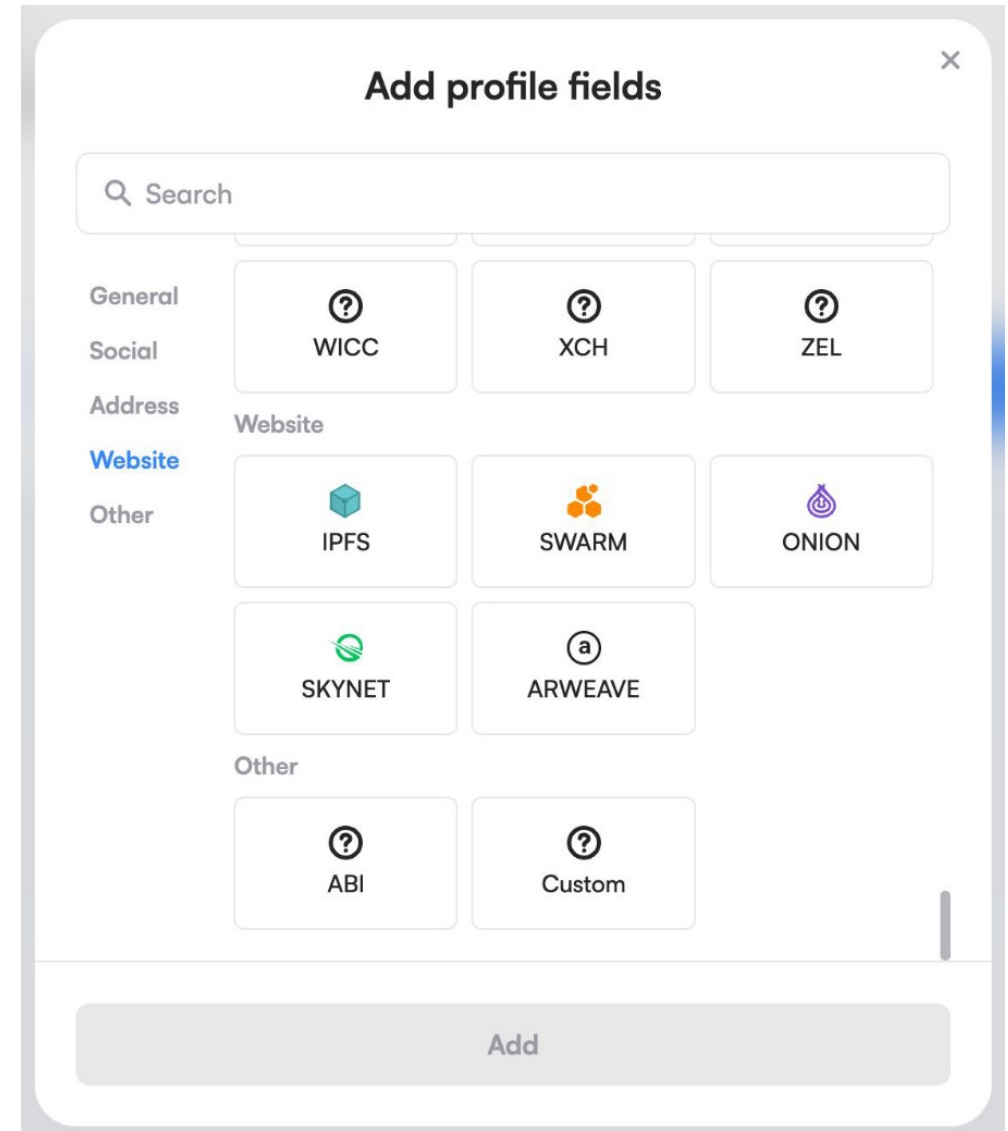
[Back](#) [Begin](#)

Source: [Federico Fontana](#)

- Prices vary (5-640US\$ per year) depending on the **length of the name one wants to register**
- Expressed in dollars in the contracts thanks to **an oracle** (you'll see it in the next lesson!)
- Can be renewed until expiration (+ a grace period)

# As a User: Fields

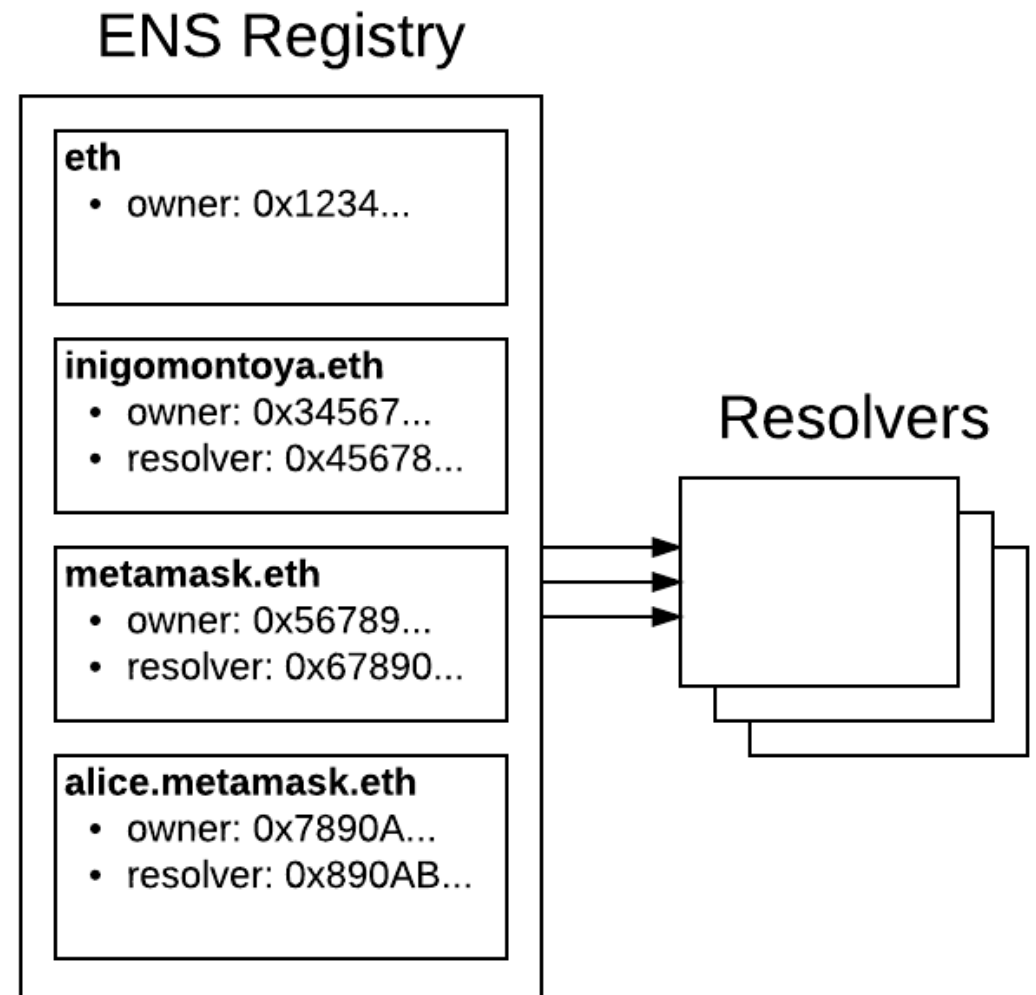
- As with DNS, you can insert **multiple fields** in your record
- Some are standardized, others just use free-form fields
  - Like DNSLink for IPFS+DNS



Source: [Federico Fontana](#)

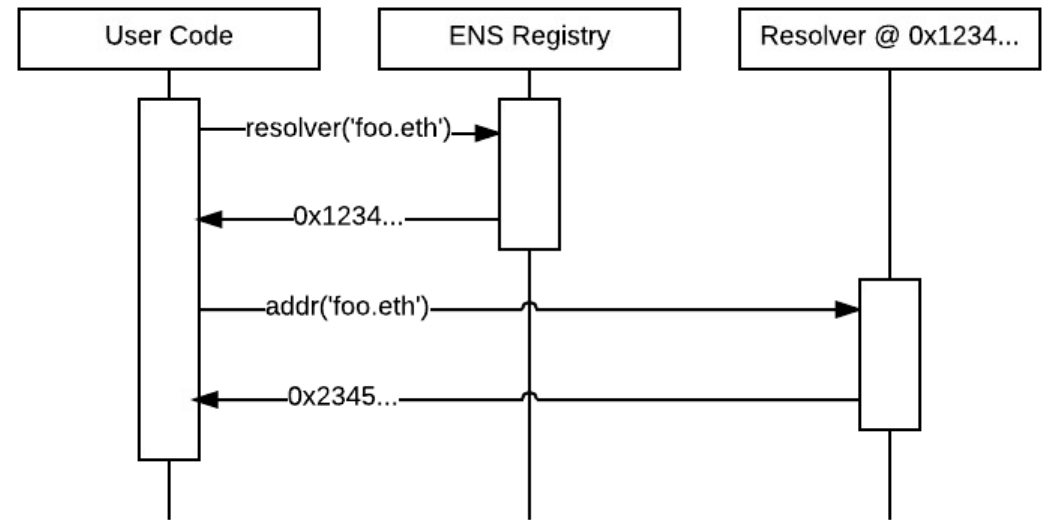
# Registry and Resolvers

- A **single registry** which maps to **resolvers** for each subdomain
- Resolvers are **smart contracts** anyone can implement
  - They must implement at least the `addr` method



# ENS Queries

- The client calls the **registry** to get the resolver's address
- The resolver answers the queries



<https://ens.readthedocs.io/en/stable/implementers.html>

# Registry

- Stores mapping between names and resolver, owner and caching time-to-live (TTL)
- Emits events such as NewOwner, Transfer, NewResolver
- Interacts with another contract, the *Controller*, responsible for registration, fees, renewal
  - Previously (up to 2020): auction model, no renewal
  - Now: commit-reveal pattern to avoid frontrunning



# Other Facilities

- NameWrapper: allows delegating a subdomain, but limiting functionalities (e.g., further transfers, change records, create subdomains)
- Reverse registrar: address `0x<addr>` can claim the `<addr>.addr.reverse` name (and optionally transfer it)
- DNS in ENS: thanks to another oracle, one can use their DNSSEC (DNS+crypto extensions) name online