





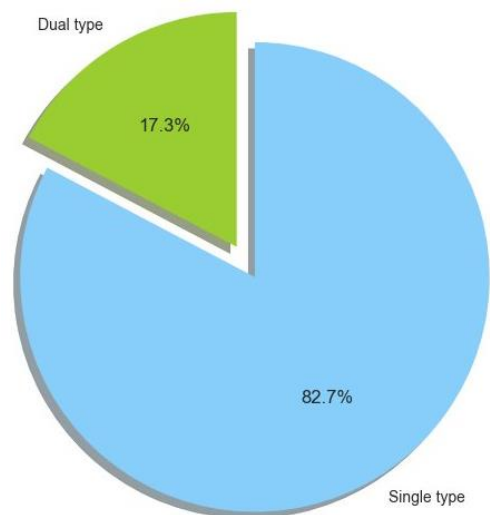
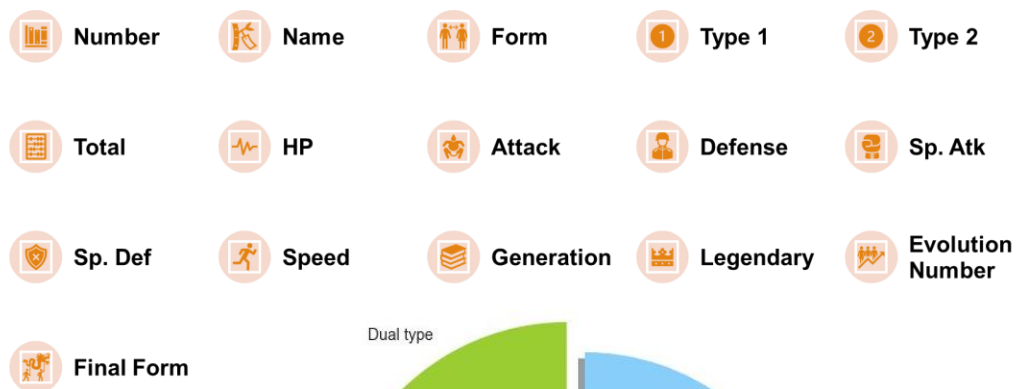
MACHINE LEARNING AND DATA ANALYSIS

MAGISTRELLO CAMILLA (4512554), ENRICO PEZZANO (4825087)

THE DATASET




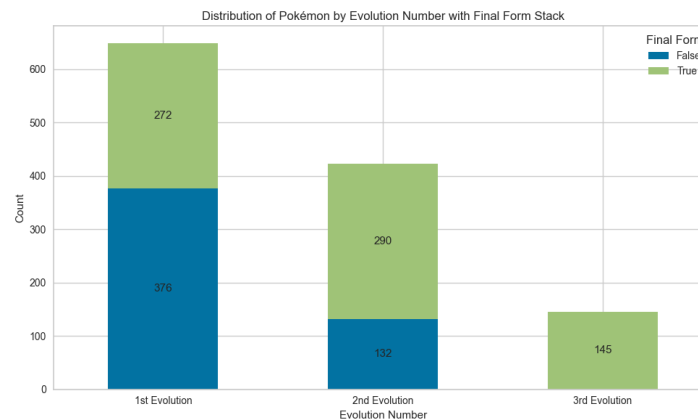



CSV file with 1200+ lines and 16 columns





HOW MANY POKEMON?

- 1088 legendary
- 793 1st or 3rd evolution
- 127 non-legendary
- 422 2nd evolution



CORRELATIONS



PROBLEMS SELECTION

Based on the features,
determine whether a
Pokémon... :

- ...is **legendary** or not
- ...is in its **second evolution**



GridSearchCsv

A fundamental tool in optimizing **hyper-parameters** of machine learning models.

It involves defining a set of values for each **hyper-parameter** of a model and testing all possible combinations of these values.



- Simplicity and clarity
- Comprehensive exploration
- Systematic evaluation




- Computational cost
- Complex models
- Choice of values

Feature Importance


They indicate how much each feature contributes to the result of a machine learning model (mainly used in models like Decision Trees and Random Forest).

KNeighborsClassifier is based on distances between points and not on coefficients or feature importance, so they cannot be calculated.

DATASET PREPARATION



Transforming the values in columns "Type1" and "Type2" from strings to numeric representations corresponding to the type.



RandomOverSampler to over-sample the legendary class by picking samples at random with replacement.

Based on the features, determine whether a Pokémon is **Legendary** or not



DATASET PREPARATION

Drop the columns:

- ☐ ID
- ☐ Name
- ☐ Form
- ☐ **Legendary**


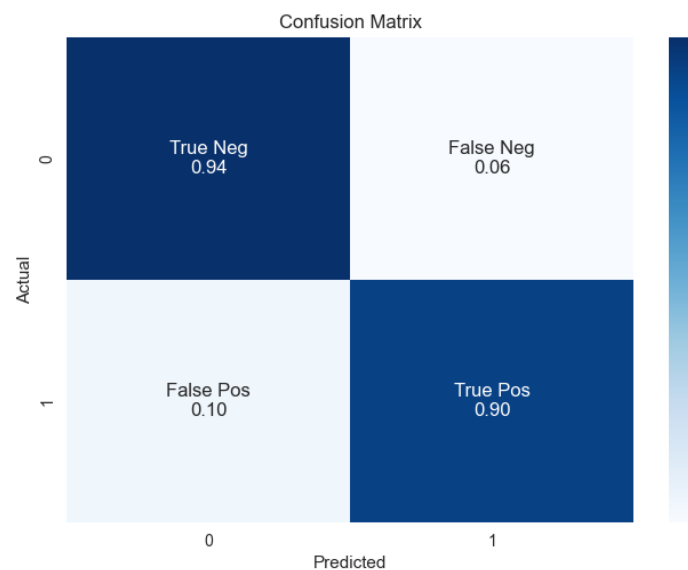



DECISION TREE CLASSIFIER WITHOUT HYPERPARAMETERS

Testing accuracy: 0.9175675675675675

Training accuracy: 1.0

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.90 | 0.94 | 0.92 | 374 |
| 1 | 0.93 | 0.90 | 0.92 | 366 |
| accuracy | | | 0.92 | 740 |
| macro avg | 0.92 | 0.92 | 0.92 | 740 |
| weighted avg | 0.92 | 0.92 | 0.92 | 740 |



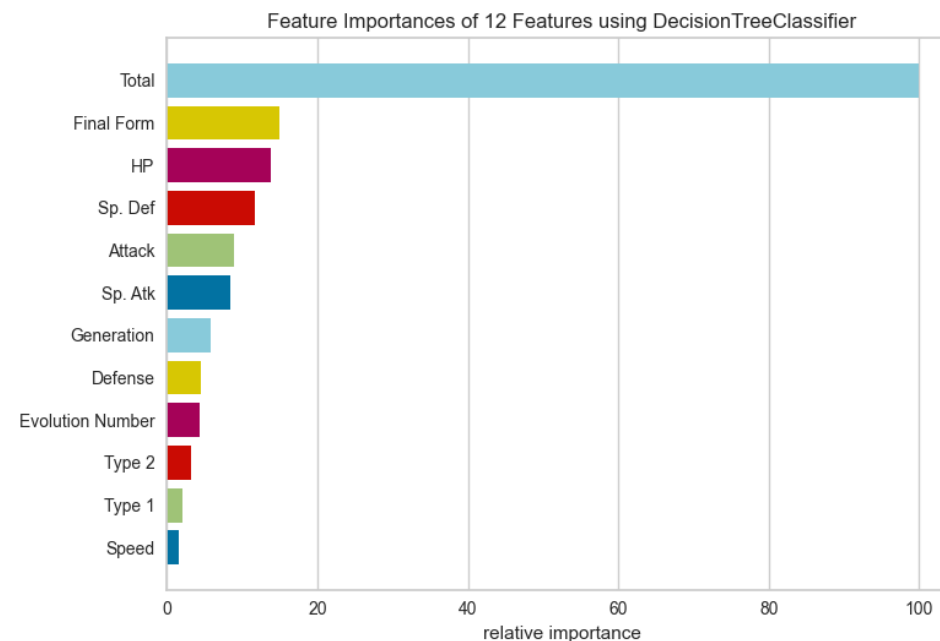
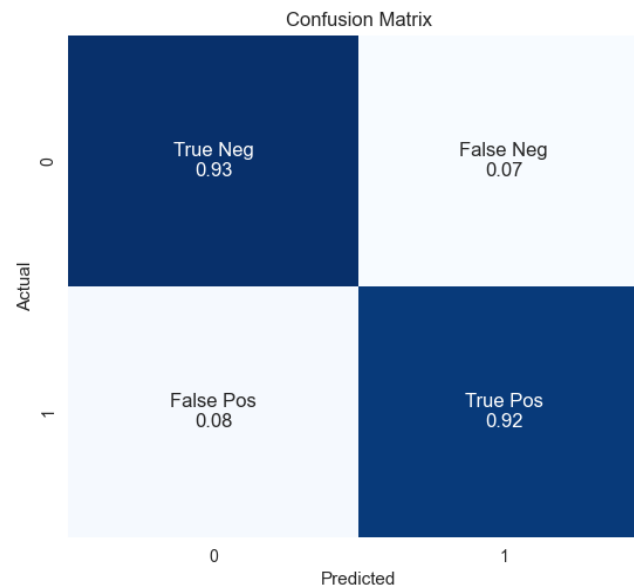



DECISION TREE CLASSIFIER WITH HYPERPARAMETERS

'criterion': 'entropy'
'max_depth': None
'min_samples_leaf': 1
'min_samples_split': 2

Testing accuracy: 0.922972972972973
Training accuracy: 1.0

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.92 | 0.93 | 0.92 | 374 |
| 1 | 0.93 | 0.92 | 0.92 | 366 |
| accuracy | | | 0.92 | 740 |
| macro avg | 0.92 | 0.92 | 0.92 | 740 |
| weighted avg | 0.92 | 0.92 | 0.92 | 740 |



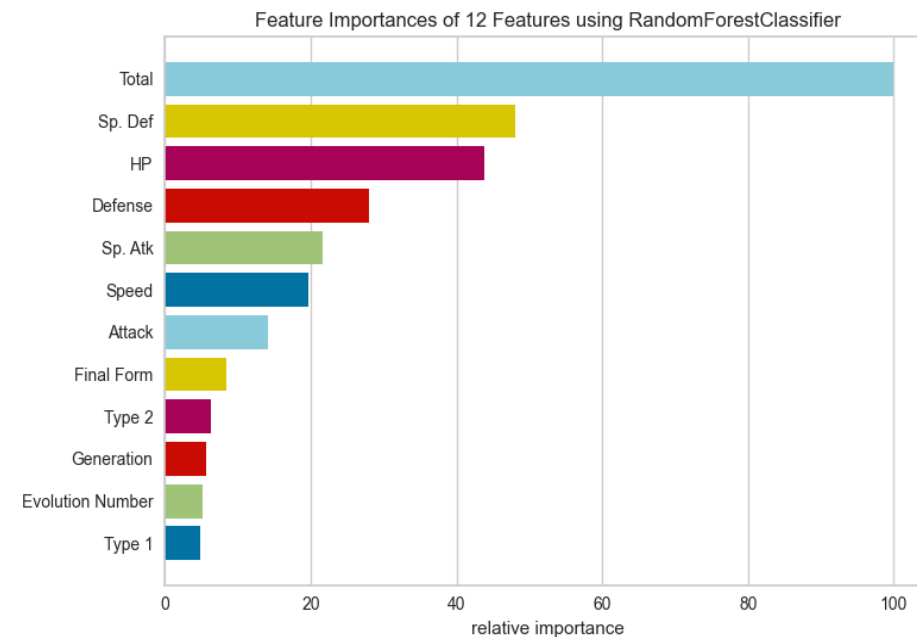
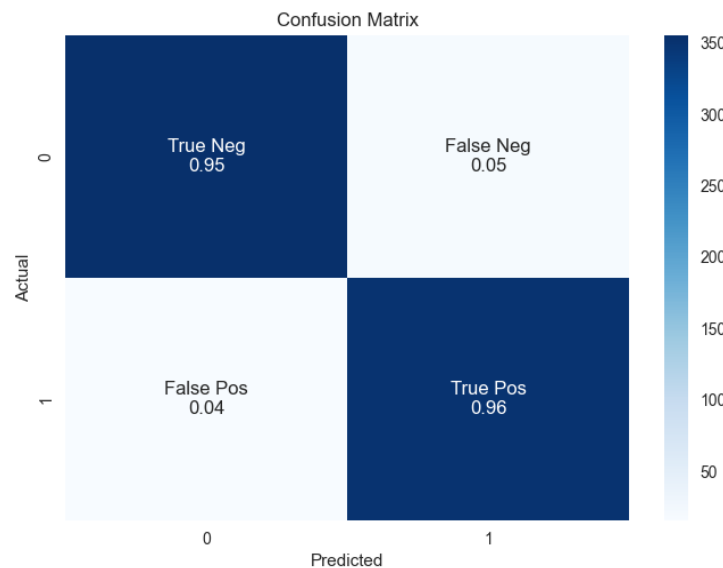



RANDOM FOREST CLASSIFIER WITHOUT HYPERPARAMETERS

Testing accuracy: 0.9527027027027027

Training accuracy: 1.0

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.96 | 0.95 | 0.95 | 374 |
| 1 | 0.95 | 0.96 | 0.95 | 366 |
| accuracy | | | 0.95 | 740 |
| macro avg | 0.95 | 0.95 | 0.95 | 740 |
| weighted avg | 0.95 | 0.95 | 0.95 | 740 |





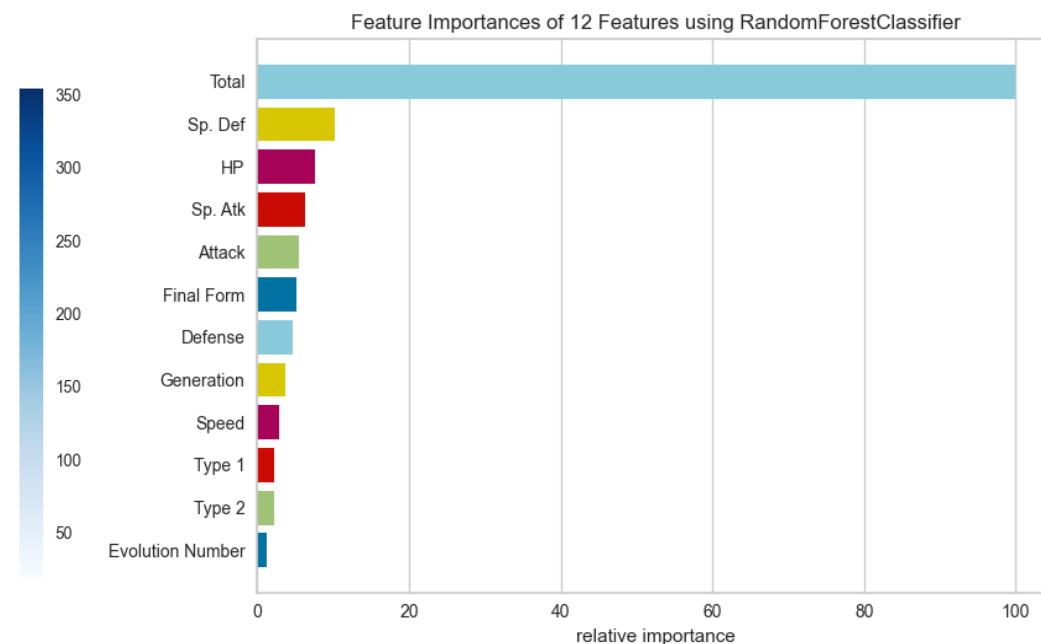

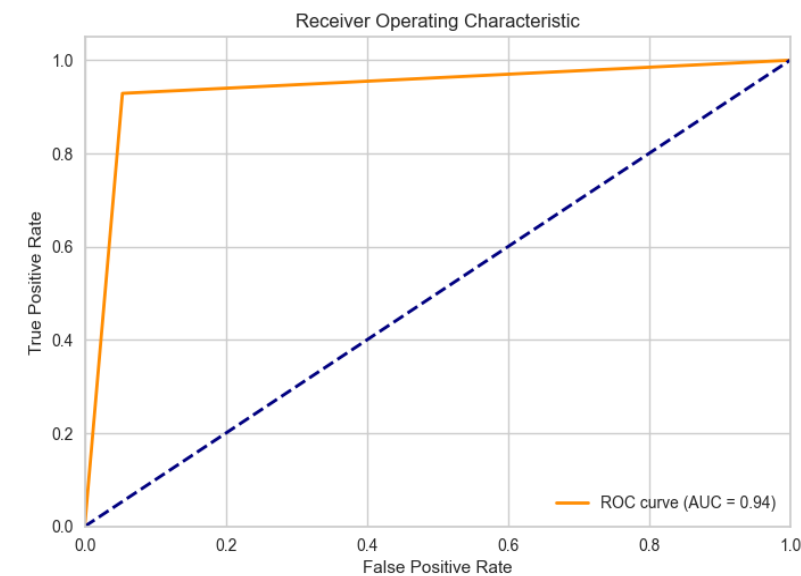
RANDOM FOREST CLASSIFIER WITH HYPERPARAMETERS


'max_depth': None
'max_features': None
'min_samples_leaf': 1
'min_samples_split': 2
'n_estimators': 70

Testing accuracy: 0.9378378378378378

Training accuracy: 1.0

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.93 | 0.95 | 0.94 | 374 |
| 1 | 0.94 | 0.93 | 0.94 | 366 |
| accuracy | | | 0.94 | 740 |
| macro avg | 0.94 | 0.94 | 0.94 | 740 |
| weighted avg | 0.94 | 0.94 | 0.94 | 740 |




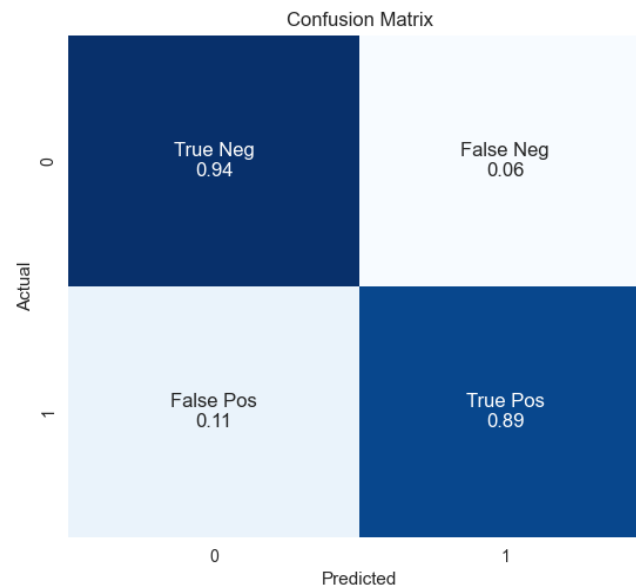




LOGISTIC REGRESSION WITHOUT HYPERPARAMETERS

Testing accuracy: 0.9148648648648648

Training accuracy: 0.915041782729805

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.89 | 0.94 | 0.92 | 374 |
| 1 | 0.94 | 0.89 | 0.91 | 366 |
| accuracy | | | 0.91 | 740 |
| macro avg | 0.92 | 0.91 | 0.91 | 740 |
| weighted avg | 0.92 | 0.91 | 0.91 | 740 |





LOGISTIC REGRESSION WITH HYPERPARAMETERS

'C': 10


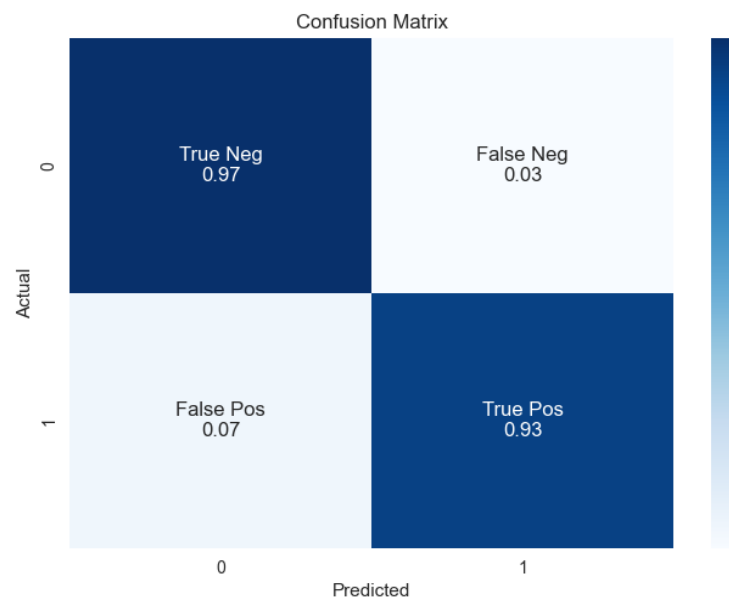

'penalty': 'l1'

'solver': 'liblinear'

Testing accuracy: 0.9472972972972973


Training accuracy: 0.9533426183844012

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.93 | 0.97 | 0.95 | 374 |
| 1 | 0.97 | 0.93 | 0.95 | 366 |
| accuracy | | | 0.95 | 740 |
| macro avg | 0.95 | 0.95 | 0.95 | 740 |
| weighted avg | 0.95 | 0.95 | 0.95 | 740 |






K-NEAREST NEIGHBORS WITHOUT HYPERPARAMETERS



Testing accuracy: 0.9216216216216216
Training accuracy: 0.9352367688022284


| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.88 | 0.98 | 0.93 | 374 |
| 1 | 0.98 | 0.86 | 0.92 | 366 |
| accuracy | | | 0.92 | 740 |
| macro avg | 0.93 | 0.92 | 0.92 | 740 |
| weighted avg | 0.93 | 0.92 | 0.92 | 740 |






K-NEAREST NEIGHBORS WITH HYPERPARAMETERS

'metric': 'manhattan',
'n_neighbors': 5,
'weights': 'distance'




Testing accuracy: 0.9094594594594595
Training accuracy: 1.0

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.86 | 0.98 | 0.92 | 374 |
| 1 | 0.98 | 0.84 | 0.90 | 366 |
| accuracy | | | 0.91 | 740 |
| macro avg | 0.92 | 0.91 | 0.91 | 740 |
| weighted avg | 0.92 | 0.91 | 0.91 | 740 |



Checking the results...

```
Type 1: KS statistic = 0.0357, p-value = 0.5459
Type 2: KS statistic = 0.0457, p-value = 0.2508
Total: KS statistic = 0.0456, p-value = 0.2526
HP: KS statistic = 0.0577, p-value = 0.0737
Attack: KS statistic = 0.0330, p-value = 0.6460
Defense: KS statistic = 0.0543, p-value = 0.1070
Sp. Atk: KS statistic = 0.0386, p-value = 0.4458
Sp. Def: KS statistic = 0.0151, p-value = 0.9998
Speed: KS statistic = 0.0263, p-value = 0.8763
Generation: KS statistic = 0.0420, p-value = 0.3439
Evolution Number: KS statistic = 0.0352, p-value = 0.5663
Final Form: KS statistic = 0.0498, p-value = 0.1695
```



- Kolmogorov–Smirnov Test:
 - **p-value** < 0.05 implies a statistically significant difference.
 - It suggests that the distributions of the attributes are very similar between the two groups, except for the Performance Status
- Legendaries totals stats too high, as we suspected
- Hence why ...

Based on the features, determine whether a Pokémon is in its **second evolution**



DATASET PREPARATION

Drop the columns:


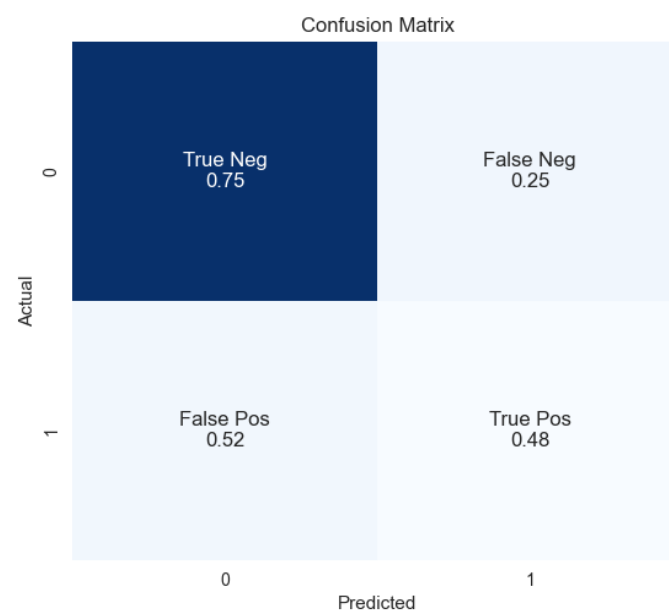

- ☐ ID
- ☐ Name
- ☐ Form
- ☐ **Evolution number**



DECISION TREE CLASSIFIER WITHOUT HYPERPARAMETERS

Testing accuracy: 0.6621621621621622
Training accuracy: 0.9993036211699164

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| False | 0.75 | 0.75 | 0.75 | 501 |
| True | 0.48 | 0.48 | 0.48 | 239 |
| accuracy | | | 0.66 | 740 |
| macro avg | 0.61 | 0.61 | 0.61 | 740 |
| weighted avg | 0.66 | 0.66 | 0.66 | 740 |





DECISION TREE CLASSIFIER WITH HYPERPARAMETERS

'criterion': 'log_loss'

'max_depth': 10


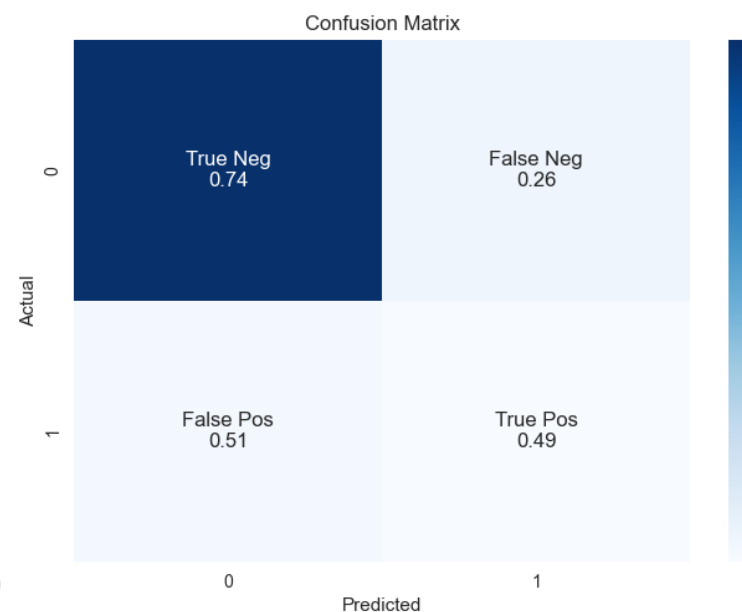

'min_samples_leaf': 1

'min_samples_split': 2

Testing accuracy: 0.6621621621621622

Training accuracy: 0.8760445682451253

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| False | 0.75 | 0.74 | 0.75 | 501 |
| True | 0.48 | 0.49 | 0.48 | 239 |
| accuracy | | | 0.66 | 740 |
| macro avg | 0.62 | 0.62 | 0.62 | 740 |
| weighted avg | 0.66 | 0.66 | 0.66 | 740 |








RANDOM FOREST CLASSIFIER WITHOUT HYPERPARAMETERS

Testing accuracy: 0.7405405405405405
Training accuracy: 0.9993036211699164

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| False | 0.77 | 0.89 | 0.82 | 501 |
| True | 0.65 | 0.43 | 0.52 | 239 |
| accuracy | | | 0.74 | 740 |
| macro avg | 0.71 | 0.66 | 0.67 | 740 |
| weighted avg | 0.73 | 0.74 | 0.72 | 740 |




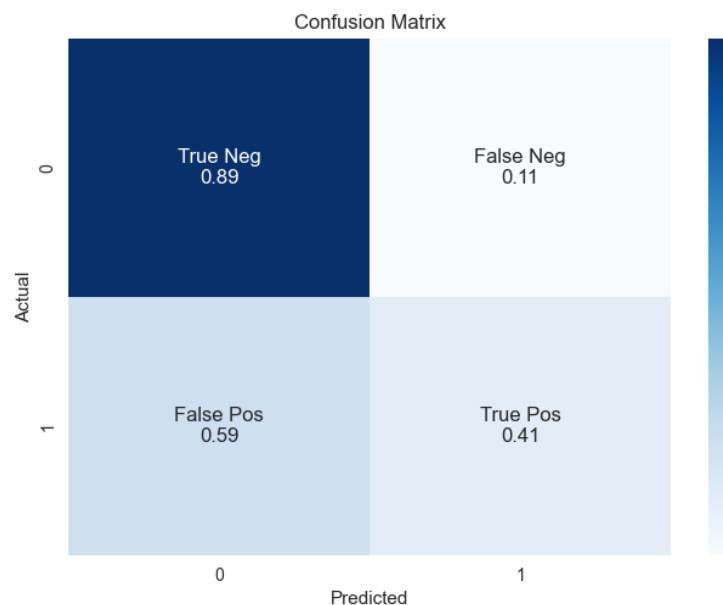



RANDOM FOREST CLASSIFIER WITH HYPERPARAMETERS

'max_depth': 10
'max_features': 'sqrt'
'min_samples_leaf': 5
'min_samples_split': 4
'n_estimators': 80

Testing accuracy: 0.7337837837837838
Training accuracy: 0.8600278551532033

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| False | 0.76 | 0.89 | 0.82 | 501 |
| True | 0.64 | 0.41 | 0.50 | 239 |
| accuracy | | | 0.73 | 740 |
| macro avg | 0.70 | 0.65 | 0.66 | 740 |
| weighted avg | 0.72 | 0.73 | 0.72 | 740 |









LOGISTIC REGRESSION WITHOUT HYPERPARAMETERS

Testing accuracy: 0.6635135135135135
Training accuracy: 0.6385793871866295

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| False | 0.68 | 0.96 | 0.80 | 501 |
| True | 0.31 | 0.03 | 0.06 | 239 |
| accuracy | | | 0.66 | 740 |
| macro avg | 0.49 | 0.50 | 0.43 | 740 |
| weighted avg | 0.56 | 0.66 | 0.56 | 740 |








LOGISTIC REGRESSION WITH HYPERPARAMETERS

'C': 0.1
'l1_ratio': 0.1,
'penalty': 'l1'
'solver': 'liblinear'


Testing accuracy: 0.6756756756756757
Training accuracy: 0.6448467966573816

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| False | 0.68 | 0.99 | 0.81 | 501 |
| True | 0.40 | 0.01 | 0.02 | 239 |
| accuracy | | | 0.68 | 740 |
| macro avg | 0.54 | 0.50 | 0.41 | 740 |
| weighted avg | 0.59 | 0.68 | 0.55 | 740 |






K-NEAREST NEIGHBORS WITHOUT HYPERPARAMETERS



Testing accuracy: 0.6797297297297298
Training accuracy: 0.7583565459610028


| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| False | 0.74 | 0.80 | 0.77 | 501 |
| True | 0.51 | 0.42 | 0.46 | 239 |
| accuracy | | | 0.68 | 740 |
| macro avg | 0.62 | 0.61 | 0.62 | 740 |
| weighted avg | 0.67 | 0.68 | 0.67 | 740 |





K-NEAREST NEIGHBORS WITH HYPERPARAMETERS


'metric': 'manhattan',
'n_neighbors': 7,
'weights': 'distance'



Testing accuracy: 0.7

Training accuracy: 0.9993036211699164

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| False | 0.76 | 0.81 | 0.78 | 501 |
| True | 0.54 | 0.48 | 0.51 | 239 |
| accuracy | | | 0.70 | 740 |
| macro avg | 0.65 | 0.64 | 0.65 | 740 |
| weighted avg | 0.69 | 0.70 | 0.69 | 740 |



CONCLUSIONS

IS LEGENDARY

- **The best model is LR with hyperparameters**, achieving a testing accuracy of 0.9473 and a training accuracy of 0.95, successfully avoiding overfitting.
- **The worst model is KNN with hyperparameters**, which suffers from overfitting and has the lowest accuracy among all models.

IS IN ITS SECOND EVOLUTION

- **The best model is the RF without hyperparameter tuning**, it achieves the best balance between testing and training performance.
- **The worst model is LR** (with or without hyperparameter tuning), which has a ROC AUC of 0.50—meaning it cannot distinguish between the positive and negative classes, equivalent to a coin toss.

CONCLUSIONS

Is Legendary? Is an easier variable to predict because Legendary Pokémon typically have much higher stats, and just a few key features are often enough to identify them.

Is Second Evolution? Is harder to predict, as it depends on the Pokémon's evolutionary structure, which is not always linear or uniform. Additionally, the data is less distinctive — for example, a second-stage Pokémon might have stats similar to those in the first or final stage.



THANK YOU