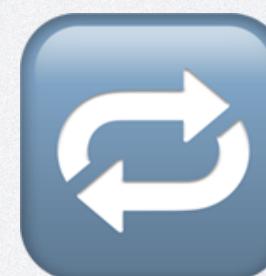


*Ritenta, sarai più fortunato.*

(Anonimo)

Baci PERUGINA.

# RETRY PATTERN



by Enrico Pezzano

# INTRODUCTION



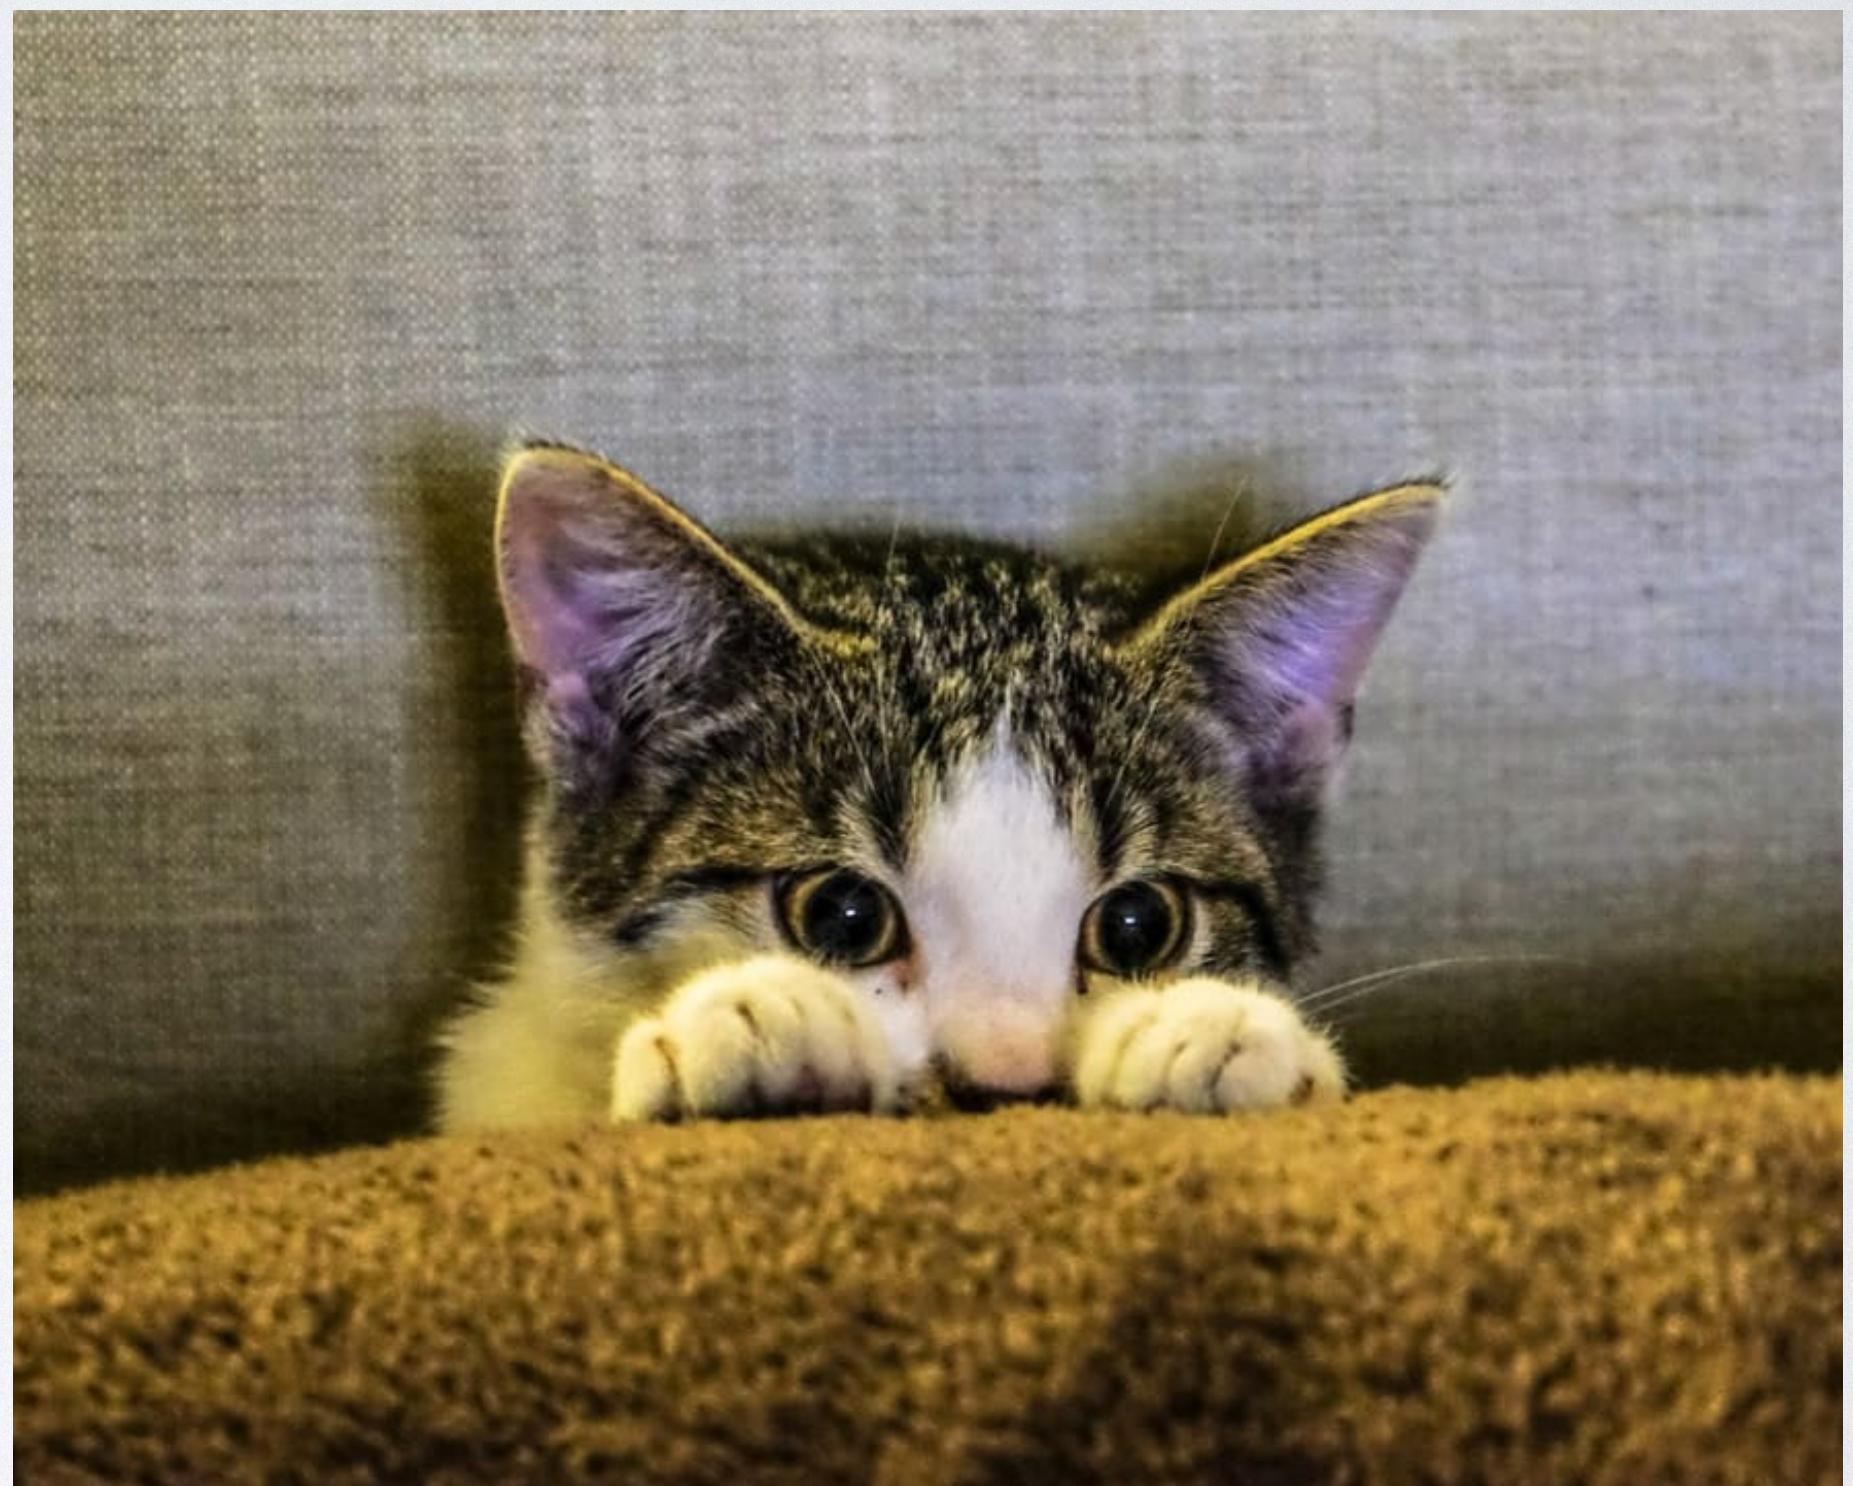
- “Resilient approach”
- Robustness & tolerance is a paramount
- Powerful & widely adopted
- Systematic & flexible
- ⚡
- Reliability & availability



# OUTLOOK

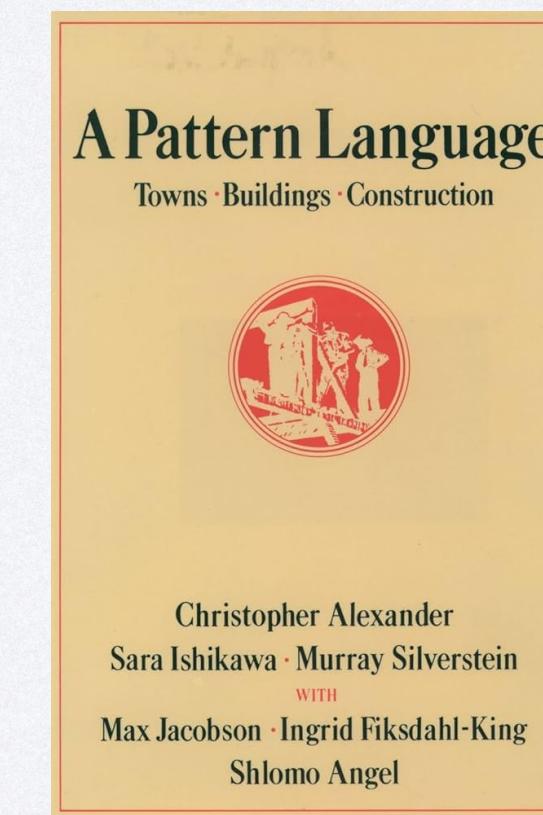


- Why do we need it ?
- What is it ?
- What is the problem?
- Pros & Cons ?
- Solutions ...
- Code examples



“Each pattern describes a problem which occurs over and over again in our environment, and then describes the core of the solution to that problem, in such a way that you can use this solution a million times over, without ever doing it the same way twice.”

—Christopher Alexander, A Pattern Language 

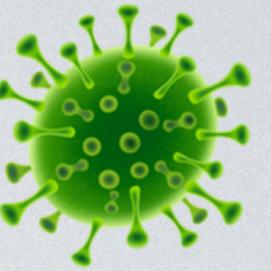


# WHAT IS IT ? 😐

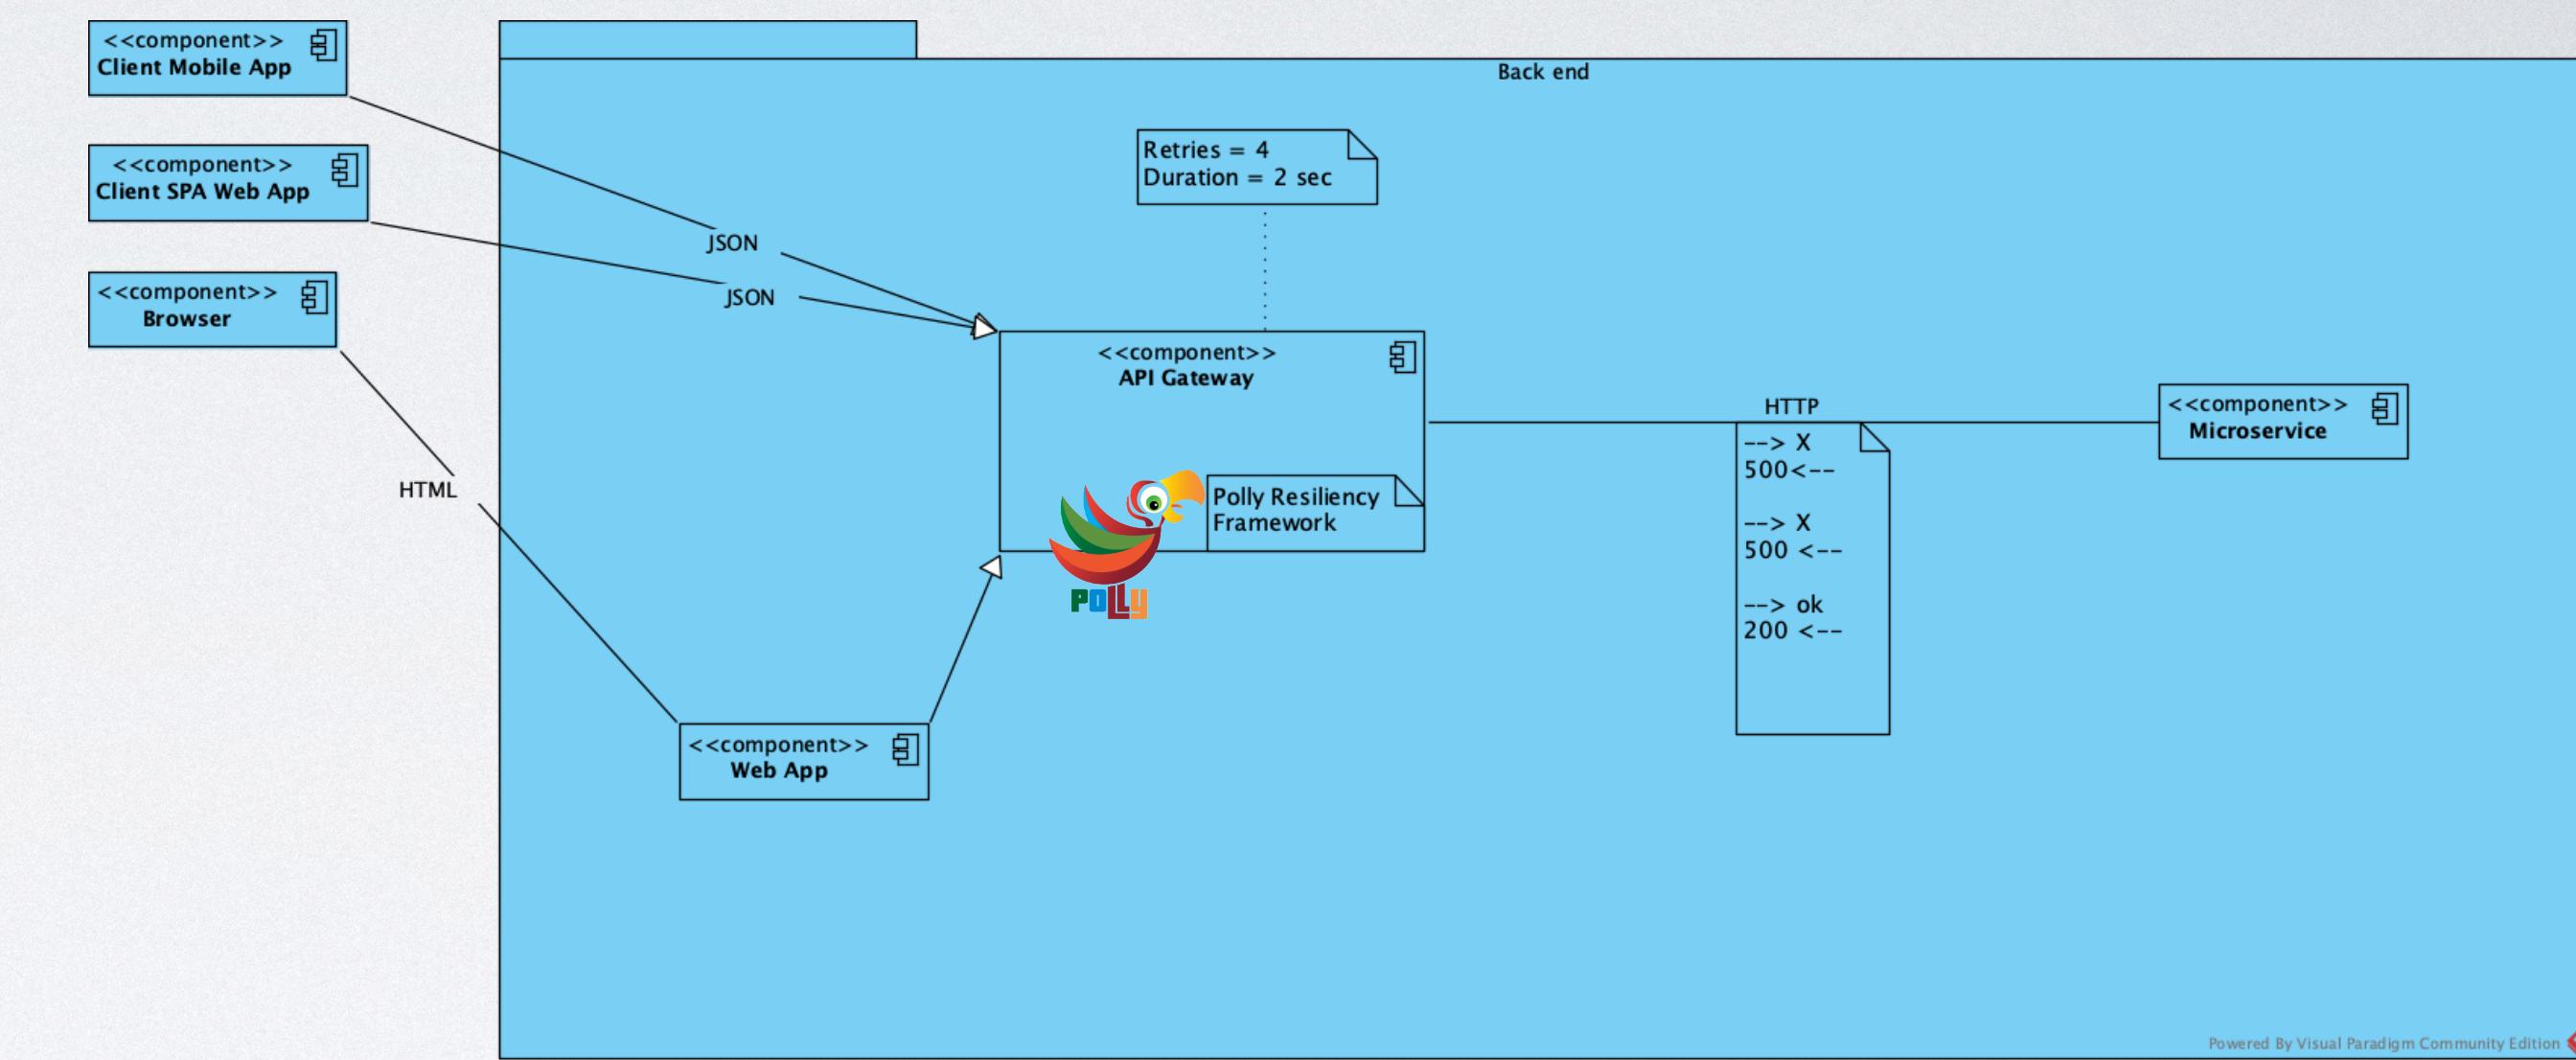
- ⚡ Specific strategy
- ⌚ Issues, external dependencies, etc
- ➡ Retry logic incorporated into the app
- ❤ Enhance system recovery from failures
- 🐆 Without jeopardising functionality



# A DESIGN PATTERN FOR μ-SERVICES



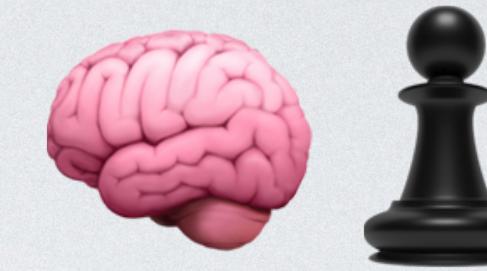
- Common pattern used also in a lot of others environments (e.g. ....)
- # of times it should retry, delay between attempts and actions to take
- Typically self-correcting → if repeated, it's likely to success



Powered By Visual Paradigm Community Edition

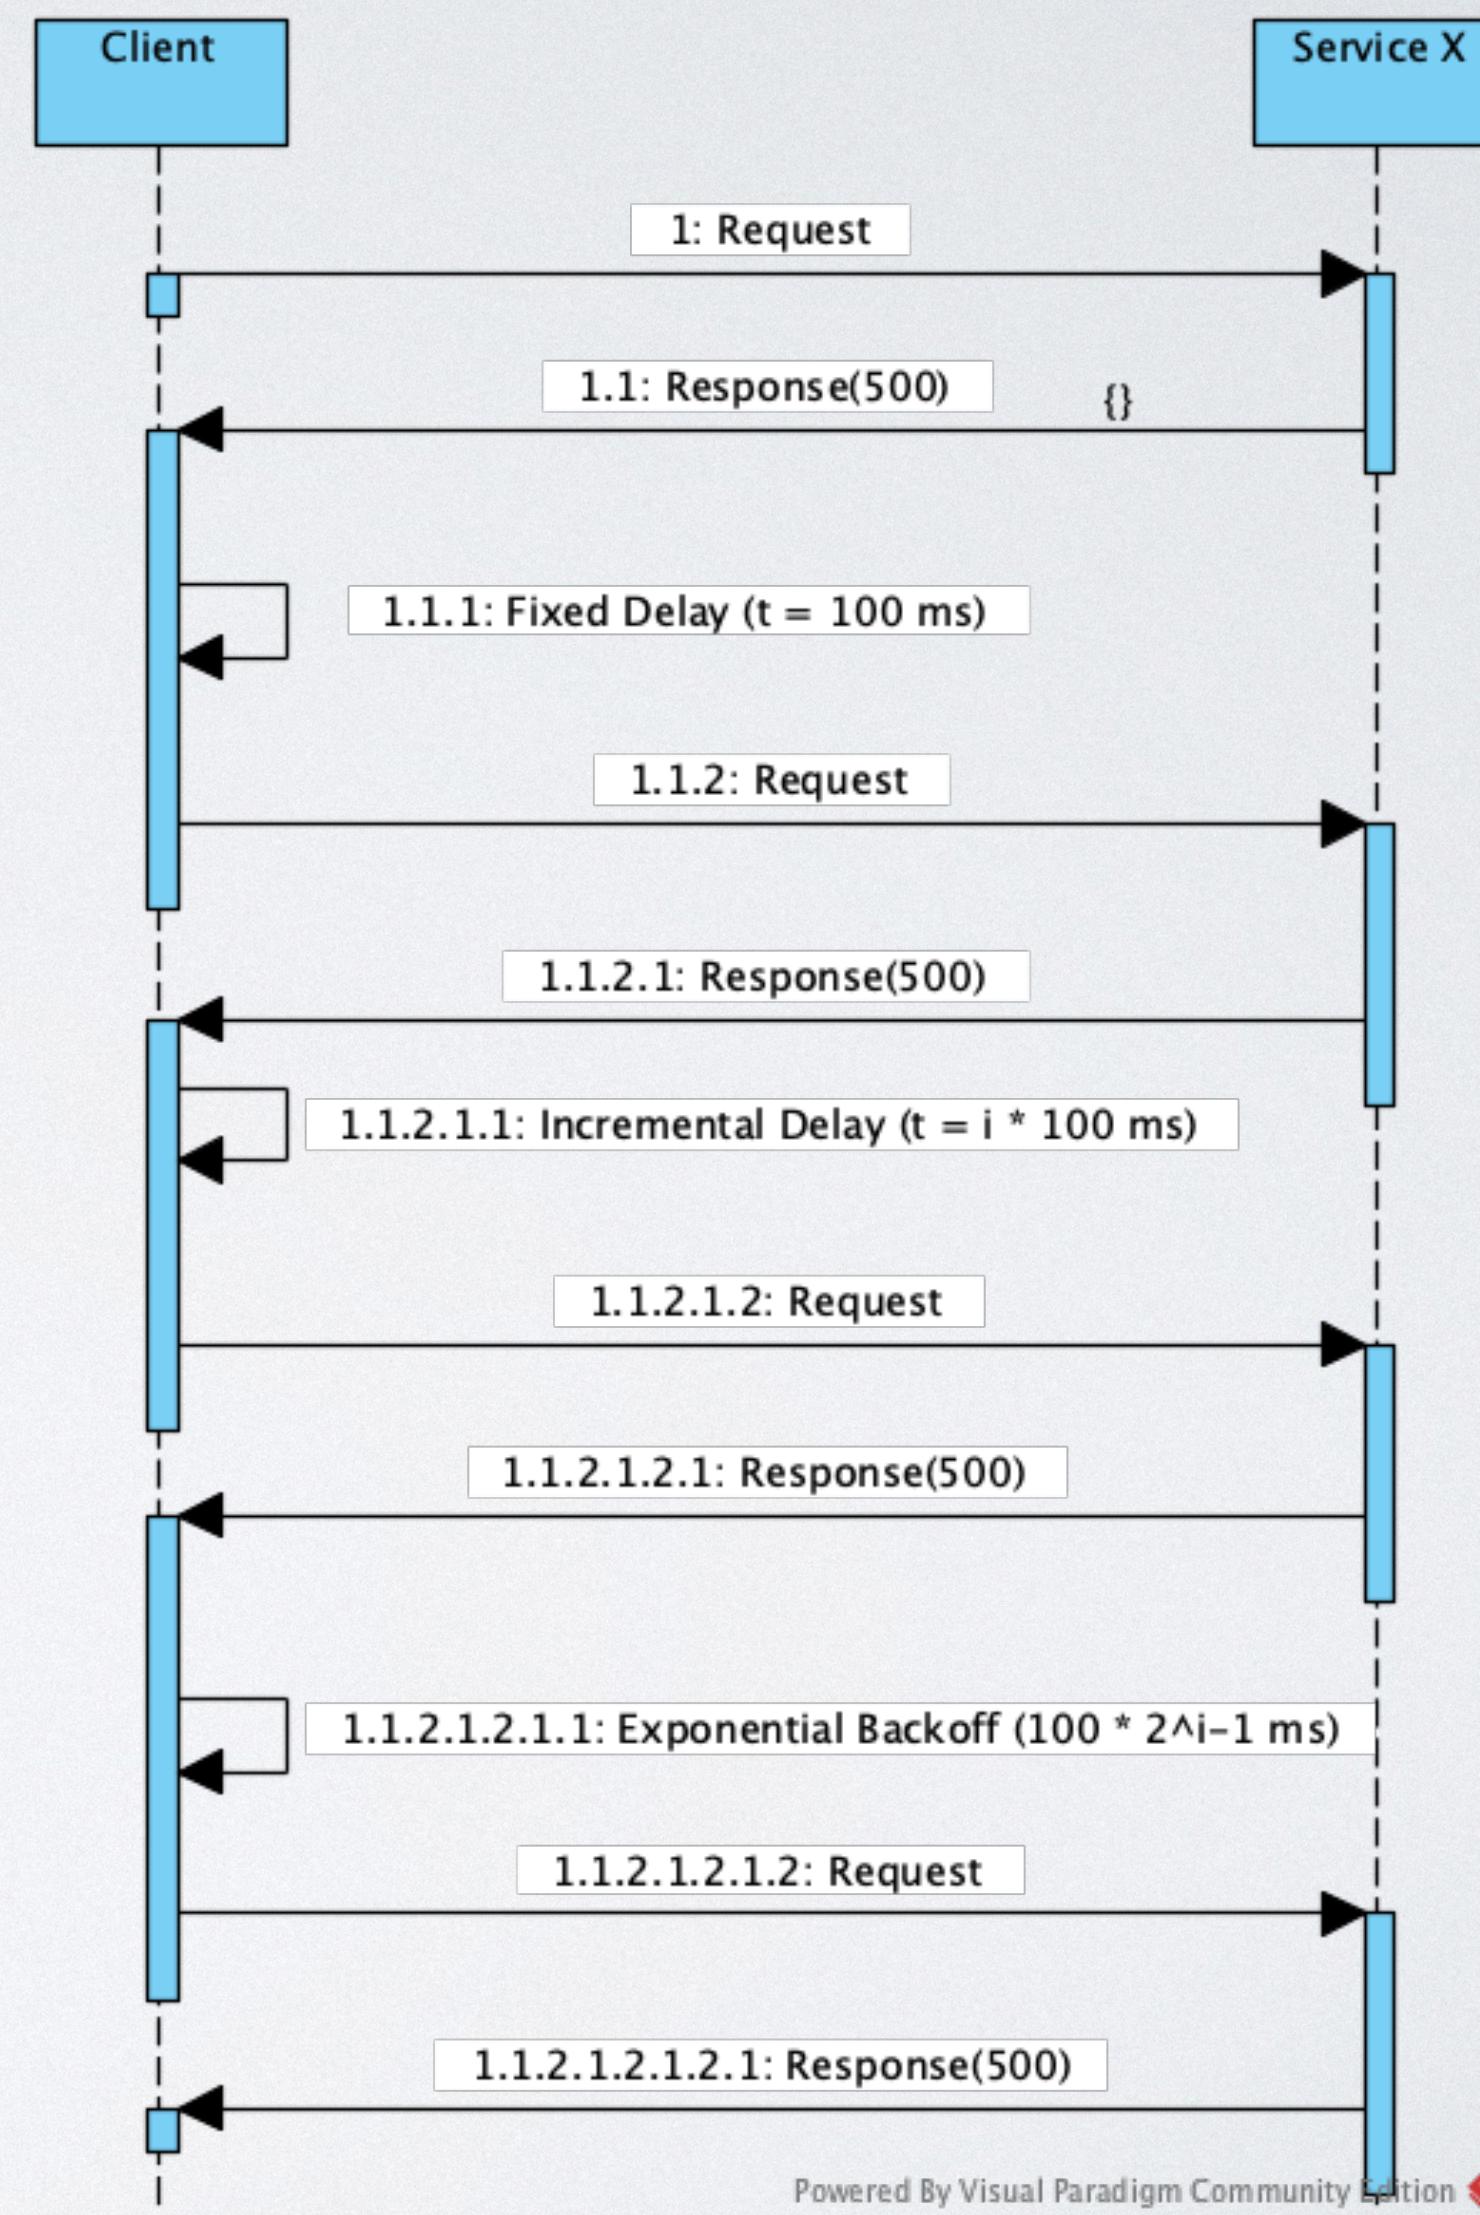
**Keep in mind that an application should be designed to handle these types of situations elegantly and transparently**

# STRATEGIES



- Retry logic:
  1. Retry
  2. Retry after delay
- Exponential backoff
- Circuit breaker
- Cancel

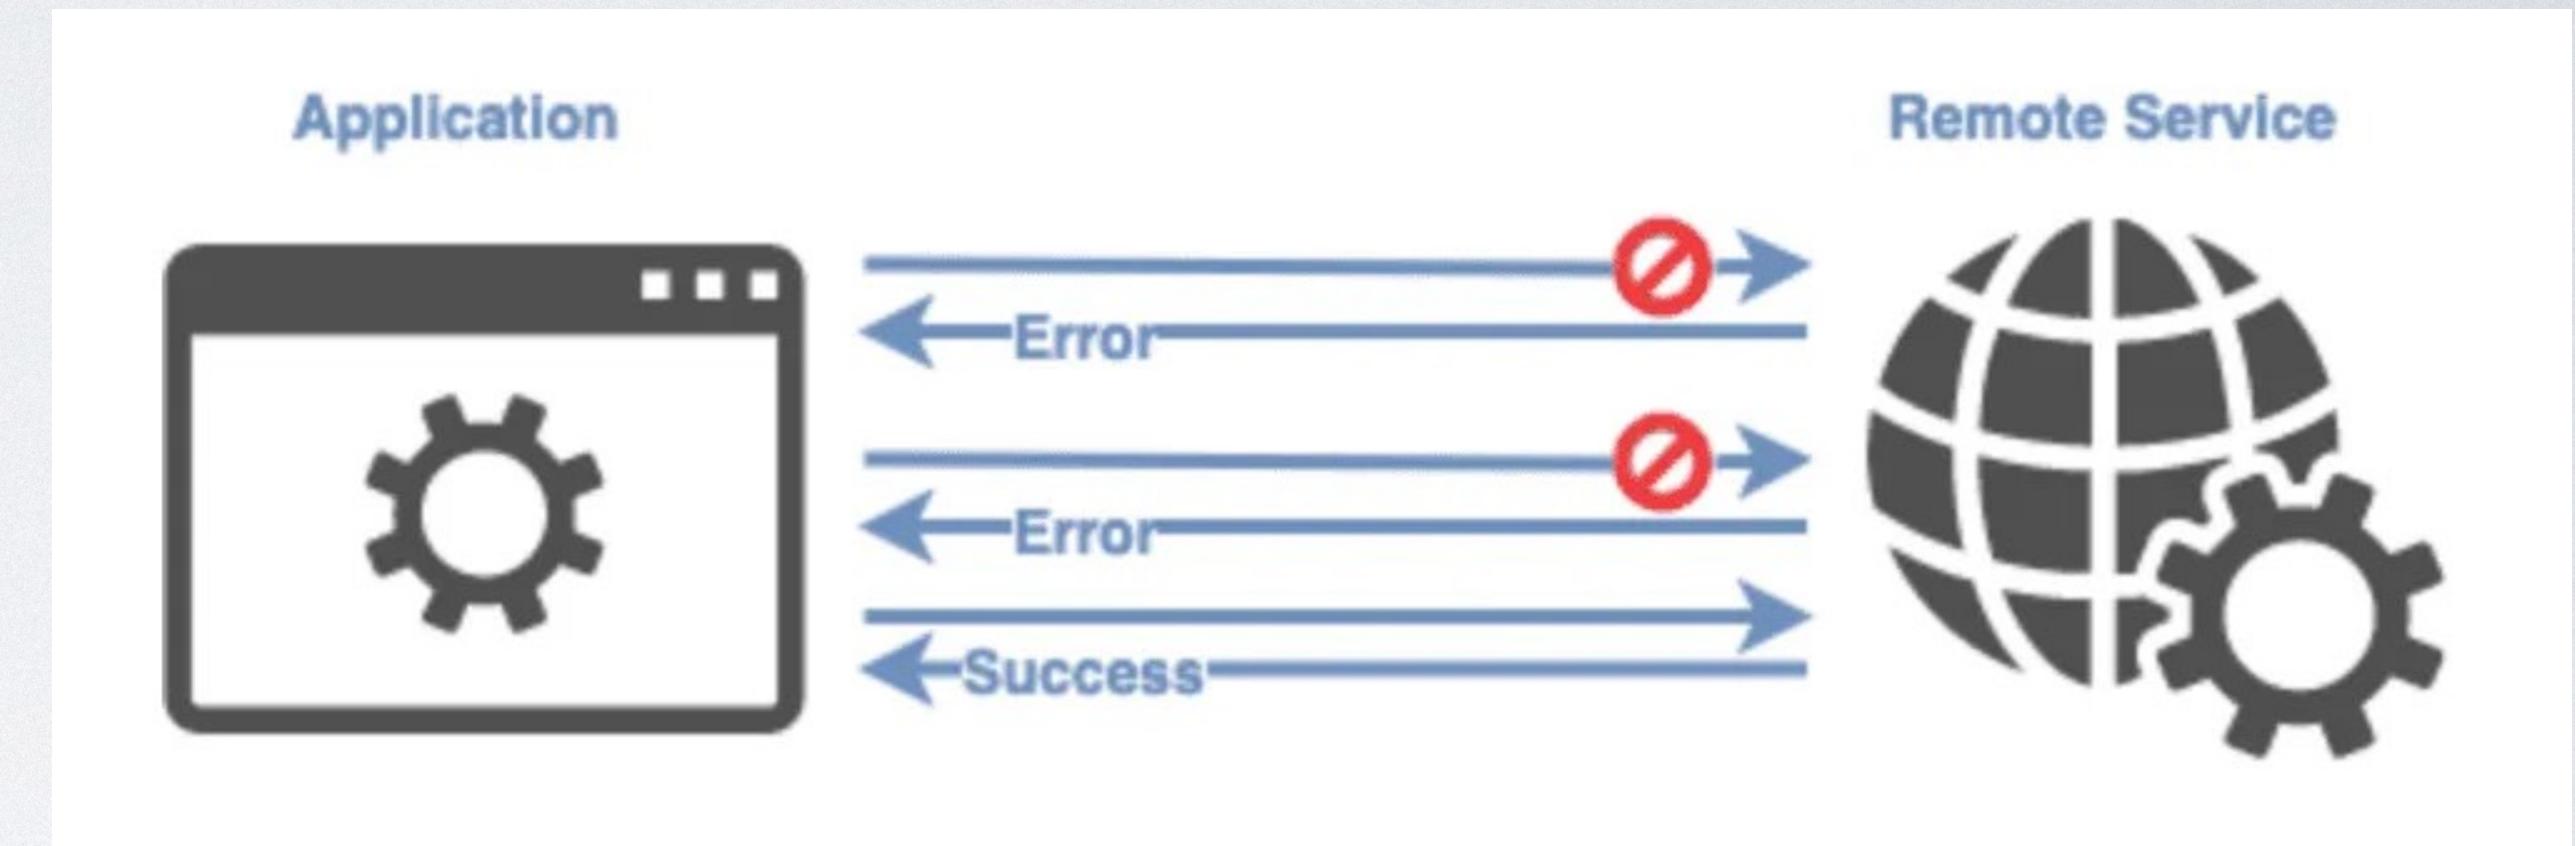
**POV:**  
**if overwhelming → longer to recover**



# ISSUES AND CONSIDERATIONS



- ✖ Without a threshold, a bigger degrade is possible
- ✖ Responsiveness is at risk, if continually trying to perform a failing operation
- Sometimes it's better to fail fast
- ✓ If exhausted resources → better solution to scale out
- ✓ If operation is idempotent, it's safe to retry.  
Otherwise...
- ★ Enhanced reliability & user experience
- ★ Adaptability to dynamic environments



**An idempotent system ensures that the actual execution on this request only happens once (doesn't matter the requests number)**

# “ONE SHOE FITS ALL” SUITABILITY?



✓ Transients faults

✓ Short lived and repeating a (failed) request

✗ Long lasting faults → retrying damage the responsiveness

✗ Non-transient faults ...

✗ Busy faults frequently → better to scale out

...if all you have is a hammer, everything looks like a nail...

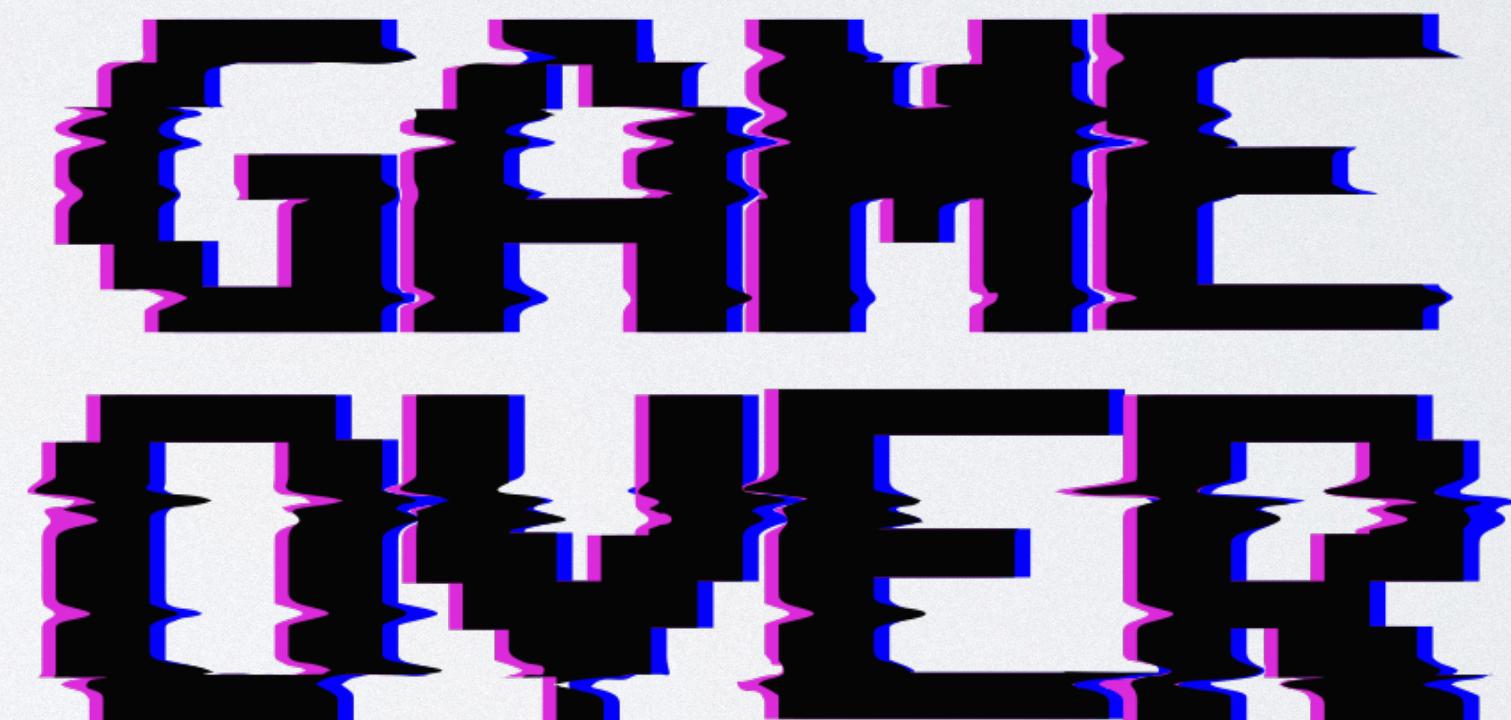
- Failures due to various reasons and various natures
  1. Some exceptions can be resolved quickly
  2. Some are longer lasting

👉 Retry policy to adjust everything

# FAQ



- How do you handle idempotence when using the retry pattern?
- What are common challenges when implementing the retry pattern?
- What monitoring and logging practices are recommended for the retry pattern?
- Are there best practices for implementing the retry pattern in **specific programming languages or frameworks?**



RETRY

```

private int retryCount = 3;
private readonly TimeSpan delay = TimeSpan.FromSeconds(5);

public async Task OperationWithBasicRetryAsync()
{
    int currentRetry = 0;

    for (;;)
    {
        try
        {
            // Call external service.
            await TransientOperationAsync();

            // Return or break.
            break;
        }
        catch (Exception ex)
        {
            Trace.TraceError("Operation Exception");

            currentRetry++;

            // Check if the exception thrown was a transient exception
            // based on the logic in the error detection strategy.
            // Determine whether to retry the operation, as well as how
            // long to wait, based on the retry strategy.
            if (currentRetry > this.retryCount || !IsTransient(ex))
            {
                // If this isn't a transient error or we shouldn't retry,
                // rethrow the exception.
                throw;
            }

            // Wait to retry the operation.
            // Consider calculating an exponential delay here and
            // using a strategy best suited for the operation and fault.
            await Task.Delay(delay);
        }
    }

    // Async method that wraps a call to a remote service (details not shown).
    private async Task TransientOperationAsync()
    {
        ...
    }
}

```

```

private bool IsTransient(Exception ex)
{
    // Determine if the exception is transient.
    // In some cases this is as simple as checking the exception type, in other
    // cases it might be necessary to inspect other properties of the exception.
    if (ex is OperationTransientException)
        return true;

    var webException = ex as WebException;
    if (webException != null)
    {
        // If the web exception contains one of the following status values
        // it might be transient.
        return new[] {WebExceptionStatus.ConnectionClosed,
                     WebExceptionStatus.Timeout,
                     WebExceptionStatus.RequestCanceled }.
            Contains(webException.Status);
    }

    // Additional exception checking logic goes here.
    return false;
}

```



We can model a recoverable scenario by instantiating `FindCustomer` like this:

```
1 final var op = new FindCustomer(  
2     "12345",  
3     new CustomerNotFoundException("not found"),  
4     new CustomerNotFoundException("still not found"),  
5     new CustomerNotFoundException("don't give up yet!")  
6 );
```

In this configuration, `FindCustomer` will throw `CustomerNotFoundException` three times, after which it will consistently return the customer's ID (`12345`).

In our hypothetical scenario, our analysts indicate that this operation typically fails 2-4 times for a given input during peak hours, and that each worker thread in the database subsystem typically needs 50ms to "recover from an error". Applying these policies would yield something like this:

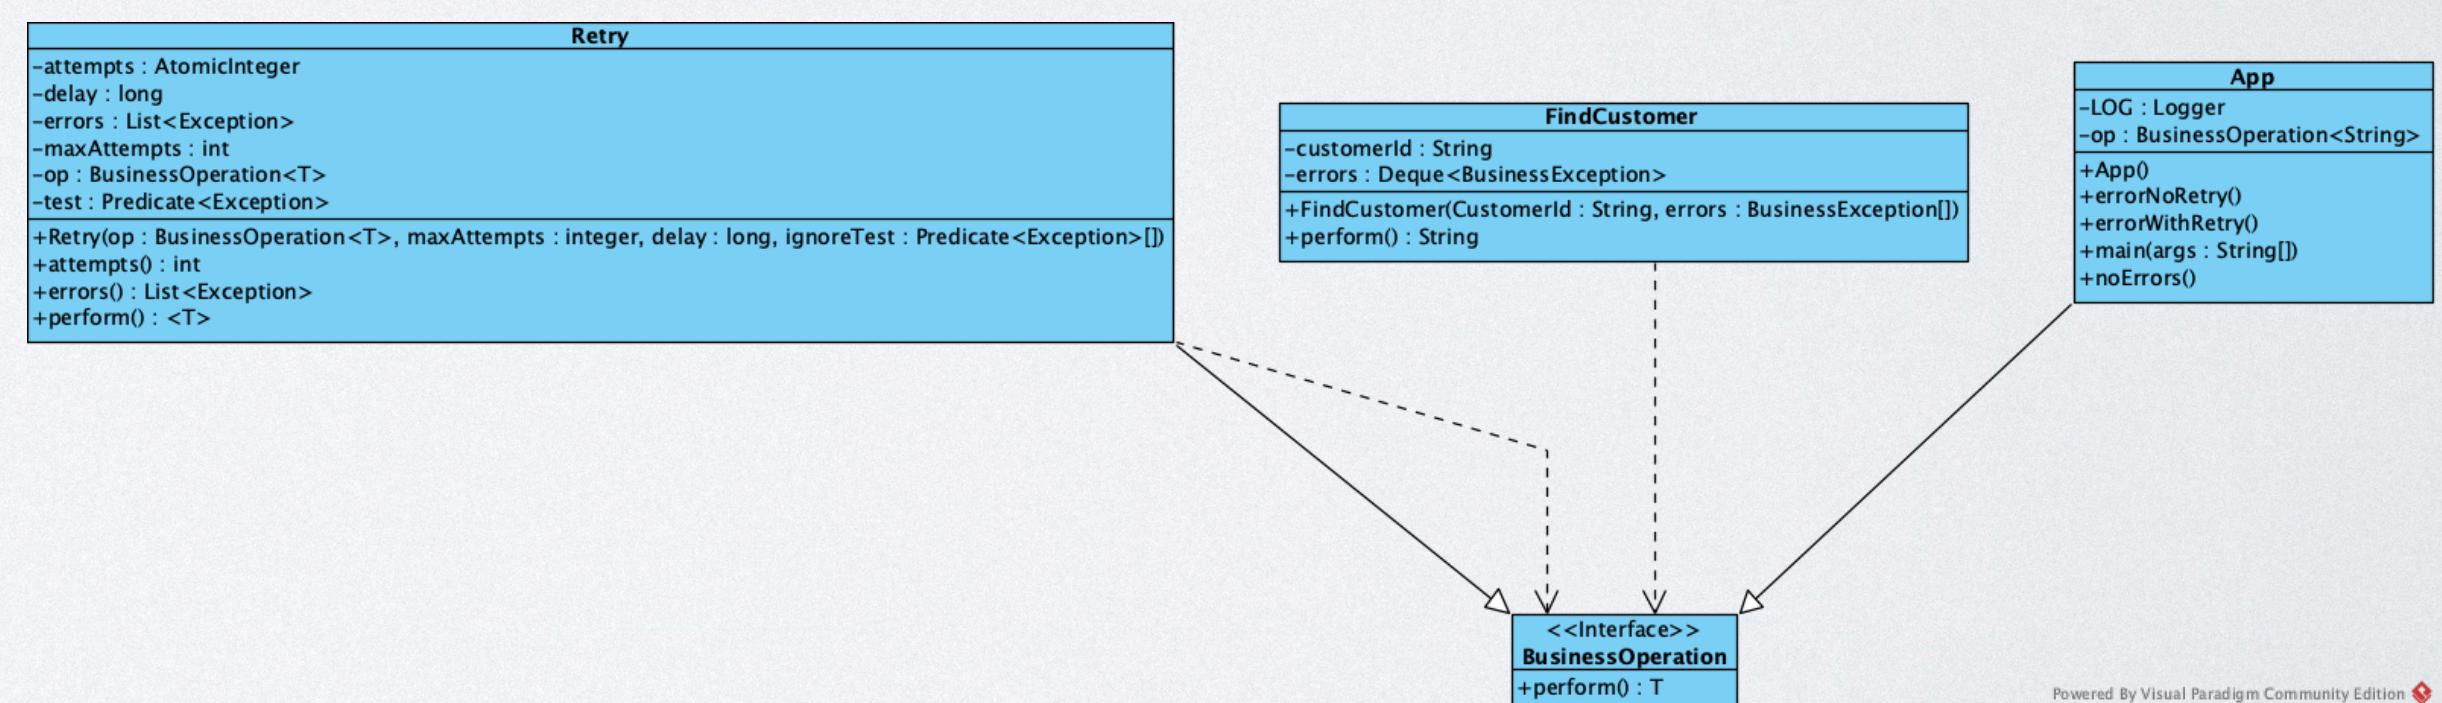
```
1 final var op = new Retry<>(  
2     new FindCustomer(  
3         "1235",  
4         new CustomerNotFoundException("not found"),  
5         new CustomerNotFoundException("still not found"),  
6         new CustomerNotFoundException("don't give up yet!")  
7     ),  
8     5,  
9     100,  
10    e -> CustomerNotFoundException.class.isAssignableFrom(e.getClass())  
11 );
```

In our hypothetical application, we have a generic interface for all operations on remote interfaces.

```
1 public interface BusinessOperation<T> {  
2     T perform() throws BusinessException;  
3 }
```

And we have an implementation of this interface that finds our customers by looking up a database.

```
1 public final class FindCustomer implements BusinessOperation<String> {  
2     @Override  
3     public String perform() throws BusinessException {  
4         ...  
5     }  
6 }
```



Powered By Visual Paradigm Community Edition

```

class RetryTest {

    private var success = "success"

    @Mock
    private lateinit var externalService: ExternalService

    private val logger: LoggerWrapper = LoggerWrapper()
    private lateinit var retry: Retry<String>

    @Before
    fun setUp() {
        MockitoAnnotations.initMocks(this)
    }

    @Test
    fun should_not_retry_when_successful() {
        retry = Retry(logger)

        val result = retry.run { success }

        assertThat(result).isEqualTo(success)
        assertThat(retry.retryCounter.get()).isEqualTo(0)
    }

    @Test
    fun should_retry_once_then_succeed_when_first_fails_and_second_successes() {
        retry = Retry(logger)
        `when`(externalService.run())
            .thenThrow(RuntimeException("Something was wrong..."))
            .thenReturn(success)

        val result = retry.run(externalService::run)

        assertThat(result).isEqualTo(success)
        assertThat(retry.retryCounter.get()).isEqualTo(1)
    }

    @Test
    fun should_throw_exception_when_max_retries() {
        retry = Retry(logger)
        `when`(externalService.run())
            .thenThrow(RuntimeException("Something was wrong..."))
            .thenThrow(RuntimeException("Something was wrong..."))
            .thenThrow(RuntimeException("Something was wrong..."))

        assertThatThrownBy {
            retry.run(externalService::run)
        }.isInstanceOf(RuntimeException::class.java)
            .hasMessage("Command fails on all of 3 retries")
    }
}

```

```

@Configuration
@EnableRetry
class Application {
    @Bean
    fun service(): Service {
        return Service()
    }
}

@Service
class Service {
    @Retryable(maxAttempts = 2, include = [RemoteAccessException::class])
    fun service() {
        // ... do something
        logger.info("Success calling external service")
    }
    @Recover
    fun recover(e: RemoteAccessException) {
        // ... do something when call to service fails
        logger.error("Recover for external service")
    }
}

```

```

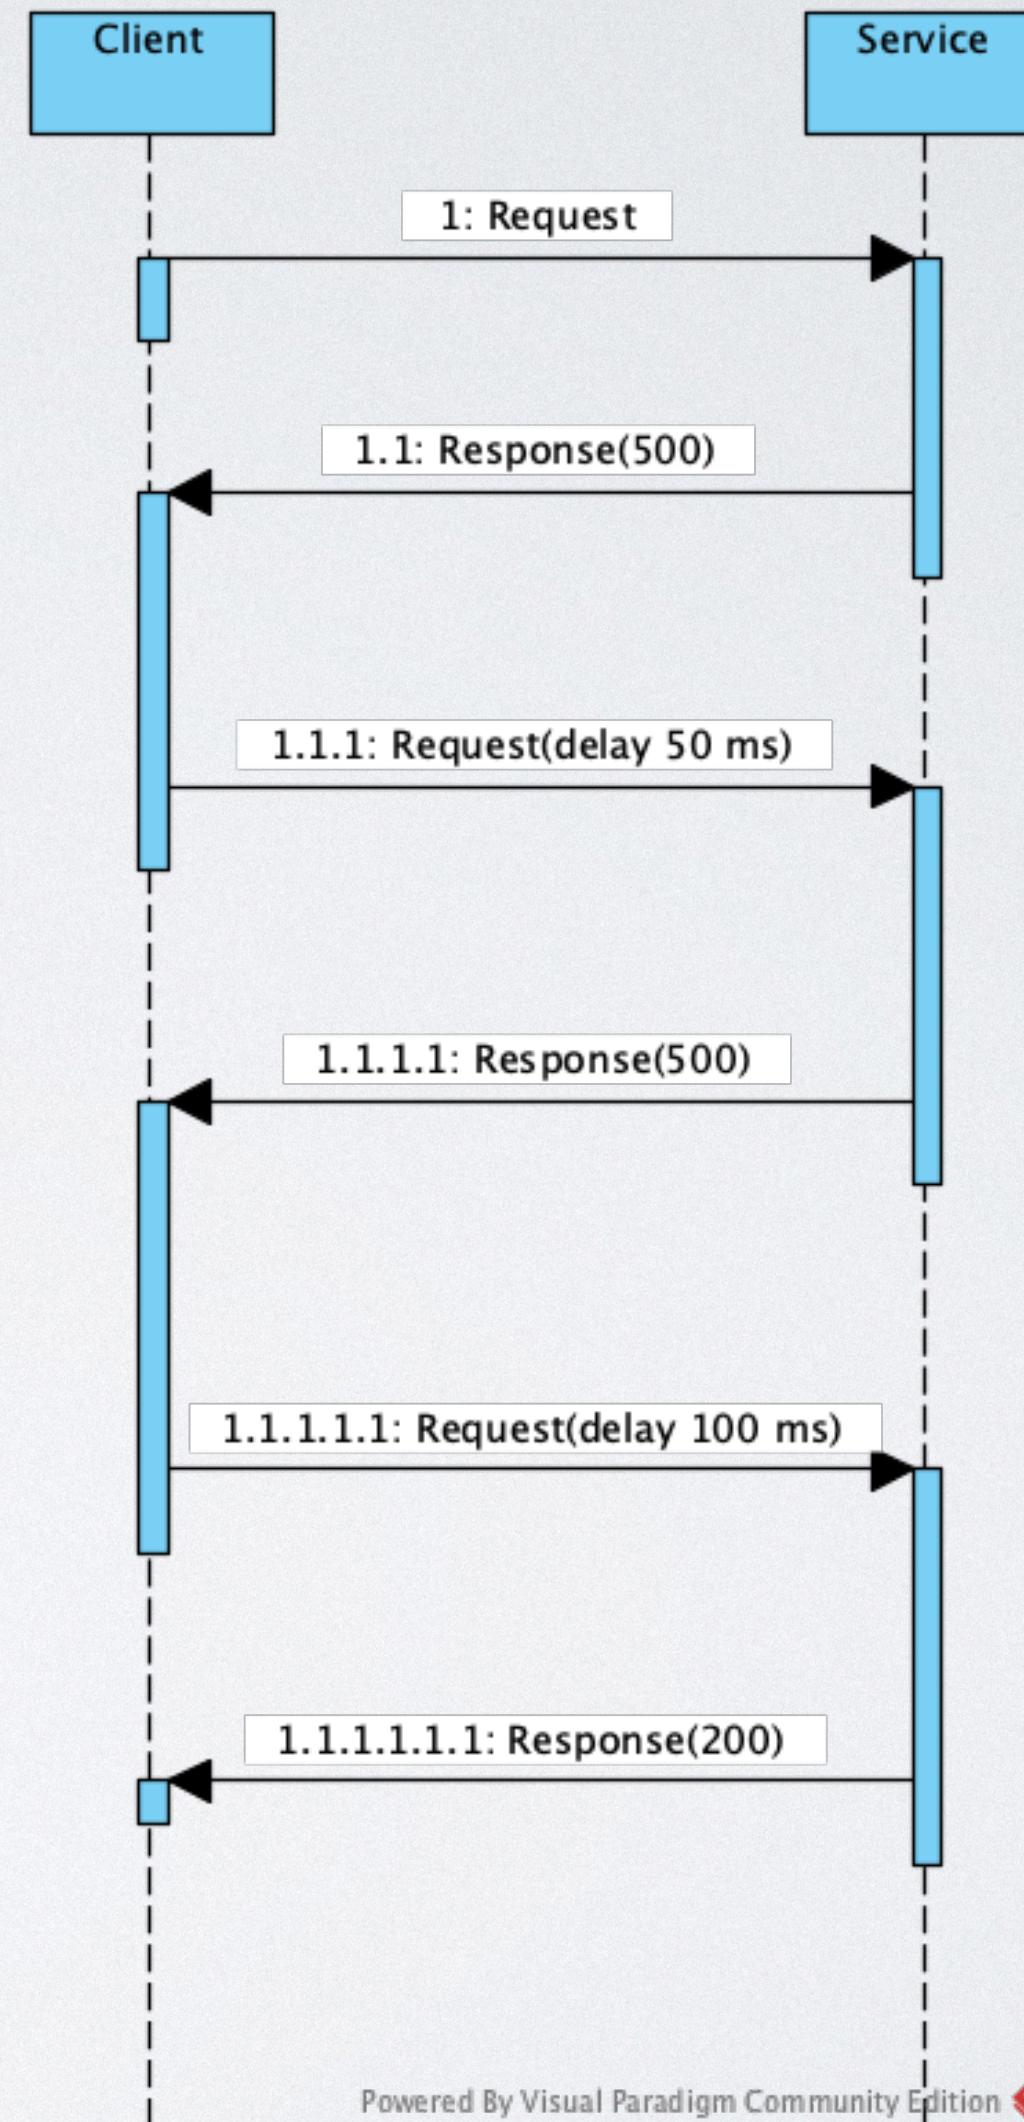
INFO 81668 --- [ Logge : RetryService ] - Calling external service...
INFO 81668 --- [ Logge : RetryService ] - Success calling external service
INFO 81668 --- [ Logge : RetryService ] - Calling external service...
INFO 81668 --- [ Logge : RetryService ] - Success calling external service
INFO 81668 --- [ Logge : RetryService ] - Calling external service...
INFO 81668 --- [ Logge : RetryService ] - Calling external service...
ERROR 81668 --- [Logger : RetryService ] - Recover output

```

# BEST PRACTICES



- ⚠ Log all connectivity failures
- ⚠ Ensure that all retry code is fully tested against a variety of failure conditions
- ⚠ Investigate the faults that are most likely to occur
- ⚠ Avoid amplifying retries at multiple levels
- ⚠ Consider a server-wide retry budget
- ⚠ Ensure retries aren't done on operations where consumers are not waiting for response
- ⚠ Never wait synchronously between retry attempts. For example, use “`Task.Delay`” in place of “`Thread.Sleep`” in .Net



# CONCLUSIONS

←  
END

- Transient faults with a remote service or resources → do take care of what the error code is before retries
- We don't want to be victim of retry anti-patterns
- Incorporating intelligent retry logic, exponential backoff and circuit breaker mechanism, developers can create robust systems capable of recovering from unexpected hiccups
- As technology continues to evolve, the RP remains a valuable tool of developers



# REFERENCES



- Online post:
  - Author: Vijaya Lakshmi, 2022
  - Title: Design patterns for micro-services
  - Website: <https://www.linkedin.com/pulse/vijaya-lakshmi-v/>
- Medium article:
  - Author: Harish Bhattbhatt, 2021
  - Title: Best practices for retry pattern
  - Website: <https://harish-bhattbhatt.medium.com/best-practices-for-retry-pattern-f29d47cd5117>
- Microsoft Article:
  - Year: 2023
  - Title: Transient fault handling
  - Website: <https://learn.microsoft.com/en-us/azure/architecture/best-practices/transient-faults>
- Google Article:
  - Author: Mike Ulrich, 2017
  - Title: Addressing Cascading Failures
  - Website: <https://sre.google/sre-book/addressing-cascading-failures/>
- Microsoft Documentation:
  - Title: Retry pattern
  - Website: <https://learn.microsoft.com/en-us/azure/architecture/patterns/retry>
- Journal article:
  - Author: Kamini Kamal, 2023
  - Title: Understanding the Retry Pattern
  - Journal: Medium
  - Website: <https://blog.bitsrc.io/retry-pattern-3f7221913e41?gi=161787a102da>
- AWS Article:
  - Author: Marc Brooker, 2015
  - Title: Addressing Cascading Failures
  - Website: <https://aws.amazon.com/blogs/architecture/exponential-backoff-and-jitter/>



